

Probabilistic Calling Context

Michael D. Bond Kathryn S. McKinley
University of Texas at Austin



Why Context Sensitivity?

- Static program location not enough

```
at com.mckoi.db.jdbcserver.JDBCInterface.executeQuery():213
```

Why Context Sensitivity?

- Static program location not enough

```
at com.mckoi.db.jdbcserver.JDBCInterface.executeQuery():213
at com.mckoi.db.jdbc.MConnection.executeQuery():348
at com.mckoi.db.jdbc.MStatement.executeQuery():110
at com.mckoi.db.jdbc.MStatement.executeQuery():127
at Test.main():48
```

Why Context Sensitivity?

- Static program location not enough

```
at com.mckoi.db.jdbcserver.JDBCInterface.executeQuery():213
at com.mckoi.db.jdbc.MConnection.executeQuery():348
at com.mckoi.db.jdbc.MStatement.executeQuery():110
at com.mckoi.db.jdbc.MStatement.executeQuery():127
at Test.main():48
```

- Motivated by

- Complex programs
- Small methods
- Virtual dispatch


Why Context Sensitivity?

■ Static program location not enough

```
at com.mckoi.db.jdbcserver.JDBCInterface.executeQuery():213
at com.mckoi.db.jdbc.MConnection.executeQuery():348
at com.mckoi.db.jdbc.MStatement.executeQuery():110
at com.mckoi.db.jdbc.MStatement.executeQuery():127
at Test.main():48
```

■ Motivated by

- ❑ Complex programs
- ❑ Small methods
- ❑ Virtual dispatch



Context Is Nontrivial

Program	API calls	
	Call sites	Distinct contexts
antlr	4,184	128,627
bloat	3,306	600,947
chart	2,335	202,603
eclipse	9,611	226,020
fop	2,225	37,710
hsqldb	947	16,050
jython	1,830	628,048
luindex	654	102,556
lusearch	507	905
pmd	1,890	847,108
xalan	1,530	17,905

Example: Residual Testing

Does behavior occur at production time that did not occur at testing time?

```
class SimpleWindow {  
    close() {  
        ...  
    }  
}
```

```
class EditorWindow {  
    close() {  
        ...  
    }  
}
```

Example: Residual Testing

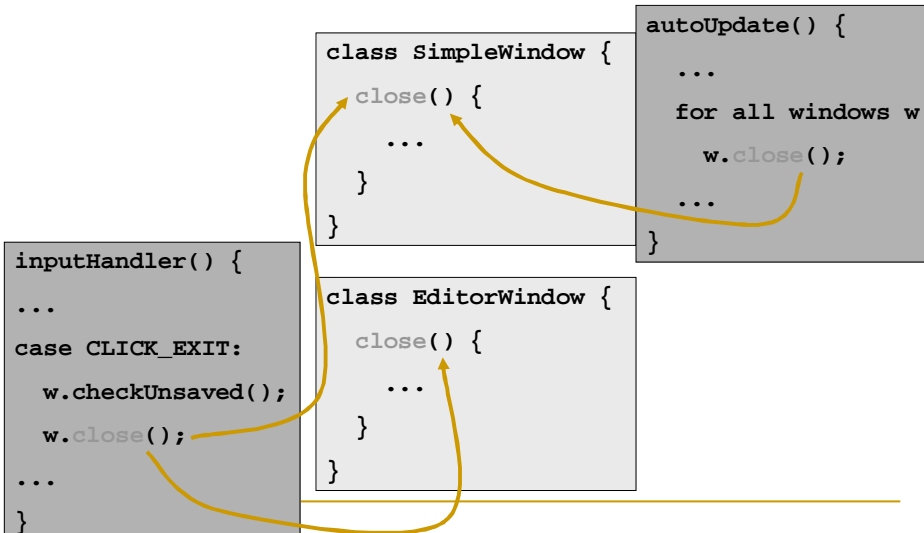
Does behavior occur at production time that did not occur at testing time?

```
inputHandler() {  
    ...  
    case CLICK_EXIT:  
        w.checkUnsaved();  
        w.close();  
    ...  
}
```

```
class SimpleWindow {  
    close() {  
        ...  
    }  
}
```

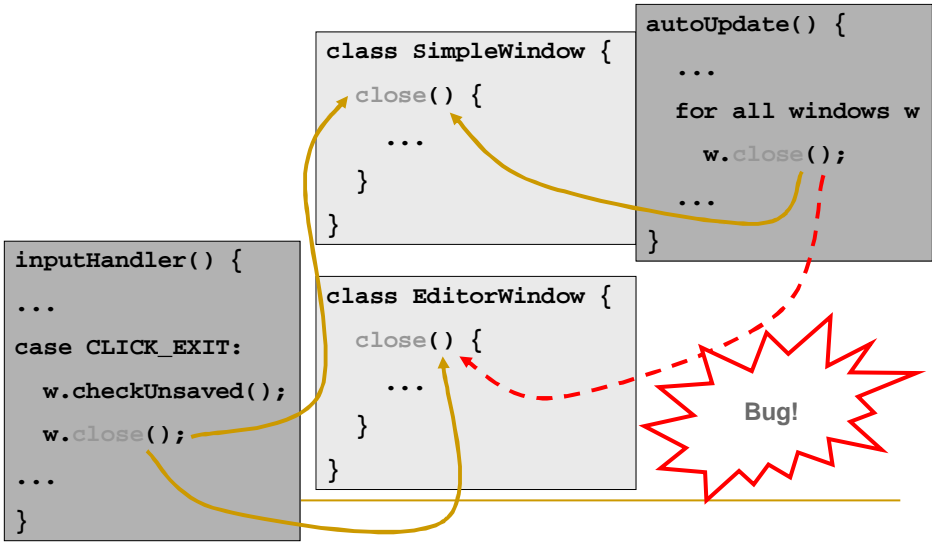
```
class EditorWindow {  
    close() {  
        ...  
    }  
}
```

```
autoUpdate() {  
    ...  
    for all windows w  
        w.close();  
    ...  
}
```



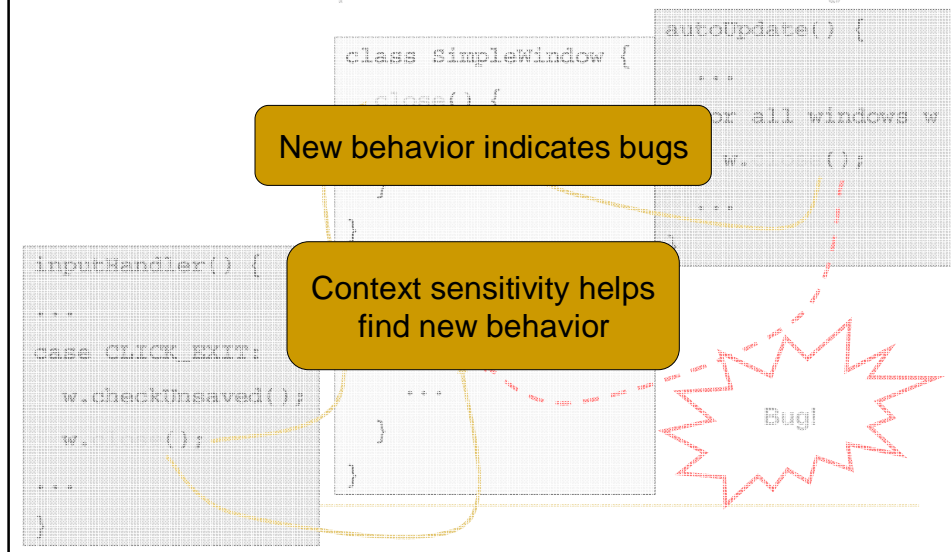
Example: Residual Testing

Does behavior occur at production time that did not occur at testing time?




Example: Residual Testing

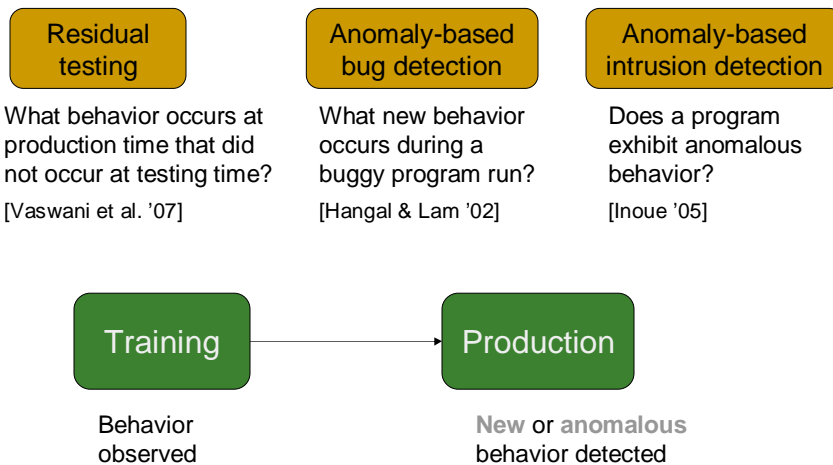
Does behavior occur at production time that did not occur at testing time?



Two-Phase Dynamic Analyses




Two-Phase Dynamic Analyses



Probabilistic Calling Context

- Adds context sensitivity to dynamic analyses
- Maintains value representing context
 - Unique with high probability
 - New value → new context → walk stack
- High accuracy: <0.1% false negatives
- Low overhead: 3% overhead, 0-8% for clients



Outline

- Introduction
- Previous approaches
- Maintaining the PCC value
- Evaluation
 - Accuracy
 - Performance

Previous Approaches

- **Tracking context** [Ammons et al. '97] [Spivey '04]
 - Maintain CCT position at each call/return
- **Walking the stack** [Nethercote & Seward '07]
- **Path profiling** [Ball & Larus '96] [Melski & Reps '99]
 - Call graphs large → path explosion
 - Virtual dispatch complicates instrumentation

Previous Approaches

- **Tracking context** [Ammons et al. '97] [Spivey '04]
 - Maintain CCT position at each call/return
- **Walking the stack** [Nethercote & Seward '07]
- **Path profiling** [Ball & Larus '96] [Melski & Reps '99]
 - Call graphs large → path explosion
 - Virtual dispatch complicates instrumentation
- **Sampling** [Zhuang et al. '06]
 - Sacrifices coverage for low overhead

Outline

- Introduction
- Previous approaches
- Maintaining the PCC value
- Evaluation
 - Accuracy
 - Performance

PCC Function

$f(V, cs)$


- **V** is PCC value
- **cs** is call site ID

PCC Function

$f(V, cs)$

□ V is PCC value

□ cs is call site ID



PCC Function

$f(V, cs) \equiv 3V + cs \pmod{2^{32}}$

□ V is PCC value


□ cs is call site ID

PCC Function


$$f(V, cs) \equiv 3V + cs \pmod{2^{32}}$$

- Motivated by MPI datatype hashing
[Langou et al. '05] [Gropp '00]
- Cheap to compute
- Desirable properties:
 - Non-commutative
 - Composition efficient to compute

Differentiating Similar Contexts




Differentiating Similar Contexts




- Non-commutative

$$f(f(V, cs_1), cs_2) \neq f(f(V, cs_2), cs_1)$$


Efficiency at Inlined Calls




Efficiency at Inlined Calls



Efficiency at Inlined Calls



Efficiency at Inlined Calls



- Composition efficient to compute

$$f^n(V, cs_i) = 3^n V + \sum_i 3^i cs_i$$

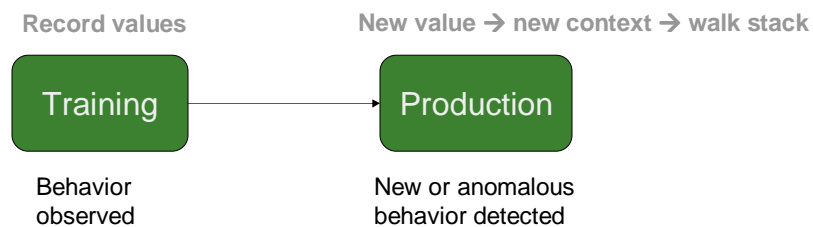
Outline

- Introduction
- Previous approaches
- Maintaining the PCC value
- Evaluation
 - Methodology
 - Evaluating potential clients
 - Accuracy
 - Performance

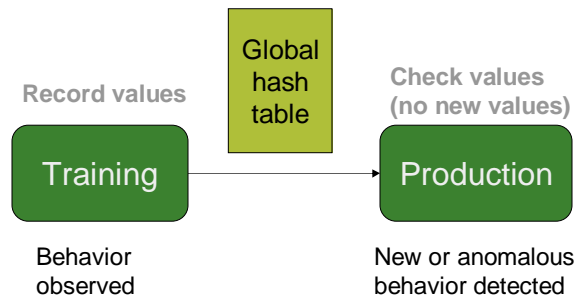
Methodology

- Implementation in Jikes RVM 2.4.6
 - Available on [Jikes RVM Research Archive](#)
- Deterministic calling context profiling
 - Maintains CCT node at each call & return
- Benchmarks: DaCapo, SPEC JBB2000, SPEC JVM98
- Platform: 3.6 GHz Pentium 4 w/Linux

How Clients Use PCC

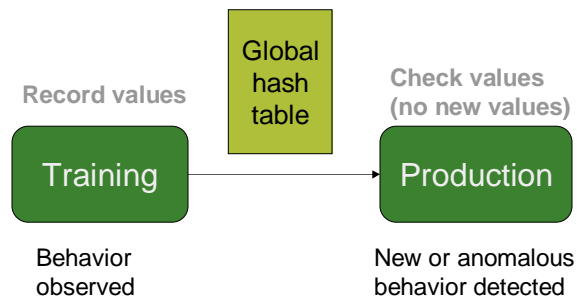


Evaluating Potential Clients



Evaluating Potential Clients

Memory overhead:
proportional to contexts



Evaluating Potential Clients

Anomaly-based intrusion detection

Check PCC value at
system calls
(Network, I/O, OS)

Residual testing

Check PCC value at
Java API calls
(calls to `java.*`)

Upper bound

Check PCC value at
all calls

Ideal Accuracy

- PCC maps context to value
 - New PCC value → new context
 - Familiar PCC value → probably familiar context

Ideal Accuracy

- PCC maps context to value
 - New PCC value → new context
 - Familiar PCC value → probably familiar context

Distinct contexts	Expected conflicts (false negatives)	
	32-bit values	64-bit values
100,000	1 (0.0%)	0 (0.0%)
1,000,000	116 (0.0%)	0 (0.0%)
10,000,000	11,632 (0.1%)	0 (0.0%)
100,000,000	1,155,170 (1.2%)	0 (0.0%)
1,000,000,000	107,882,641 (10.8%)	0 (0.0%)
10,000,000,000	6,123,623,065 (61.2%)	3 (0.0%)

Ideal Accuracy

- PCC maps context to value
 - New PCC value → new context
 - Familiar PCC value → probably familiar context

Distinct contexts	Expected conflicts (false negatives)	
	32-bit values	64-bit values
100,000	1 (0.0%)	0 (0.0%)
1,000,000	116 (0.0%)	0 (0.0%)
10,000,000	11,632 (0.1%)	0 (0.0%)
100,000,000	1,155,170 (1.2%)	0 (0.0%)
1,000,000,000	107,882,641 (10.8%)	0 (0.0%)
10,000,000,000	6,123,623,065 (61.2%)	3 (0.0%)

Ideal Accuracy

- PCC maps context to value
 - New PCC value → new context
 - Familiar PCC value → probably familiar context

	Expected conflicts (false negatives)	
Distinct contexts	32-bit values	64-bit values
100,000	1 (0.0%)	0 (0.0%)
1,000,000	116 (0.0%)	0 (0.0%)
10,000,000	11,632 (0.1%)	All calls %
100,000,000	1,155,170 (1.2%)	0 (0.0%)
1,000,000,000	107,882,641 (10.8%)	0 (0.0%)
10,000,000,000	6,123,623,065 (61.2%)	3 (0.0%)

Ideal Accuracy

- PCC maps context to value
 - New PCC value → new context
 - Familiar PCC value → probably familiar context

	Expected conflicts (false negatives)	
Distinct contexts	32-bit values	64-bit values
100,000	1 (0.0%)	0 (0.0%)
1,000,000	116 (0.0%)	0 (0.0%)
10,000,000	11,632 (0.1%)	0 (0.0%)
100,000,000	1,155,170 (1.2%)	0 (0.0%)
1,000,000,000	107,882,641 (10.8%)	0 (0.0%)
10,000,000,000	6,123,623,065 (61.2%)	3 (0.0%)

PCC's Accuracy

Program	System calls		
	Dynamic	Distinct	Conf.
antlr	211,490	1,567	0
bloat	12	10	0
chart	63	62	0
eclipse	14,110	197	0
fop	18	17	0
hsqldb	12	12	0
jython	5,929	4,289	0
luindex	2,615	14	0
lusearch	141	11	0
pmd	1,045	25	0
xalan	137,895	59	0


PCC's Accuracy

Program	System calls			Java API calls		
	Dynamic	Distinct	Conf.	Dynamic	Distinct	Conf.
antlr	211,490	1,567	0	24,422,013	128,627	3
bloat	12	10	0	1,159,281,573	600,947	40
chart	63	62	0	258,891,525	202,603	4
eclipse	14,110	197	0	132,507,343	226,020	5
fop	18	17	0	9,918,275	37,710	0
hsqldb	12	12	0	81,161,541	16,050	0
jython	5,929	4,289	0	543,845,772	628,048	48
luindex	2,615	14	0	39,733,214	102,556	0
lusearch	141	11	0	113,511,311	905	0
pmd	1,045	25	0	537,017,118	847,108	79
xalan	137,895	59	0	2,105,838,670	17,905	0

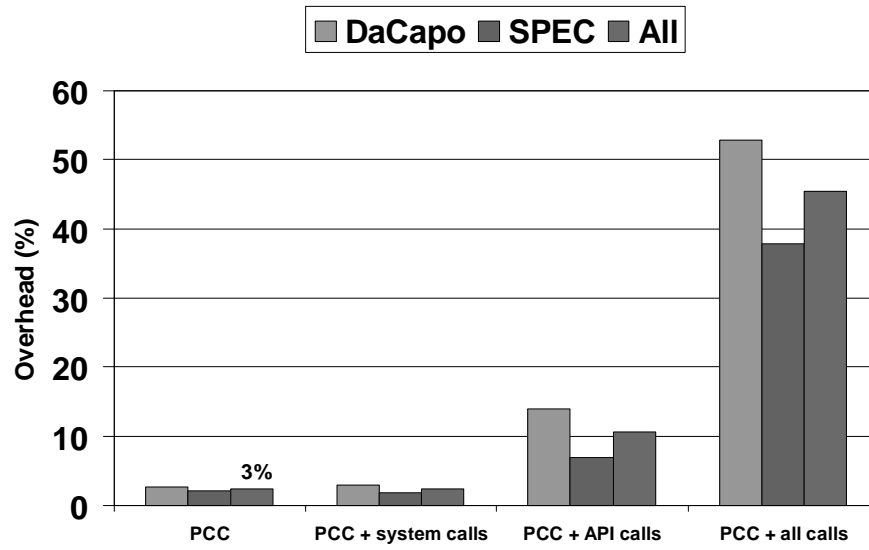
PCC's Accuracy

Program	All calls		
	Dynamic	Distinct	Conf.
antlr	490,363,211	1,006,578	118
bloat	6,276,446,059	1,980,205	453
chart	908,459,469	845,432	91
eclipse	1,266,810,504	4,815,901	2,652
fop	44,200,446	174,955	2
hsqldb	877,680,667	110,795	1
jython	5,326,949,158	3,859,545	1,738
luindex	740,053,104	374,201	12
lusearch	1,439,034,336	6,039	0
pmd	2,726,876,957	8,043,096	7,653
xalan	10,083,858,546	163,205	6

PCC's Execution Time Overhead



PCC's Execution Time Overhead



Summary

- PCC maintains calling context value
 - New value indicates new behavior
- Low overhead
 - Maintaining PCC value adds 3%
 - Checking PCC value 0-8%
 - Memory overhead proportional to contexts
- High accuracy
 - Less than 0.1% false negative rate
- PCC adds context sensitivity to clients that detect anomalous behavior

Summary

Thank you!

- PCC maintains calling context value
 - New value indicates new behavior
- Low overhead
 - Maintaining PCC value adds 3%
 - Checking PCC value 0-8%
 - Memory overhead proportional to contexts
- High accuracy
 - Less than 0.1% false negative rate
- PCC adds context sensitivity to clients that detect anomalous behavior

Extra slides

Context Sensitivity Mostly Unused

- Do paths capture enough behavior?

