


# Efficient Deterministic Replay of Multithreaded Executions in a Managed Language Virtual Machine



# Nondeterminism is problematic

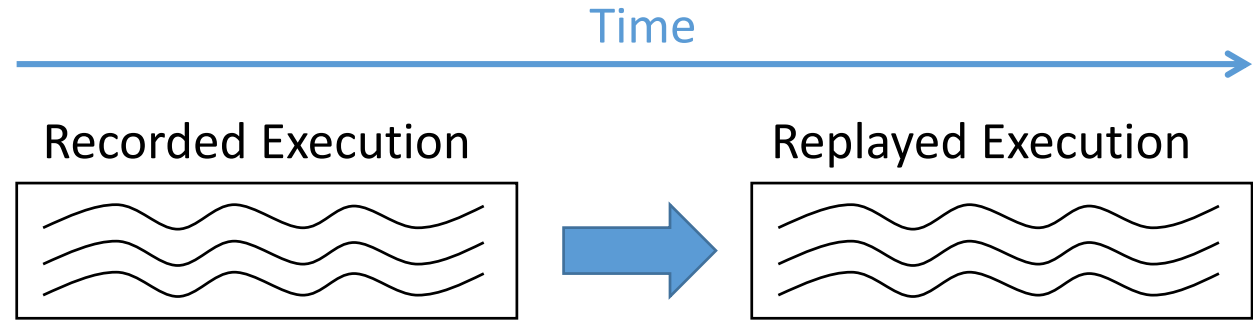
- Reproduce and Debug
  - Replication for fault-tolerance
  - Record and Replay!
    - RecPlay, M. Ronsse, et al, 1999
    - Respec, D. Lee, et al, 2010
    - DoublePlay, K. Veeraraghavan, et al, 2011
    - Chimera, D. Lee, et al, 2012
    - CLAP, J. Huang, et al, 2013
- And many others...



# Record and Replay Types


- Offline

- Debugging



- Online

- Fault tolerance
- Distribute dynamic analysis



# Record and Replay Challenges

- Single-thread, *external* sources of nondeterminism
  - I/O, time, SysCall, etc.
  - Garbage collection, hash code
  - Adaptive compilation and dynamic classloading
    - Even harder for metacircular JVM (Jikes RVM)
- Multithreaded, *internal* nondeterminism
  - Thread interleaving
  - Hard and expensive to capture or control

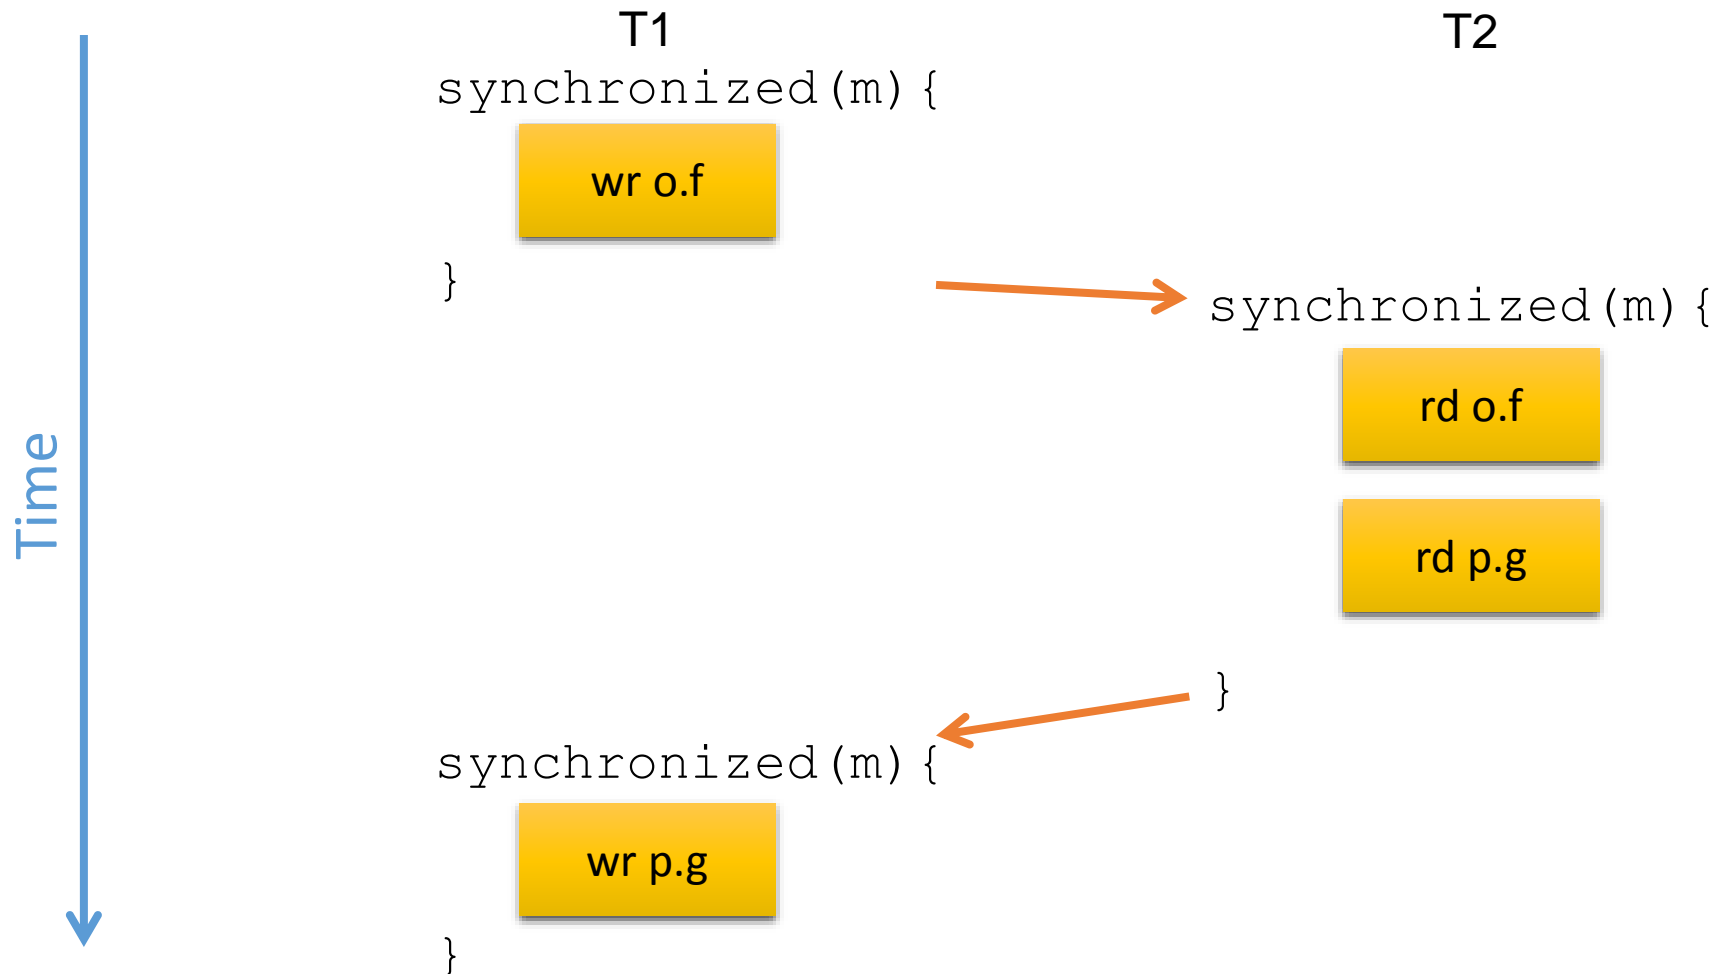
# Contributions

Handle both nondeterminisms

- External (VM, system, I/O)
  - Non-trivial engineering effort
  - Novel methodology to sidestep nondeterminisms
    - Fork-and-recompile
- Internal (multithreading)
  - Two dynamic analyses
    - RECORD
    - REPLAY
  - Low overhead
  - Fewer limitations

# Handling internal nondeterminism

# Easy case: data-race-free execution

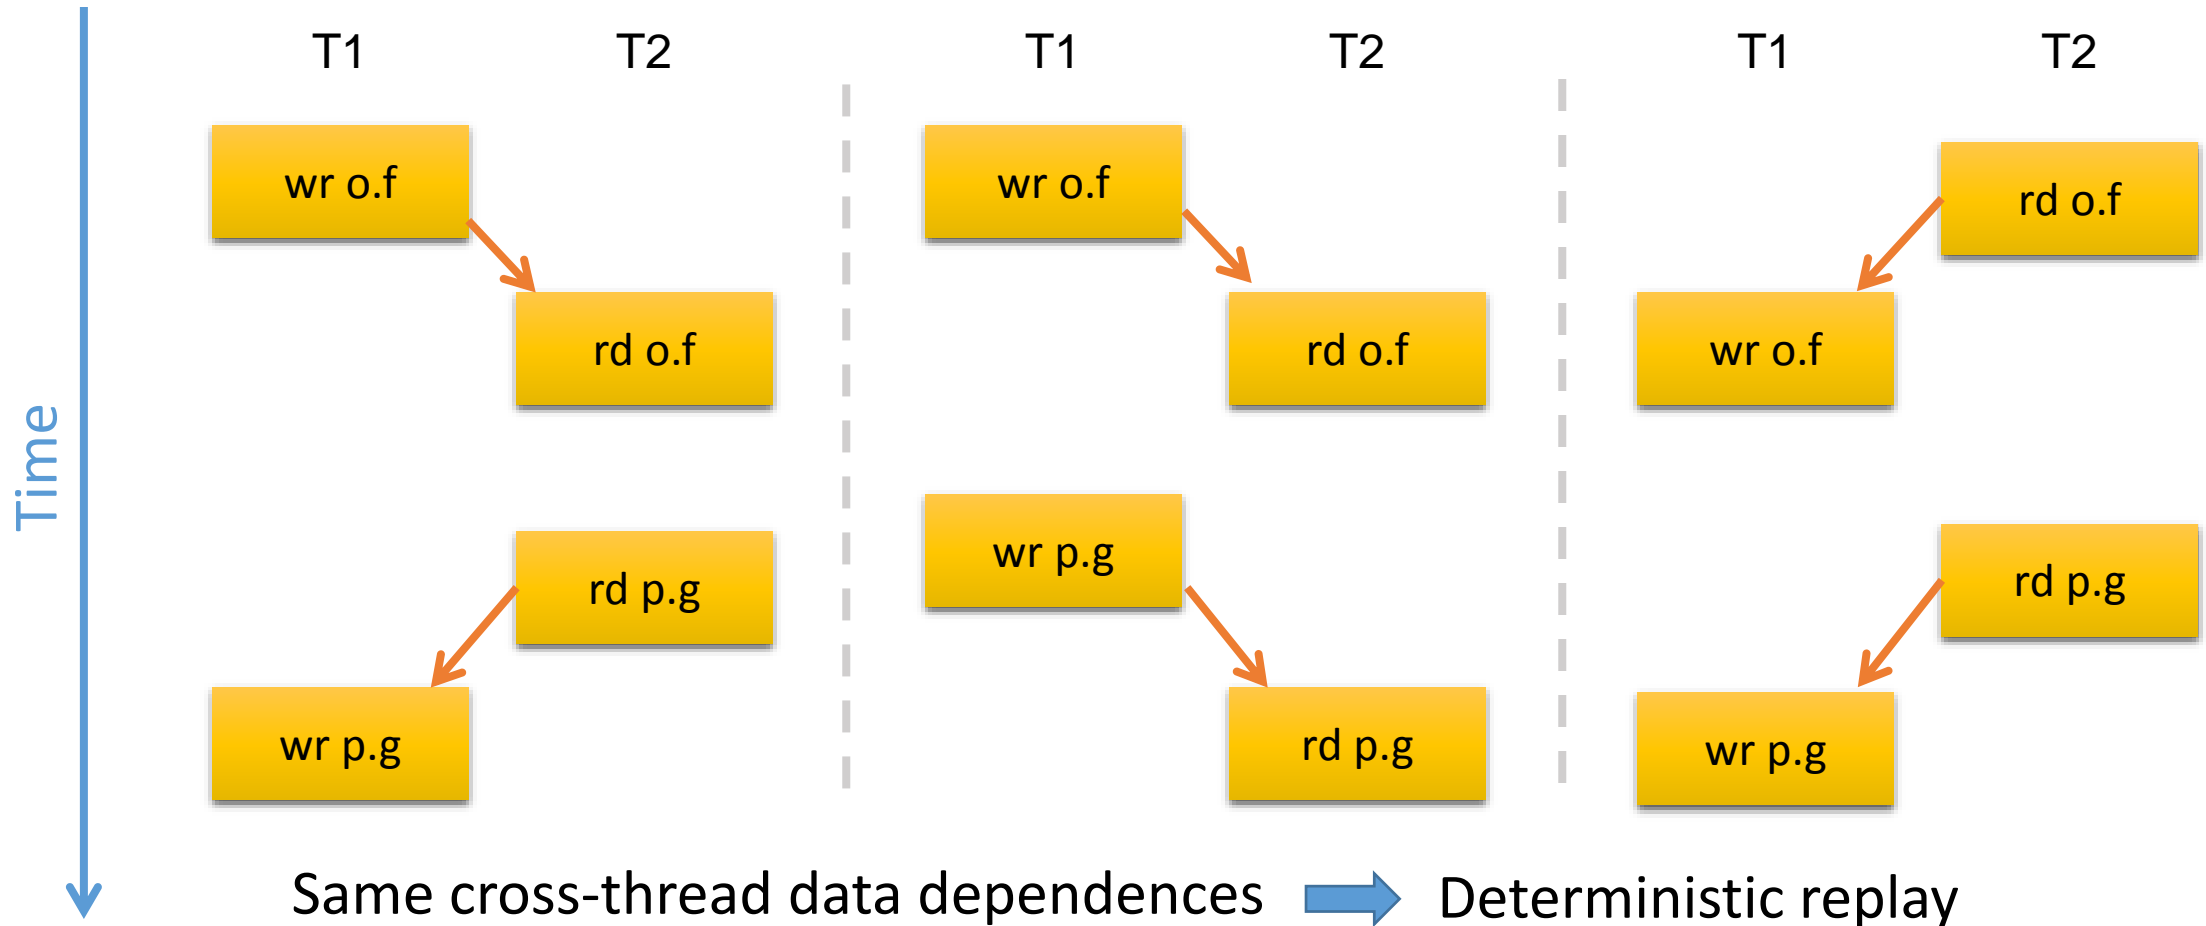


# Hard case: executions with data races

- Unfortunately, most real-world programs have data races
  - Instrument all potentially racy memory accesses to capture *cross-thread data dependences*
  - Add many synchronization operations, very expensive!



# Cross-thread data dependences



# Limitations of existing multithreaded record and replay approaches

- High overhead for recording cross-thread dependences
- OR do not handle racy execution
- OR support only offline or online replay, but not both
- OR rely on speculation and extra cores
- OR need custom hardware

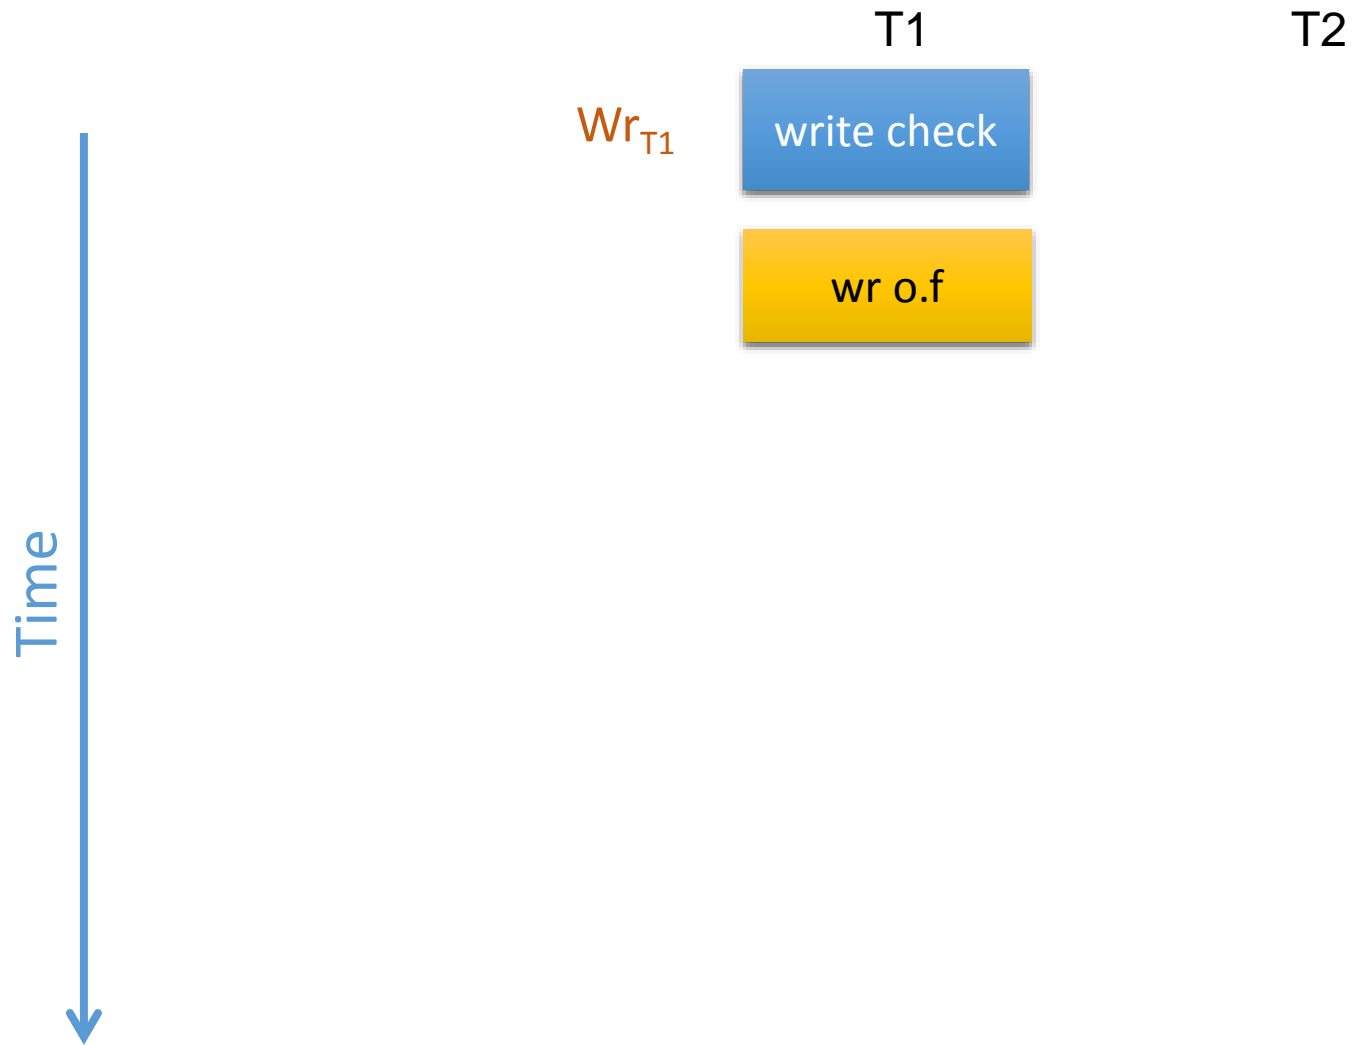
*Our approach overcomes all of these limitations at the same time!*

# RECORD

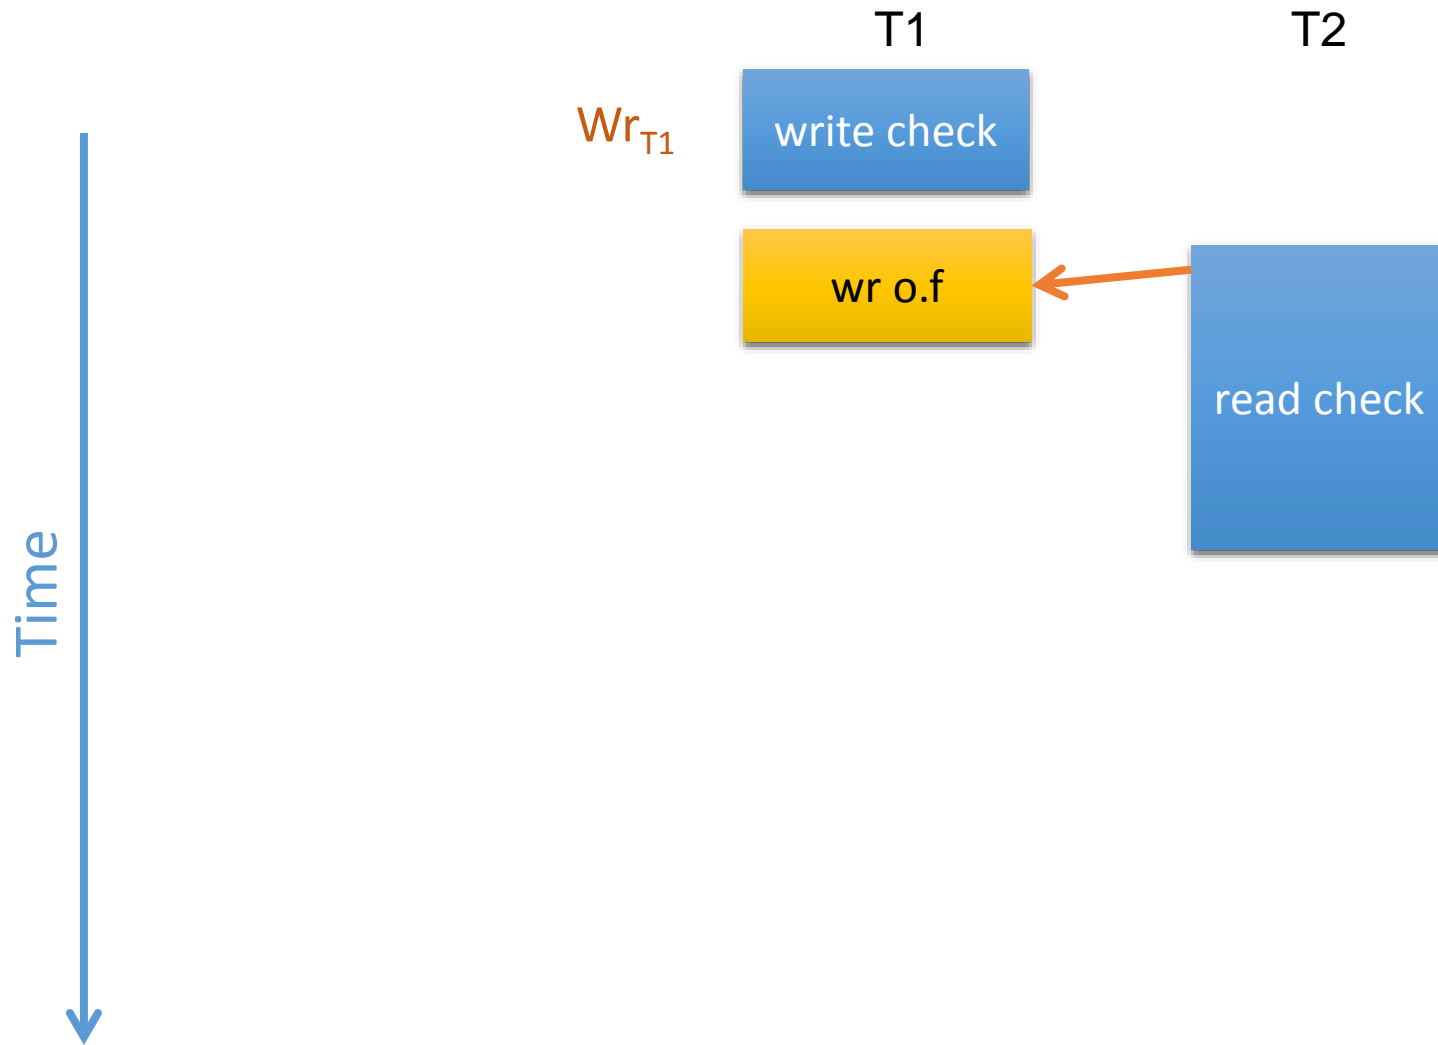
- Builds on our prior work “Octet”
- Octet tracks cross-thread dependences at object granularity
- Each object has an ownership state
  - Analogous to cache coherence protocol
- For simplicity, consider
$$o.state \in \{ Wr_T, Rd_T \}$$
- Cross-thread dependence  $\Rightarrow$  state transition

# State transition example

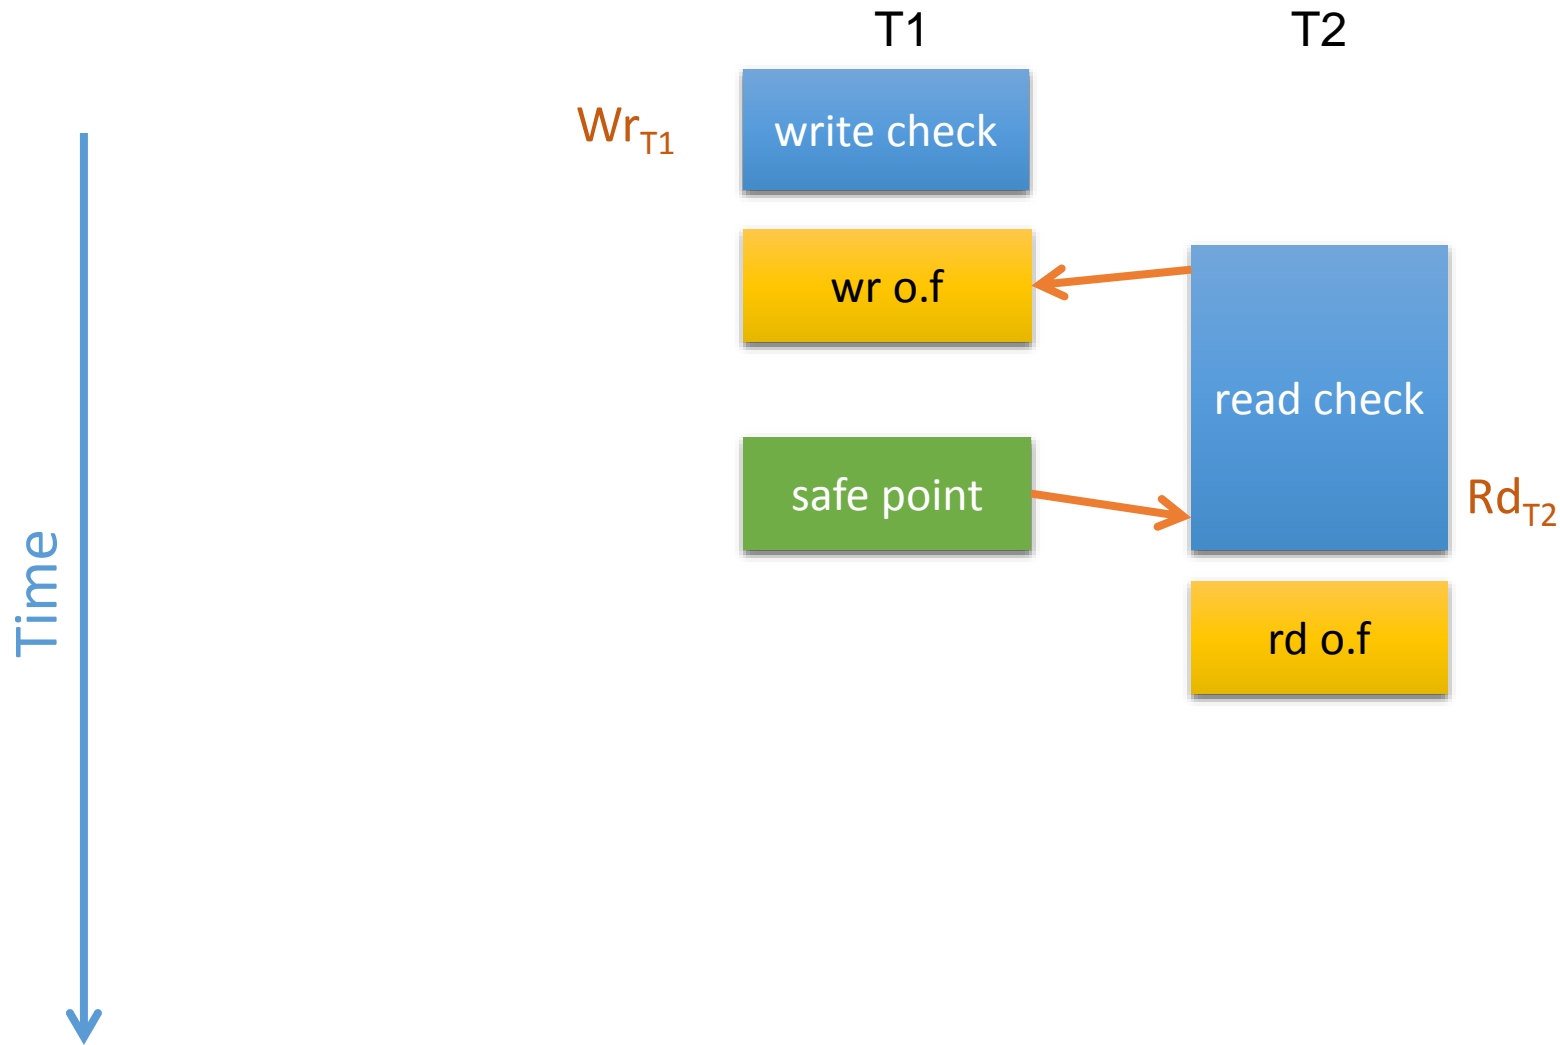
Initially o.state =  $Wr_{T1}$ , p.state =  $Rd_{T2}$



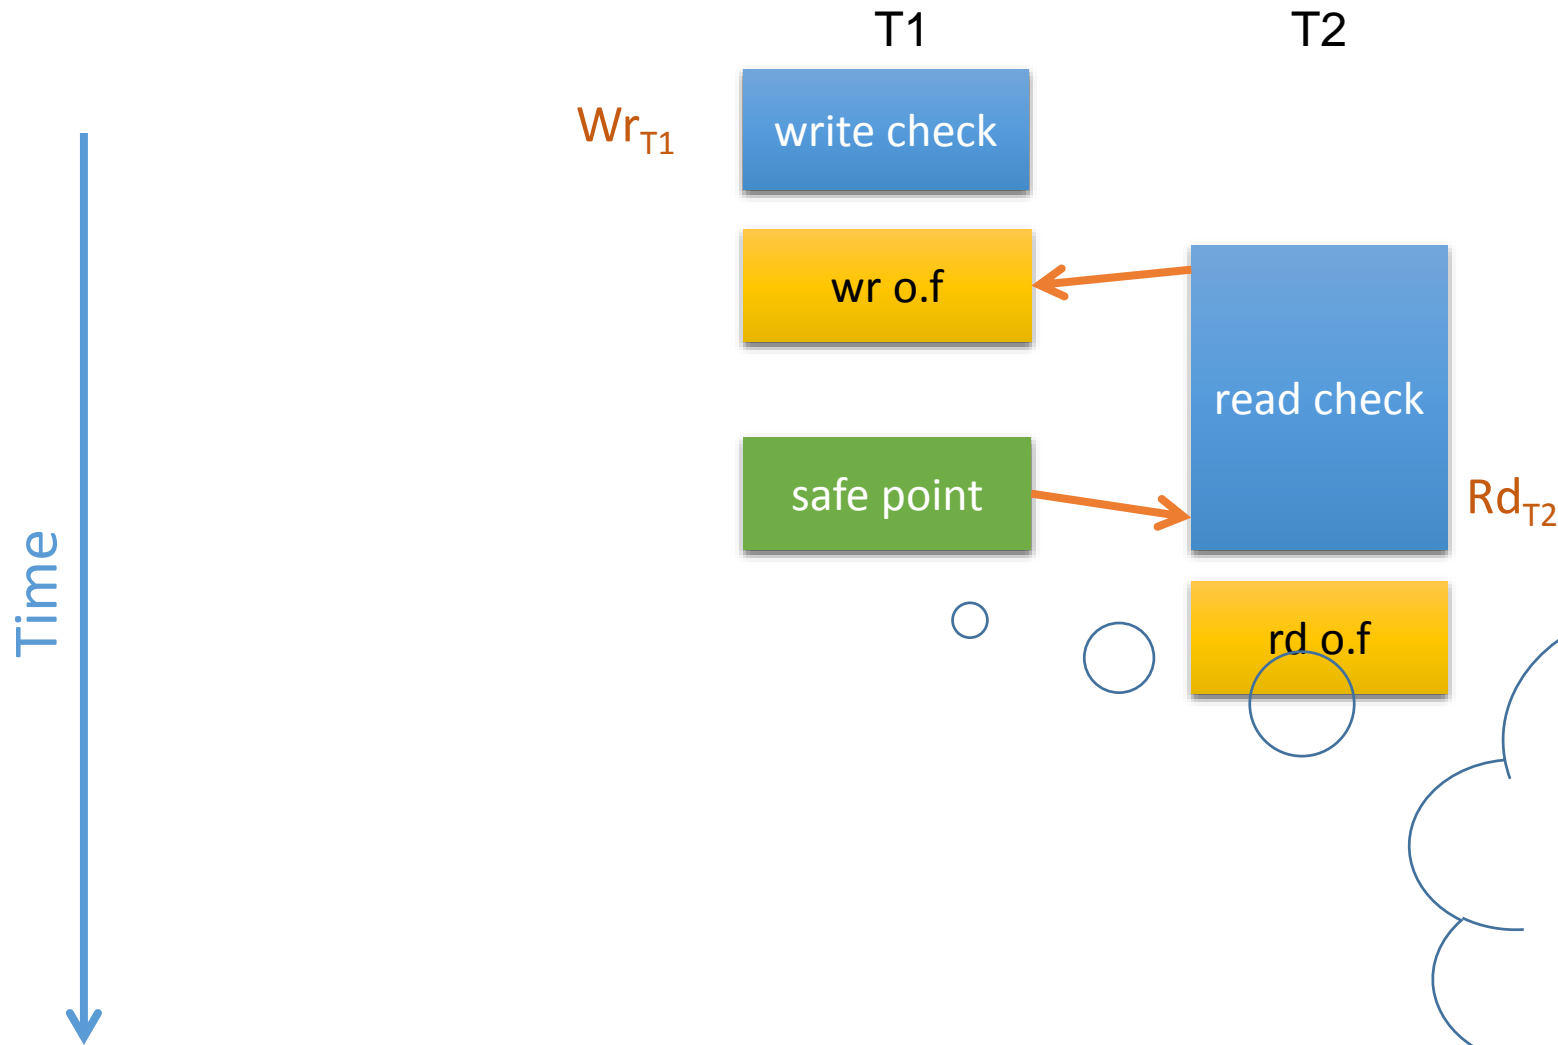
# State transition example



# State transition example

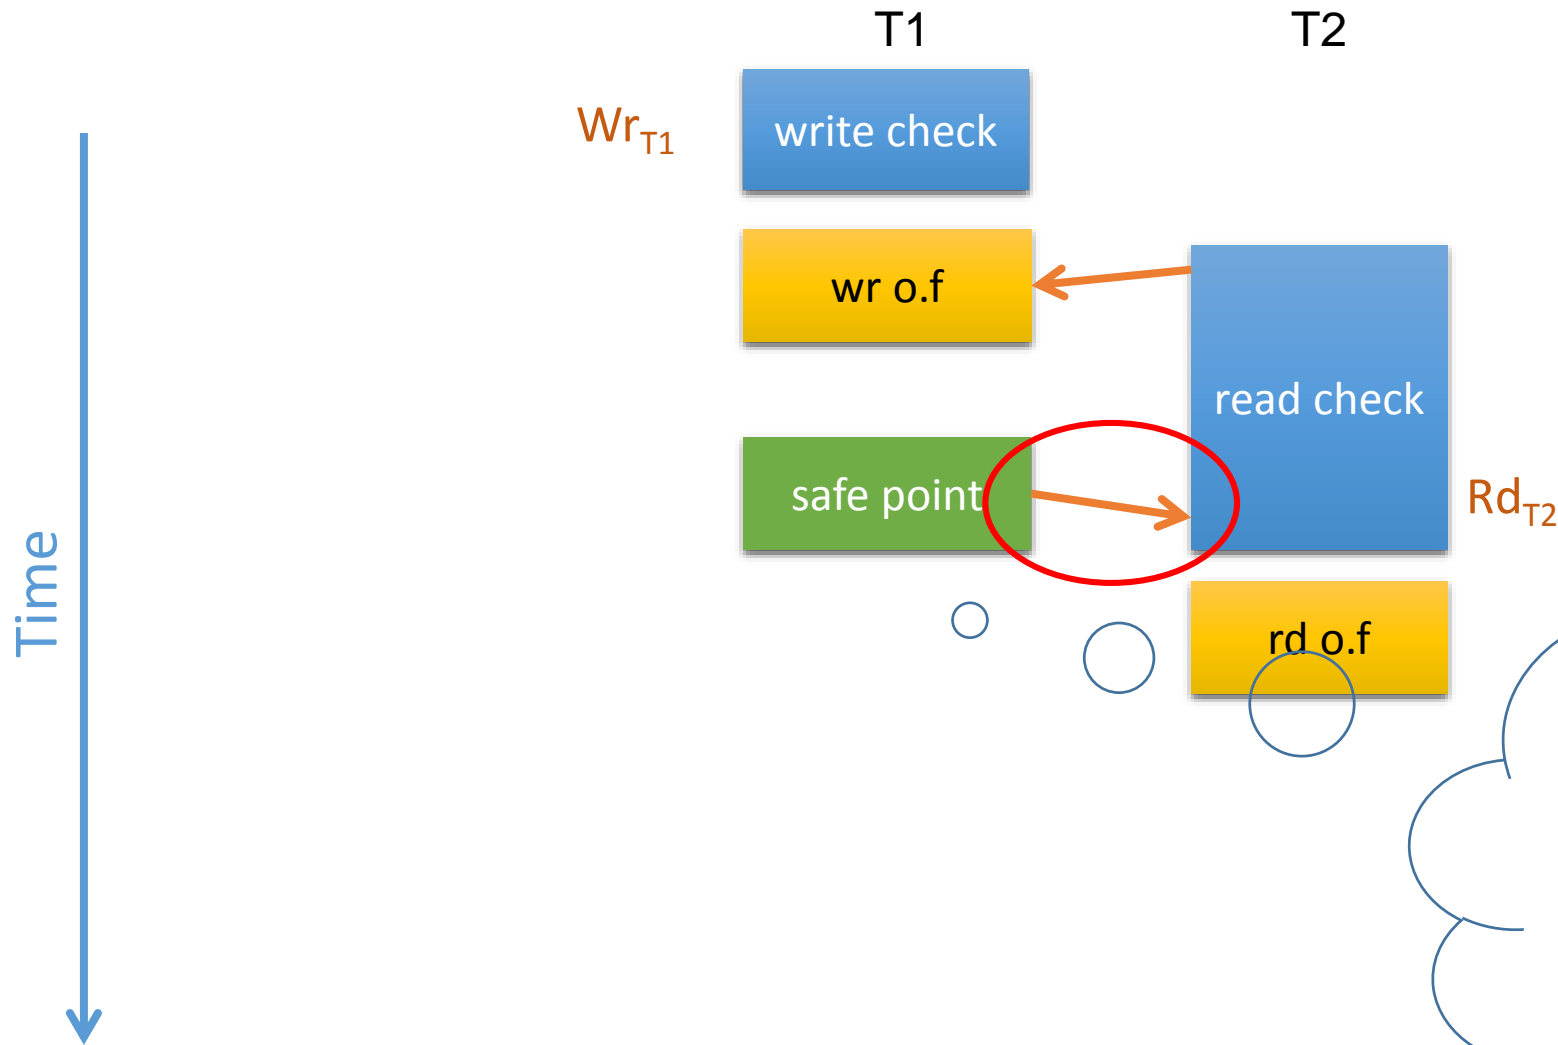


# State transition example



How to record this cross-thread dependence?

# State transition example



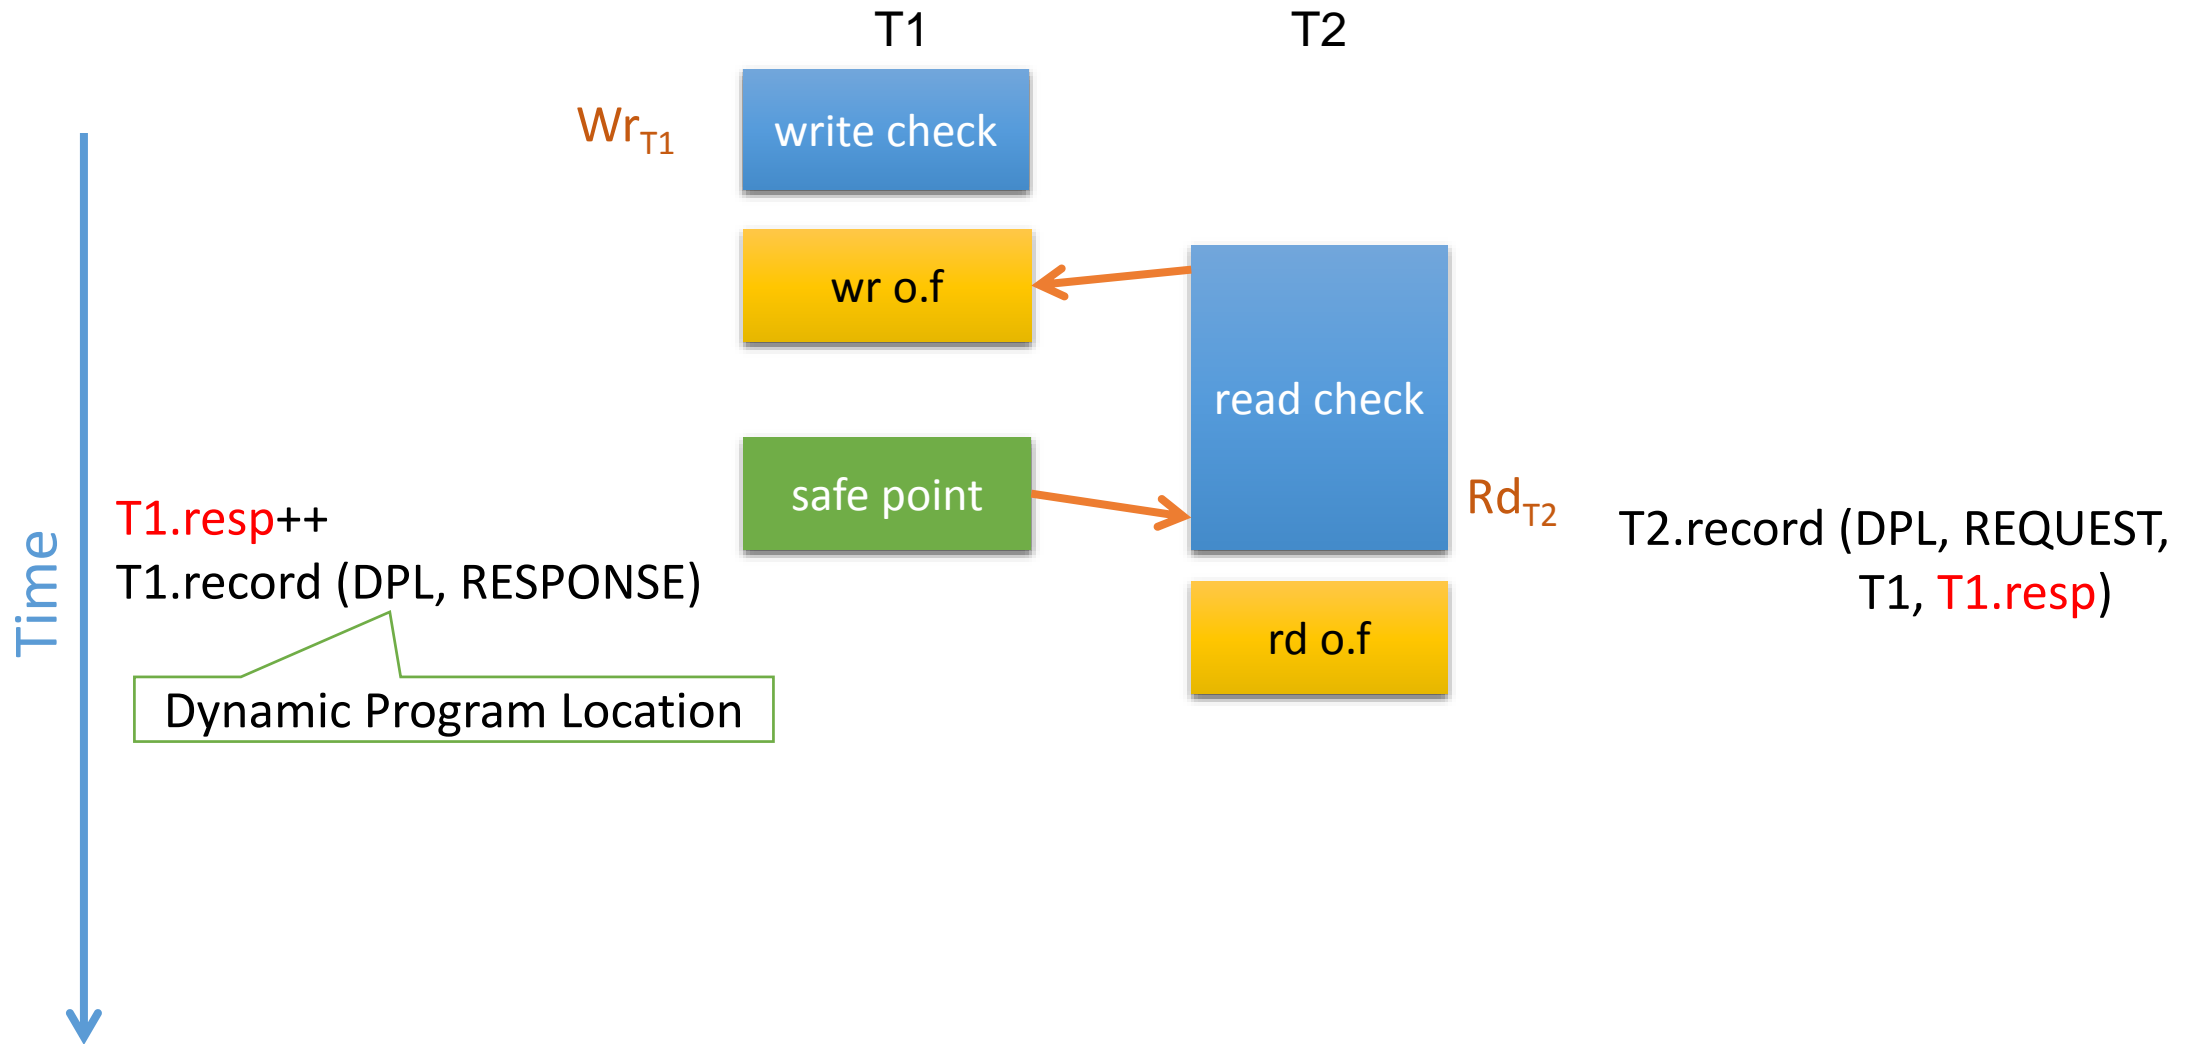
How to record this cross-thread dependence?



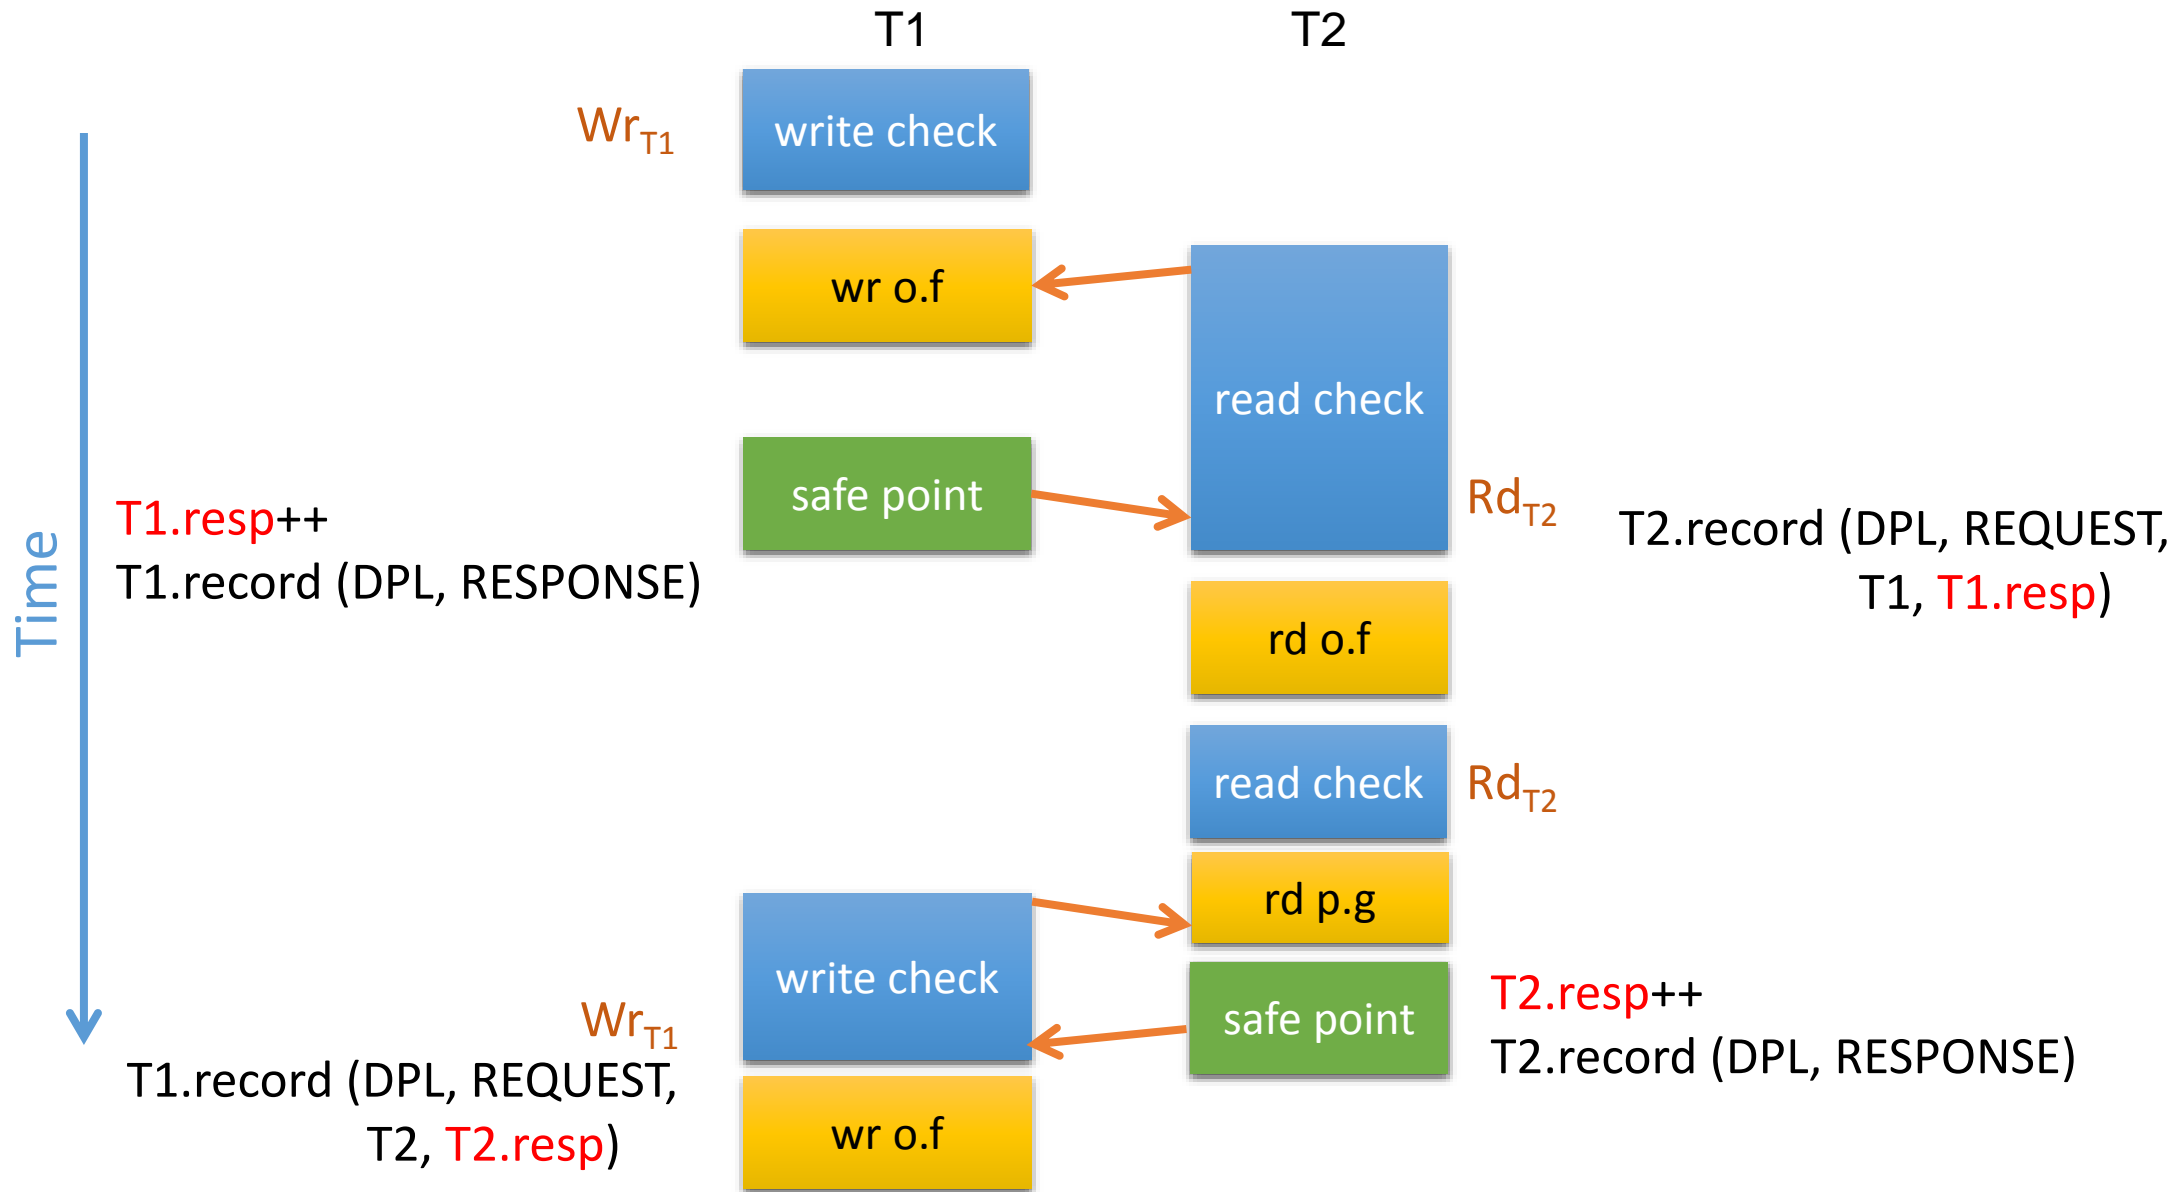
# RECORD Design

- What?
  - Response edge
  - Per-thread response counter
- Where in execution?
  - Dynamic Program Location (DPL)
- Per-thread log file

# RECORD example



# RECORD example



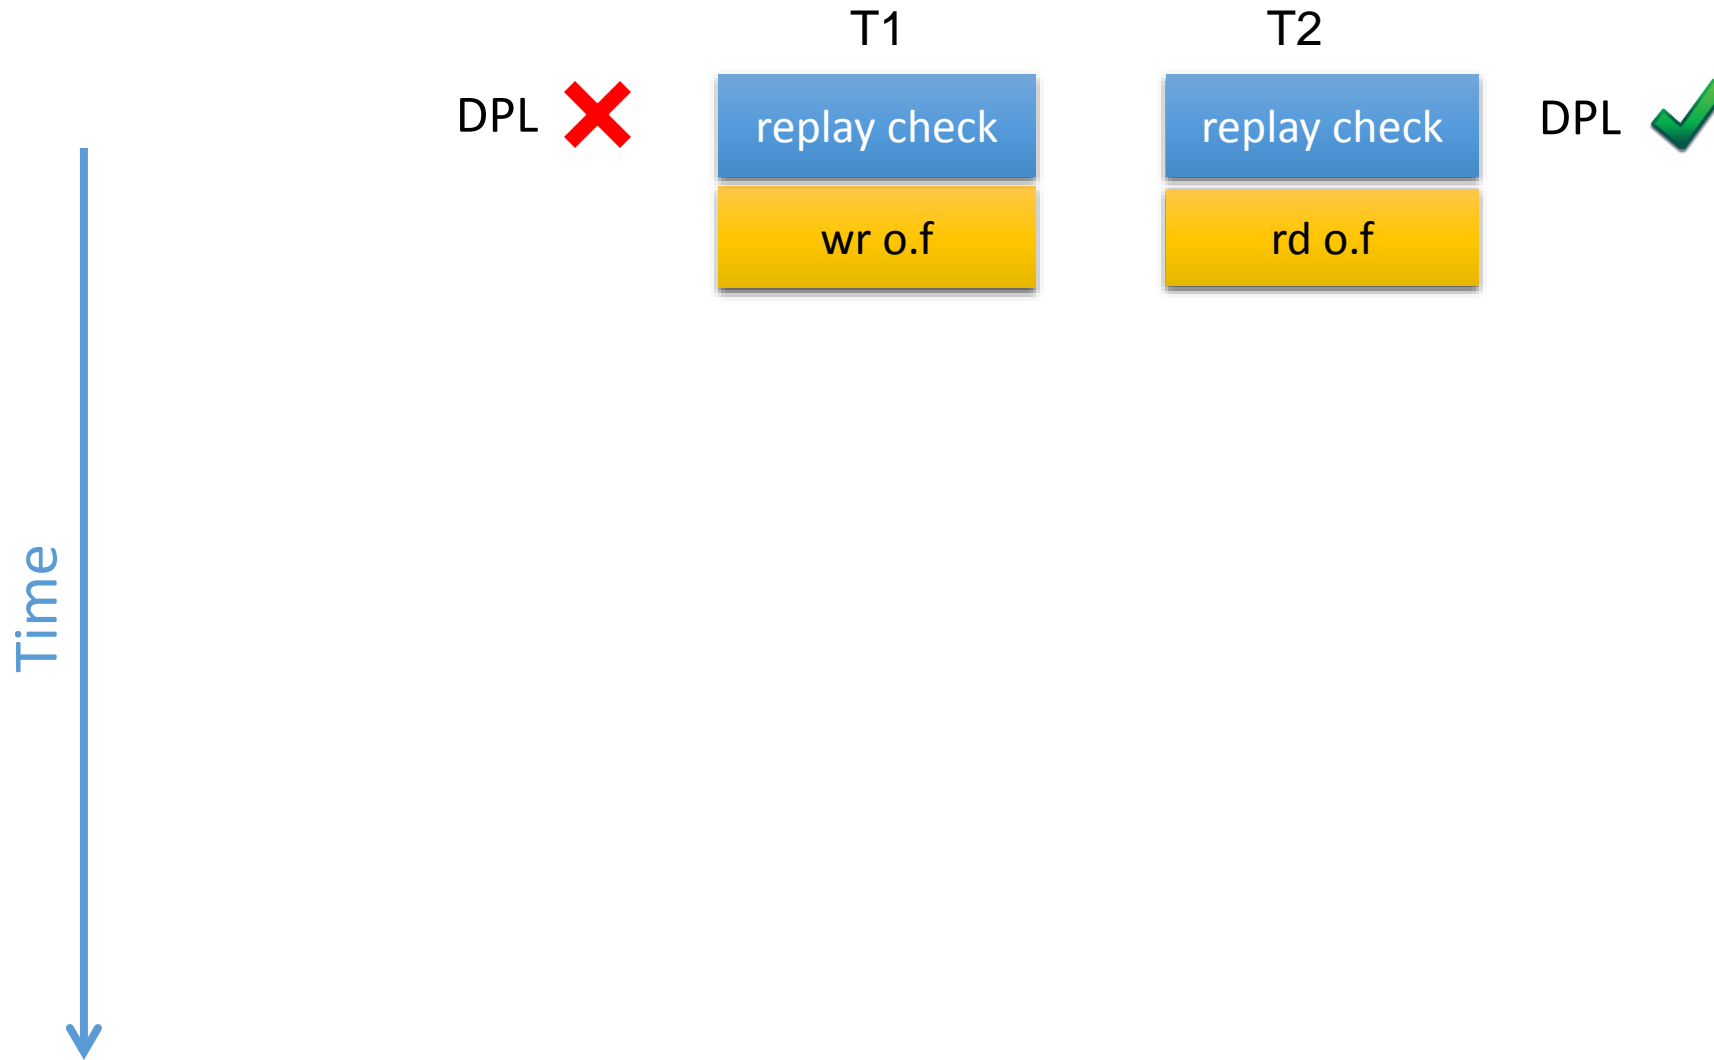
# RECORD observation

- Adds low overhead
  - Most accesses (>99%) are same-state
  - Only one load and one if check

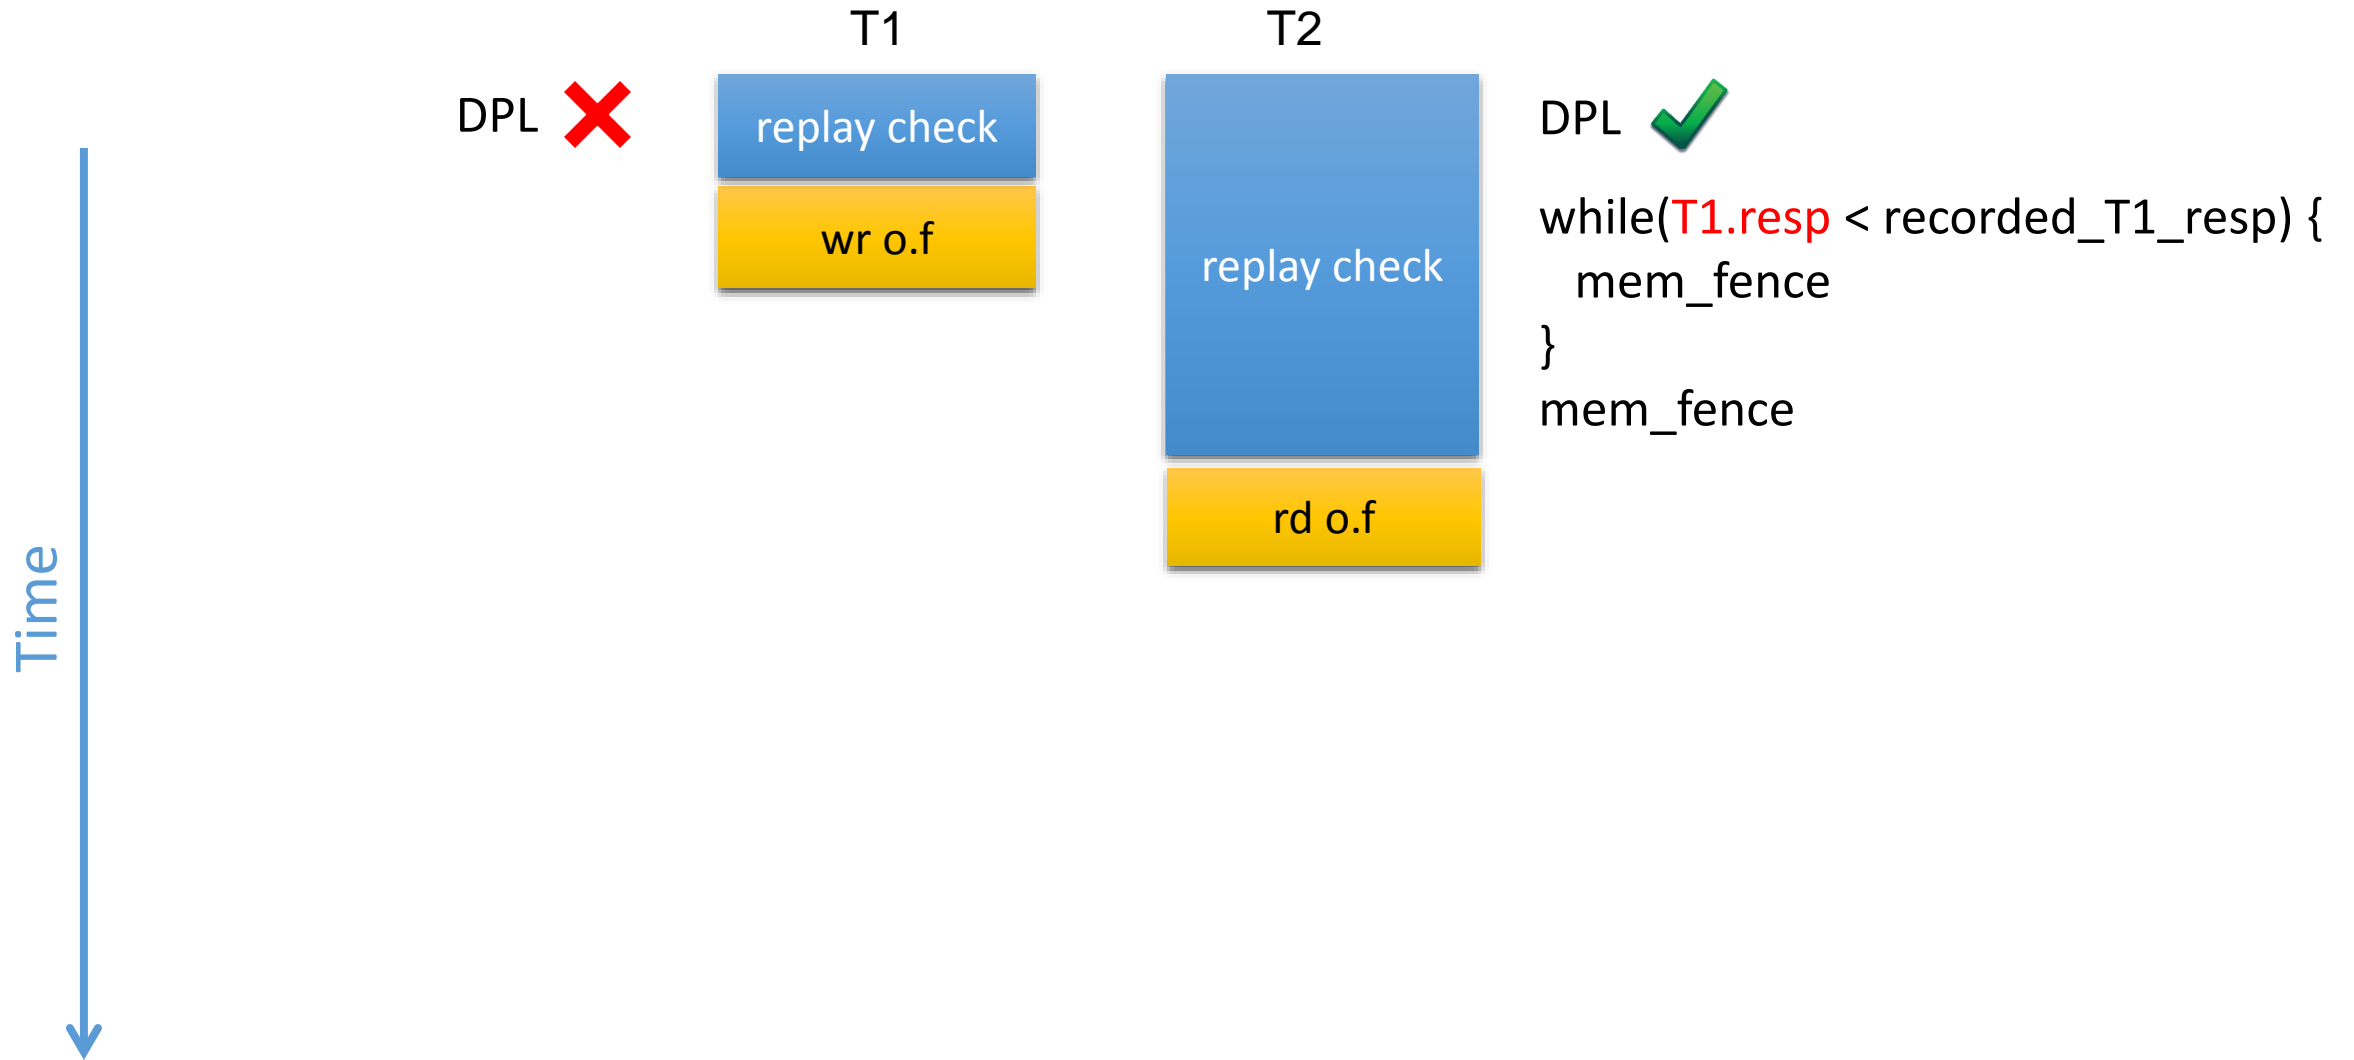
# REPLAY

- Goal: enforce recorded edges
- Instrument every possible edge source and sink
  - Check if DPL matches
- Edge source
  - Increment counter
- Edge sink
  - Wait for counter to reach recorded value

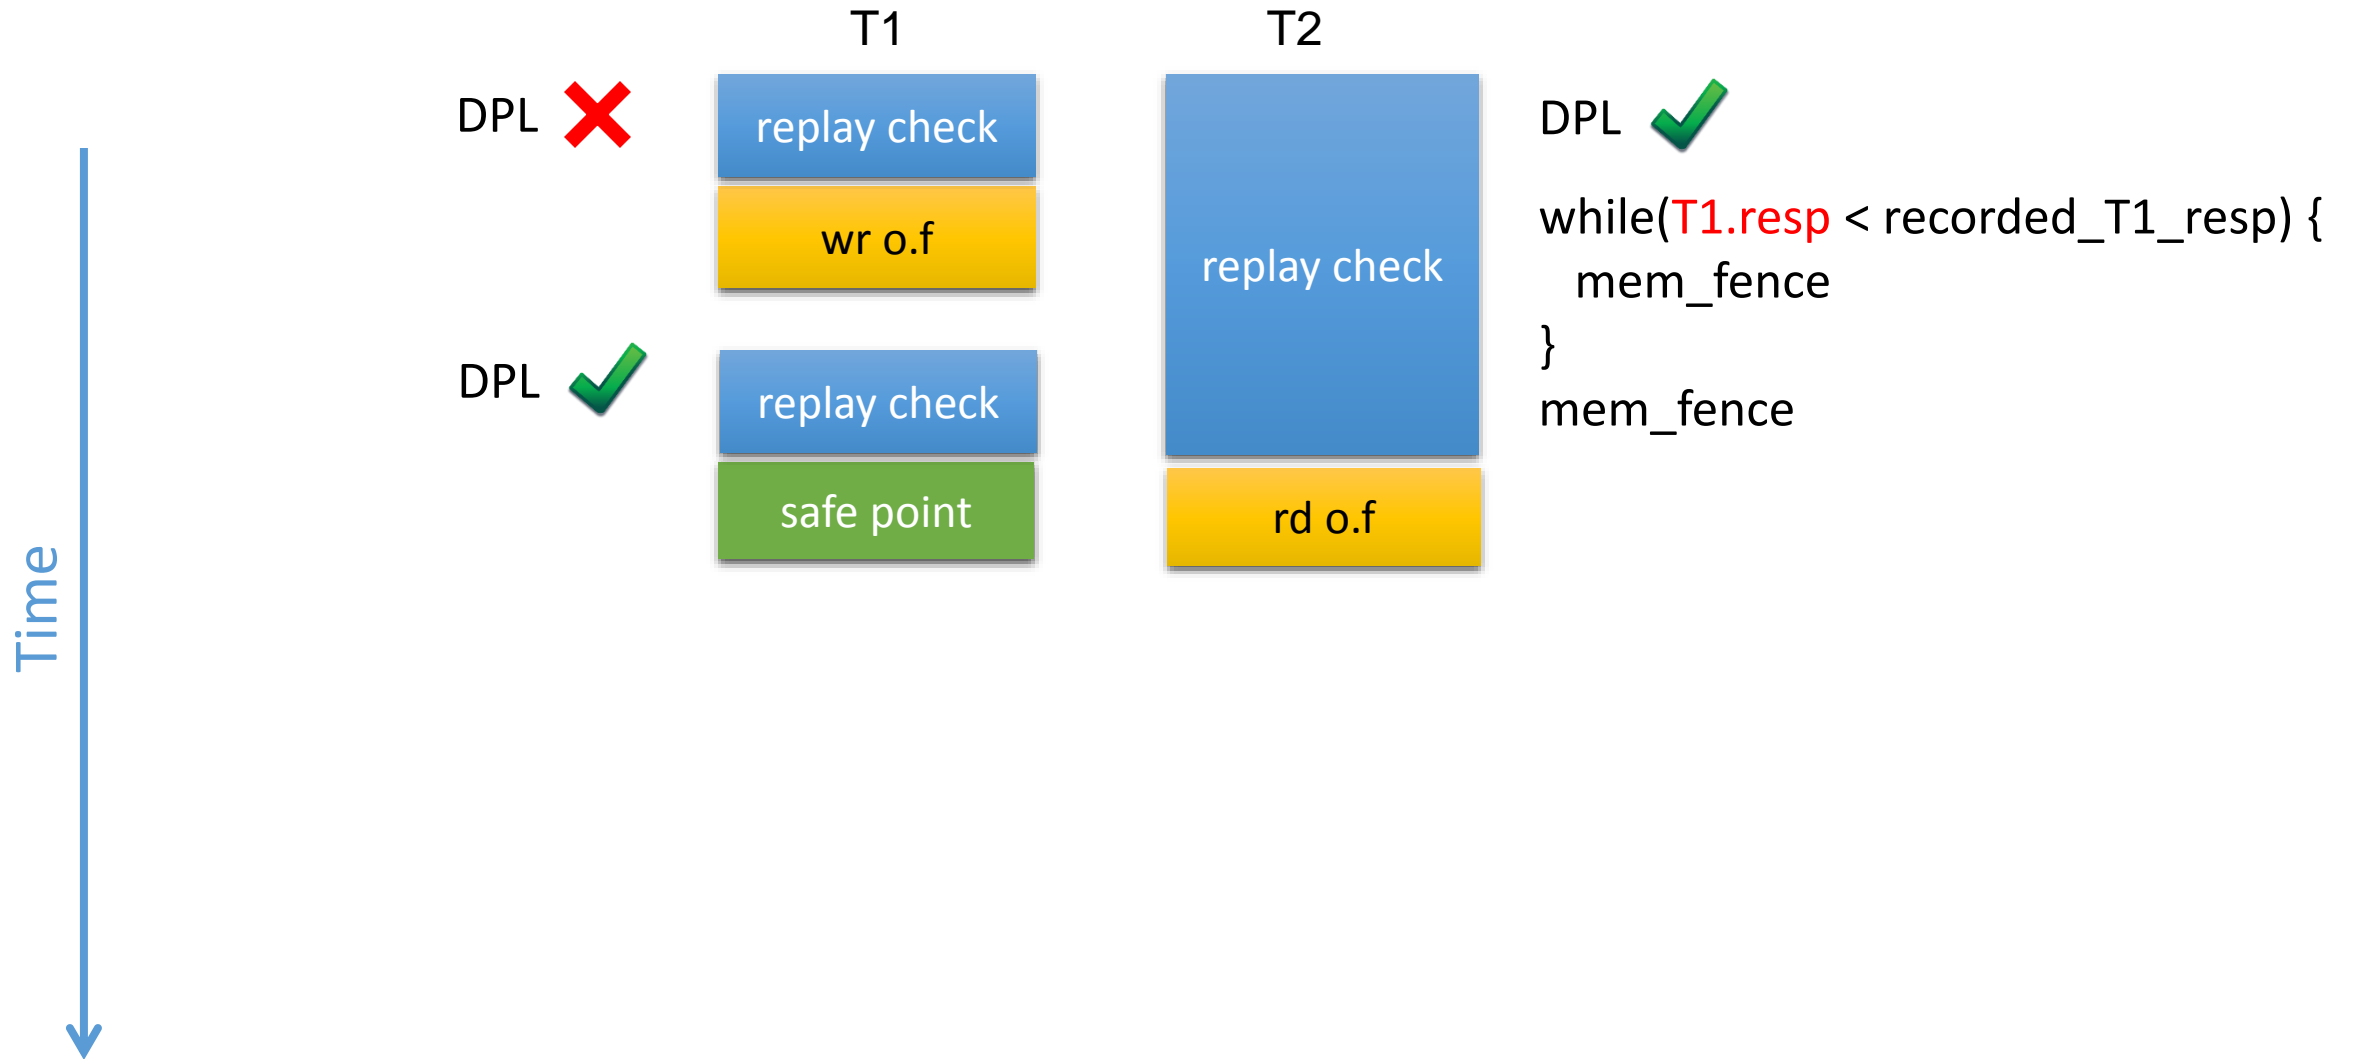
# REPLAY example



# REPLAY example

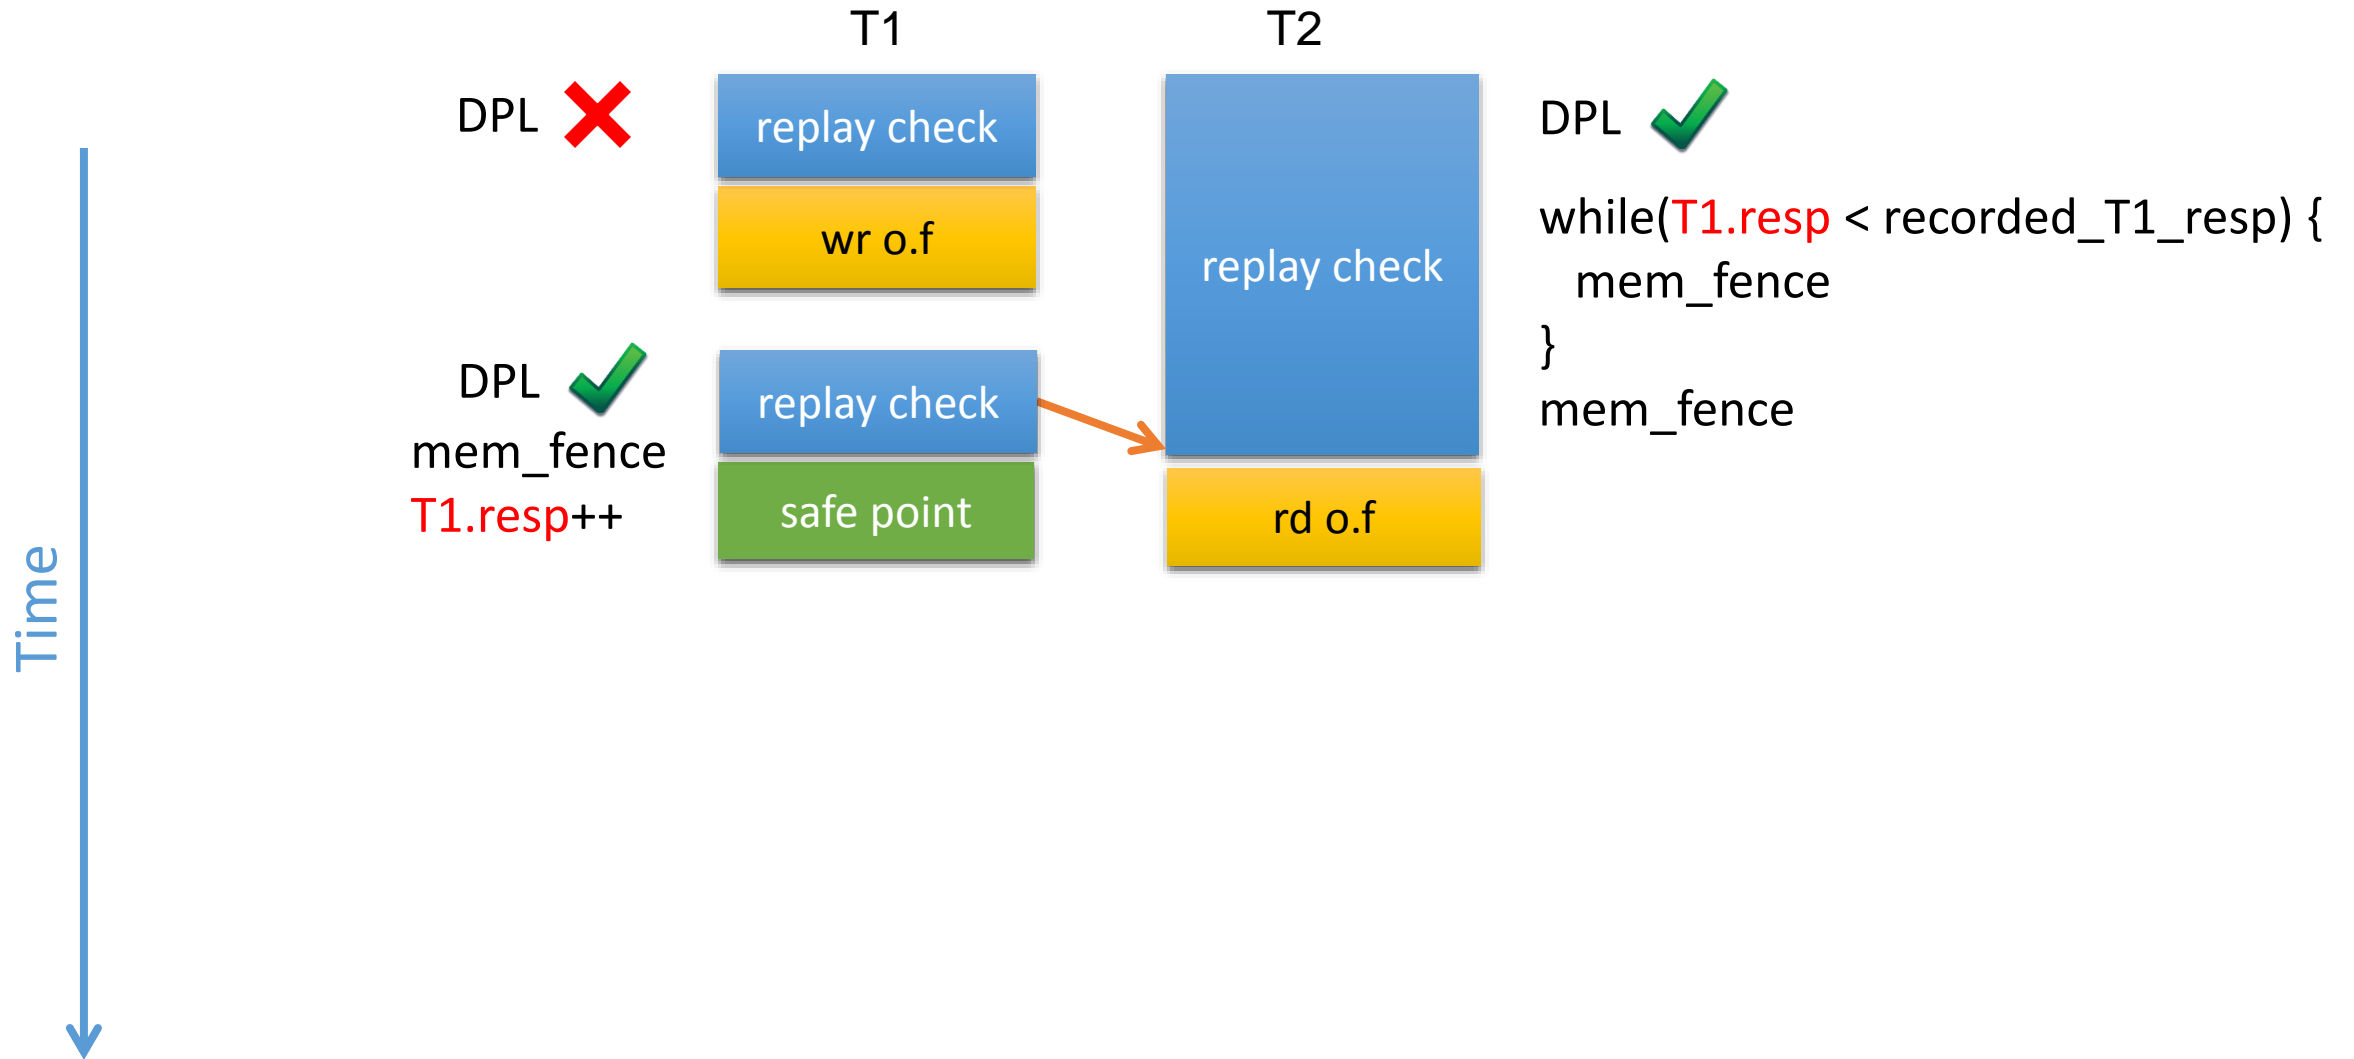


# REPLAY example

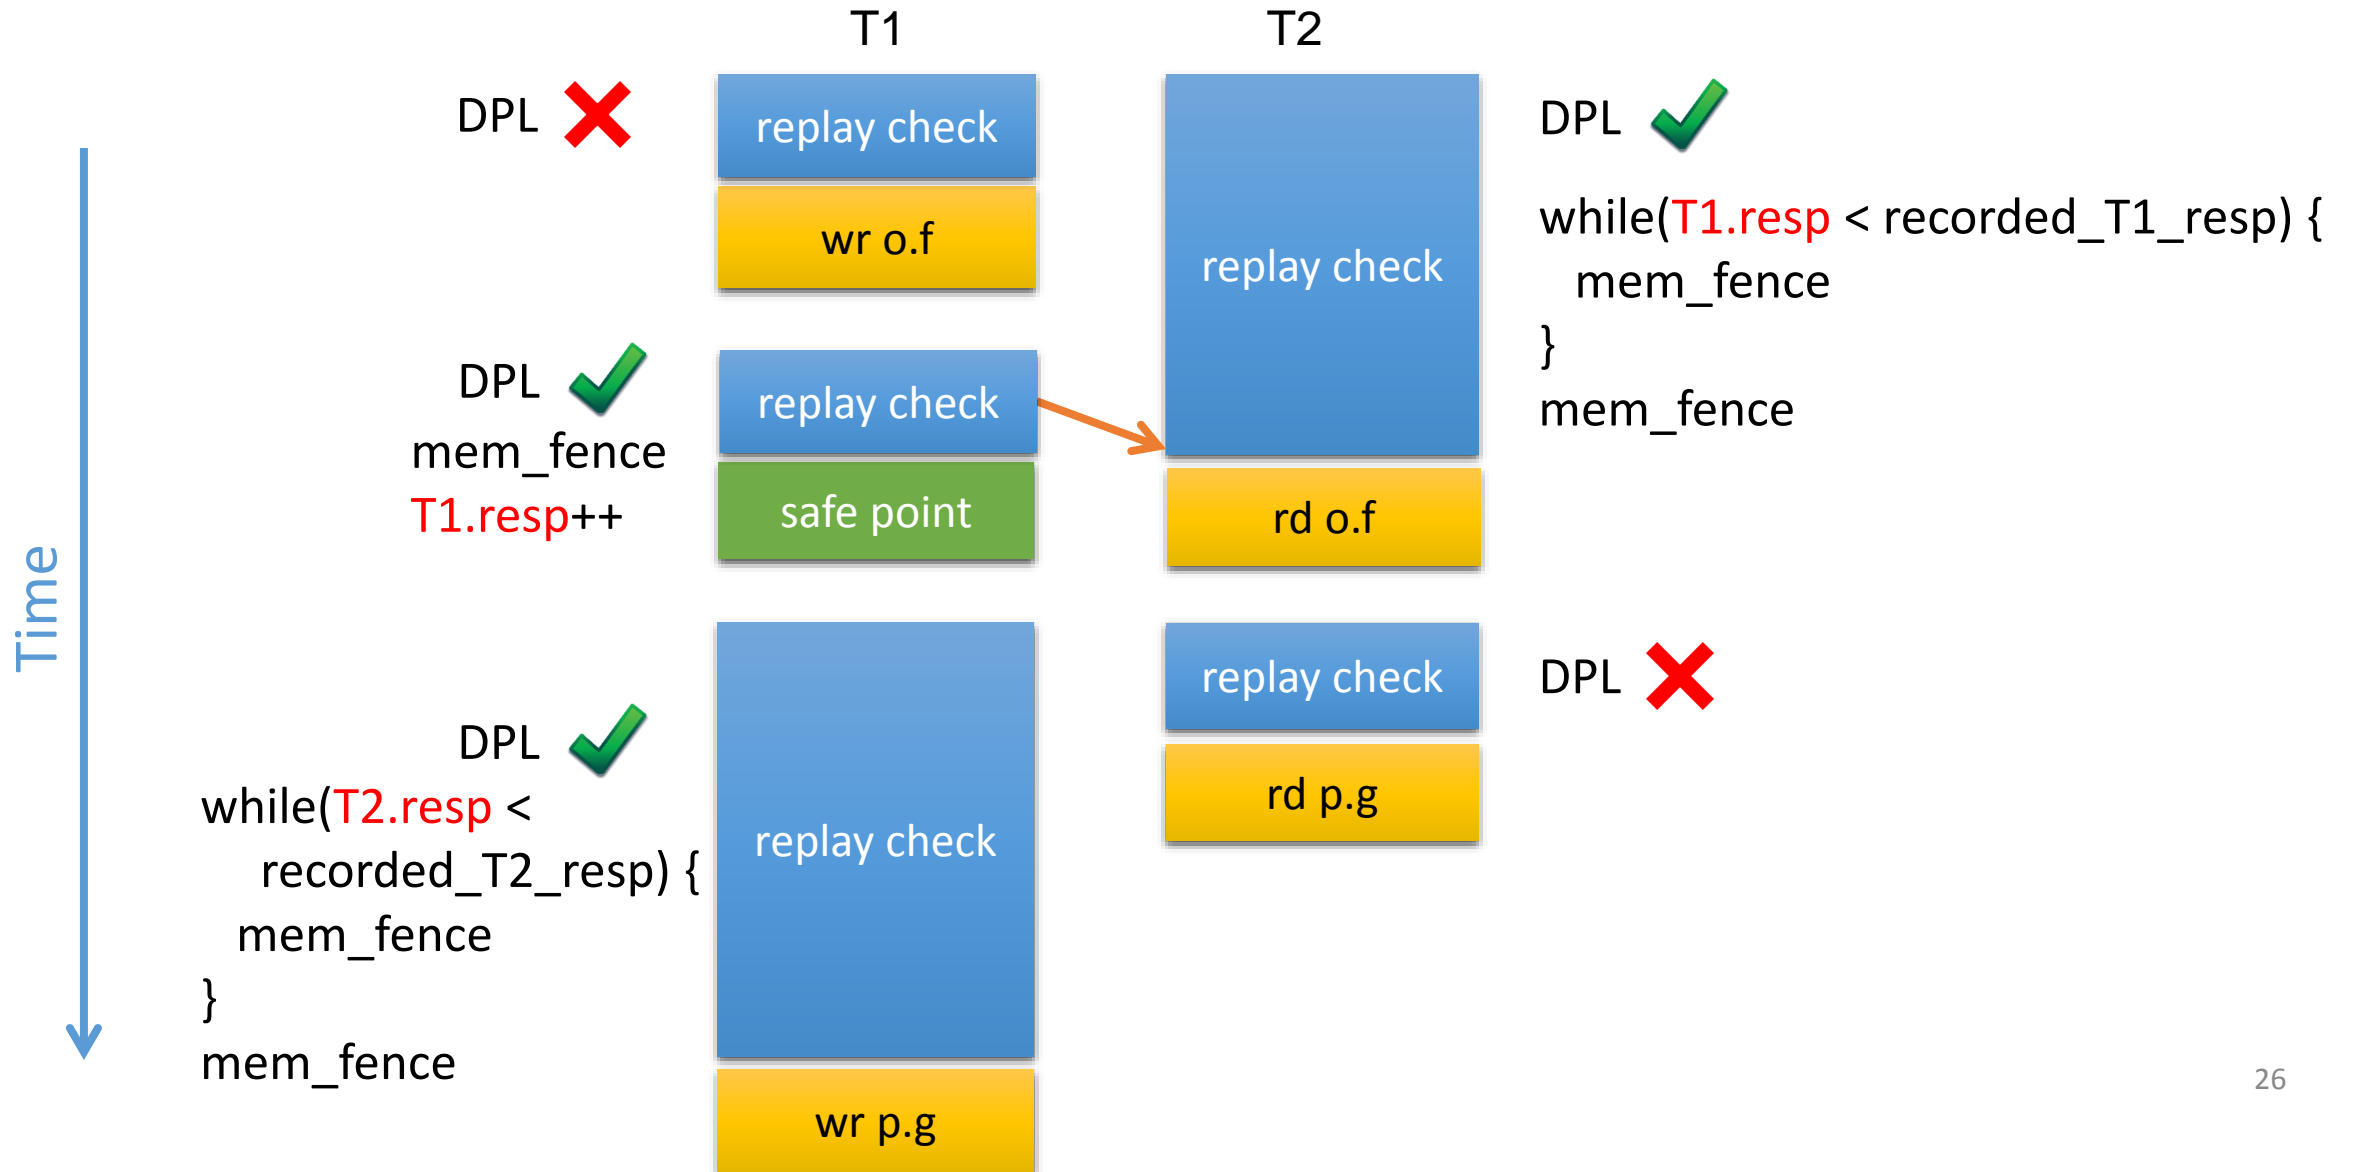




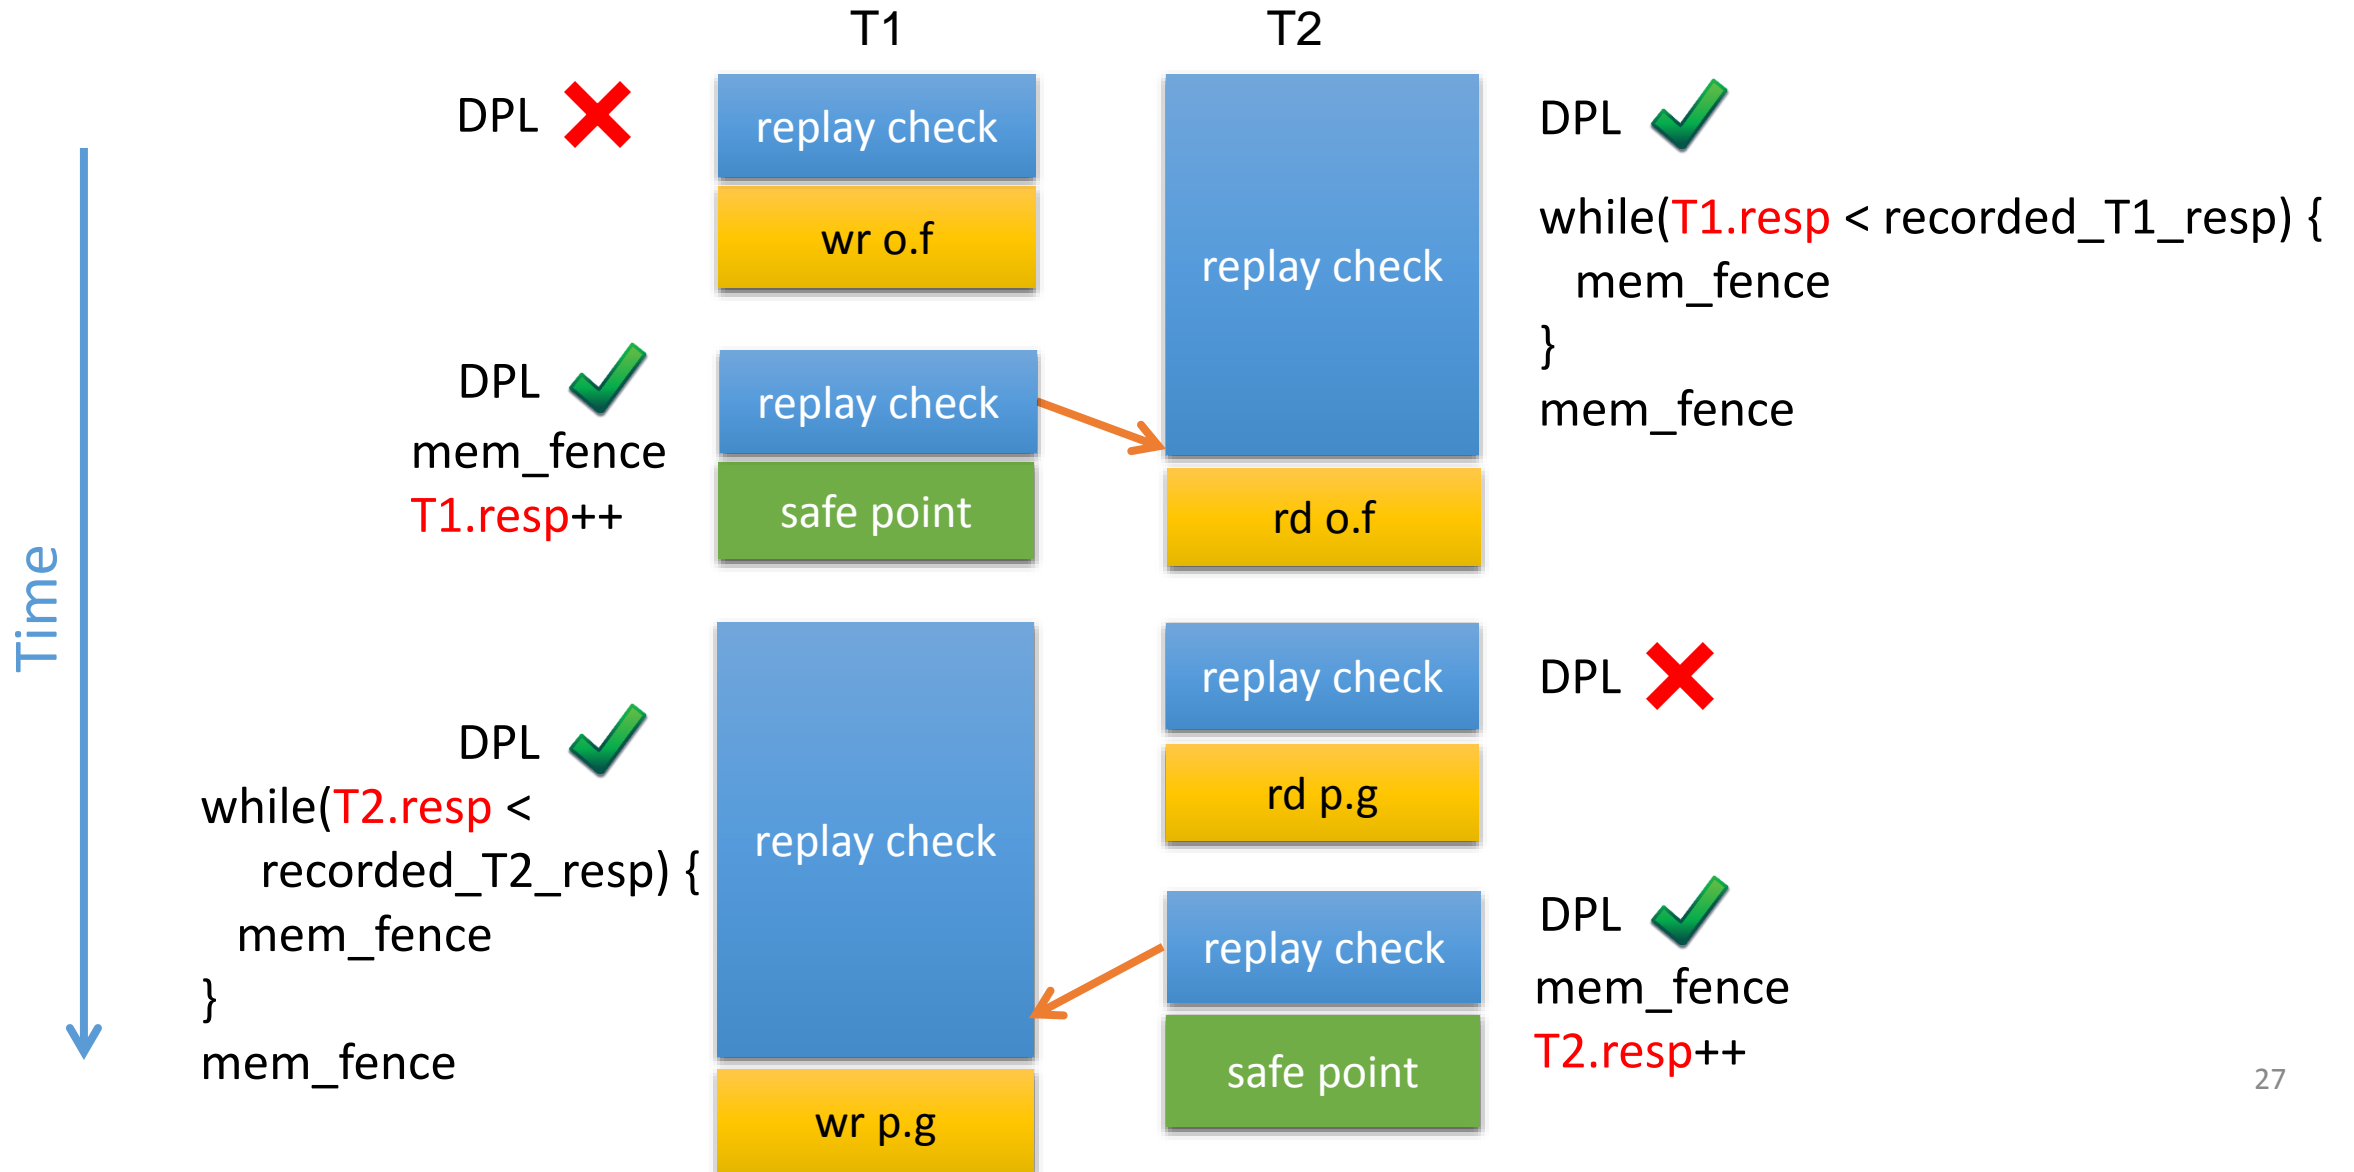
# REPLAY example



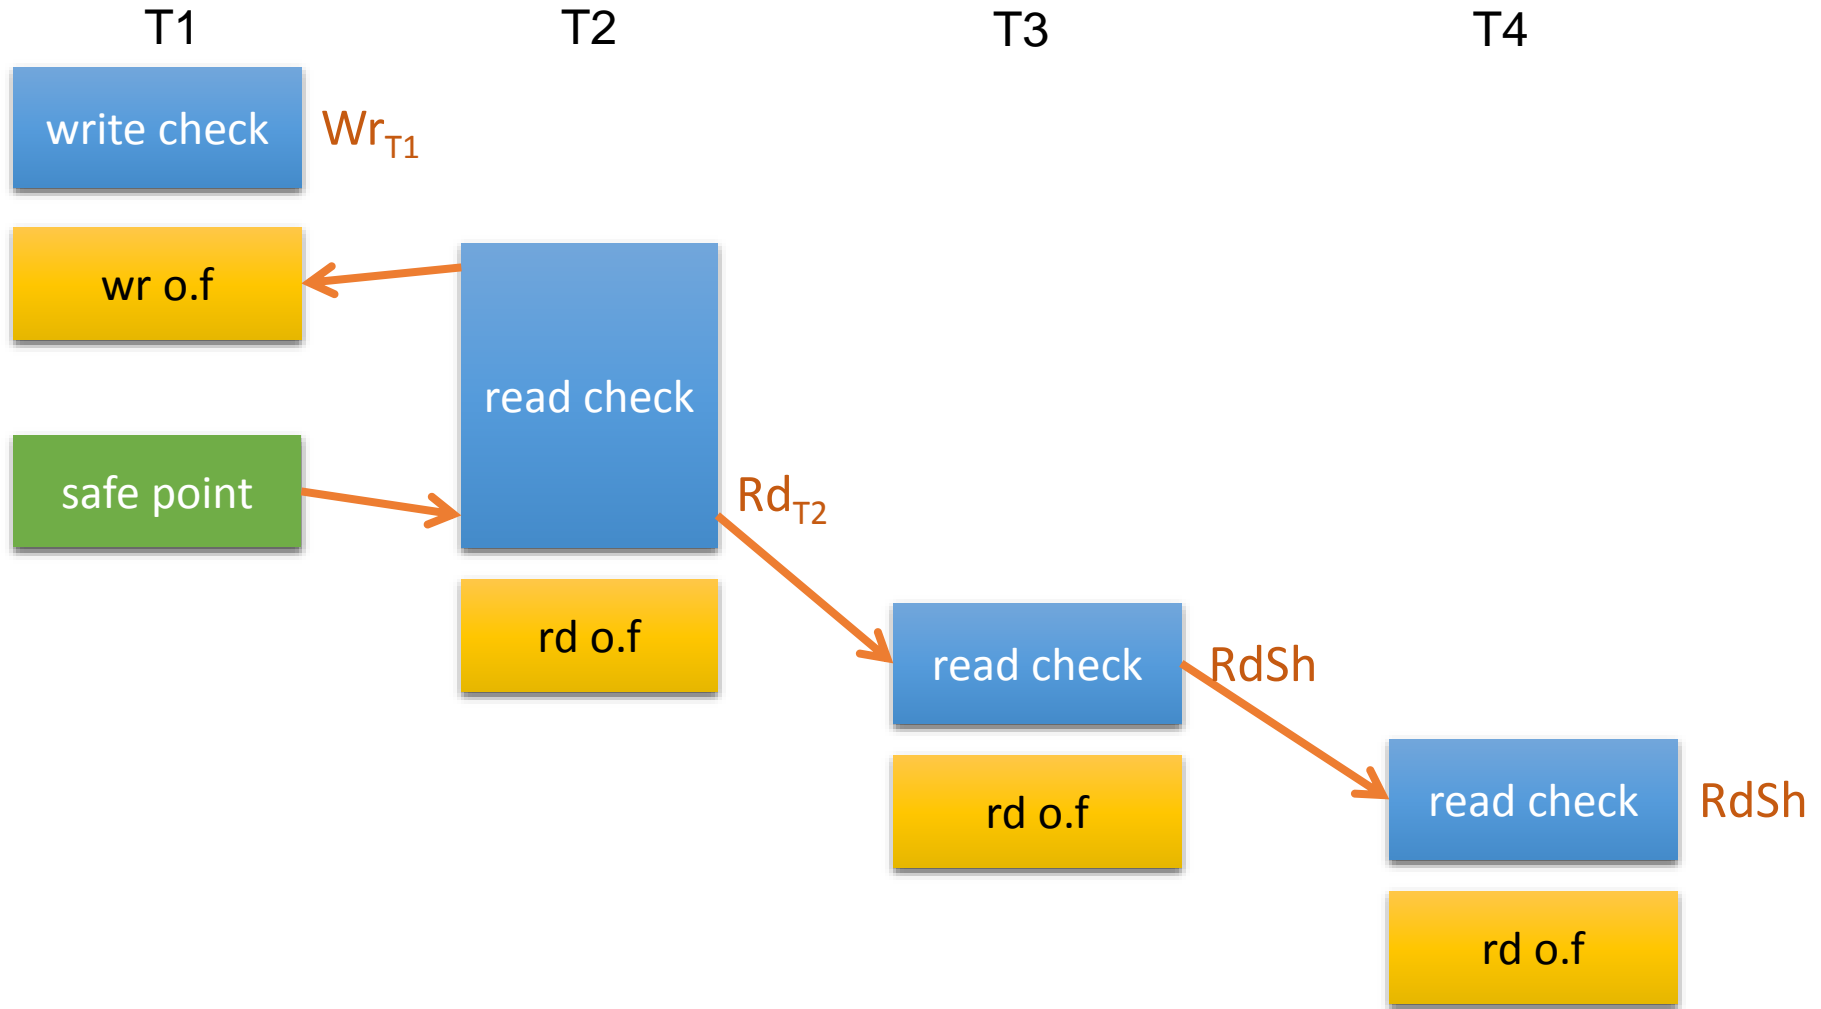
# REPLAY example



# REPLAY example

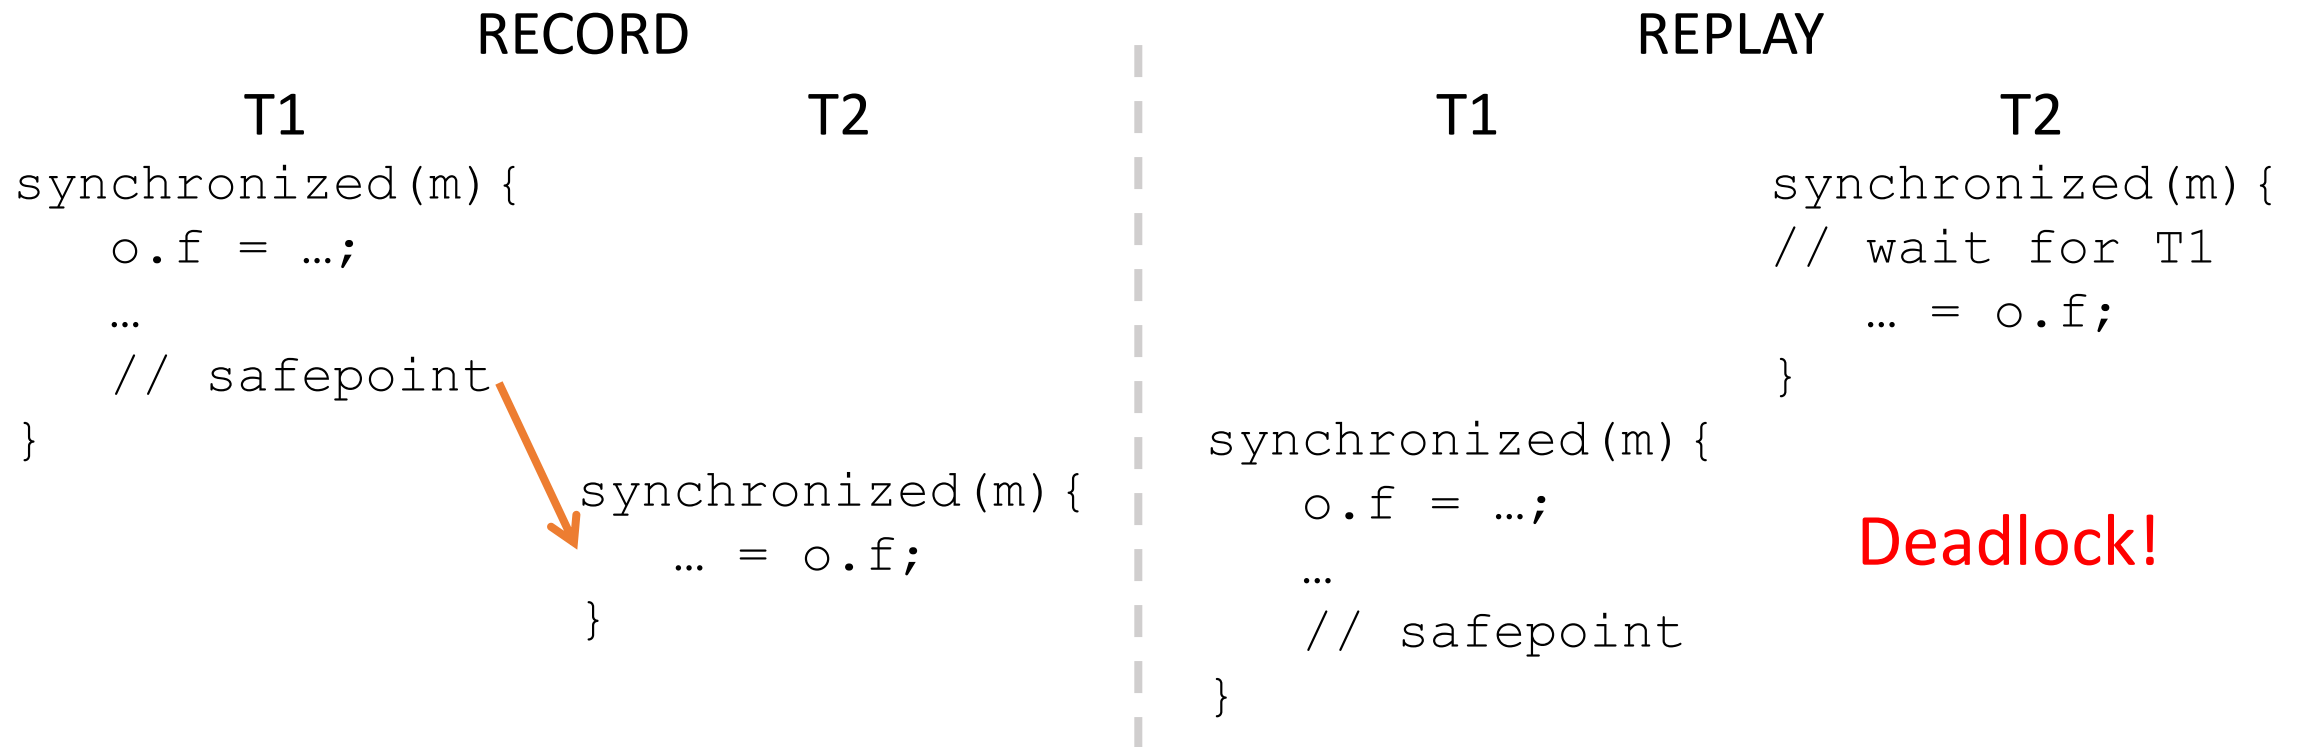


# A more complicated case



# REPLAY observations

- Do not track object's state
  - Only per-thread or global counters
- How about program synchronization?





# Elide program synchronization

- Necessary
  - RECORD does not track synchronization
  - Otherwise deadlock
- Side effect: more parallelism (see Evaluation)

# Handling external nondeterminism

# Goal: Application-level determinism

- No need to track JVM's cross-thread dependences!
  - Jikes RVM
- Contexts for compiled code

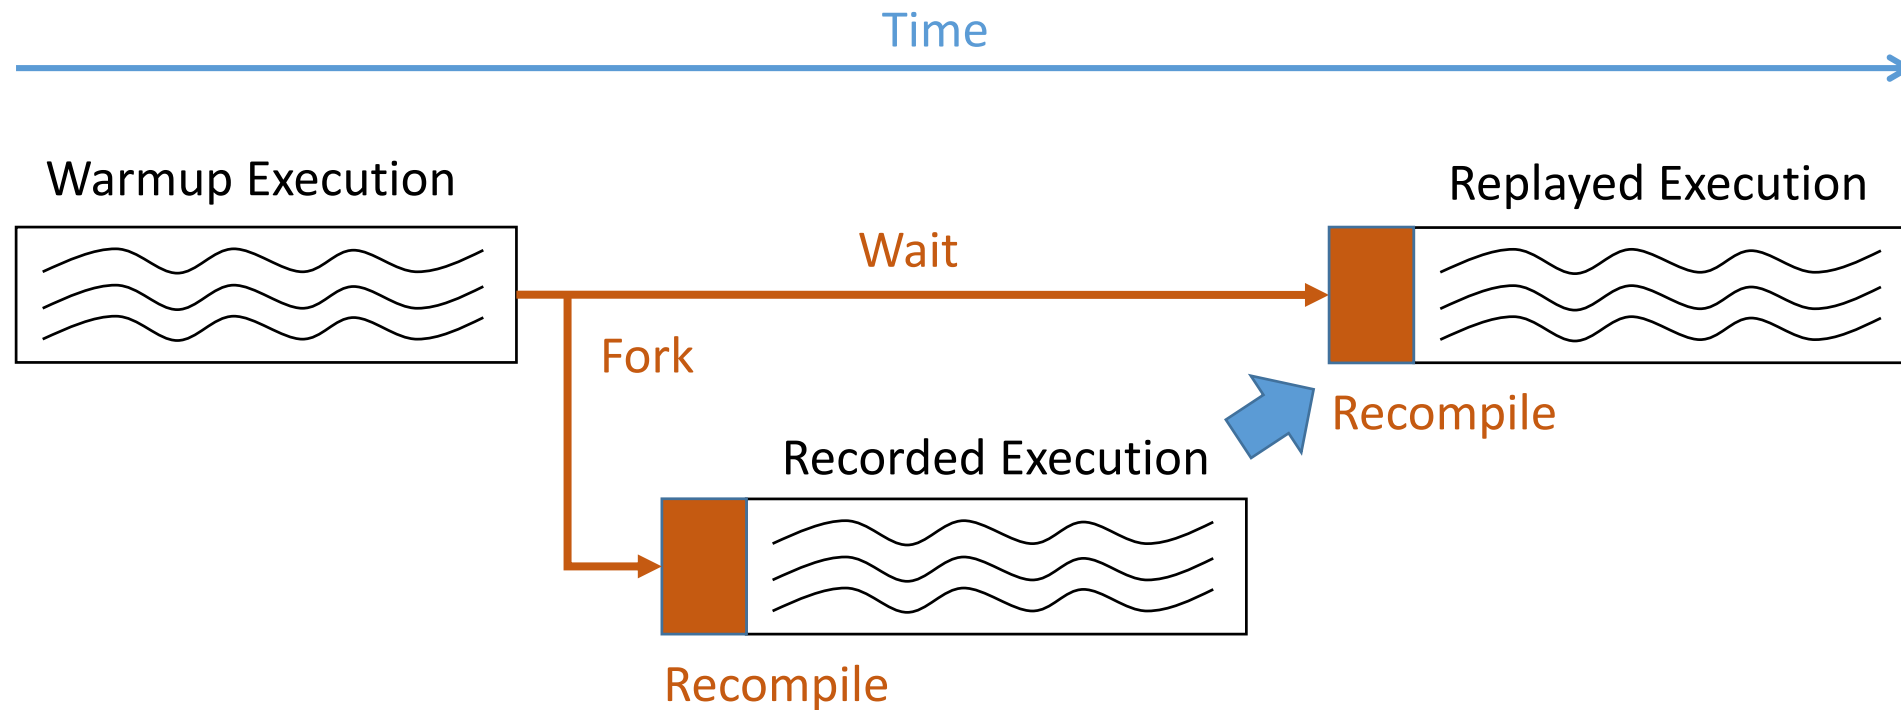




# External nondeterminisms

- Handled nondeterminisms
  - Stop-the-world GC, record and replay DPLs of GC points
  - Deterministic hash code
  - Deterministic “logical time”
  - Deterministic I/O
  - etc.
- Adaptive compilation and dynamic classloading
  - Most challenging (esp. in Jikes)!
  - Fork-and-recompile

# Fork-and-recompile



# Evaluation

- Implementation in Jikes RVM
  - Publicly available (<http://sourceforge.net/p/jikesrvm/research-archive/49/>)
- DaCapo 2006 & 2009, SPEC JBB 2000 & 2005
- 64 cores (AMD Opteron 6272)

# REPLAY Efficacy

Benchmark	REPLAY Success Rate
	Default w/ value logging
hsqldb6	100%
lusearch6	100%
xalan6	100%
avrorra9	100%
jython9	100%
luindex9	100%
lusearch9	100%
pmd9	100%
sunflow9	100%
xalan9	60%
pjbb2000	100%
pjbb2005	100%

# REPLAY Efficacy

Benchmark	REPLAY Success Rate	
	Default w/ value logging	Ignore HB edge w/ value logging
hsqldb6	100%	0%
lusearch6	100%	0%
xalan6	100%	0%
avrorra9	100%	0%
jython9	100%	0%
luindex9	100%	0%
lusearch9	100%	0%
pmd9	100%	0%
sunflow9	100%	0%
xalan9	60%	0%
pjbb2000	100%	0%
pjbb2005	100%	0%


# RECORD logging throughput

Benchmark	RECORD Logging MB/s
hsqldb6	0.7
lusearch6	<0.1
xalan6	7.7
avrorra9	2.5
gython9	<0.1
luindex9	<0.1
lusearch9	<0.1
pmd9	0.1
sunflow9	0.1
xalan9	9.7
pjbb2000	1.1
pjbb2005	4.8

# RECORD logging throughput


Benchmark	RECORD Logging MB/s
hsqldb6	0.7
lusearch6	<0.1
xalan6	7.7
avrorra9	2.5
gython9	<0.1
luindex9	<0.1
lusearch9	<0.1
pmd9	0.1
sunflow9	0.1
xalan9	9.7
pjbb2000	1.1
pjbb2005	4.8

# Performance





# Performance



# Conclusion

- Handle both external and internal nondeterminisms
  - Metacircular JVM
  - fork-and-recompile
- Efficient record and replay
  - Low overhead in RECORD
  - More parallelism in REPLAY
- Overcome many limitations simultaneously
  - Online/offline, software-only, no speculation, etc.