Leak Pruning

Michael Bond Kathryn McKinley

The University of Texas at Austin

Deployed Software Fails

Deployed Software Fails

- Driverless truck
  - 10,000 lines of C#

http://www.codeproject.com/KB/showcase/IfOnlyWedUsedANTSProfiler.aspx

Deployed Software Fails

- Driverless truck
  - 10,000 lines of C#
- Leak: past obstacles remained reachable

http://www.codeproject.com/KB/showcase/IfOnlyWedUsedANTSProfiler.aspx

Deployed Software Fails

- Driverless truck
  - 10,000 lines of C#
- Leak: past obstacles remained reachable
- No immediate symptoms

"This problem was pernicious because it only showed up after 40 minutes to an hour of driving around and collecting obstacles."

http://www.codeproject.com/KB/showcase/IfOnlyWedUsedANTSProfiler.aspx

Deployed Software Fails

- Driverless truck
  - 10,000 lines of C#
- Leak: past obstacles remained reachable
- No immediate symptoms

"This problem was pernicious because it only showed up after 40 minutes to an hour of driving around and collecting obstacles."

- Quick "fix": restart after 40 minutes

http://www.codeproject.com/KB/showcase/IfOnlyWedUsedANTSProfiler.aspx

Deployed Software Fails

- Driverless truck
  - 10,000 lines of C#
- Leak: past obstacles remained reachable
- No immediate symptoms

  "This problem was pernicious because it only showed up after 40 minutes to an hour of driving around and collecting obstacles."
- Quick "fix": restart after 40 minutes
- Environment sensitive
  - More obstacles in deployed setting
  - Unresponsive after 28 minutes

http://www.codeproject.com/KB/showcase/IfOnlyWedUsedANTSProfiler.aspx

Deployed Software Fails

- Driverless truck
  - 10,000
- Leak: pa hable
- No immediate symptoms

  "This problem was p wed up after 40 minutes to an hour of driving cles."
- Quick "fix": nutes
- Environment sensitive
  - More obst
  - Unrespon

Different environments & inputs → different behavior

Hard to fix before deployment

Uncertainty in deployed systems

Tolerating Memory Leaks

- Deployed systems have leaks
  - Critical systems need immediate help
- **Leak pruning** tolerates bad effects
  - Reclaims memory automatically
  - High precision & low overhead
  - Bounds resources
  - Preserves semantics

Outline

- Why tolerate leaks
- Why leaks are so bad
- How leak pruning works
- How leak pruning predicts leaked objects
- Evaluation

Memory Leaks in Deployed Systems

- Memory leaks are a real problem
  - Managed languages do not eliminate them

Memory Leaks in Deployed Systems

- Memory leaks are a real problem
  - Managed languages do not eliminate them

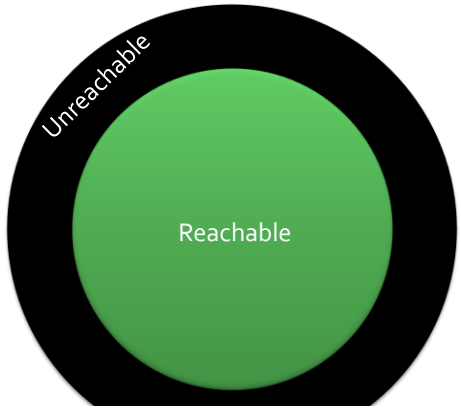

Memory Leaks in Deployed Systems

- Memory leaks are a real problem
  - Managed languages do not eliminate them

Memory Leaks in Deployed Systems

- Memory leaks are a real problem
  - Managed languages do not eliminate them

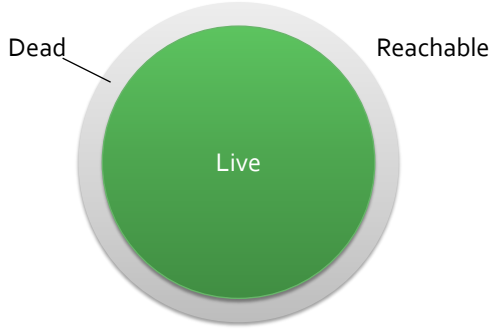

Memory Leaks in Deployed Systems

- Memory leaks are a real problem
  - Managed languages do not eliminate them

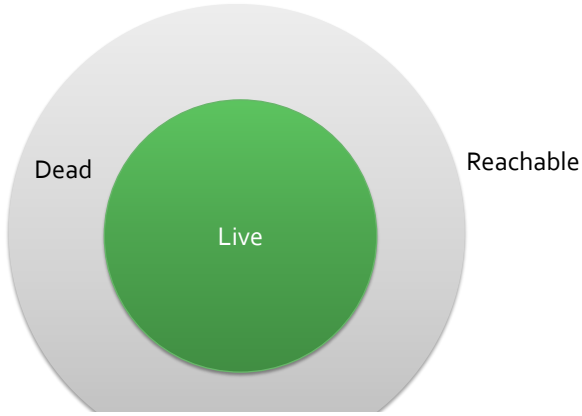

Memory Leaks in Deployed Systems

- Memory leaks are a real problem
  - Managed languages do not eliminate them

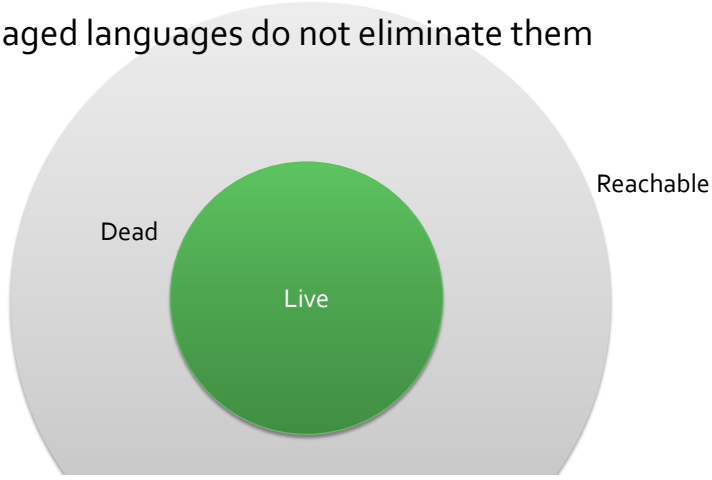

Memory Leaks in Deployed Systems

- Memory leaks are a real problem
  - Managed languages do not eliminate them

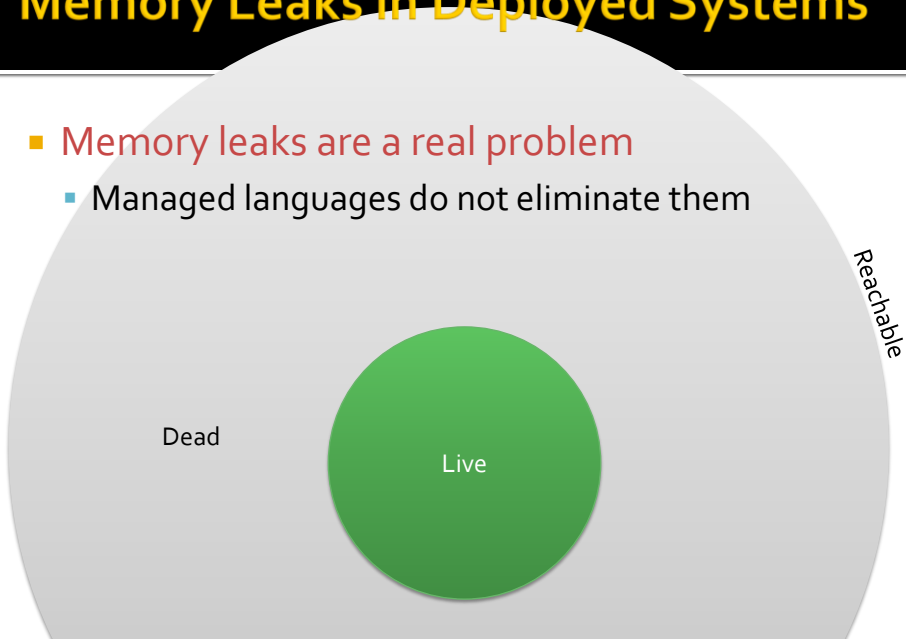

Memory Leaks in Deployed Systems

- Memory leaks are a real problem
  - Managed languages do not eliminate them
  - Slow & crash real programs

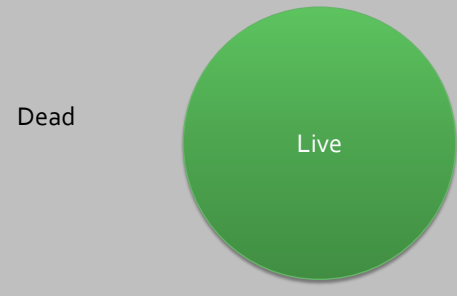

Memory Leaks in Deployed Systems

- Memory leaks are a real problem
  - Managed languages do not eliminate them
  - Slow & crash real programs
  - Unacceptable for some applications

Memory Leaks in Deployed Systems

- Memory leaks are a real problem
  - Managed languages do not eliminate them
  - Slow & crash real programs
  - Unacceptable for some applications

- Fixing leaks is hard
  - Leaks take time to materialize
  - Failure far from cause

Outline

- Why tolerate leaks
- Why leaks are so bad
- How leak pruning works
- How leak pruning predicts leaked objects
- Evaluation

Reclaiming Memory while Preserving Semantics

- Garbage collection based on liveness

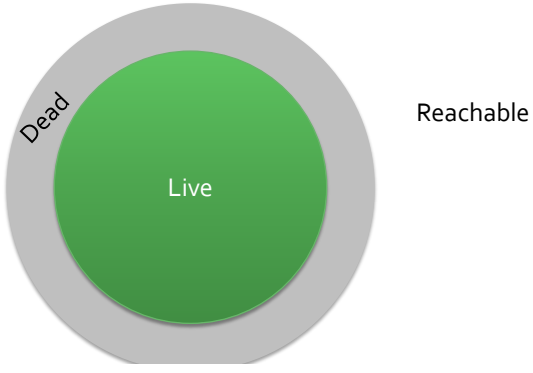

Reclaiming Memory while Preserving Semantics

- Garbage collection based on liveness
  - Take action when memory exhausted

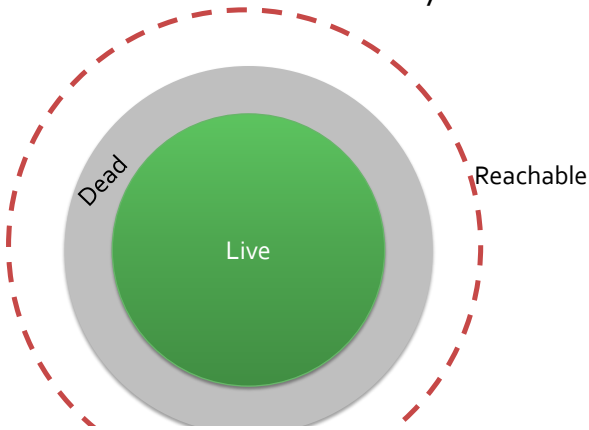

Reclaiming Memory while Preserving Semantics

- Garbage collection based on liveness
  - Take action when memory exhausted

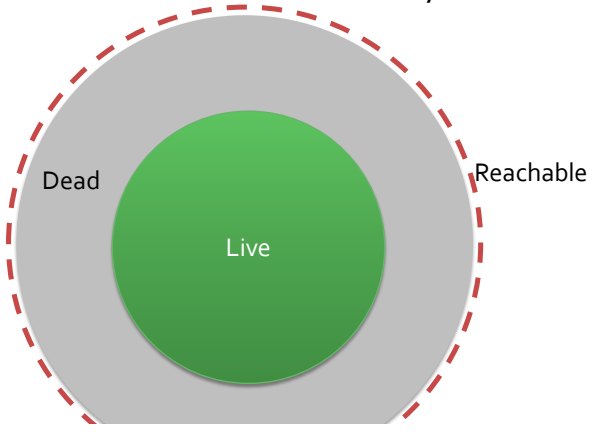

Reclaiming Memory while Preserving Semantics

- Garbage collection based on liveness
  - Take action when memory exhausted

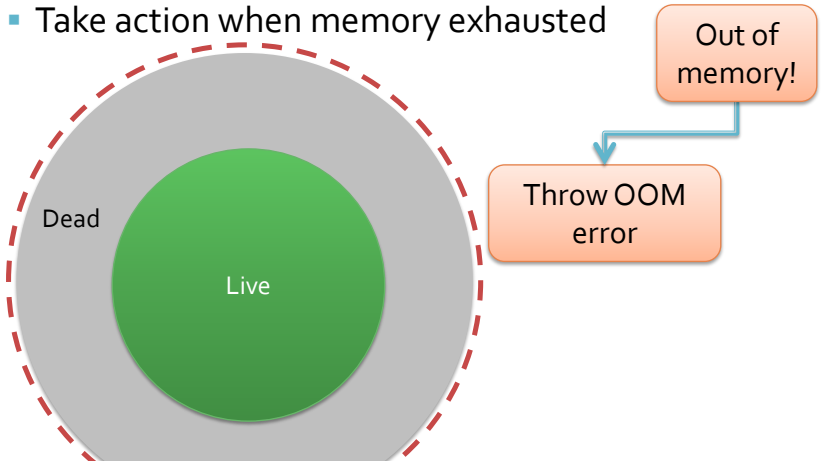

Reclaiming Memory while Preserving Semantics

- Garbage collection based on liveness
  - Take action when memory exhausted

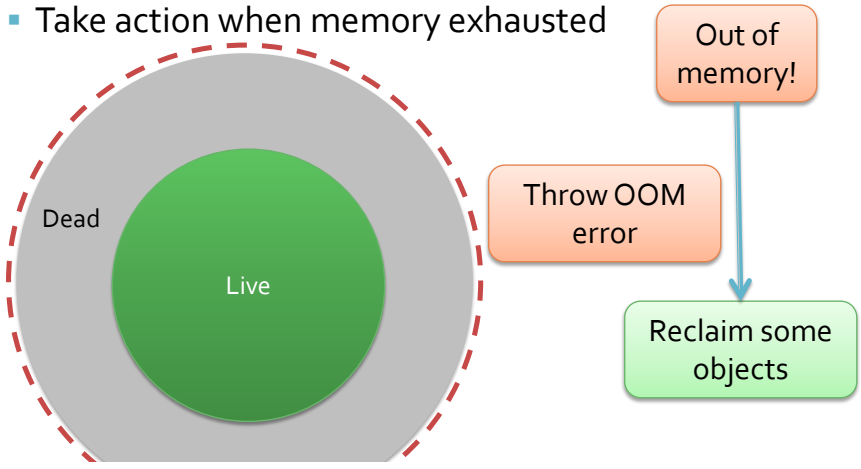

Reclaiming Memory while Preserving Semantics

- Garbage collection based on liveness
  - Reclaim predicted dead objects

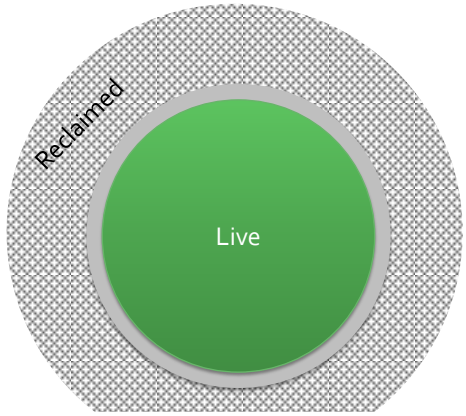

Reclaiming Memory while Preserving Semantics

- Garbage collection based on liveness
  - Reclaim predicted dead objects

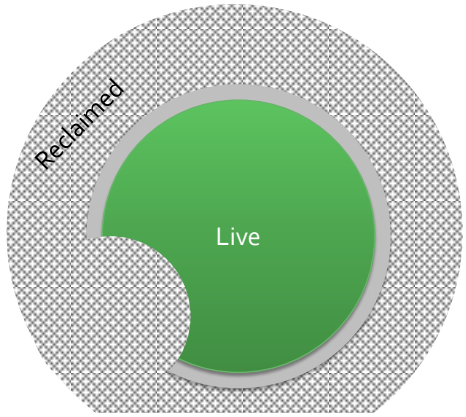

Reclaiming Memory while Preserving Semantics

- Garbage collection based on liveness
  - Reclaim predicted dead objects

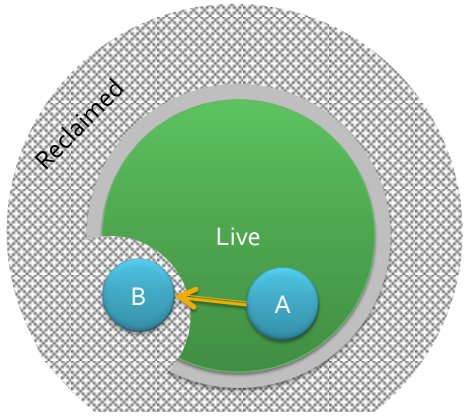

Reclaiming Memory while Preserving Semantics

- Garbage collection based on liveness
  - Reclaim predicted dead objects

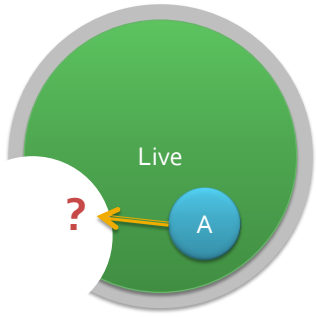

Reclaiming Memory while Preserving Semantics

- Garbage collection based on liveness
  - Reclaim predicted dead objects
  - *Poison* references to reclaimed objects

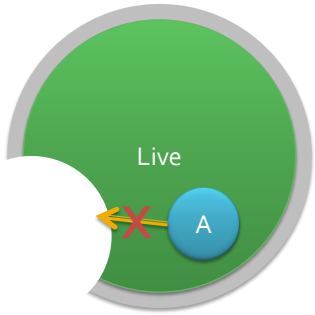

Reclaiming Memory while Preserving Semantics

- Garbage collection based on liveness
  - Reclaim predicted dead objects
  - *Poison* references to reclaimed objects

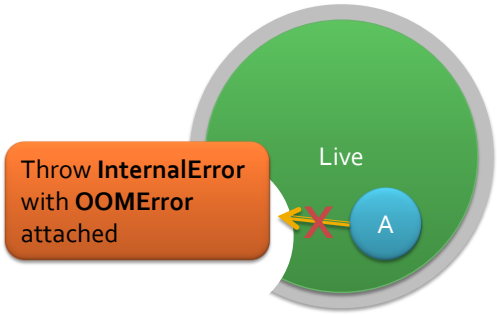

Reclaiming Memory while Preserving Semantics

- Garbage collection based on liveness
  - Reclaim predicted dead objects
  - *Poison* references to reclaimed objects

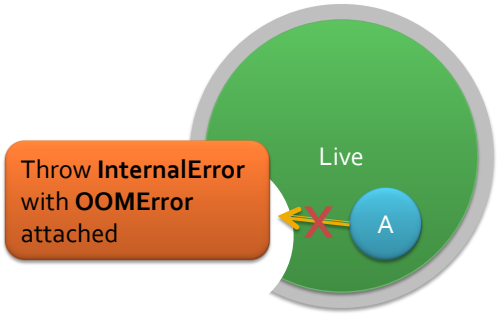

Worst case: defers fatal errors

Best case: keeps leaky programs running much longer or indefinitely

Outline

- Why tolerate leaks
- Why leaks are so bad
- How leak pruning works
- How leak pruning predicts leaked objects
- Evaluation

Predicting Dead Objects

- Predicting the future
  - Leaked objects → not used again
  - Highly *stale* objects → *likely* leaked

[Chilimbi & Hauswirth '04] [Qin et al. '05] [Bond & McKinley '06]

Predicting Dead Objects

- Predicting the future
  - Leaked objects → not used again
  - Highly *stale* objects → *likely* leaked
- Alternative: offload to disk

[**Melt,** Bond & McKinley '08]
[**LeakSurvivor,** Tang et al. '08]
[**Panacea,** Goldstein et al. '07, Breitgand et al. '07]
  - Tolerates mispredictions
  - Eventually exhausts disk (if disk at all)

Predicting Dead Objects

- Predicting the future
  - Leaked objects → not used again
  - Highly *stale* objects → *likely* leaked
- Alternat
  [Melt, Bo
  [LeakSur
  [Panacea '07]
  - Tolera
  - Eventu

Need high precision

One misprediction: program terminates

Predicting Dead Data Structures

- Identify references to prune
  - Roots of leaked data structures
  - Categorize by *reference type*

Predicting Dead Data Structures

- Identify references to prune
  - Roots of leaked data structures
  - Categorize by *reference type*
- Criteria
  - Highly stale references
    - More stale than *most stale instance accessed previously*

Predicting Dead Data Structures

- Identify references to prune
  - Roots of leaked data structures
  - Categorize by *reference type*
- Criteria
  - Highly stale references
    - More stale than *most stale instance accessed previously*
  - Reference type keeping most bytes reachable

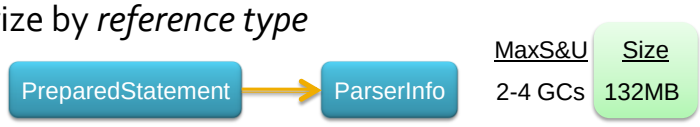

Predicting Dead Data Structures

- Identify references to prune
  - Roots of leaked data structures
  - Categorize by *reference type*
- Criteria

MaxS&U	Size
2-4 GCs	132MB

  - Highly stale references
    - More stale than *most stale instance accessed previously*
  - Reference type keeping most bytes reachable

Piggyback on GC: two-phase transitive closure

MySQL Leak

- Based on colleague's JDBC application
- Leak: SQL statements remain in set

MySQL Leak

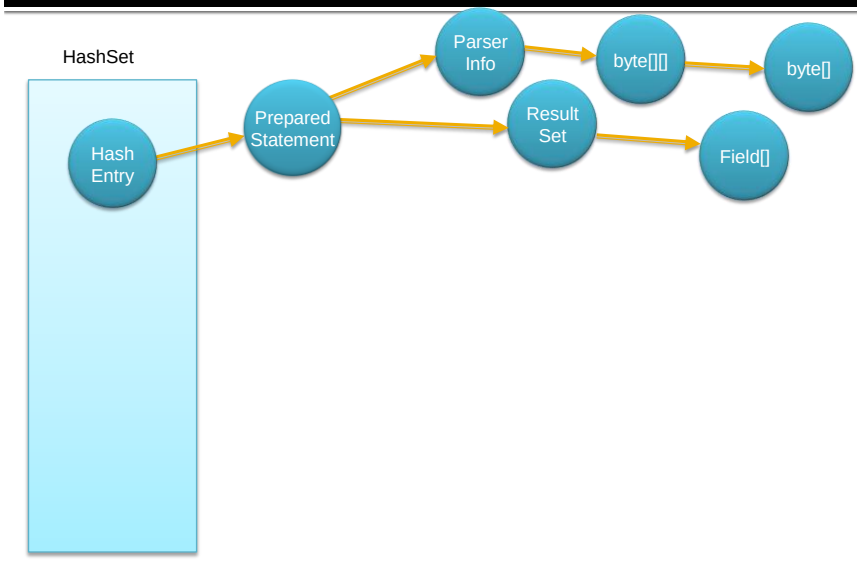

MySQL Leak

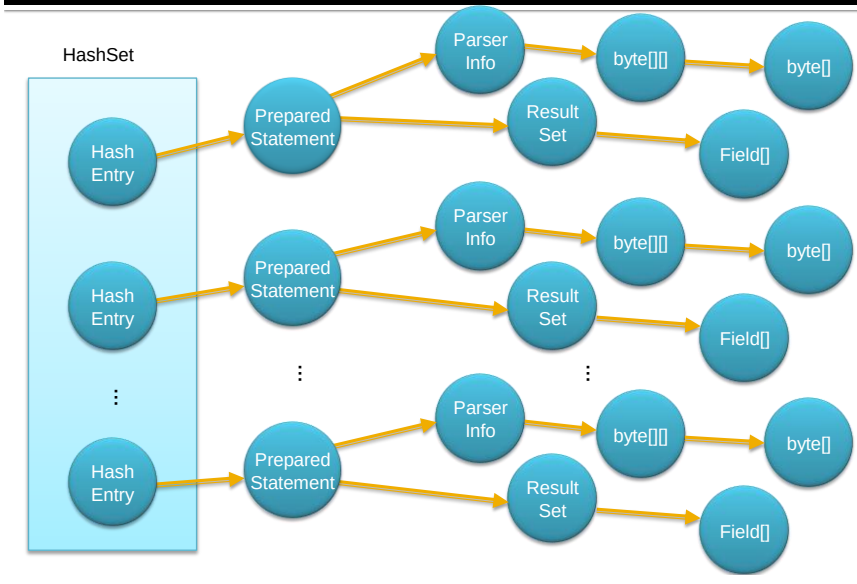

Transitive Closure: Phase One

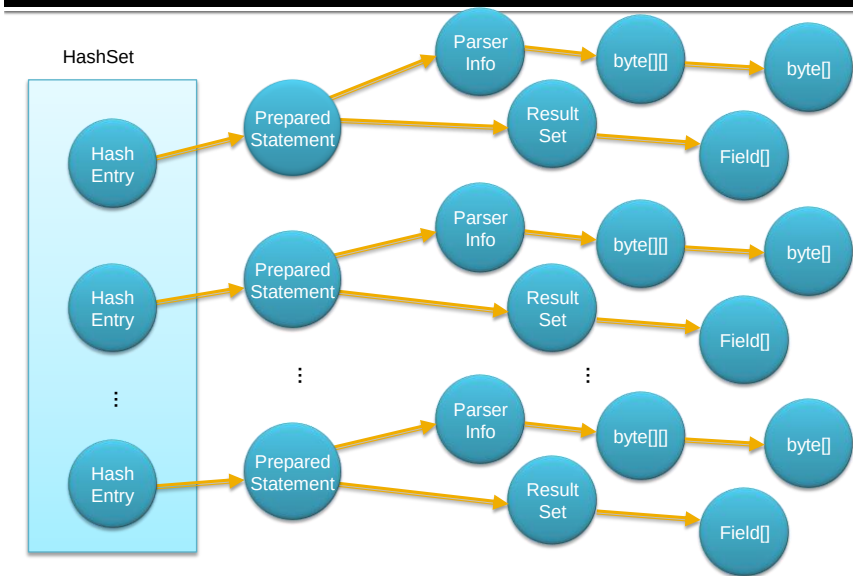

Transitive Closure: Phase One

		MaxS&U
HashEntry	PreparedStatement	16-32
PreparedStatement	ParserInfo	0-1
PreparedStatement	ResultSet	0-1

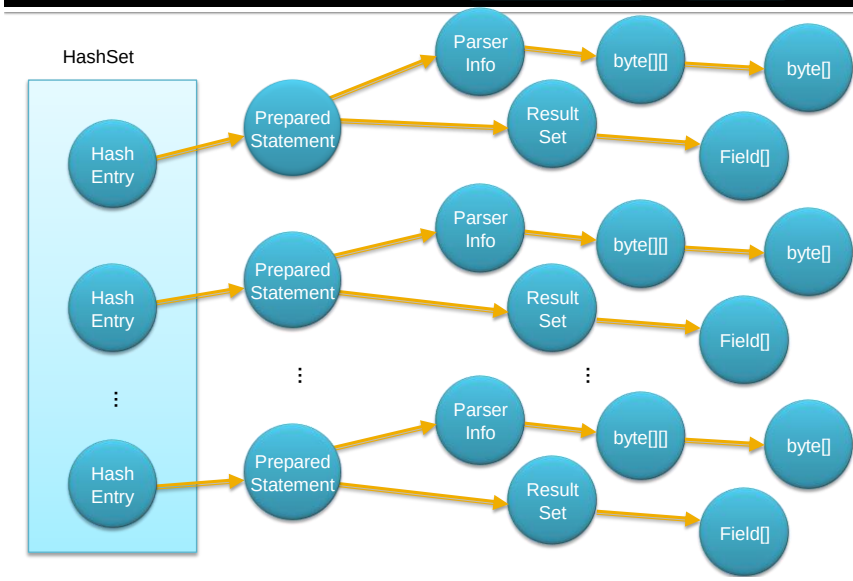

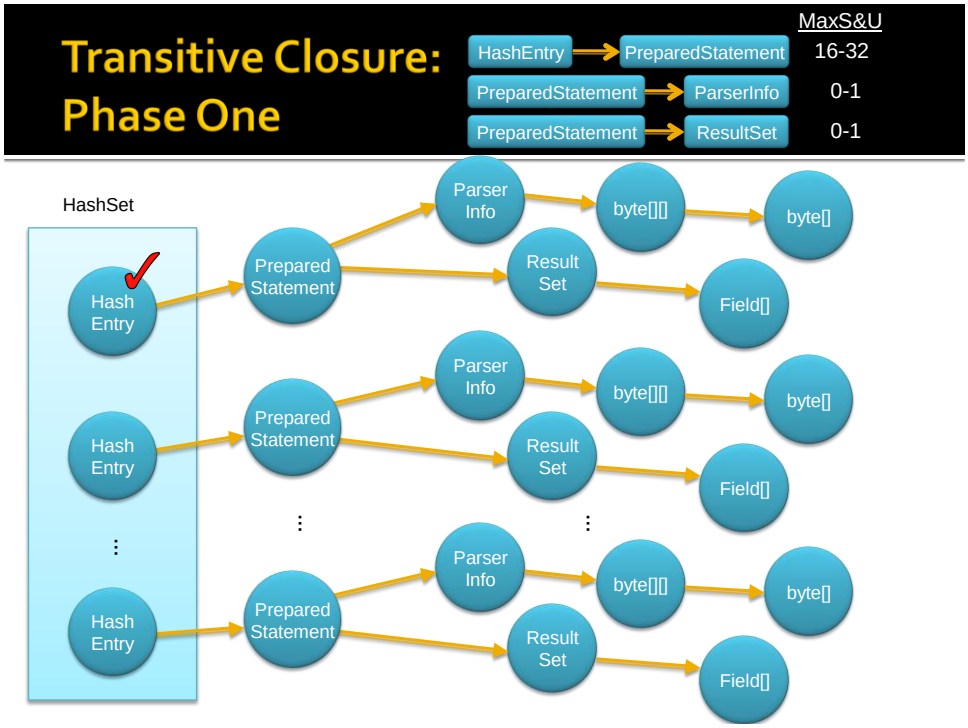
Transitive Closure: Phase One
MaxS&U
HashEntry
PreparedStatement
16-32
PreparedStatement
ParserInfo
0-1
PreparedStatement
ResultSet
0-1
HashSet
Hash Entry
Hash Entry
Hash Entry
Prepared Statement
Prepared Statement
Prepared Statement
Parser Info
Parser Info
Parser Info
Result Set
Result Set
Result Set
byte[][]
byte[][]
byte[][]
byte[]
byte[]
byte[]
Field[]
Field[]
Field[]

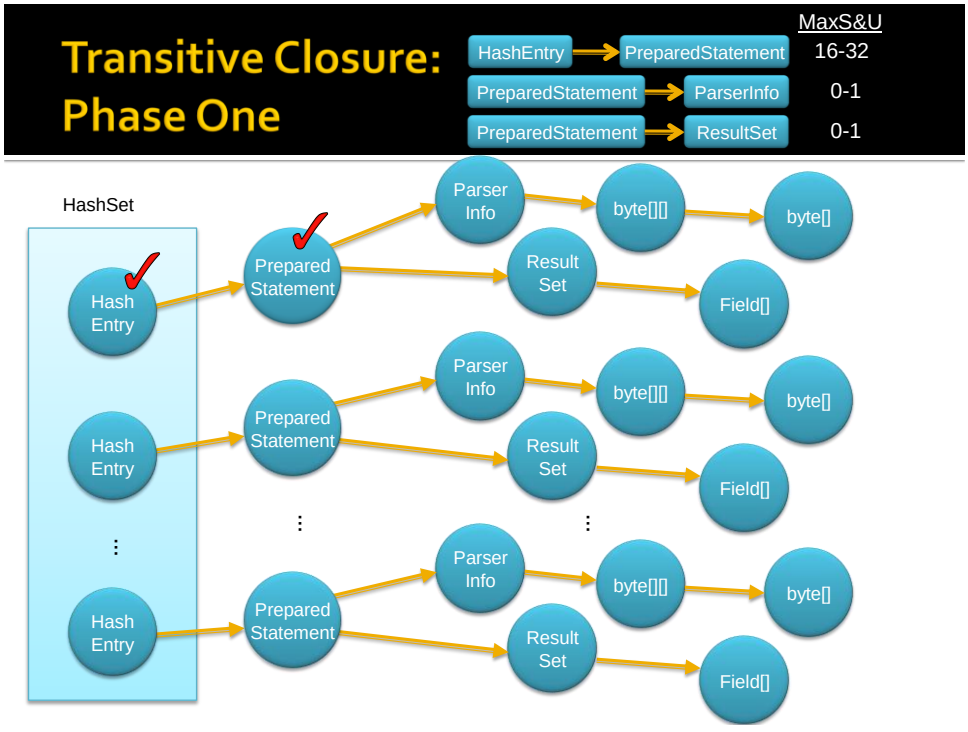
Transitive Closure: Phase One
MaxS&U
HashEntry
PreparedStatement
16-32
PreparedStatement
ParserInfo
0-1
PreparedStatement
ResultSet
0-1
HashSet
Hash Entry
Hash Entry
Hash Entry
Prepared Statement
Prepared Statement
Prepared Statement
Parser Info
Parser Info
Parser Info
Result Set
Result Set
Result Set
byte[][]
byte[][]
byte[][]
byte[]
byte[]
byte[]
Field[]
Field[]
Field[]

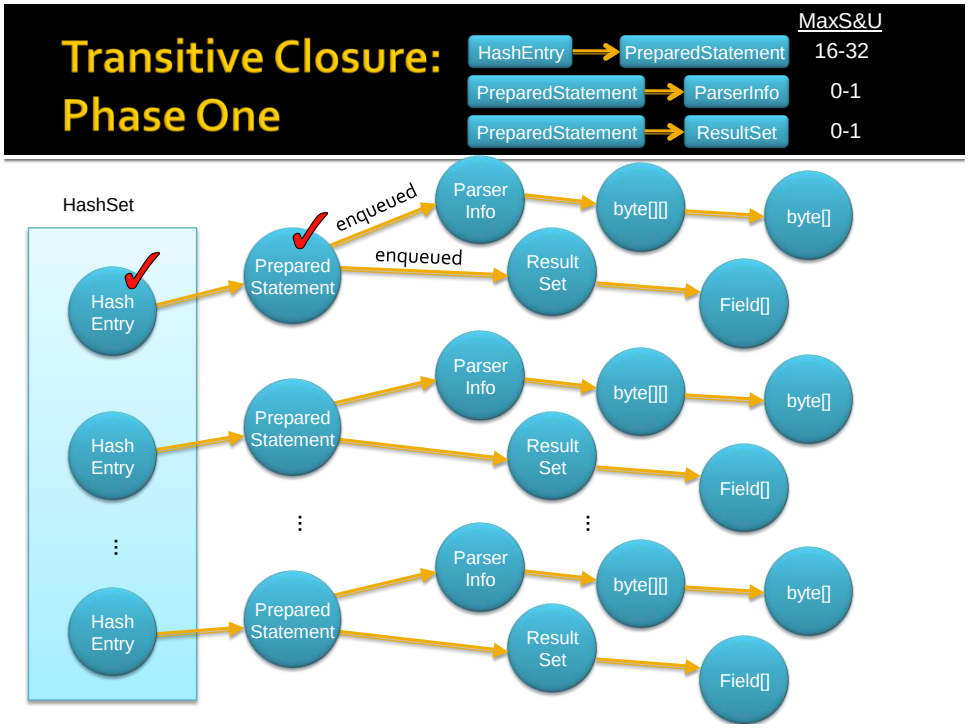
Transitive Closure:
Phase One
MaxS&U
HashEntry
PreparedStatement
16-32
PreparedStatement
ParserInfo
0-1
PreparedStatement
ResultSet
0-1
HashSet
Hash Entry
Prepared Statement
enqueued
enqueued
Parser Info
Result Set
byte[][]
byte[]
Field[]

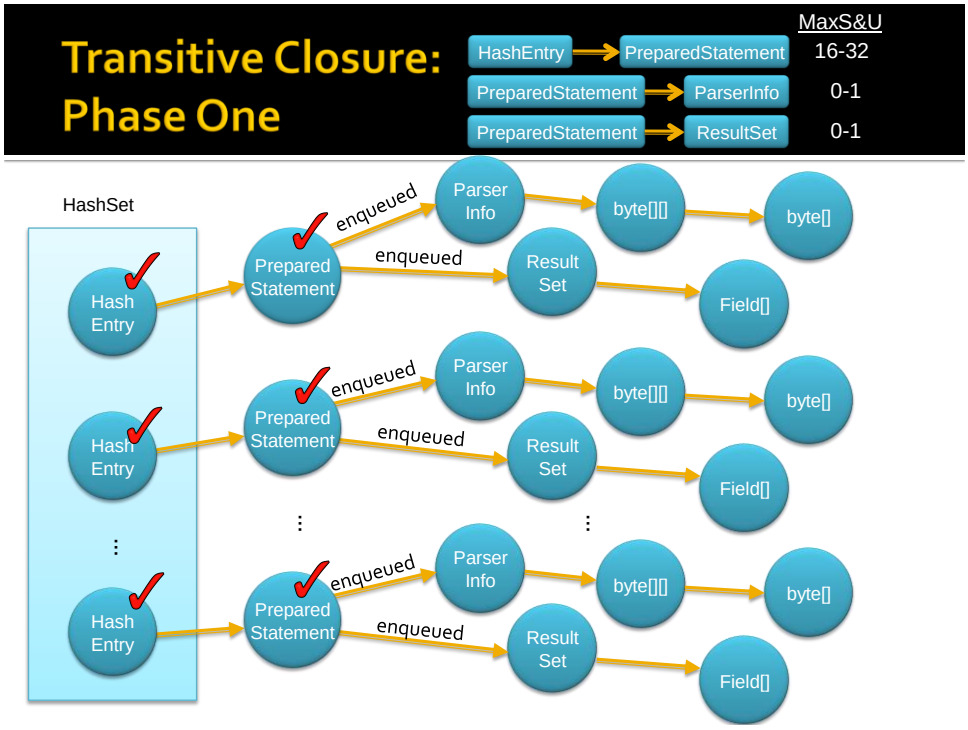
Transitive Closure:
Phase One
MaxS&U
HashEntry
PreparedStatement
16-32
PreparedStatement
ParserInfo
0-1
PreparedStatement
ResultSet
0-1
HashSet
Hash Entry
Prepared Statement
enqueued
enqueued
Parser Info
Result Set
byte[][]
byte[]
Field[]

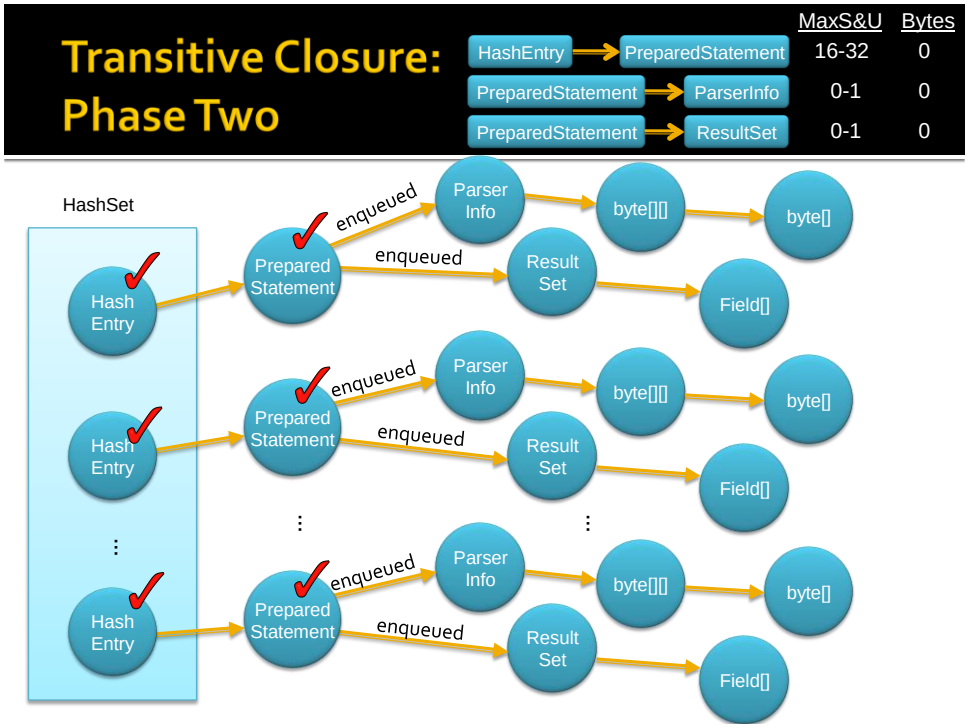
Transitive Closure: Phase Two
HashEntry → PreparedStatement
PreparedStatement → ParserInfo
PreparedStatement → ResultSet
MaxS&U
Bytes
16-32
0
0-1
0
0-1
0
HashSet
Hash Entry
Prepared Statement
enqueued
Parser Info
Result Set
byte[][]
byte[]
Field[]

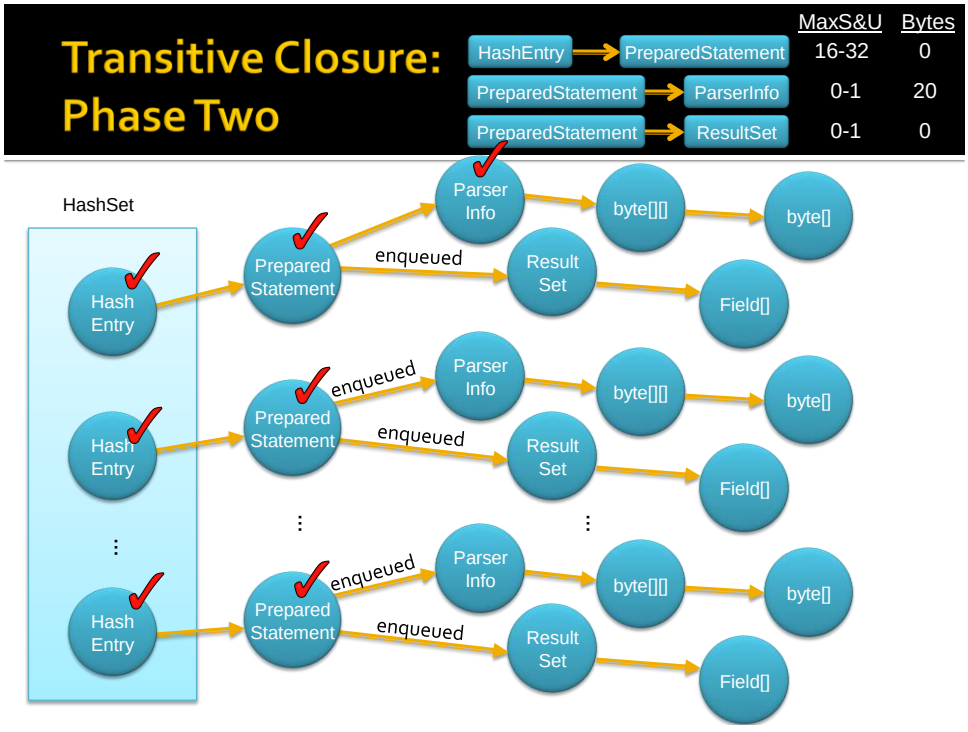
Transitive Closure: Phase Two
HashEntry → PreparedStatement
PreparedStatement → ParserInfo
PreparedStatement → ResultSet
MaxS&U
Bytes
16-32
0
0-1
20
0-1
0
HashSet
Hash Entry
Prepared Statement
enqueued
Parser Info
Result Set
byte[][]
byte[]
Field[]

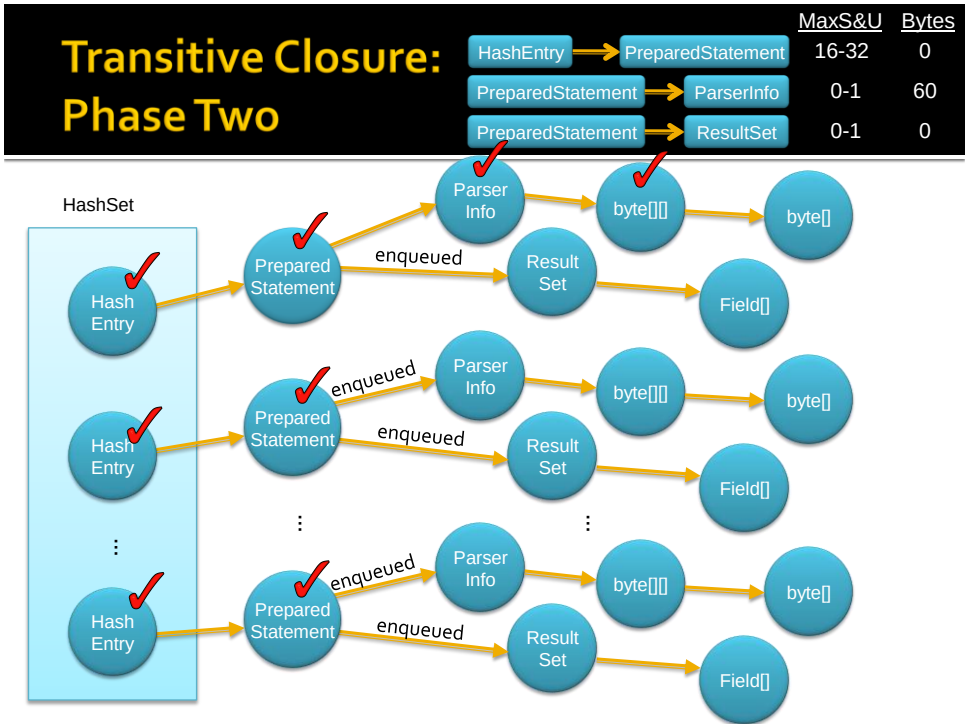
Transitive Closure:
Phase Two
HashEntry
PreparedStatement
PreparedStatement
ParserInfo
PreparedStatement
ResultSet
MaxS&U
Bytes
16-32
0
0-1
60
0-1
0
HashSet
Hash Entry
Prepared Statement
Parser Info
Result Set
byte[][]
byte[]
Field[]
enqueued

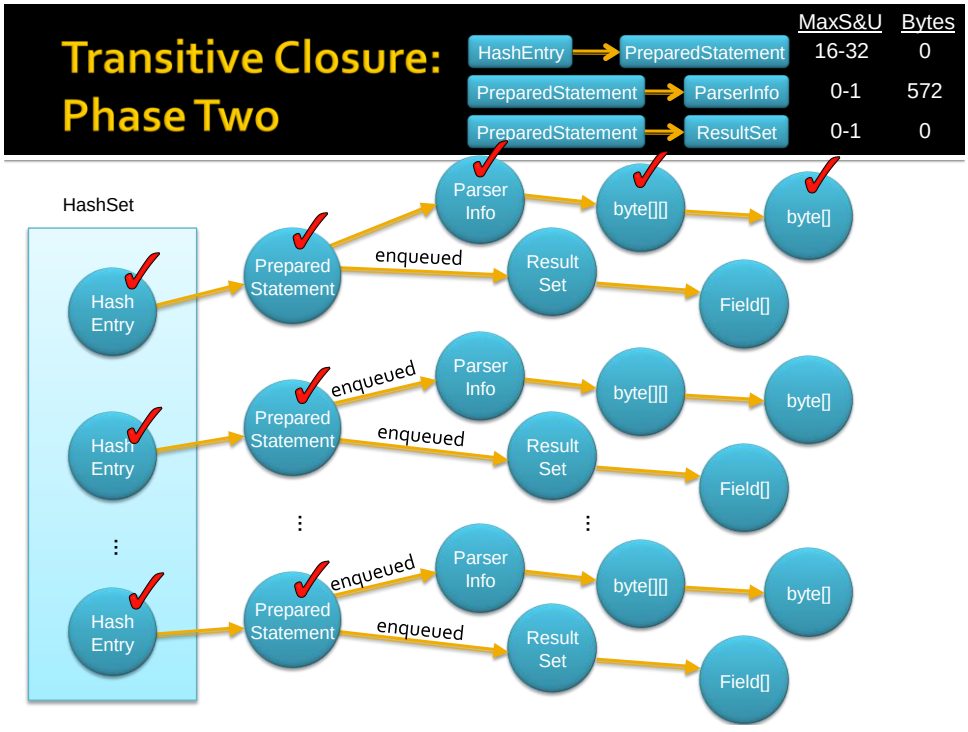
Transitive Closure:
Phase Two
HashEntry
PreparedStatement
PreparedStatement
ParserInfo
PreparedStatement
ResultSet
MaxS&U
Bytes
16-32
0
0-1
572
0-1
0
HashSet
Hash Entry
Prepared Statement
Parser Info
Result Set
byte[][]
byte[]
Field[]
enqueued

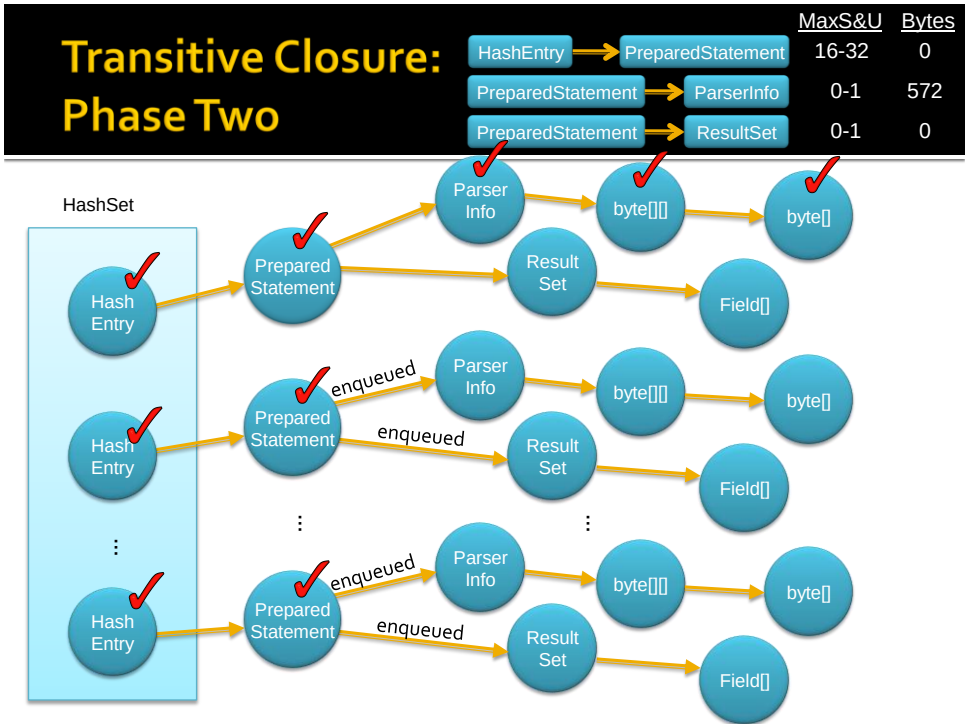
Transitive Closure:
Phase Two
HashEntry
PreparedStatement
MaxS&U
16-32
Bytes
0
PreparedStatement
ParserInfo
0-1
572
PreparedStatement
ResultSet
0-1
0
HashSet
Hash Entry
Prepared Statement
Parser Info
Result Set
byte[][]
byte[]
Field[]
enqueued
enqueued
enqueued
enqueued

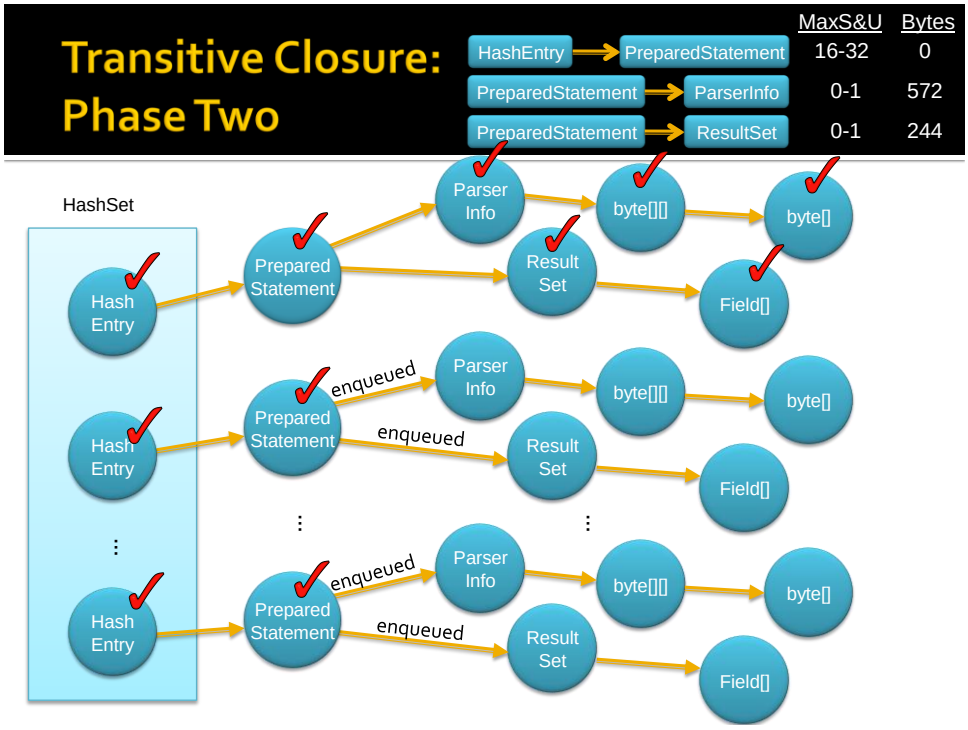
Transitive Closure:
Phase Two
HashEntry
PreparedStatement
MaxS&U
16-32
Bytes
0
PreparedStatement
ParserInfo
0-1
572
PreparedStatement
ResultSet
0-1
244
HashSet
Hash Entry
Prepared Statement
Parser Info
Result Set
byte[][]
byte[]
Field[]
enqueued
enqueued
enqueued
enqueued

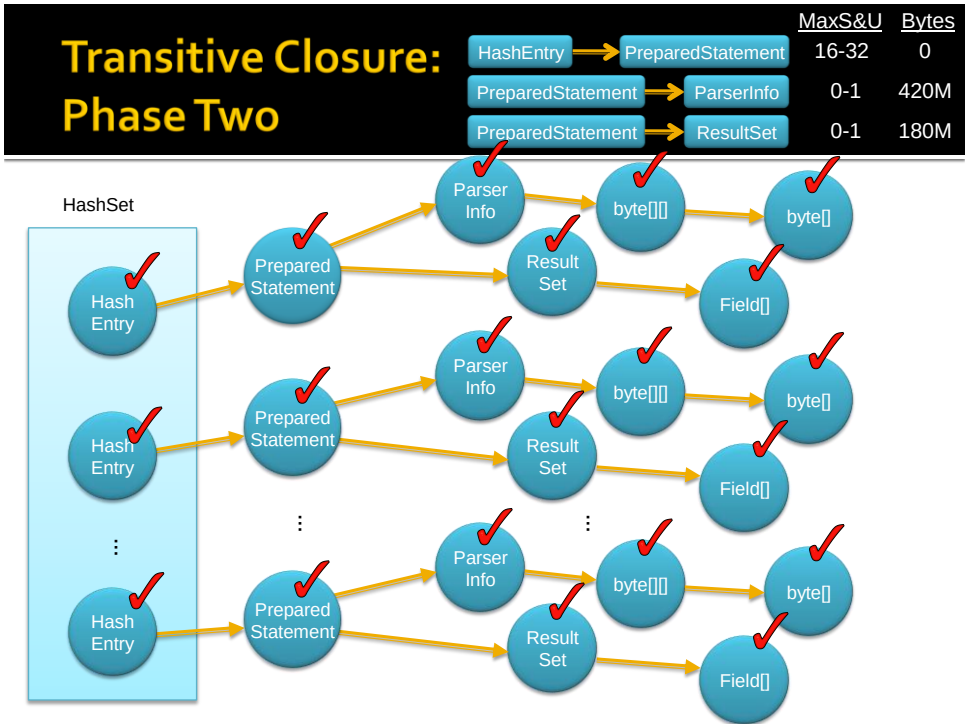
Transitive Closure: Phase Two
MaxS&U
Bytes
HashEntry
PreparedStatement
16-32
0
PreparedStatement
ParserInfo
0-1
420M
PreparedStatement
ResultSet
0-1
180M
HashSet
Hash Entry
Prepared Statement
Parser Info
Result Set
byte[][]
byte[]
Field[]

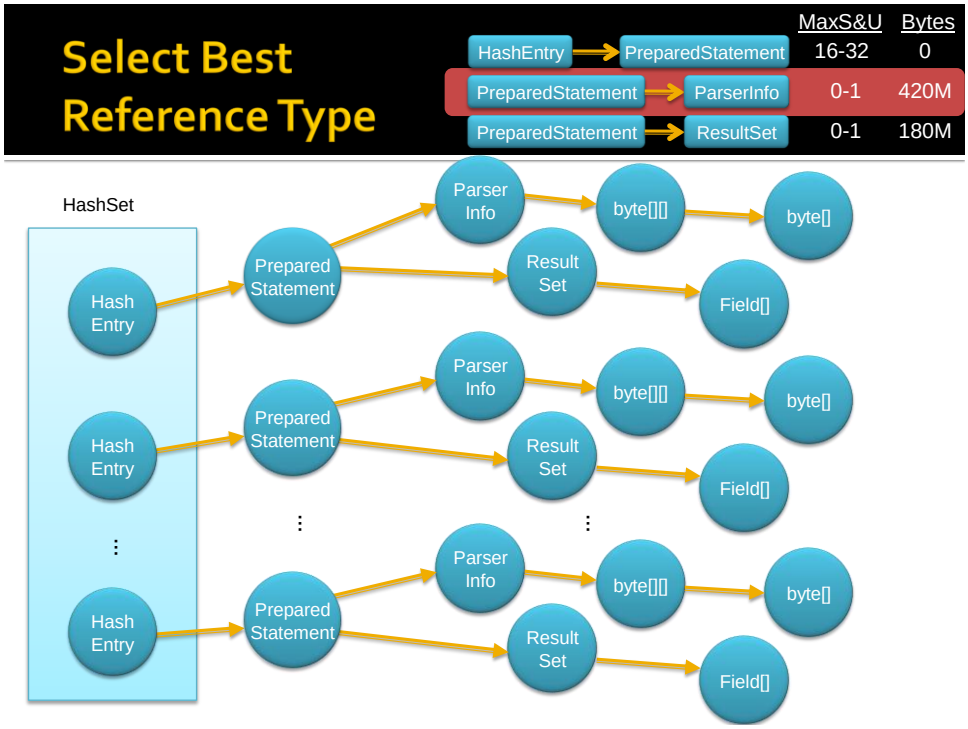
Select Best Reference Type
MaxS&U
Bytes
HashEntry
PreparedStatement
16-32
0
PreparedStatement
ParserInfo
0-1
420M
PreparedStatement
ResultSet
0-1
180M
HashSet
Hash Entry
Prepared Statement
Parser Info
Result Set
byte[][]
byte[]
Field[]

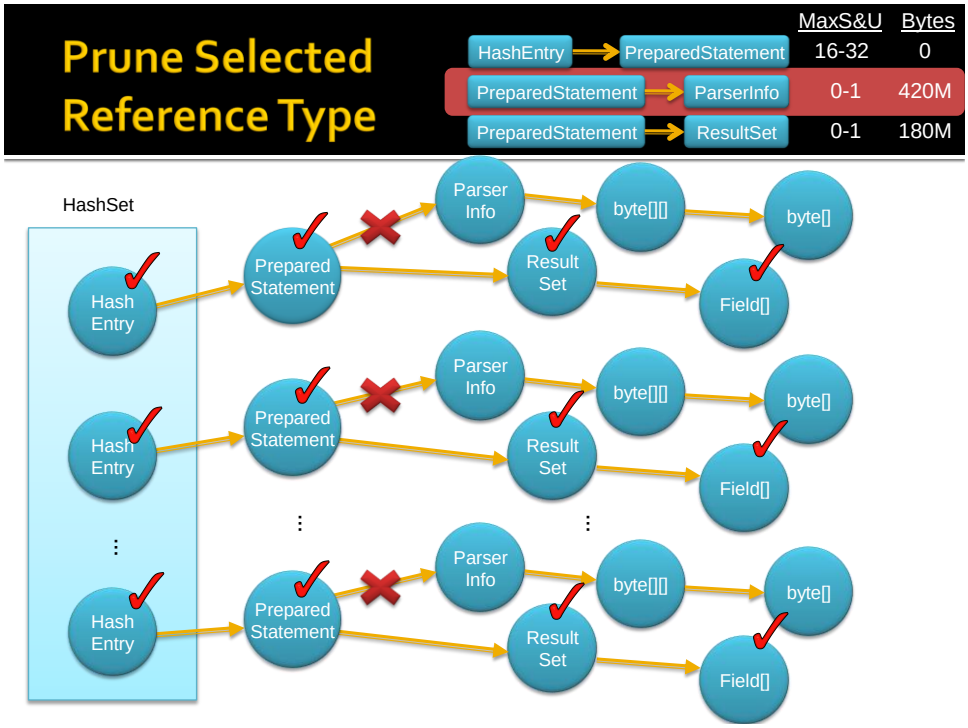
Prune Selected Reference Type
MaxS&U
Bytes
HashEntry
PreparedStatement
16-32
0
PreparedStatement
ParserInfo
0-1
420M
PreparedStatement
ResultSet
0-1
180M
HashSet
Hash Entry
Prepared Statement
Parser Info
Result Set
byte[][]
byte[]
Field[]

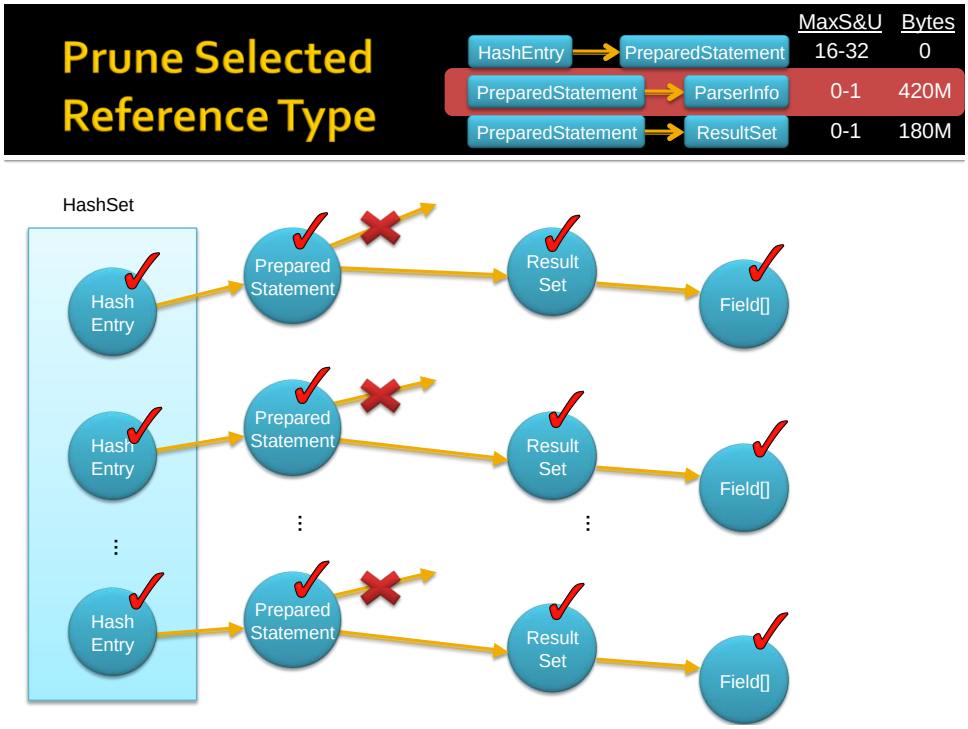
Prune Selected Reference Type
MaxS&U
Bytes
HashEntry
PreparedStatement
16-32
0
PreparedStatement
ParserInfo
0-1
420M
PreparedStatement
ResultSet
0-1
180M
HashSet
Hash Entry
Prepared Statement
Result Set
Field[]

Leak Partially Eliminated

		MaxS&U	Bytes
HashEntry	PreparedStatement	16-32	0
PreparedStatement	ParserInfo	0-1	0
PreparedStatement	ResultSet	0-1	0

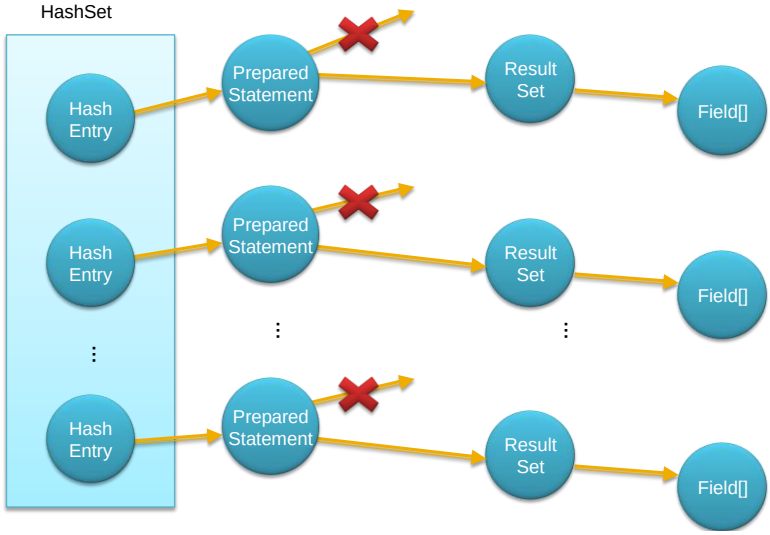

Transitive Closure: Phase One

		MaxS&U	Bytes
HashEntry	PreparedStatement	16-32	0
PreparedStatement	ParserInfo	0-1	0
PreparedStatement	ResultSet	0-1	0

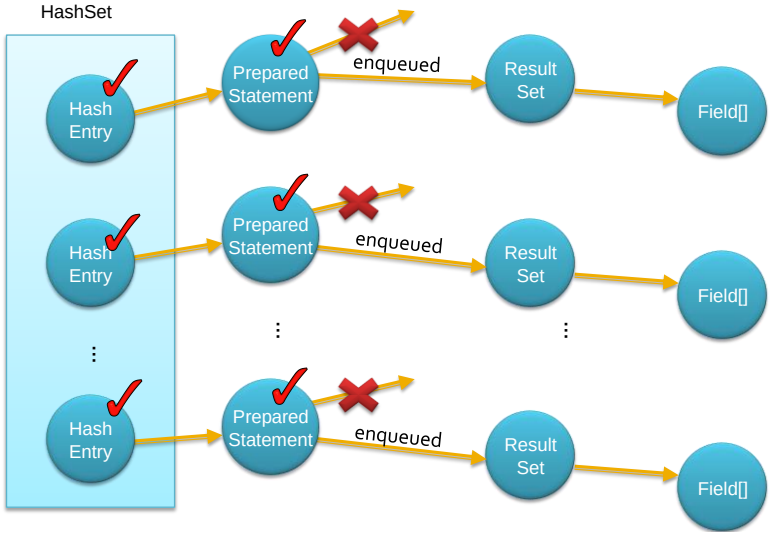

Transitive Closure: Phase Two

	MaxS&U	Bytes
HashEntry → PreparedStatement	16-32	0
PreparedStatement → ParserInfo	0-1	50M
PreparedStatement → ResultSet	0-1	220M

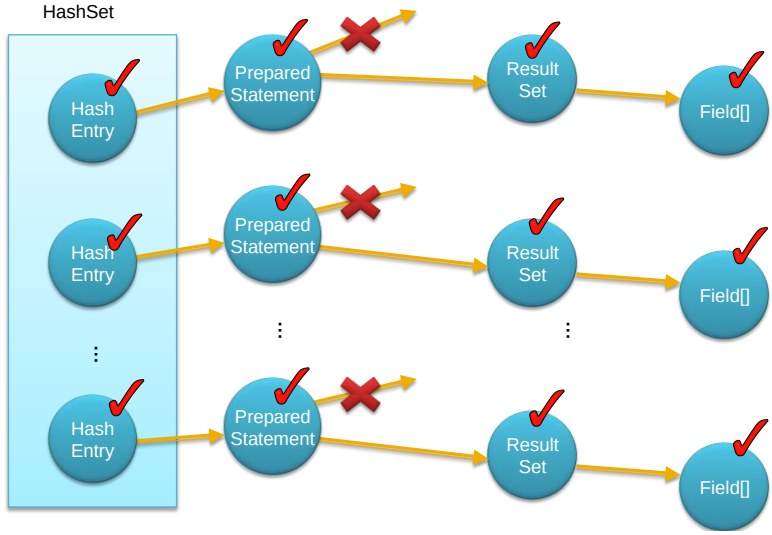

Select Best Reference Type

	MaxS&U	Bytes
HashEntry → PreparedStatement	16-32	0
PreparedStatement → ParserInfo	0-1	50M
PreparedStatement → ResultSet	0-1	220M

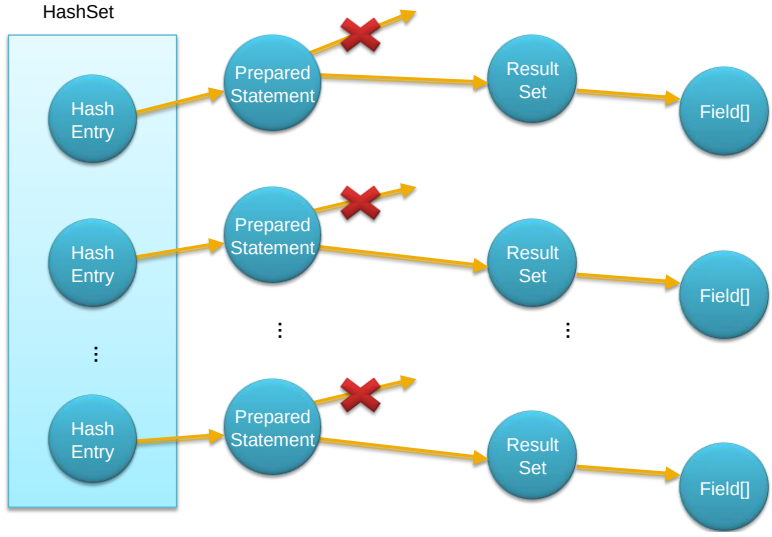

Prune Selected Reference Type

	MaxS&U	Bytes
HashEntry → PreparedStatement	16-32	0
PreparedStatement → ParserInfo	0-1	50M
PreparedStatement → ResultSet	0-1	220M

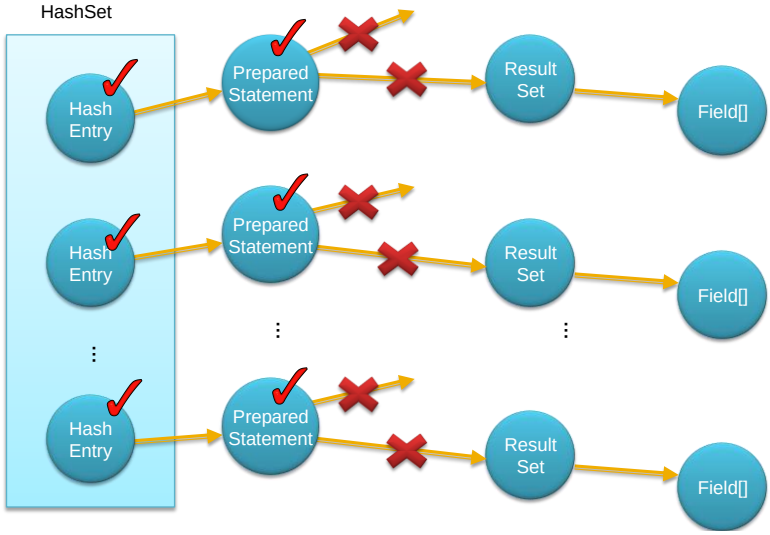

Prune Selected Reference Type

	MaxS&U	Bytes
HashEntry → PreparedStatement	16-32	0
PreparedStatement → ParserInfo	0-1	50M
PreparedStatement → ResultSet	0-1	220M

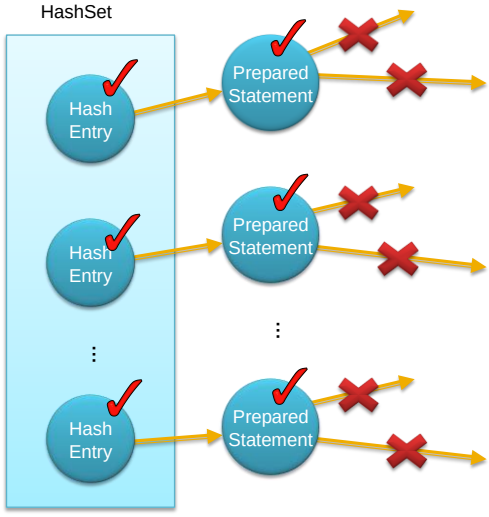

Prediction Summary

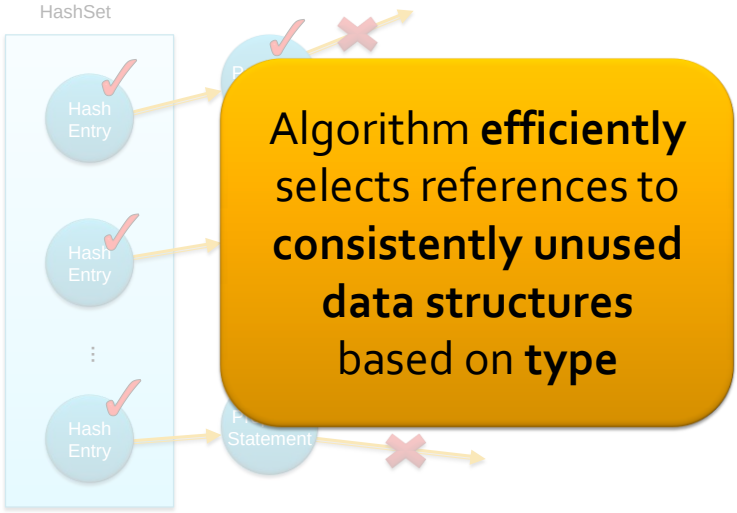

Outline

- Why tolerate leaks
- Why leaks are so bad
- How leak pruning works
- How leak pruning predicts leaked objects
- Evaluation

Implementation

- Leak pruning added to Jikes RVM 2.9.2

  `http://www.jikesrvm.org/Research+Archive`

Implementation

- Leak pruning added to Jikes RVM 2.9.2

  `http://www.jikesrvm.org/Research+Archive`

- Performance stress test
  - Non-leaking programs: DaCapo & SPEC benchmarks
  - Low overhead
    - 3% (Core 2) or 5% (Pentium 4)
    - Primarily read barriers

Tolerating Leaks

Leak	Leak pruning's effect
Eclipse "Diff"	Tolerates until 24-hr limit (>200X longer)
ListLeak	Tolerates until 24-hr limit (>25,000X longer)
SwapLeak	Tolerates until 24-hr limit (>2,200X longer)
Eclipse "Copy-Paste"	Most dead but some live (81X longer)
MySQL	Most dead but some live (35X longer)
JbbMod	All dead but pruning misses some (21X longer)
SPECjbb2000	Heap growth is mostly live (4.7X longer)
Mckoi Database	Thread leak: extra support needed (1.6X longer)
DualLeak	Heap growth is live (No help)

Tolerating Leaks

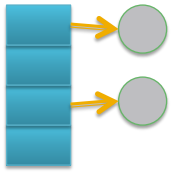

Eclipse "Copy-Paste"	Most dead but some live (81X longer)
MySQL	Most dead but some live (35X longer)

SPECjbb2000	Heap growth is mostly live (4.7X longer)

DualLeak	Heap growth is live (No help)

Tolerating Leaks

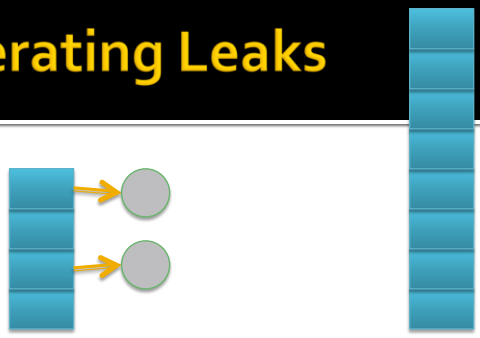

Eclipse “Copy-Paste”	Most dead but some live (81X longer)
MySQL	Most dead but some live (35X longer)
SPECjbb2000	Heap growth is mostly live (4.7X longer)
DualLeak	Heap growth is live (No help)

Tolerating Leaks

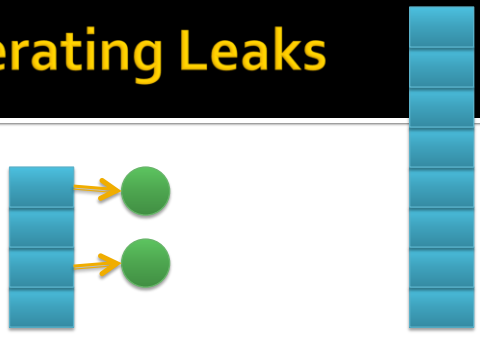

Eclipse “Copy-Paste”	Most dead but some live (81X longer)
MySQL	Most dead but some live (35X longer)
SPECjbb2000	Heap growth is mostly live (4.7X longer)
DualLeak	Heap growth is live (No help)

Tolerating Leaks

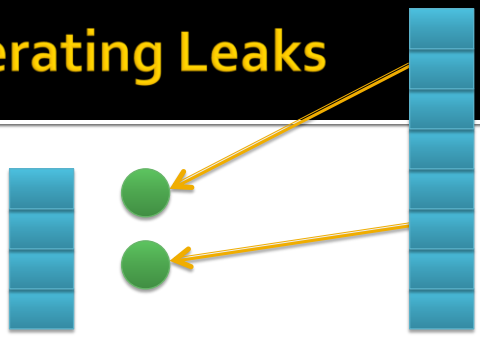

Eclipse “Copy-Paste”	Most dead but some live (81X longer)
MySQL	Most dead but some live (35X longer)

SPECjbb2000	Heap growth is mostly live (4.7X longer)

DualLeak	Heap growth is live (No help)

Tolerating Leaks

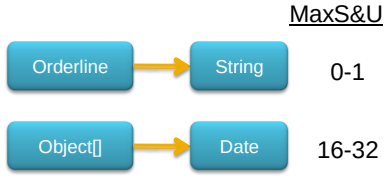

JbbMod	All dead but pruning misses some (21X longer)

Tolerating Leaks

Leak	Leak pruning's effect
Eclipse "Diff"	Tolerates until 24-hr limit (>200X longer)
ListLeak	Tolerates until 24-hr limit (>25,000X longer)
SwapLeak	Tolerates until 24-hr limit (>2,200X longer)
Eclipse "Copy-Paste"	Most dead but some live (81X longer)
MySQL	Most dead but some live (35X longer)
JbbMod	All dead but pruning misses some (21X longer)
SPECjbb2000	Heap growth is mostly live (4.7X longer)
Mckoi Database	Thread leak: extra support needed (1.6X longer)
DualLeak	Heap growth is live (No help)

Tolerating Leaks

Eclipse "Diff"	Tolerates until 24-hr limit (>200X longer)

- Reported on Eclipse Bugzilla
- Leak: recursive difference
- Automated with Eclipse Plugin

Eclipse Diff: Reachable Memory

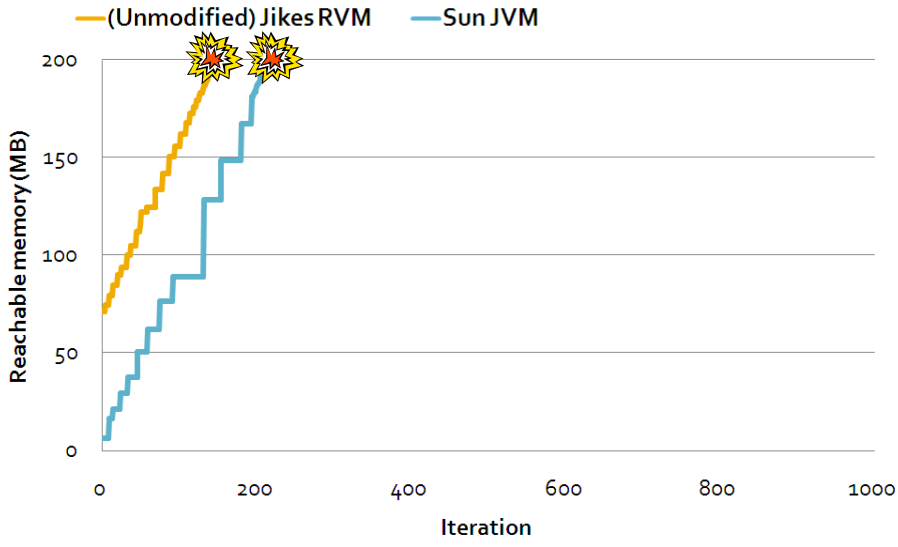

Eclipse Diff: Reachable Memory

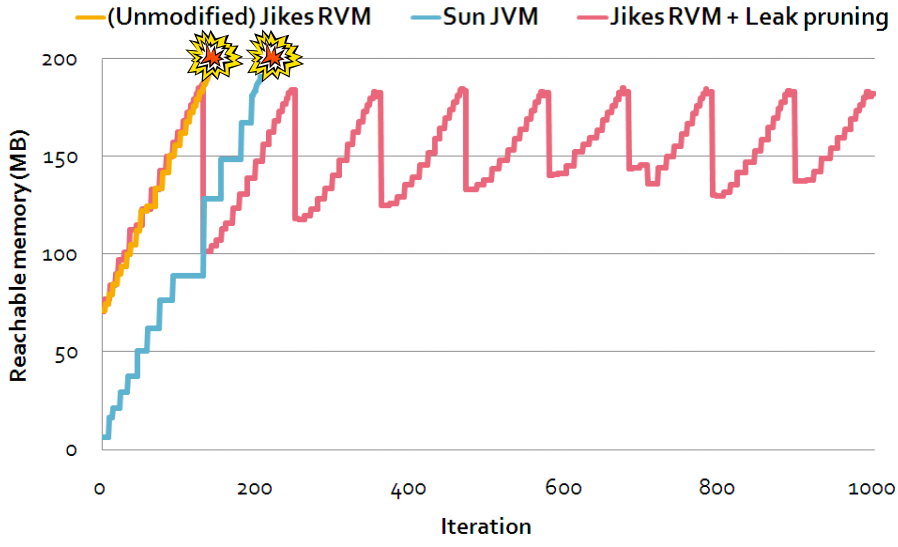

Eclipse Diff: Reachable Memory

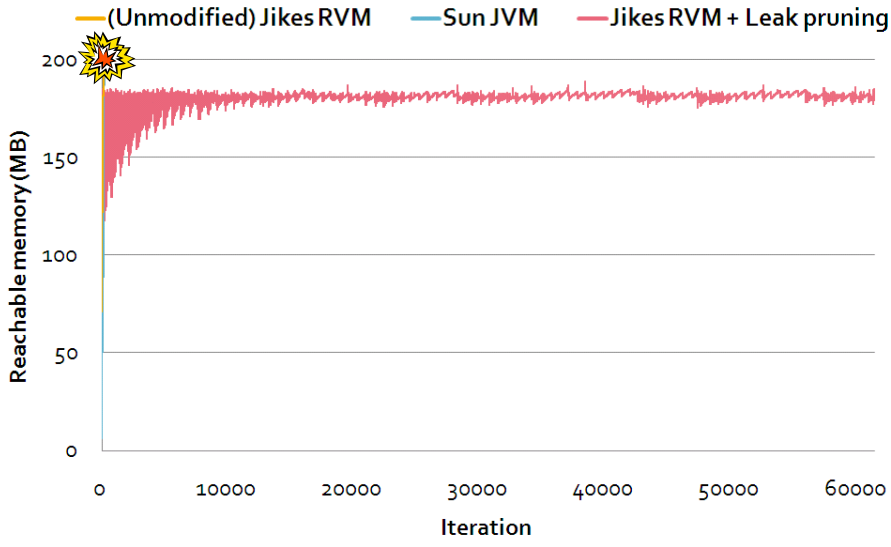

Eclipse Diff: Performance

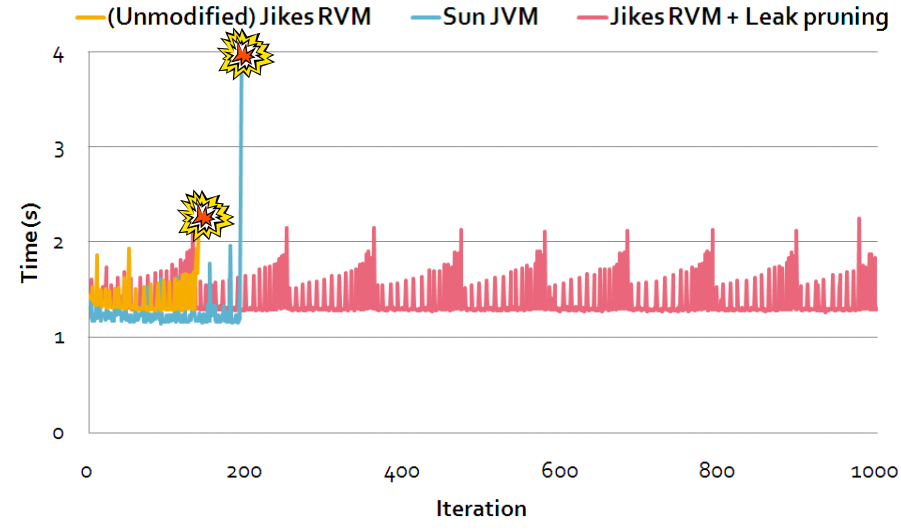

Eclipse Diff: Performance

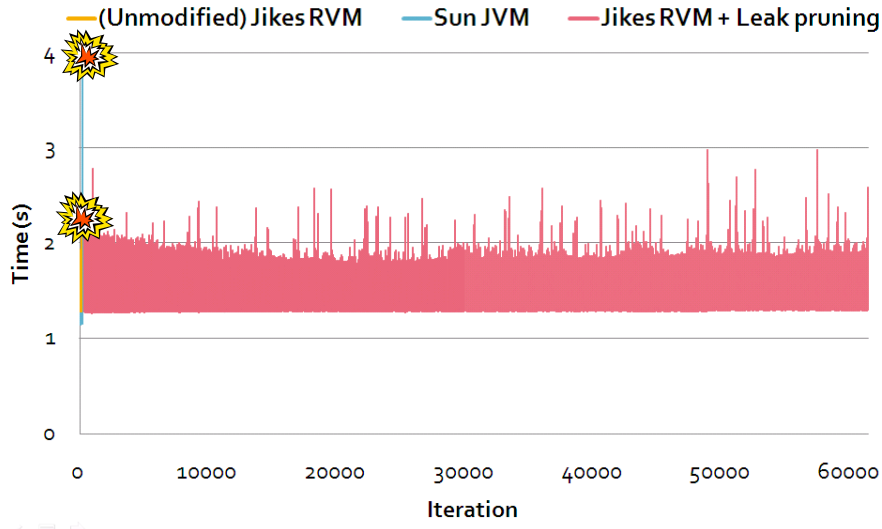

Performance Summary

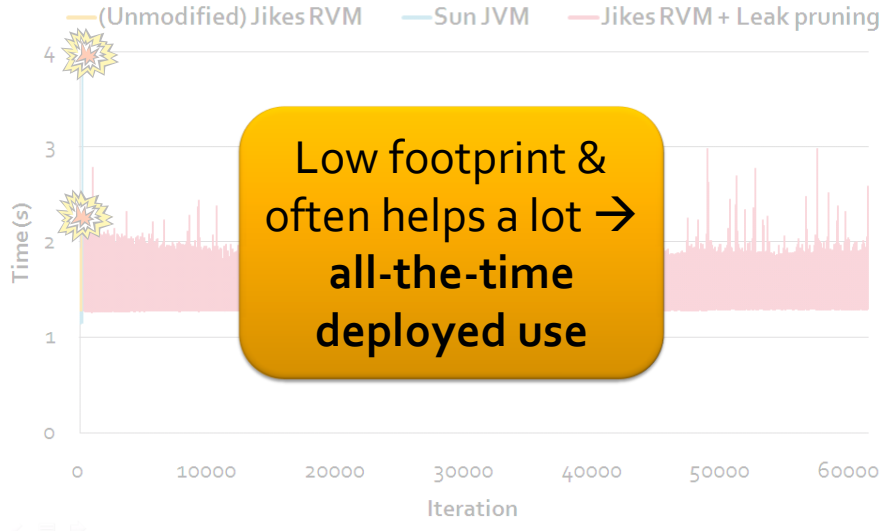

Related Work

- Memory corruption & concurrency bugs
  - Perturb layout & scheduling [**Rx,** Qin et al. '05] [**DieHard,** Berger & Zorn '06] [**Atom-Aid,** Lucia et al. '08] [**Grace,** Berger et al. '08]
- Fatal errors
  - Ignore errors [**Failure-oblivious computing,** Rinard et al. '04] [**Append,** Dobolyi & Weimer '08]

Related Work

- Memory corruption & concurrency bugs
  - Perturb layout & scheduling [**Rx,** Qin et al. '05] [**DieHard,** Berger & Zorn '06] [**Atom-Aid,** Lucia et al. '08] [**Grace,** Berger et al. '08]
- Fatal errors
  - Ignore errors [**Failure-oblivious computing,** Rinard et al. '04] [**Append,** Dobolyi & Weimer '08]
- Memory leaks in unmanaged languages
  - Leak-friendly layout [**Plug,** Novark et al. '08]
  - Bound allocation sites [**Cyclic alloc,** Nguyen & Rinard '07]

Tolerating Memory Leaks in Managed Languages

- Finding leaks before deployment is hard
  - Deployed systems need immediate help
- Leak pruning: GC based on liveness
  - High precision & low overhead
  - Bounded resources
  - Preserves semantics

Tolerating Memory Leaks in Managed Languages

- Finding leaks before deployment is hard
  - Deployed systems need immediate help
- Leak pruning: GC based on liveness
  - High precision & low overhead
  - Bounded resources
  - Preserves semantics

Thank you!

Backup

Eclipse "Diff" Leak

- Reported on Eclipse Bugzilla
- Recursive difference leaks memory

Eclipse "Diff" Leak

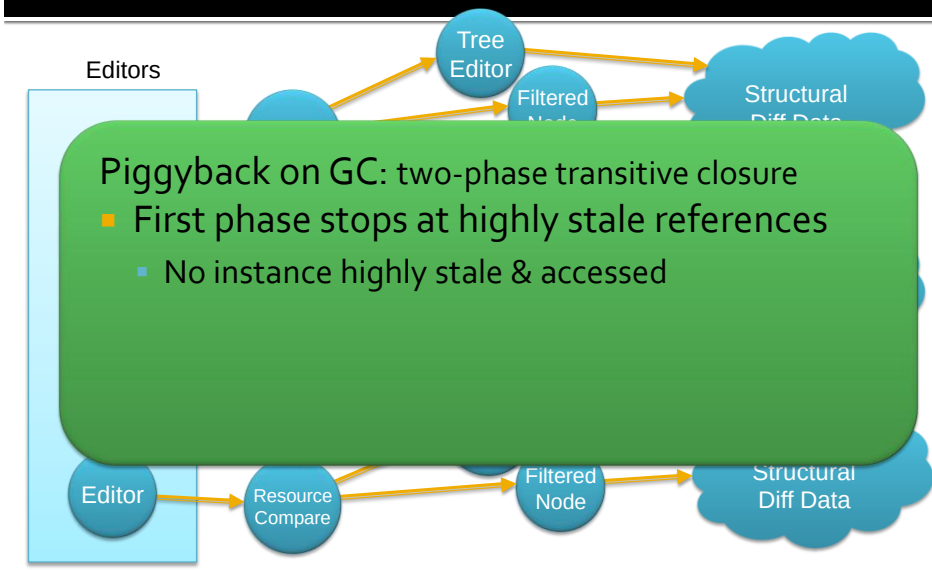

Fine-Grained Usage Tracking

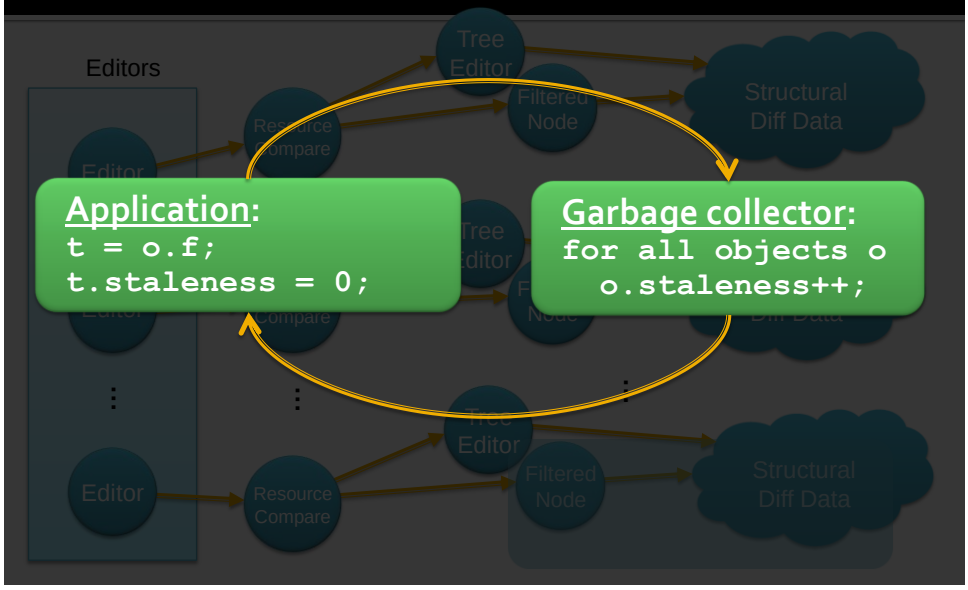

Fine-Grained Usage Tracking

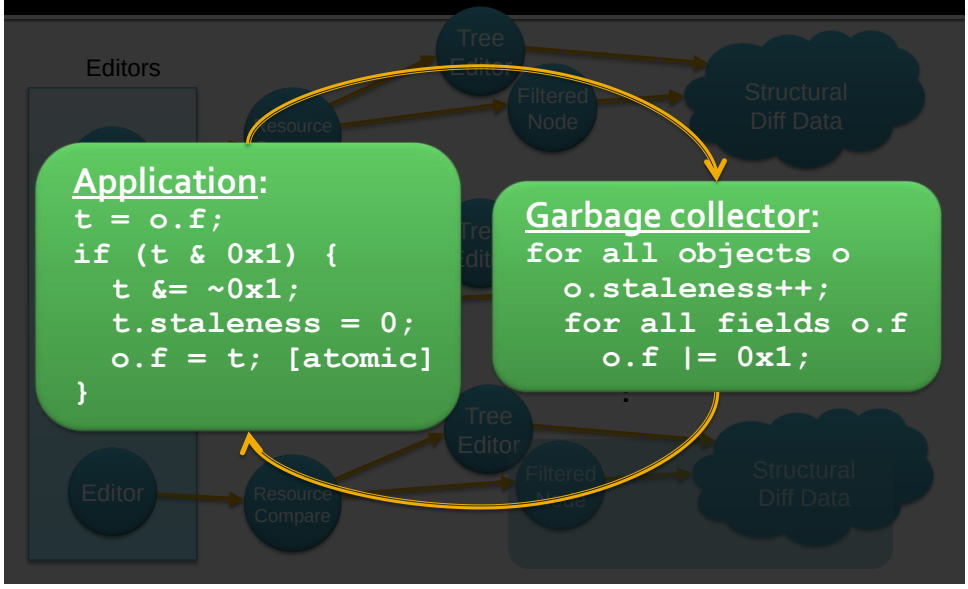

Fine-Grained Usage Tracking

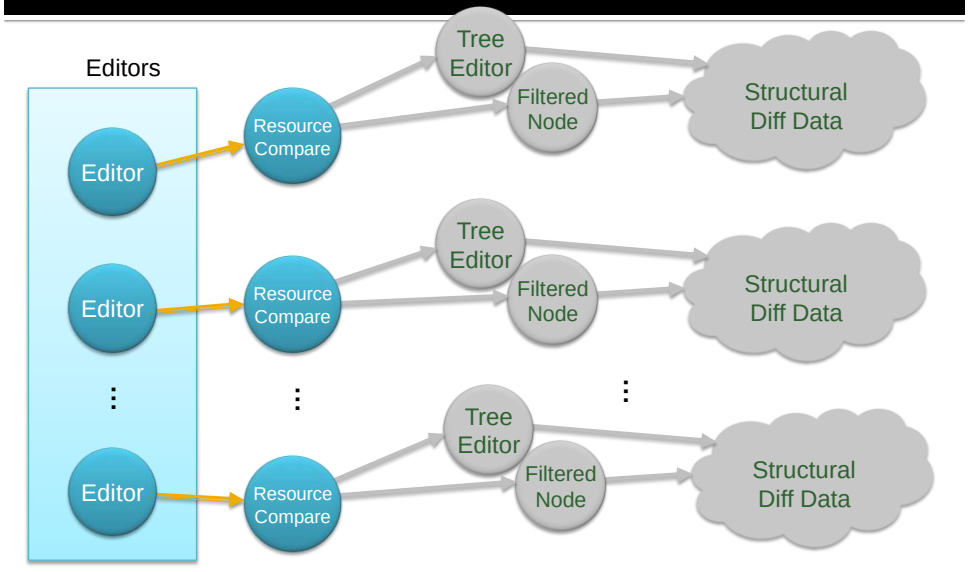

Summarize Stale Data Structures

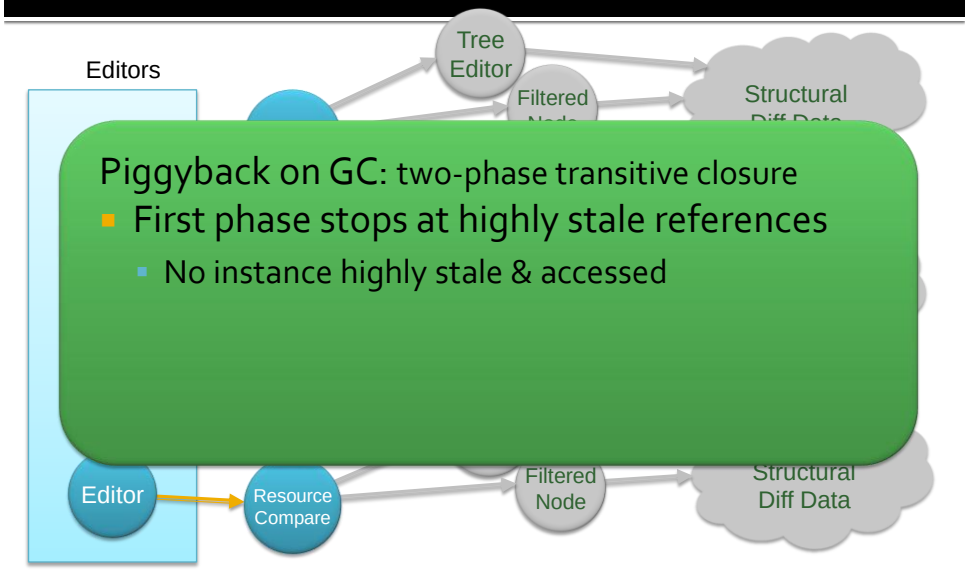

Phase One of Transitive Closure

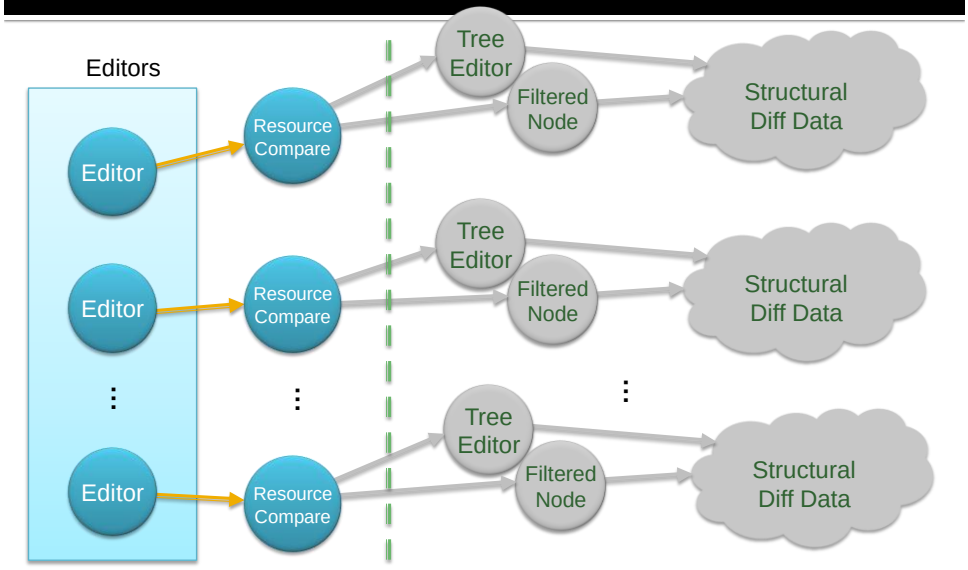

Phase One of Transitive Closure

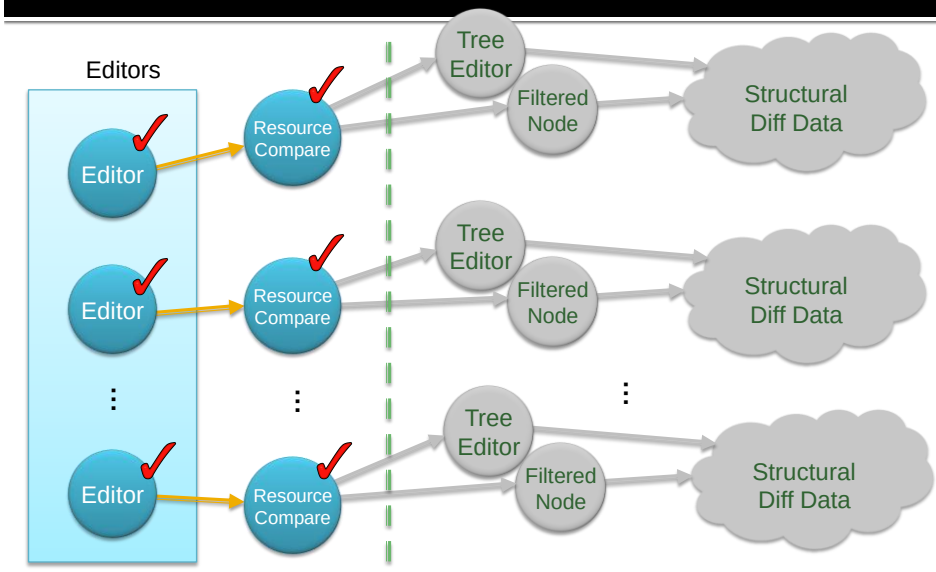

Phase Two of Transitive Closure

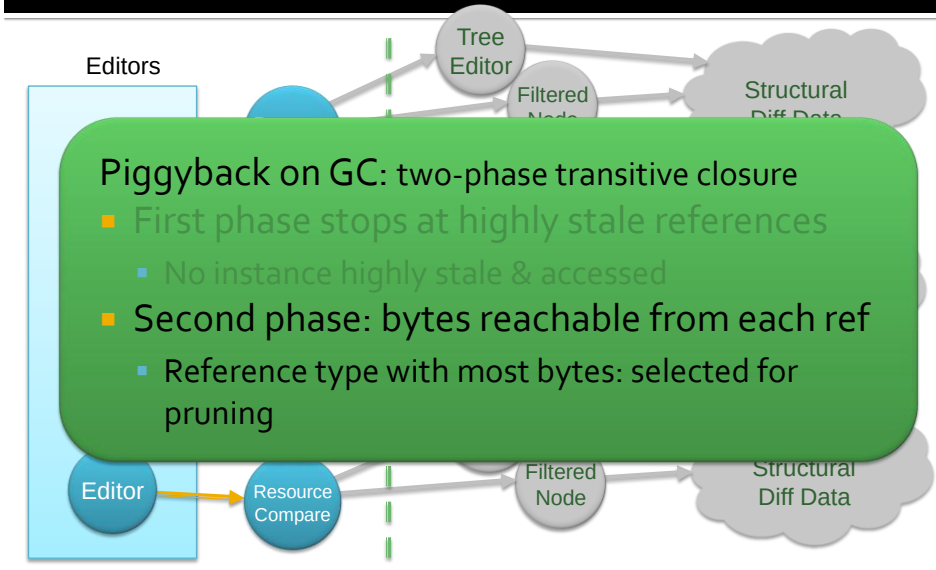

Piggyback on GC: two-phase transitive closure

- First phase stops at highly stale references
  - No instance highly stale & accessed
- Second phase: bytes reachable from each ref
  - Reference type with most bytes: selected for pruning

Phase Two of Transitive Closure

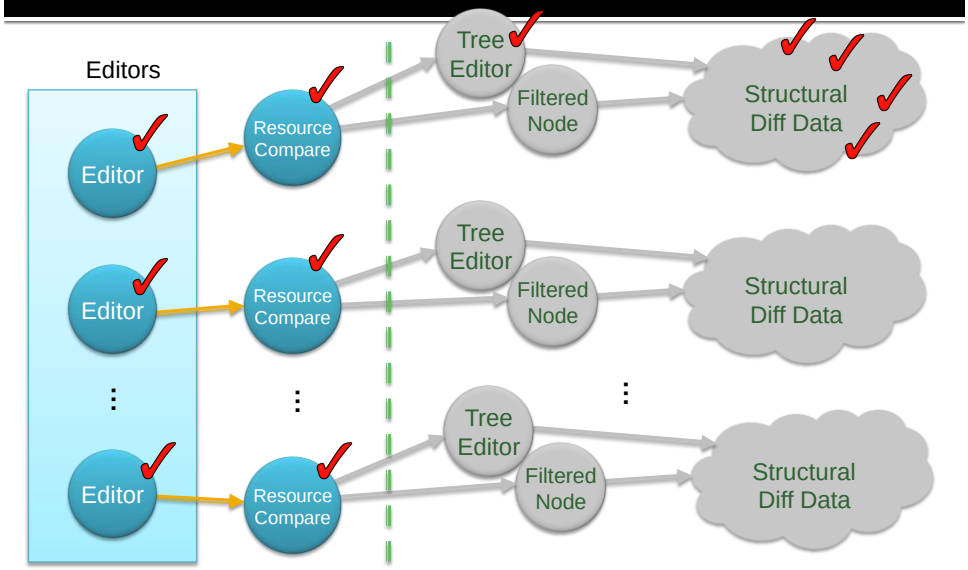

Phase Two of Transitive Closure

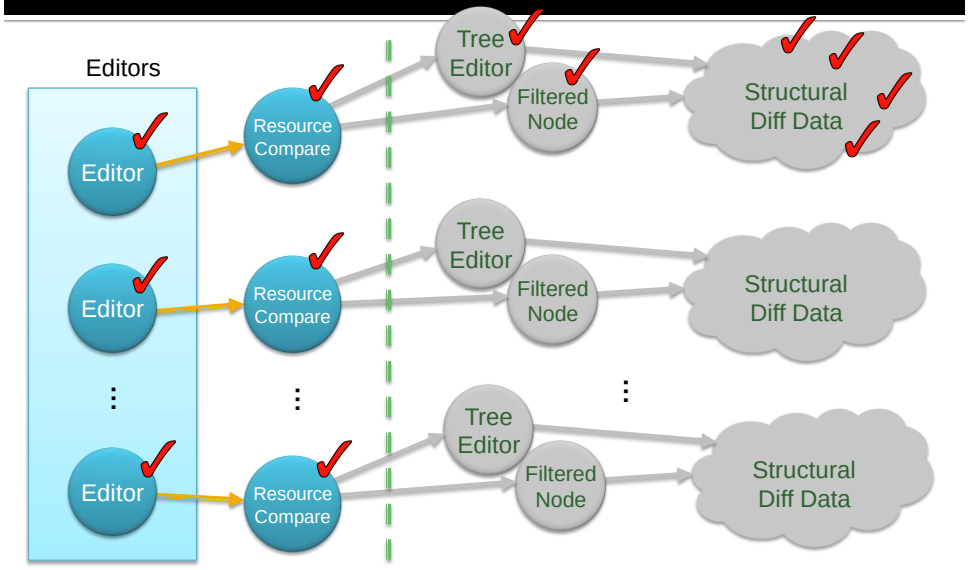

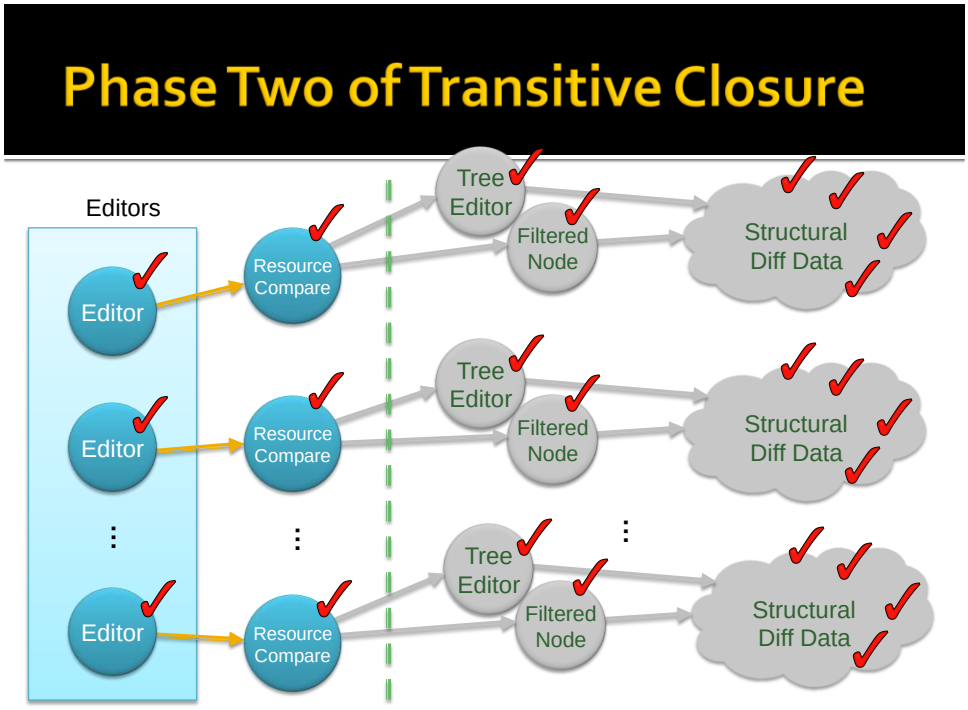
Phase Two of Transitive Closure
Editors
Editor
Editor
Editor
Resource Compare
Resource Compare
Resource Compare
Tree Editor
Filtered Node
Tree Editor
Filtered Node
Tree Editor
Filtered Node
Structural Diff Data
Structural Diff Data
Structural Diff Data

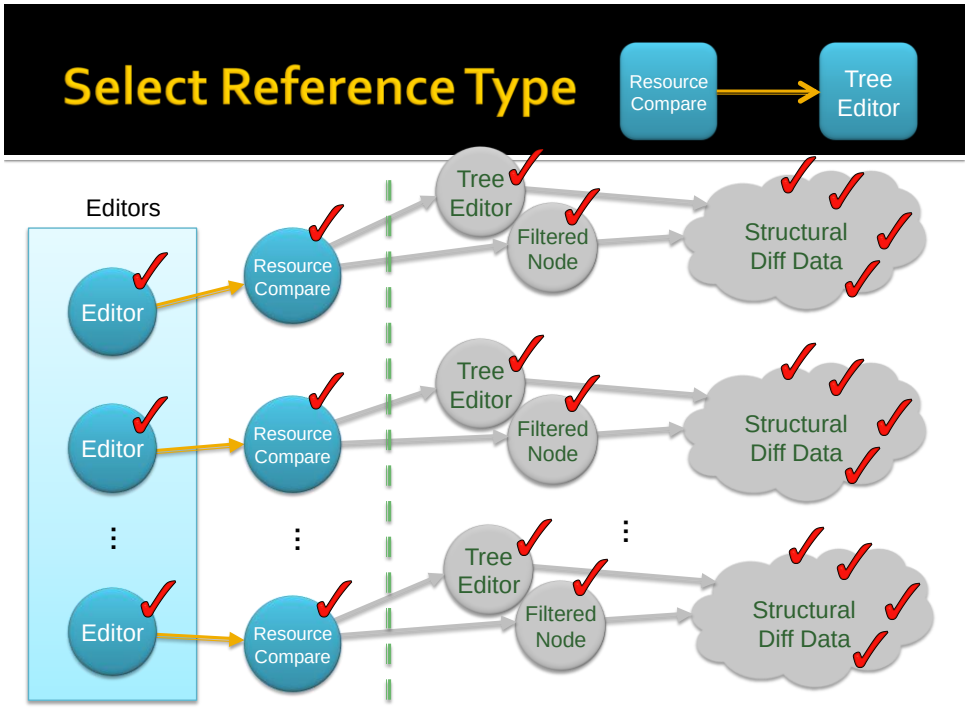
Select Reference Type
Resource Compare
Tree Editor
Editors
Editor
Editor
Editor
Resource Compare
Resource Compare
Resource Compare
Tree Editor
Filtered Node
Tree Editor
Filtered Node
Tree Editor
Filtered Node
Structural Diff Data
Structural Diff Data
Structural Diff Data

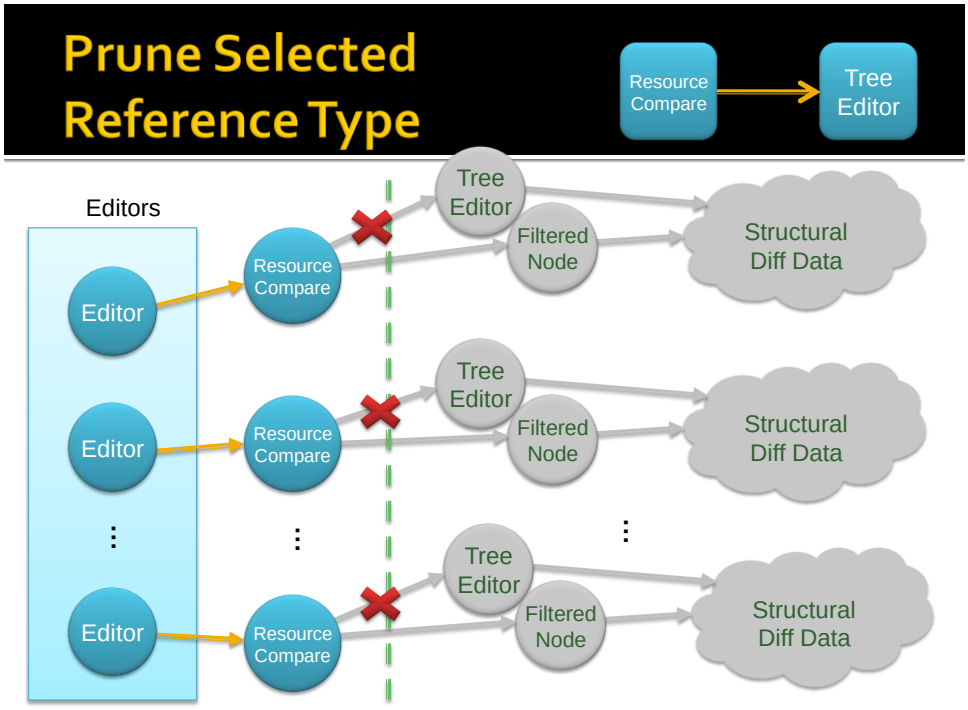
Prune Selected Reference Type
Resource Compare
Tree Editor
Editors
Editor
Editor
Editor
Resource Compare
Resource Compare
Resource Compare
Tree Editor
Tree Editor
Tree Editor
Filtered Node
Filtered Node
Filtered Node
Structural Diff Data
Structural Diff Data
Structural Diff Data

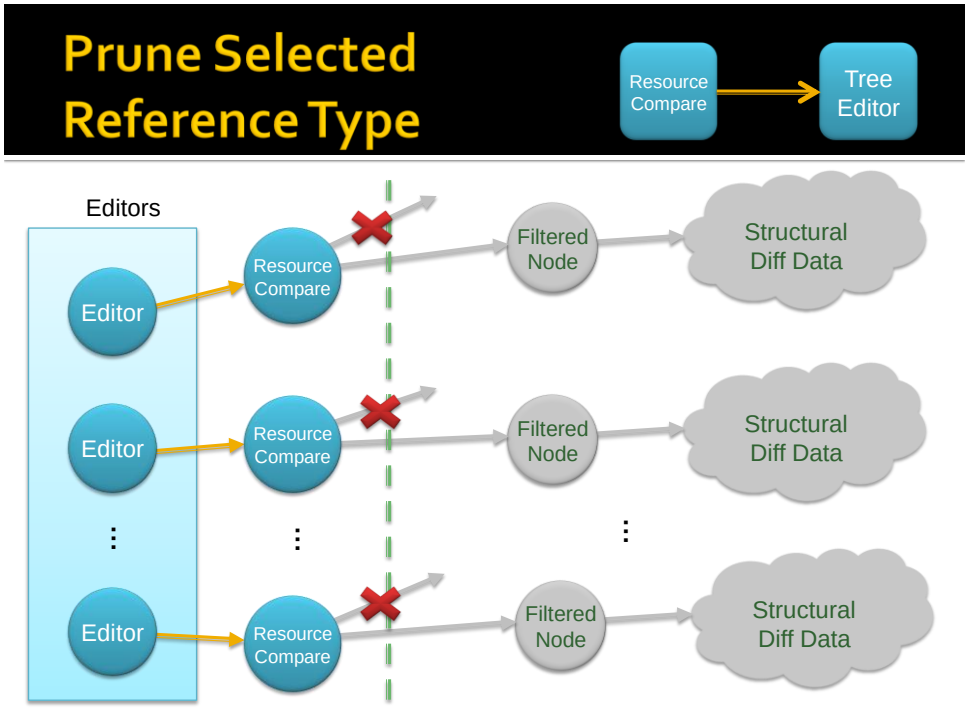
Prune Selected Reference Type
Resource Compare
Tree Editor
Editors
Editor
Editor
Editor
Resource Compare
Resource Compare
Resource Compare
Filtered Node
Filtered Node
Filtered Node
Structural Diff Data
Structural Diff Data
Structural Diff Data

Phase One of Transitive Closure

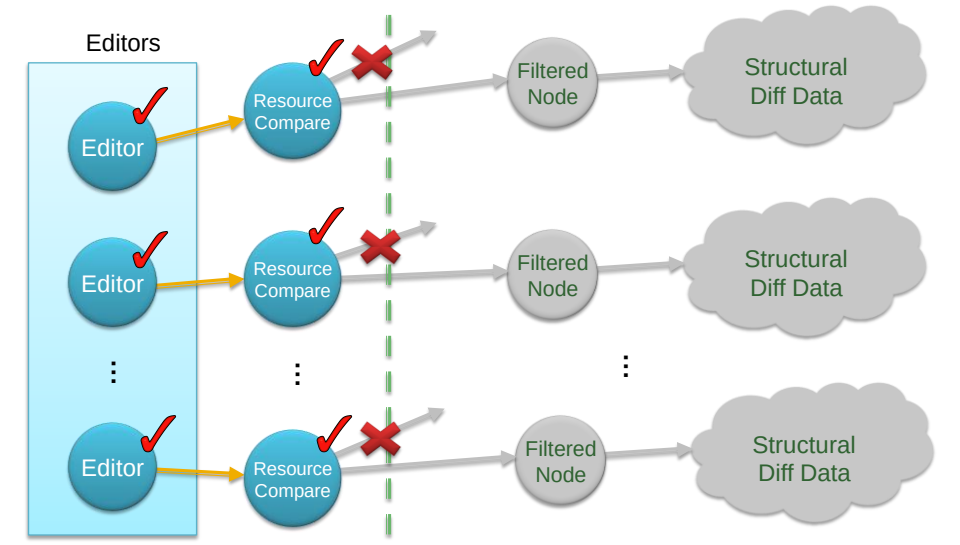

Select Another Reference Type

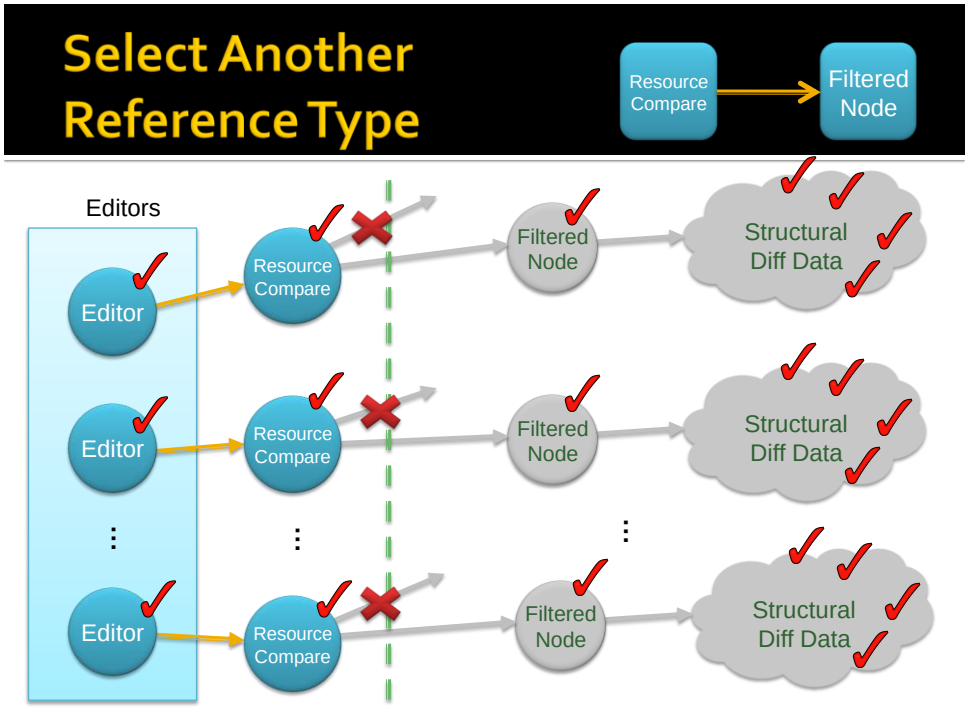

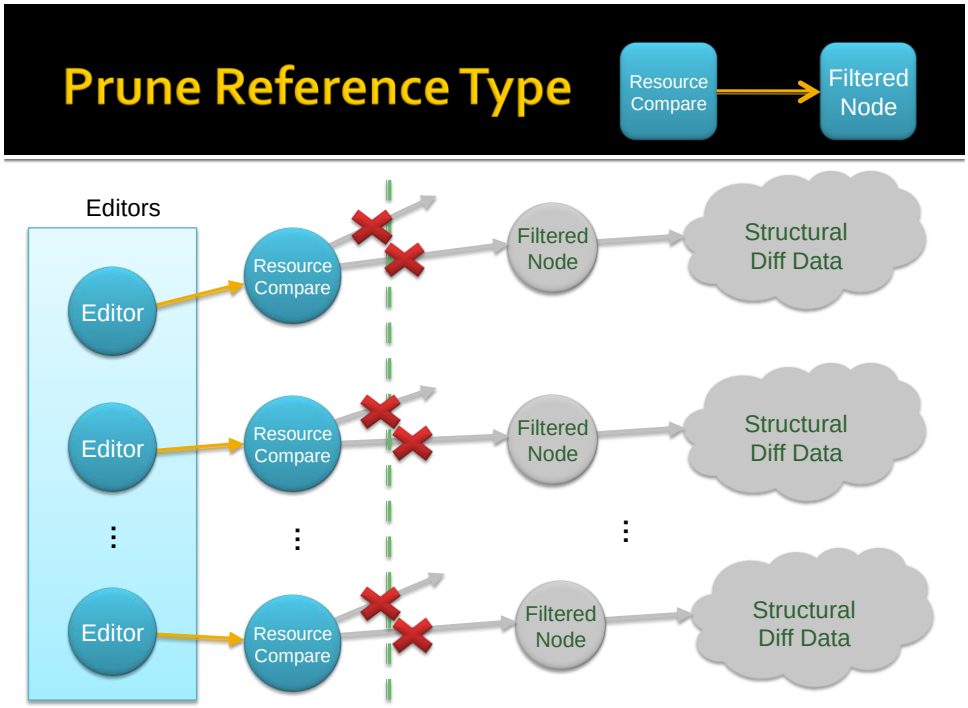
Prune Reference Type
Resource Compare
Filtered Node
Editors
Editor
Editor
Editor
Resource Compare
Resource Compare
Resource Compare
Filtered Node
Filtered Node
Filtered Node
Structural Diff Data
Structural Diff Data
Structural Diff Data

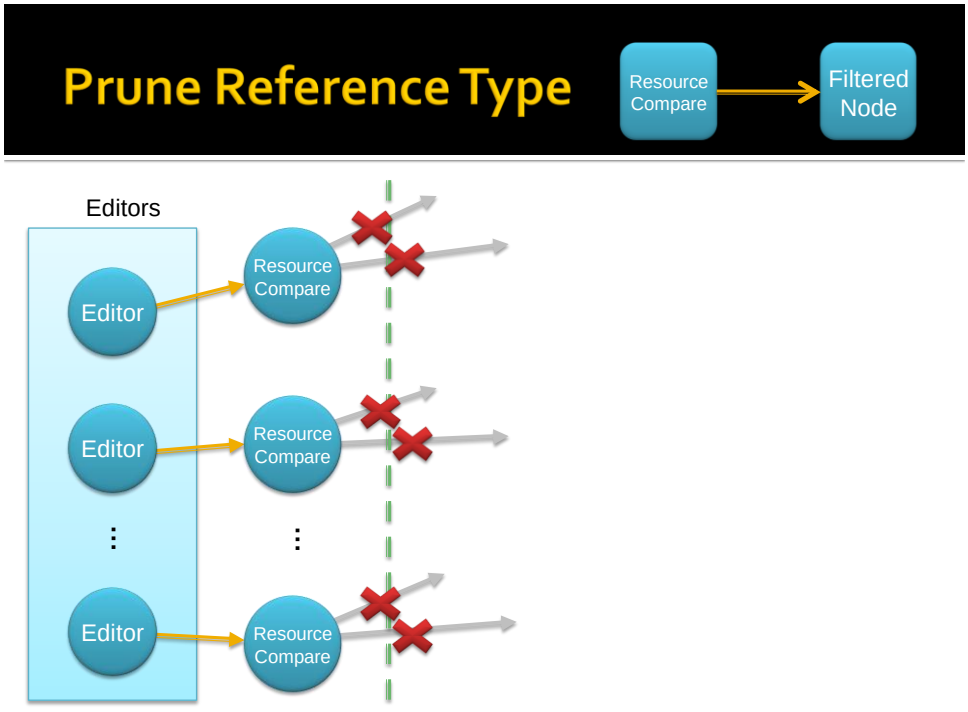
Prune Reference Type
Resource Compare
Filtered Node
Editors
Editor
Editor
Editor
Resource Compare
Resource Compare
Resource Compare