

Legato: Bounded Region Serializability Using Commodity Hardware Transactional Memory

Aritra Sengupta

Man Cao

Michael D. Bond

and

Milind Kulkarni



Programming Language Semantics?



Data Races

C++ no guarantee of semantics – “catch-fire” semantics

Java provides weak semantics

Weak Semantics

```
Data data = null;  
boolean done = false;
```

T1

```
data = new Data();  
done = true;
```

T2

```
if (done)  
    data.foo();
```

Weak Semantics

```
Data data = null;  
boolean done = false;
```

T1

```
data = new Data();  
done = true;
```

T2

```
if (done)  
    data.foo();
```



Null Pointer
Exception

Weak Semantics

```
Data data = null;  
boolean done = false;
```

T1

```
data = new Data();  
done = true;
```

No data
dependence:
Reordering effect

T2

```
if (done)  
data.foo();
```

Null Pointer
Exception

Need for Stronger Memory Models

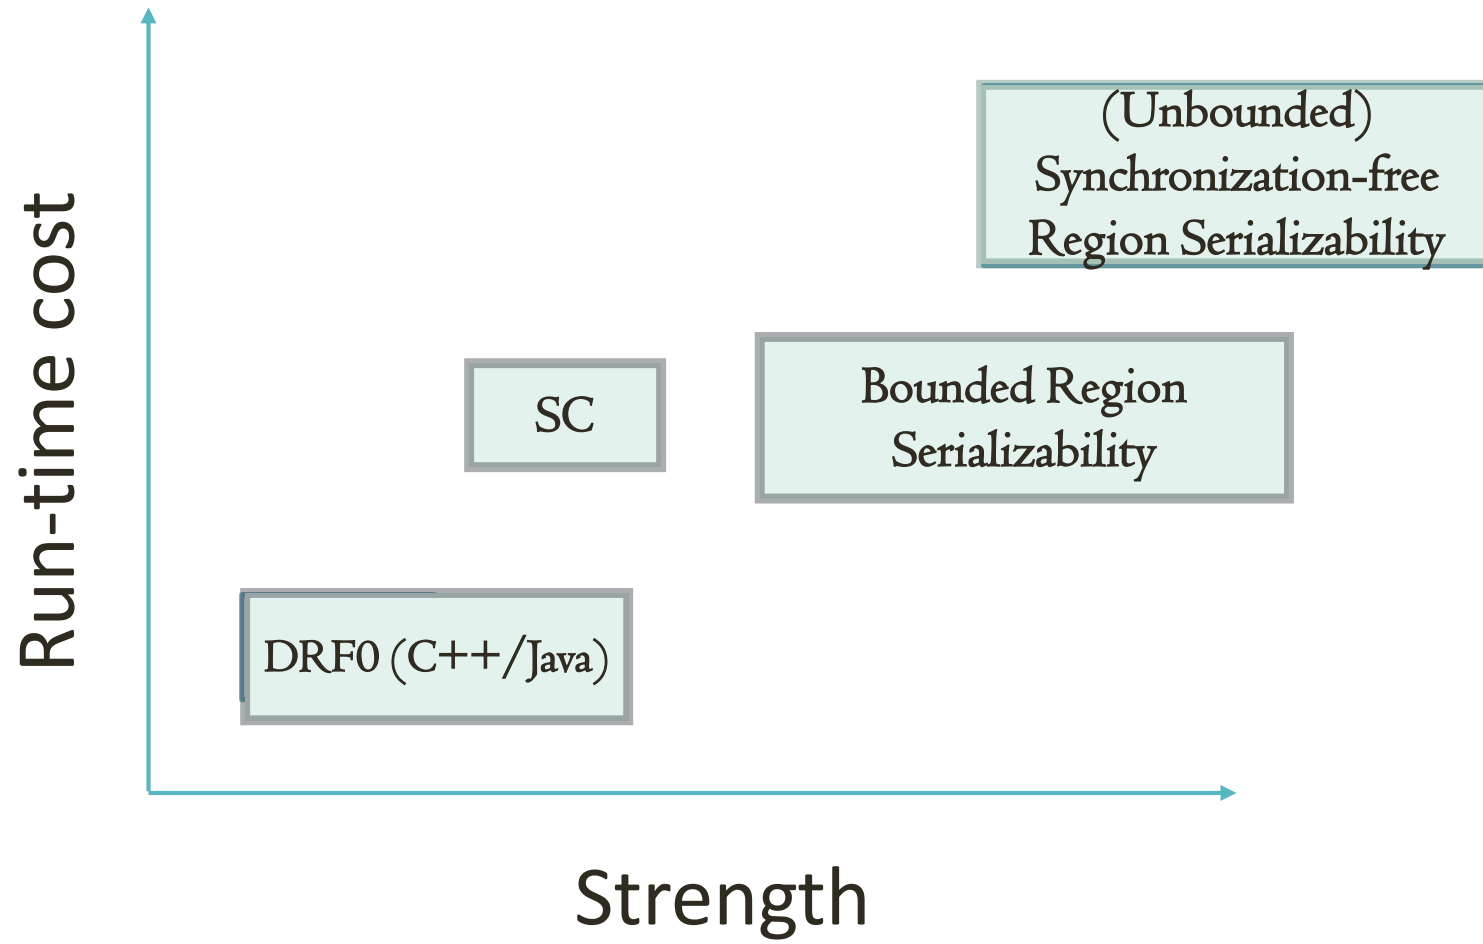
“The inability to define reasonable semantics for programs with data races is not just a theoretical shortcoming, but a fundamental hole in the foundation of our languages and systems...”

Give better semantics to programs with data races

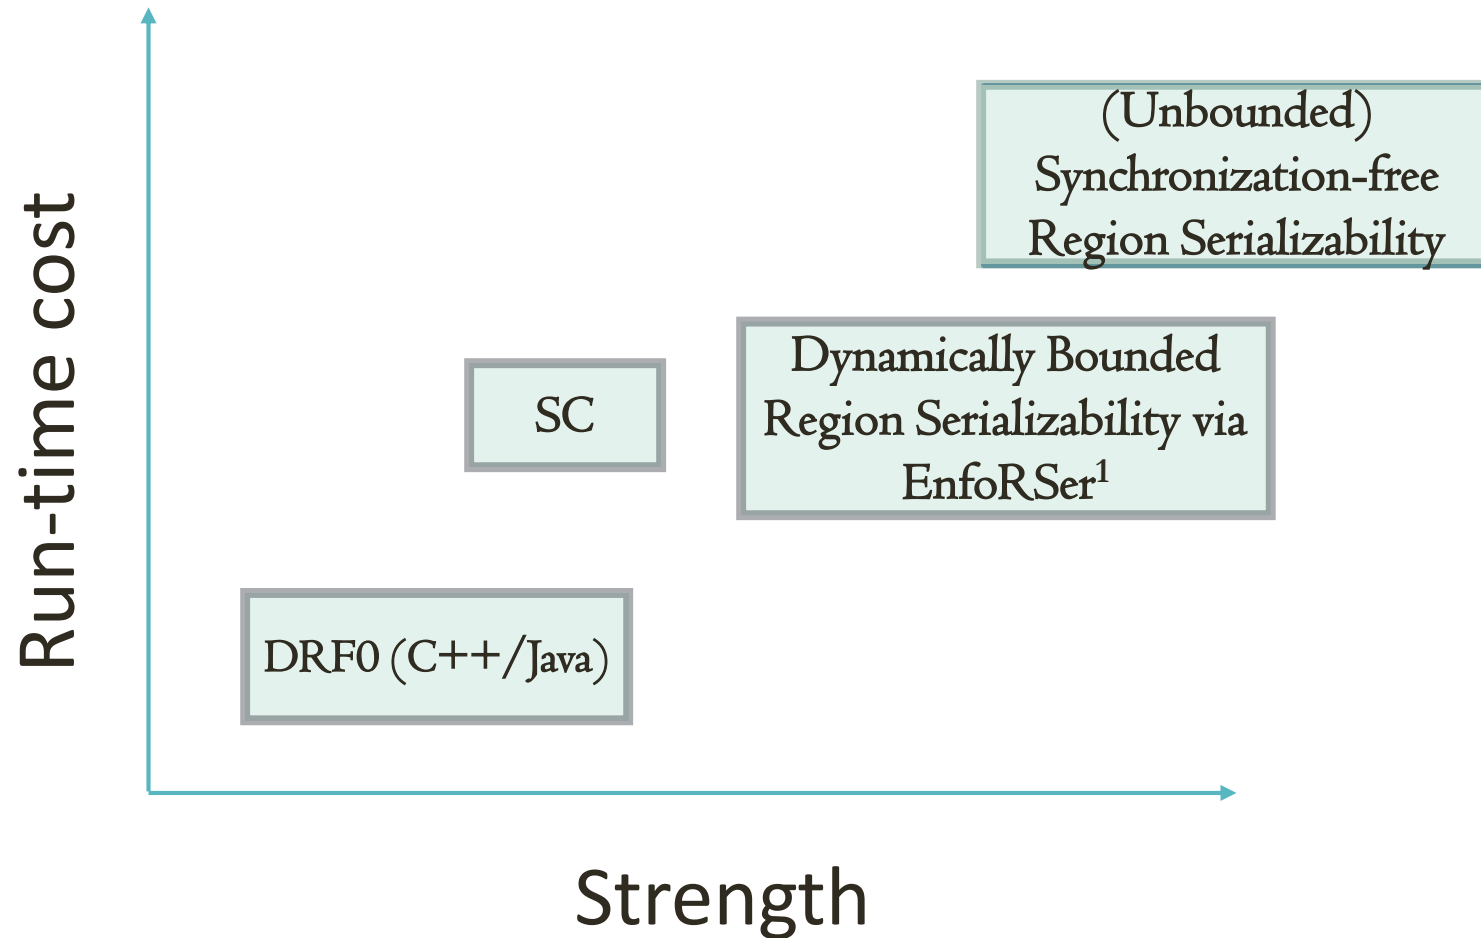
Stronger memory models

– Adve and Boehm, CACM, 2010

End-to-End Memory Models: Run-time cost vs Strength




End-to-End Memory Models: Run-time cost vs Strength



1. Sengupta et al. Hybrid Static-Dynamic Analysis for Statically Bounded Region Serializability. ASPLOS, 2015.

Dynamically Bounded Region Serializability (DBRS)



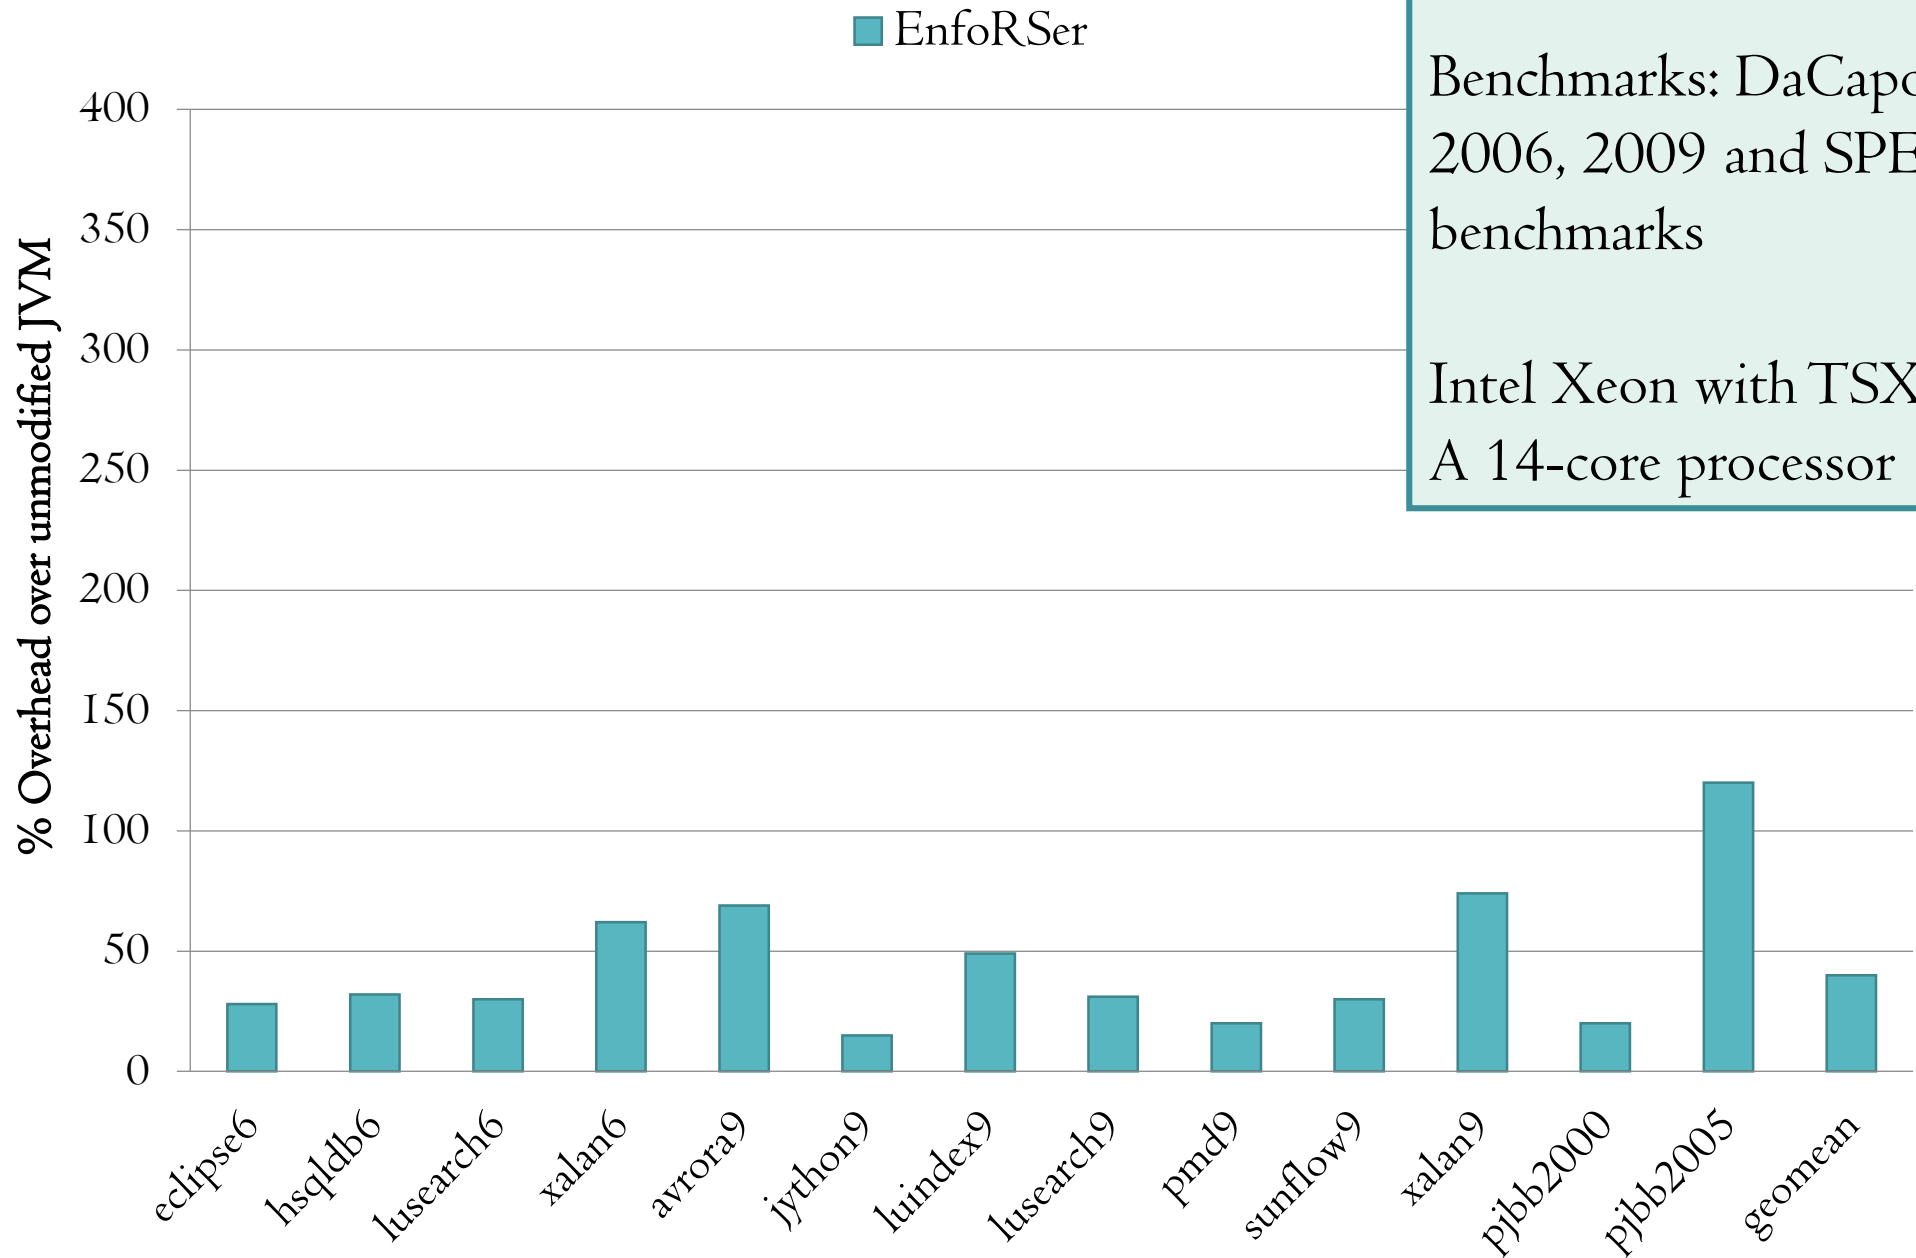
Dynamically Bounded Region Serializability

SC+Atomicity of
bounded regions¹

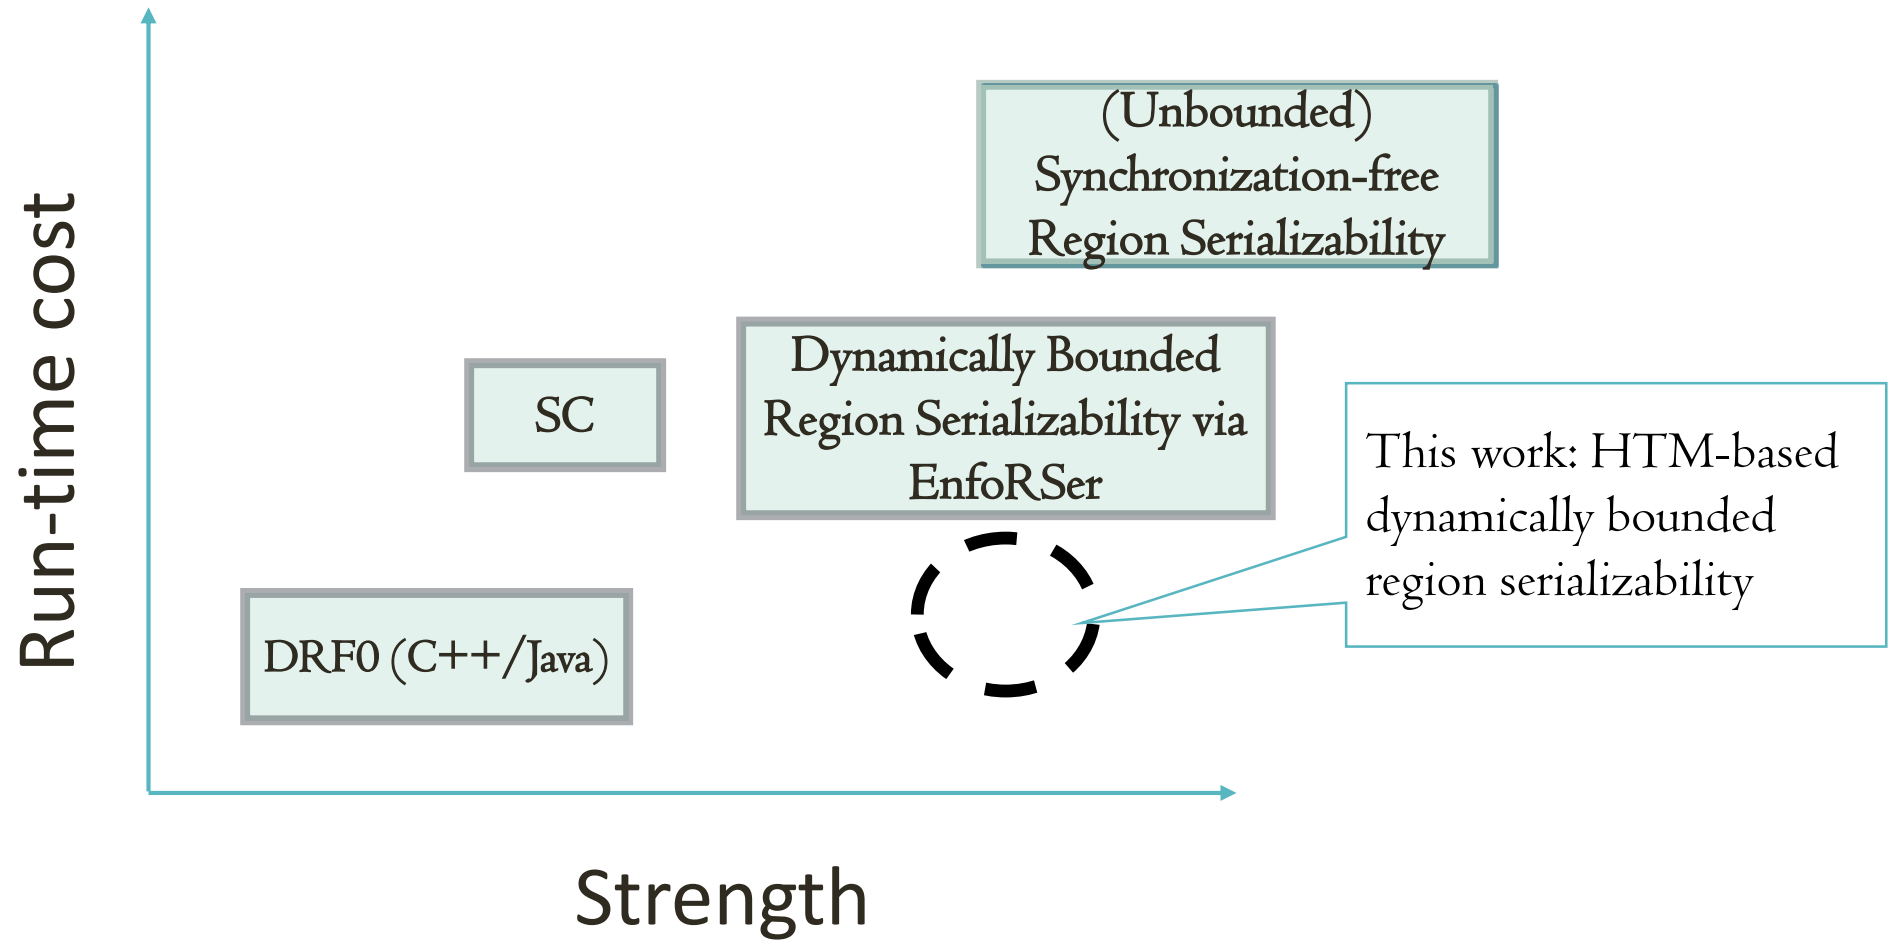
- Eliminates SC violation
- Eliminates some atomicity violations

	DRF0	SC	DBRS
hsqldb6	Infinite loop	Correct	Correct
sunflow9	Null pointer exception	Correct	Correct
jbb2000	Corrupt output	Corrupt output	Correct
jbb2000	Infinite loop	Correct	Correct
sor	Infinite loop	Correct	Correct
lufact	Infinite loop	Correct	Correct
moldyn	Infinite loop	Correct	Correct
raytracer	Fails validation	Fails validation	Correct

1. Sengupta et al. Hybrid Static-Dynamic Analysis for Statically Bounded Region Serializability. ASPLOS, 2015.




End-to-End Memory Models: Run-time cost vs Strength




Outline

- Challenges
 - Naïve implementation with hardware transactional memory
 - Limitations of using HTM for DBRS
- Approach
 - Overcoming limitations
 - Our approach to DBRS enforcement: **Legato**
- Evaluation

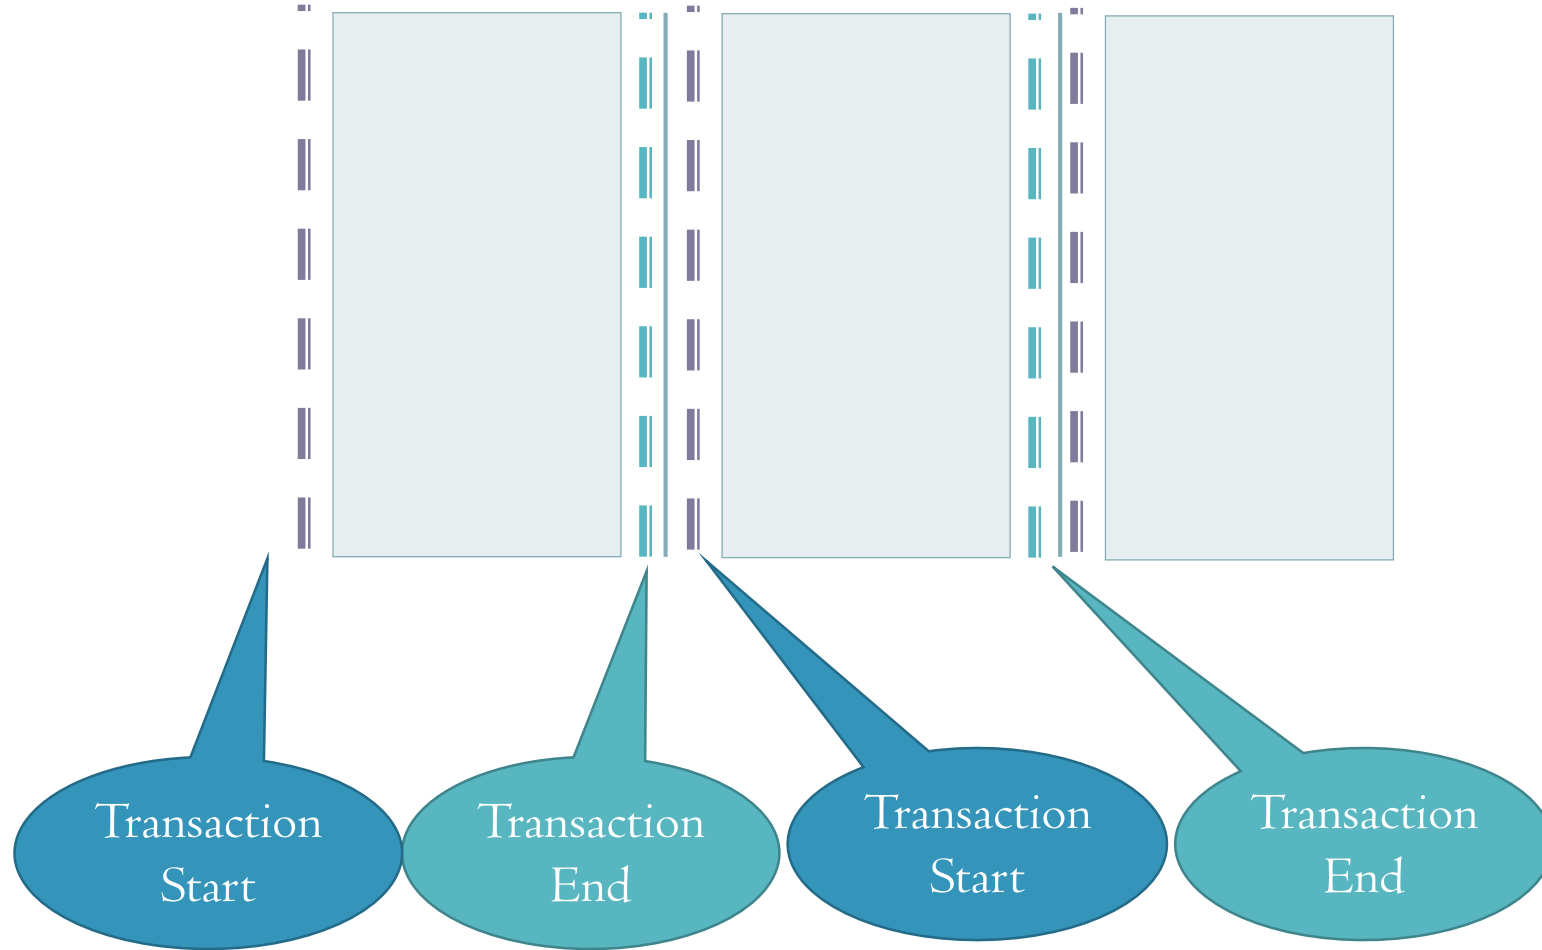
Enforcing DBRS with Commodity Hardware Transactional Memory (HTM)

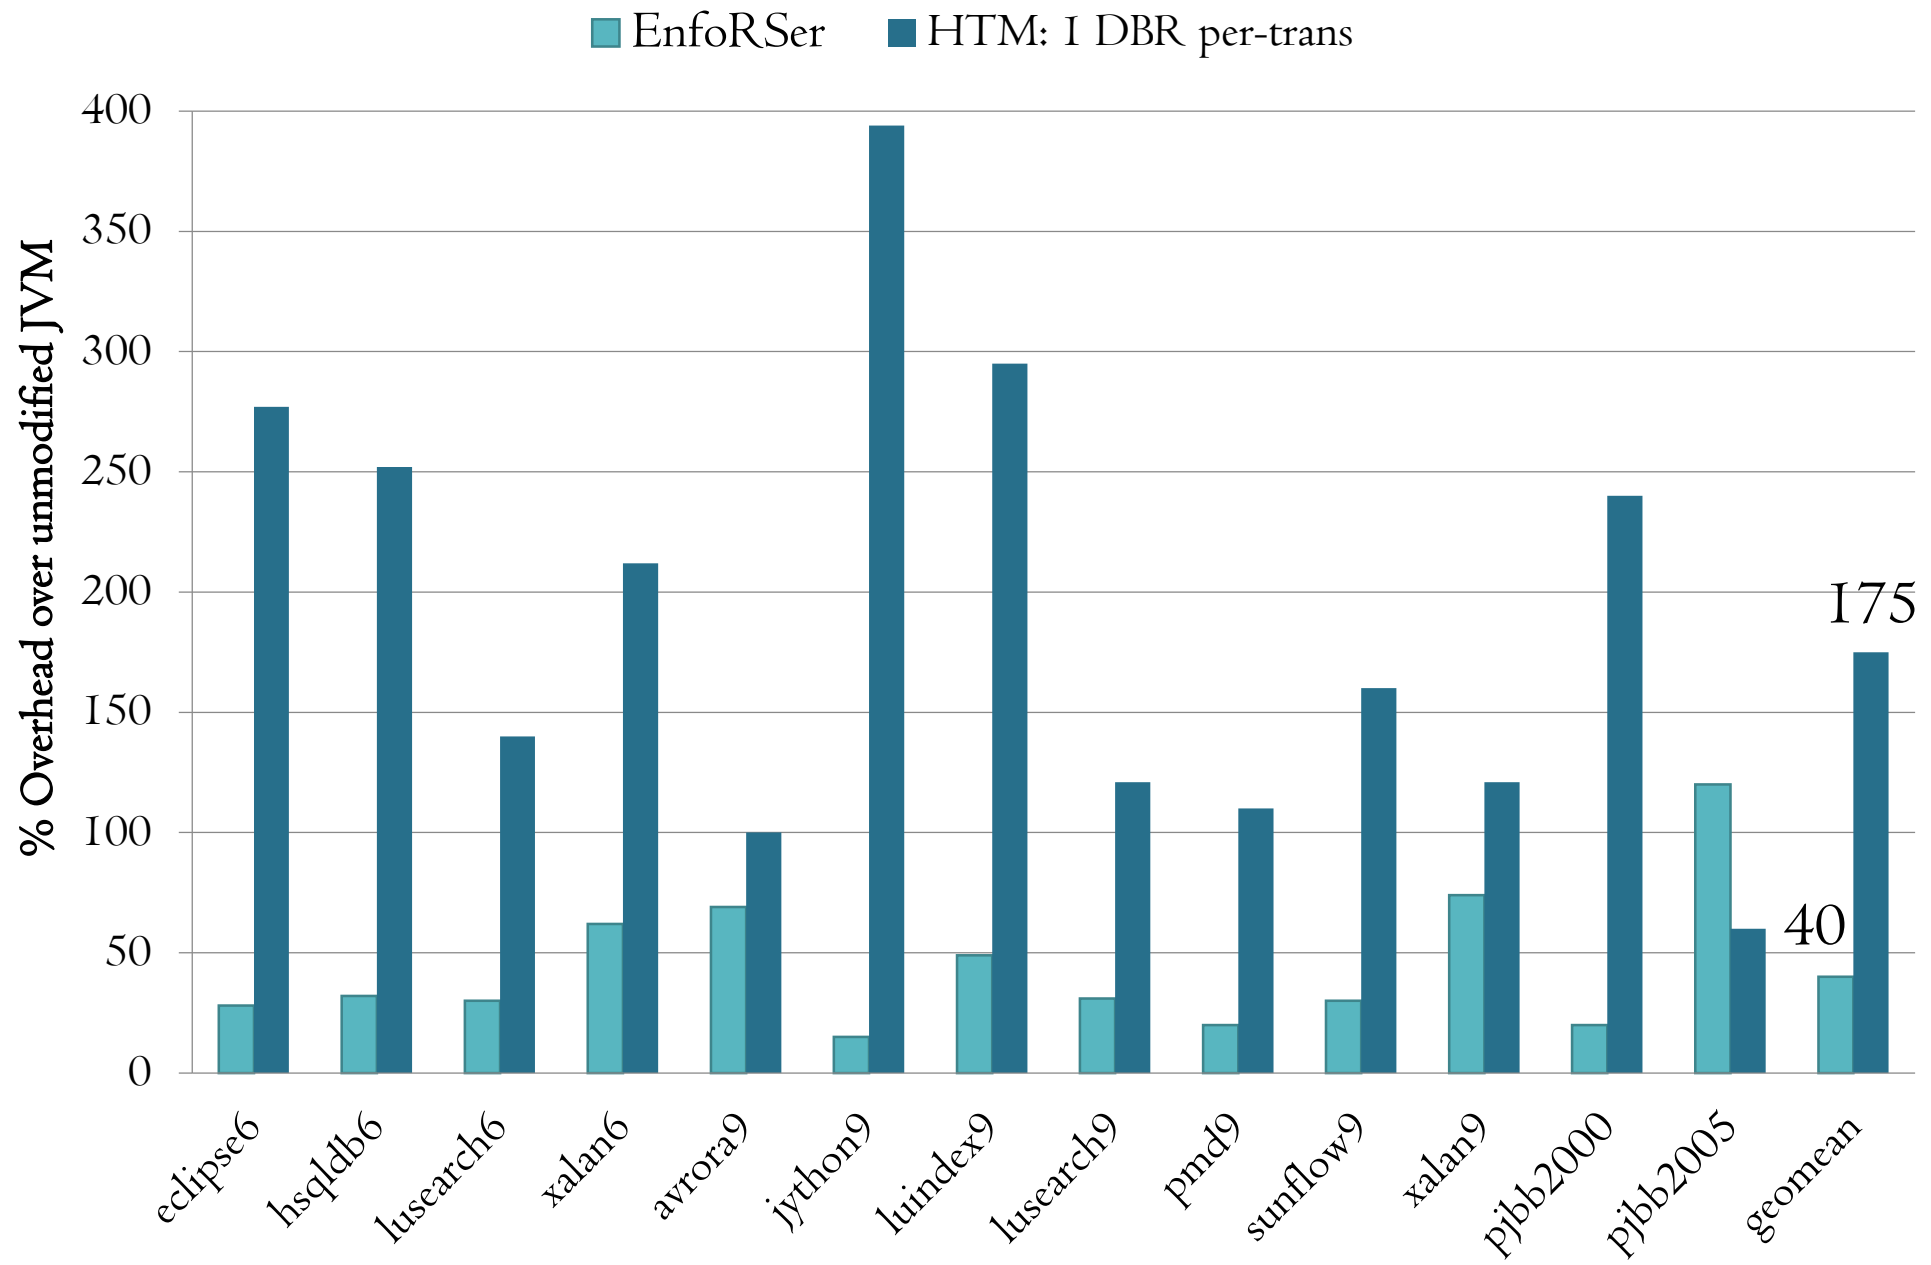



Enforcing DBRS with Commodity Hardware Transactional Memory (HTM)



Enforcing DBRS with Commodity Hardware Transactional Memory (HTM)









Legato: Key Idea

Merge several regions into a single transaction:
amortize the cost of starting and stopping a transaction


Legato: Key Idea




Legato: Key Idea



Legato: Key Idea



Legato: Key Idea




Challenges

Conflicts abort transactions: wasted work


Capacity aborts: larger transactions have **larger footprint**,
unknown *a priori*

HTM-unfriendly operations: hardware interrupts, page faults etc.

Per-transaction cost vs Abort cost



Per-transaction cost vs Abort cost



Transient cause: conflicts,
hardware interrupts

=>


Retry transaction

Capacity abort: end
before culprit region


=>

Start new transaction

Per-transaction cost vs Abort cost



Per-transaction cost vs Abort cost



Transient cause: conflicts,
hardware interrupts

=>


Retry transaction

Capacity abort: end
before culprit region

=>

Start new transaction

Per-transaction cost vs Abort cost



Transient cause: conflicts,
hardware interrupts

=>

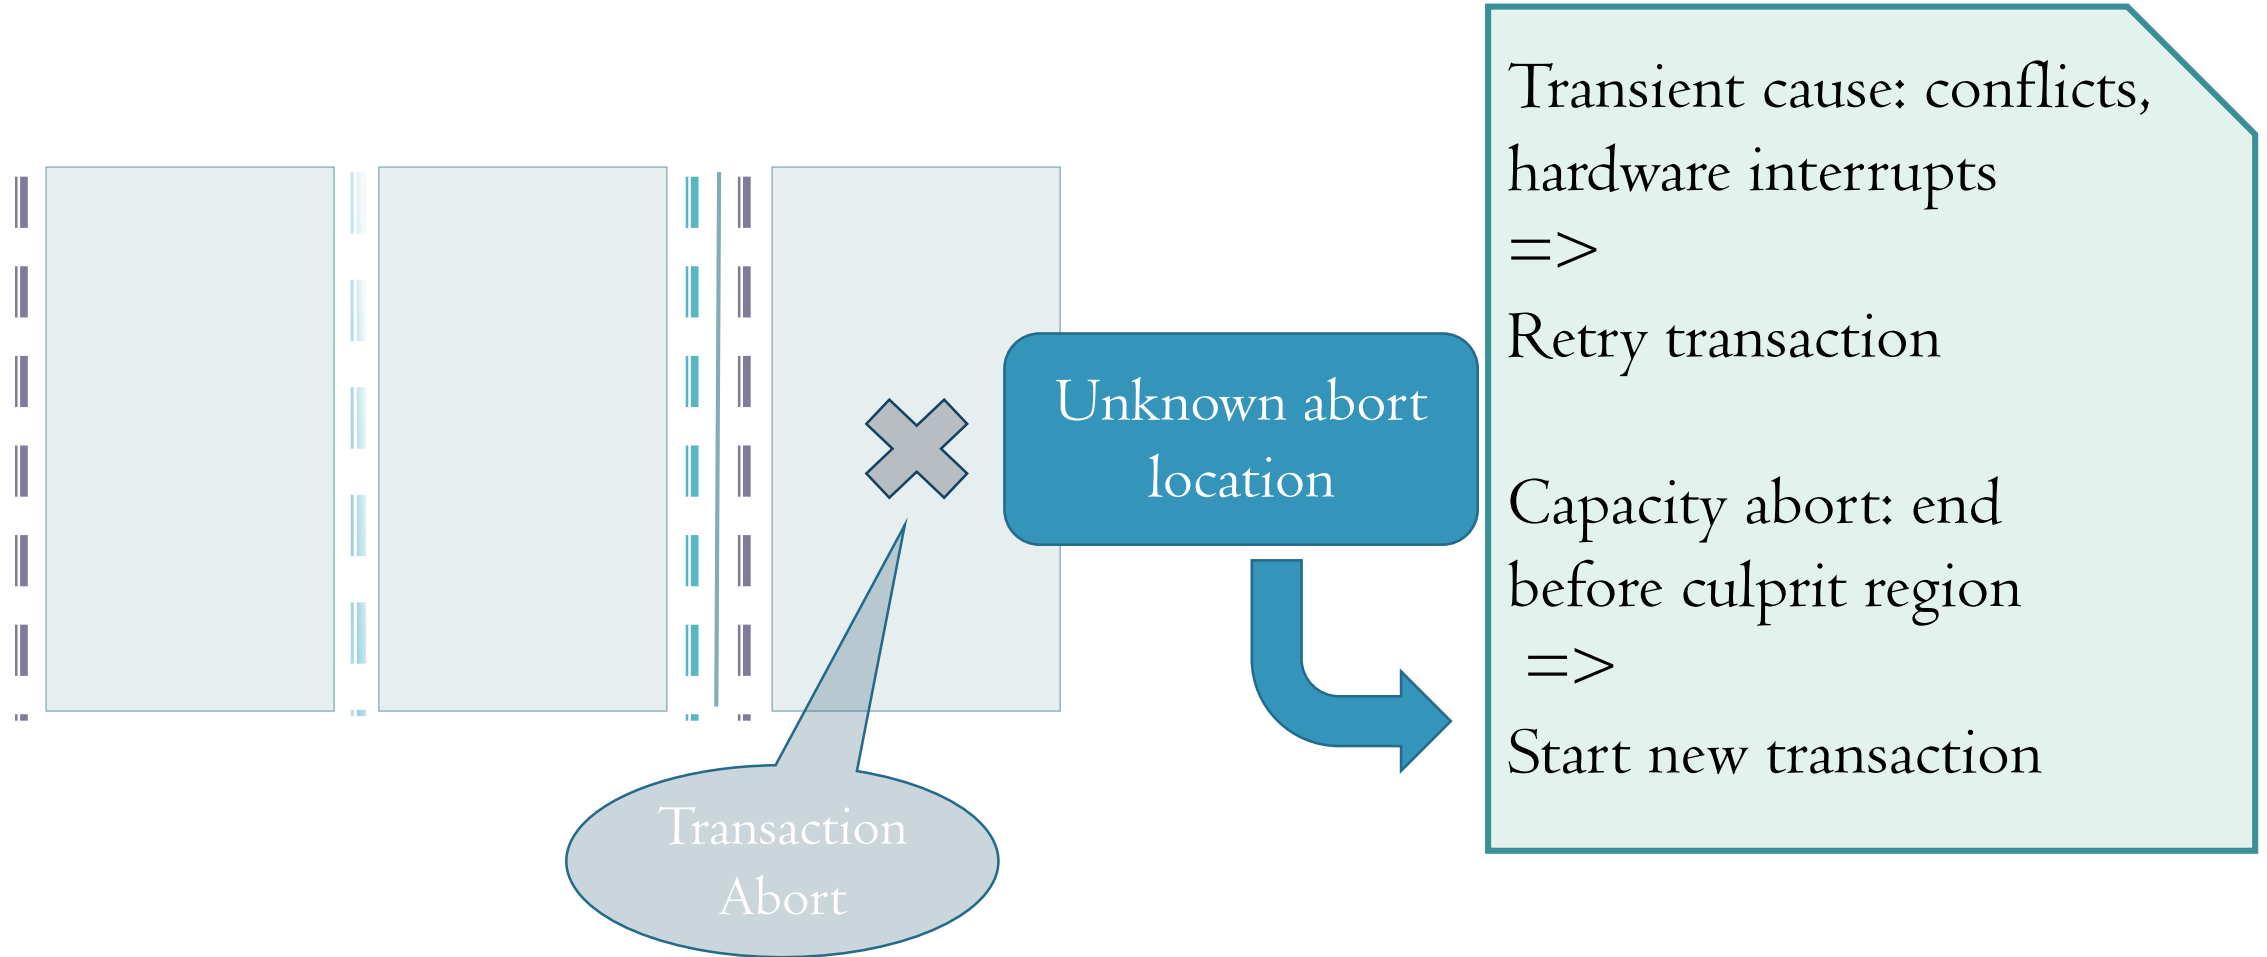
Retry transaction

Capacity abort: end
before culprit region


=>

Start new transaction

Per-transaction cost vs Abort cost



Per-transaction cost vs Abort cost



Transient cause: conflicts,
hardware interrupts

=>


Retry transaction

Capacity abort: end
before culprit region

=>

Start new transaction

Per-transaction cost vs Abort cost



Transaction
End

Minimal
Learning



Transient cause: conflicts,
hardware interrupts

=>

Retry transaction

Capacity abort: end
before culprit region

=>

Start new transaction

Future transaction
behavior unknown

Legato: Our Approach


Decide on a merge target: use **history** of previous transaction

1. **Temporary target** changes **rapidly**: capture transient effects

2. **Setpoint** or “**steady state**” target changes **slowly**: capture program phases


Merging Algorithm

Curr Target: 8
Setpoint: 8



Merging Algorithm: Initial Phase


Curr Target: 8
Setpoint: 8



Transaction
Start

Merging Algorithm: Initial Phase


Curr Target: 8
Setpoint: 8



Transaction
Merge


Merging Algorithm: Initial Phase

Curr Target: 8
Setpoint: 8



Transaction
Merge


Merging Algorithm: Recover from Transient Error



Curr Target: $8/2 = 4$
Setpoint: 8

Rapid action:


Merging Algorithm: Recover from Transient Error



Transaction
Start

Curr Target: 4
Setpoint: 8


Merging Algorithm: Recover from Transient Error



Transaction
Merge


Curr Target: 4
Setpoint: 8

Merging Algorithm: Recover from Transient Error



Curr Target: $4/2 = 2$
Setpoint: 8


Merging Algorithm: Recover from Transient Error



Transaction
Start

Curr Target: 2
Setpoint: 8

Merging Algorithm: Recover from Transient Error




Transaction
End

Curr Target: 2
Setpoint: 8


No Action:
Skeptical

Merging Algorithm: Build up to Setpoint




Curr Target: $2 * 2 = 4$
Setpoint: 8

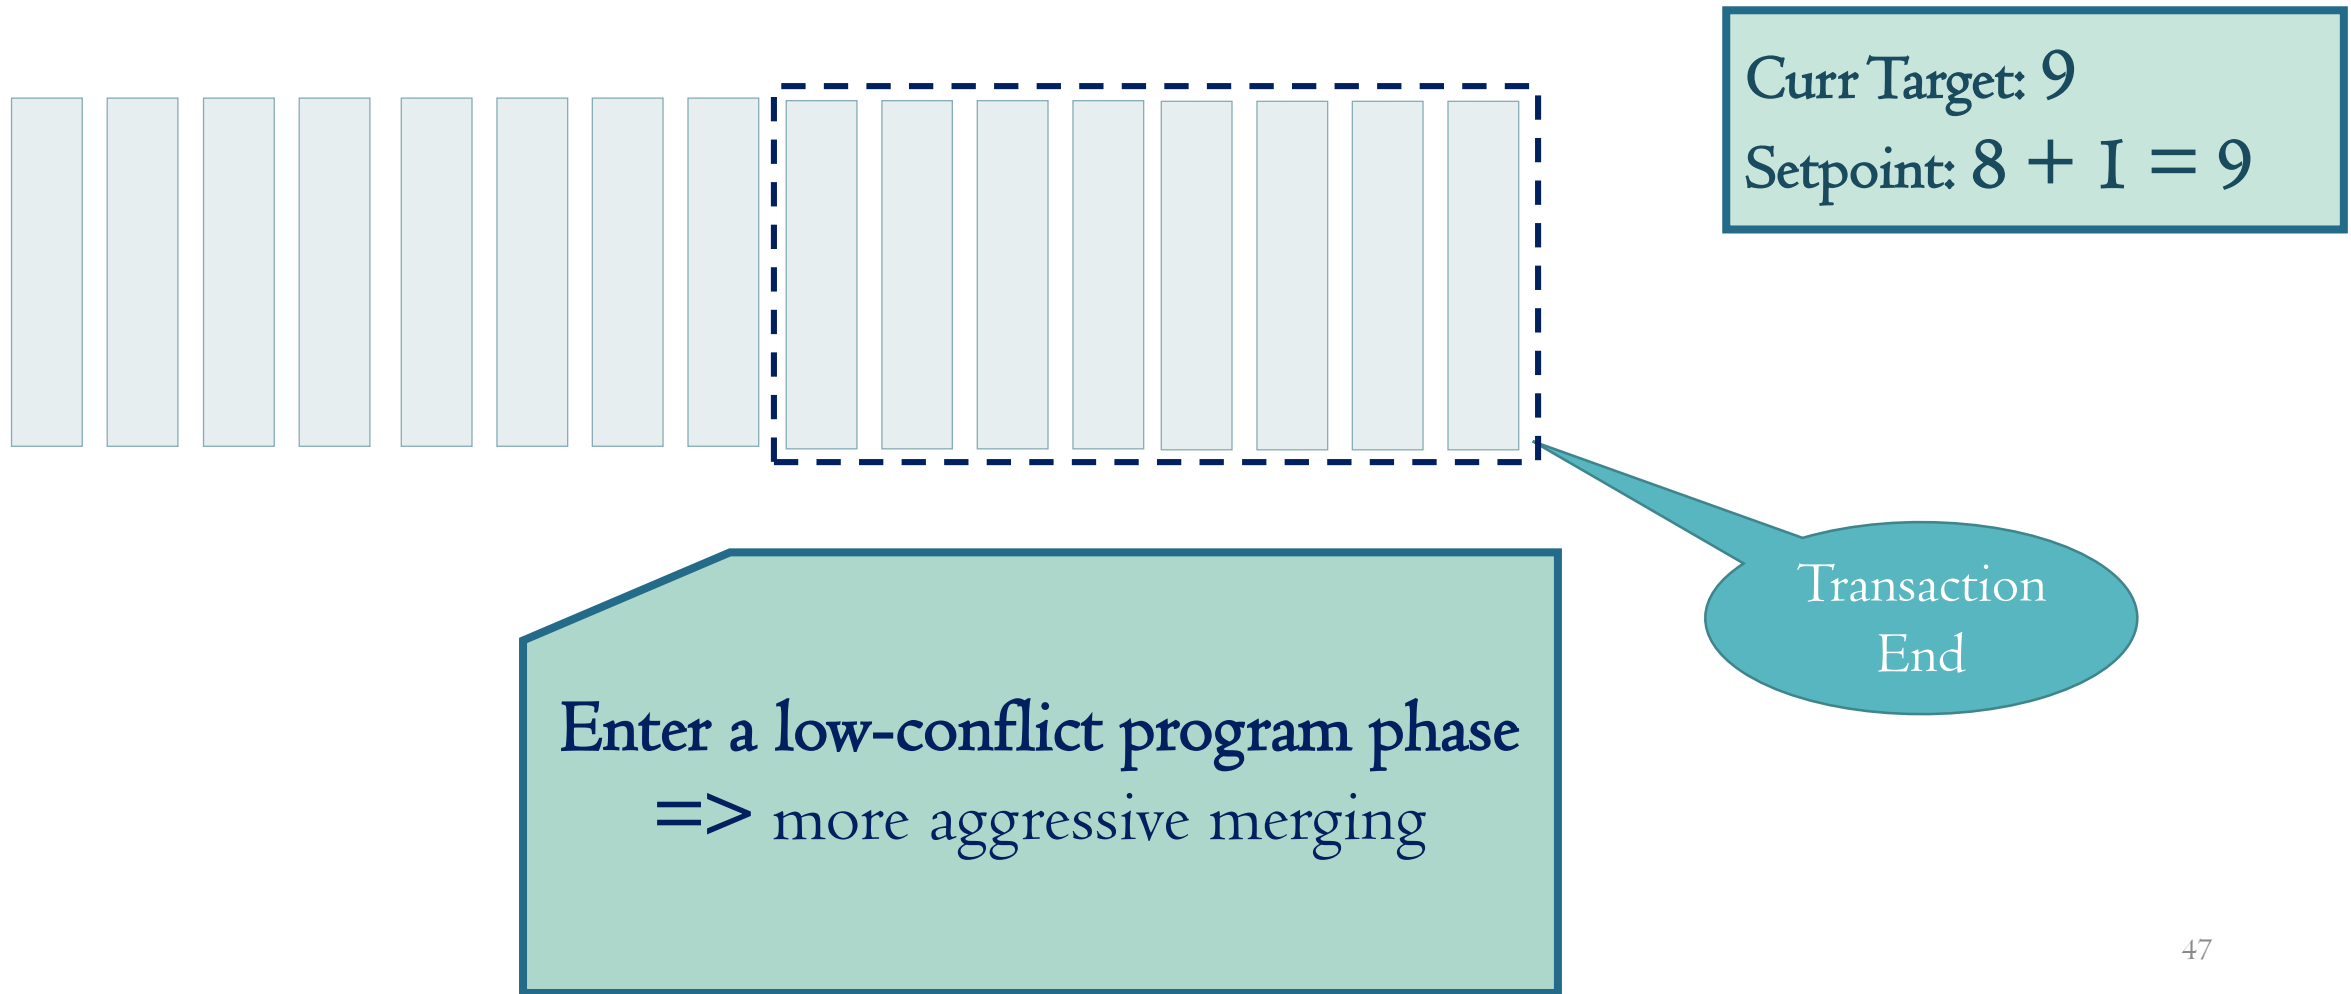
Merging Algorithm: Build up to Setpoint




Merging Algorithm: Change of Program Phase



Merging Algorithm: Change of Program Phase



Merging Algorithm: Change of Program Phase



Implementation and Evaluation



Implementation


Developed in Jikes RVM 3.1.3

Code publicly available in Jikes RVM Research Archive


Run-time Performance




Run-time Performance




Run-time Performance



Run-time Performance




Run-time Performance




JIT Compilation Time




JIT Compilation Time



JIT Compilation Time



JIT Compilation Time



Related Work

- *Checks conflicts in bounded region*
DRFx, Marino et al., PLDI 2010
- *Checks conflicts in synchronization-free regions*
Conflict Exceptions, Lucia et al., ISCA 2010
- *Enforces atomicity of bounded regions*
BulkCompiler, Ahn et al., MICRO 2009
Atom-Aid, Lucia et al., ISCA 2008
- *Reducing dynamic misspeculations*
BlockChop, Mars and Kumar, ISCA 2012

Memory Models: Run-time cost vs Strength

