

Tolerating Memory Leaks

Michael D. Bond Kathryn S. McKinley



Bugs in Deployed Software

- Deployed software fails
 - Different environment and inputs → different behaviors
- Greater complexity & reliance



Bugs in Deployed Software

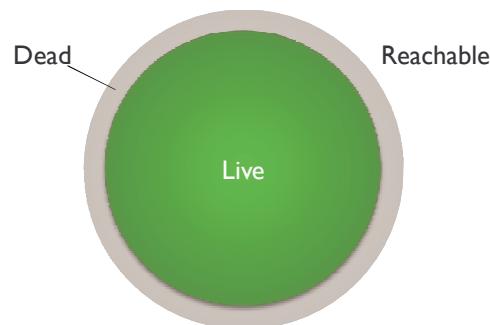
- Deployed software fails
 - Different environment and inputs → different behaviors
- Greater complexity & reliance
- Memory leaks are a real problem
- Fixing leaks is hard

Memory Leaks in Deployed Systems

- Memory leaks are a real problem
 - Managed languages do not eliminate them

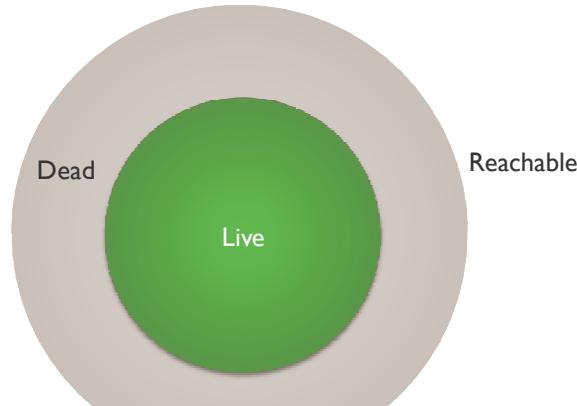
Memory Leaks in Deployed Systems

- Memory leaks are a real problem
 - Managed languages do not eliminate them



Memory Leaks in Deployed Systems

- Memory leaks are a real problem
 - Managed languages do not eliminate them



Memory Leaks in Deployed Systems

- Memory leaks are a real problem
 - Managed languages do not eliminate them

The diagram consists of three concentric circles. The innermost circle is green and labeled 'Live'. The middle ring is light gray and labeled 'Dead' on its left side. The outermost ring is also light gray and labeled 'Reachable' on its right side.

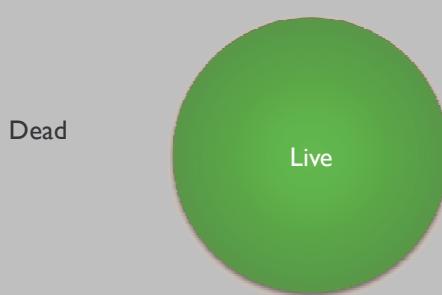
Memory Leaks in Deployed Systems

- Memory leaks are a real problem
 - Managed languages do not eliminate them

The diagram consists of three concentric circles. The innermost circle is green and labeled 'Live'. The middle ring is light gray and labeled 'Dead' on its left side. The outermost ring is also light gray and labeled 'Reachable' on its right side, with the word 'Reachable' also written vertically along the top edge of the outer ring.

Memory Leaks in Deployed Systems

- **Memory leaks are a real problem**
 - Managed languages do not eliminate them
 - Slow & crash real programs



Memory Leaks in Deployed Systems

- **Memory leaks are a real problem**
 - Managed languages do not eliminate them
 - Slow & crash real programs
 - Unacceptable for some applications

Memory Leaks in Deployed Systems

- **Memory leaks are a real problem**
 - Managed languages do not eliminate them
 - Slow & crash real programs
 - Unacceptable for some applications

- **Fixing leaks is hard**
 - Leaks take time to materialize
 - Failure far from cause

Example

- Driverless truck
 - 10,000 lines of C#
- Leak: past obstacles remained reachable



<http://www.codeproject.com/KB/showcase/IfOnlyWeUsedANTSProfiler.aspx>

Example

- Driverless truck
 - 10,000 lines of C#
- Leak: past obstacles remained reachable
- No immediate symptoms

"This problem was pernicious because it only showed up after 40 minutes to an hour of driving around and collecting obstacles."

<http://www.codeproject.com/KB/showcase/IfOnlyWeUsedANTSProfiler.aspx>

Example

- Driverless truck
 - 10,000 lines of C#
- Leak: past obstacles remained reachable
- No immediate symptoms

"This problem was pernicious because it only showed up after 40 minutes to an hour of driving around and collecting obstacles."

- Quick "fix": restart after 40 minutes

<http://www.codeproject.com/KB/showcase/IfOnlyWeUsedANTSProfiler.aspx>

Example

- Driverless truck
 - 10,000 lines of C#
- Leak: past obstacles remained reachable
- No immediate symptoms

"This problem was pernicious because it only showed up after 40 minutes to an hour of driving around and collecting obstacles."
- Quick "fix": restart after 40 minutes
- Environment sensitive
 - More obstacles in deployment
 - Unresponsive after 28 minutes



<http://www.codeproject.com/KB/showcase/IfOnlyWeUsedANTSProfiler.aspx>

Uncertainty in Deployed Software

- Unknown leaks; unexpected failures
- Online leak diagnosis helps
 - Too late to help failing systems

Uncertainty in Deployed Software

- Unknown leaks; unexpected failures
- Online leak diagnosis helps
 - Too late to help failing systems
- Also tolerate leaks

Uncertainty in Deployed Software

- Unknown leaks; unexpected failures
- Online leak diagnosis helps
 - Too late to help failing systems
- Also tolerate leaks

Illusion of fix

Uncertainty in Deployed Software

- Unknown leaks; unexpected failures
- Online leak diagnosis helps
 - Too late to help failing systems
- Also tolerate leaks

Illusion of fix

Eliminate bad effects

Don't slow Don't crash

Uncertainty in Deployed Software

- Unknown leaks; unexpected failures
- Online leak diagnosis helps
 - Too late to help failing systems
- Also tolerate leaks

Illusion of fix

Eliminate bad effects

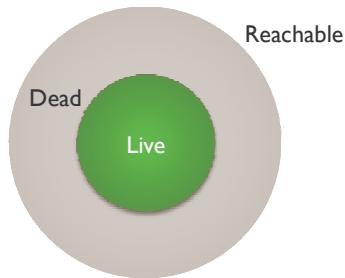
Don't slow Don't crash

Preserve semantics

Defer OOM errors

Predicting the Future

- Dead objects \Leftrightarrow not used again
- Highly stale objects \Rightarrow likely leaked



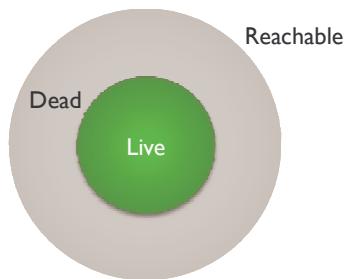
Predicting the Future

- Dead objects \Leftrightarrow not used again
- Highly stale objects \Rightarrow likely leaked

[Chilimbi & Hauswirth '04]

[Qin et al. '05]

[Bond & McKinley '06]

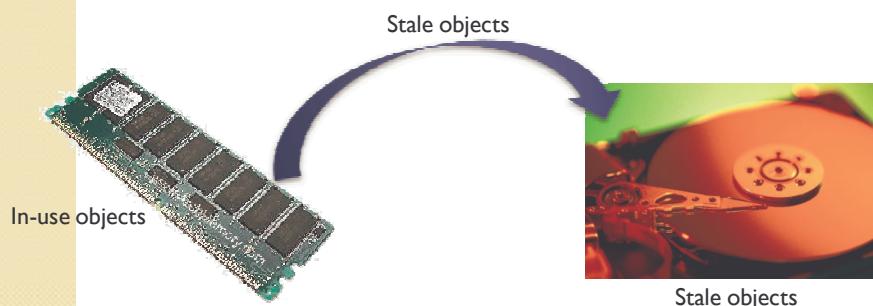


Tolerating Leaks with Melt

- Move **highly stale** objects to disk
 - Much larger than memory
 - Time & space proportional to live memory
 - Preserve semantics



Sounds like Paging!

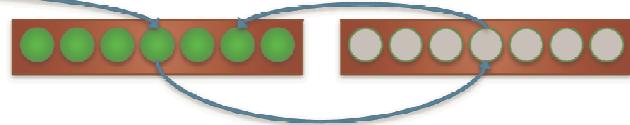


Sounds like Paging!

- Paging insufficient for managed languages
 - Need object granularity



- GC's working set is all reachable objects

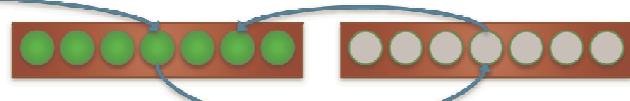


Sounds like Paging!

- Paging insufficient for managed languages
 - Need object granularity

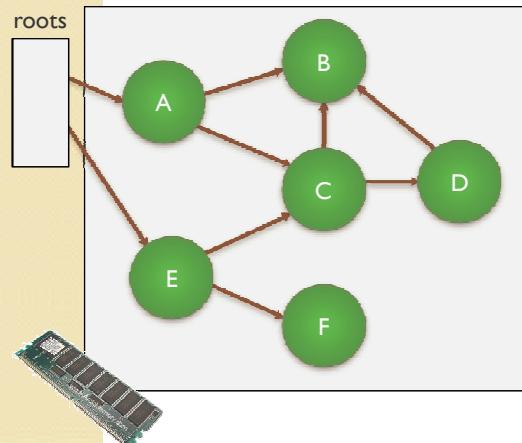


- GC's working set is all reachable objects

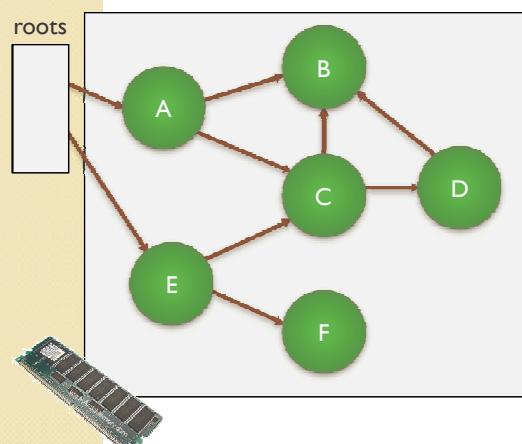


Bookmarking collection [Hertz et al.'05]

Challenge #1: How does Melt identify stale objects?



Challenge #1: How does Melt identify stale objects?



GC:
for all fields a.f
a.f |= 0x1;

Challenge #1: How does Melt identify stale objects?

The diagram shows a memory graph with six nodes labeled A through F. Nodes A, B, C, D, E, and F are represented by green circles. Node A is connected to B, C, and E. Node B is connected to A, C, and D. Node C is connected to A, B, D, and E. Node D is connected to C. Node E is connected to A, C, and F. Node F is connected to E. A white rectangular box labeled "roots" contains node A. Below the graph is a small image of a RAM chip.

GC:

```
for all fields a.f  
  a.f |= 0x1;
```

Challenge #1: How does Melt identify stale objects?

The diagram shows a memory graph with six nodes labeled A through F. Nodes A, B, C, D, E, and F are represented by green circles. Node A is connected to B, C, and E. Node B is connected to A, C, and D. Node C is connected to A, B, D, and E. Node D is connected to C. Node E is connected to A, C, and F. Node F is connected to E. A white rectangular box labeled "roots" contains node A. Below the graph is a small image of a RAM chip.

GC:

```
for all fields a.f  
  a.f |= 0x1;
```

Application:

```
b = a.f;  
if (b & 0x1) {  
  b &= ~0x1;  
  a.f = b; [atomic]  
}
```

Challenge #1: How does Melt identify stale objects?

GC:

```
for all fields a.f
  a.f |= 0x1;
```

Application:

```
b = a.f;
if (b & 0x1) {
  b &= ~0x1;
  a.f = b; [atomic]
}
```

Add 6% to application time

Challenge #1: How does Melt identify stale objects?

GC:

```
for all fields a.f
  a.f |= 0x1;
```

Application:

```
b = a.f;
if (b & 0x1) {
  b &= ~0x1;
  a.f = b; [atomic]
}
```

Challenge #1: How does Melt identify stale objects?

roots

```

GC:
for all fields a.f
  a.f |= 0x1;

Application:
b = a.f;
if (b & 0x1) {
  b &= ~0x1;
  a.f = b; [atomic]
}

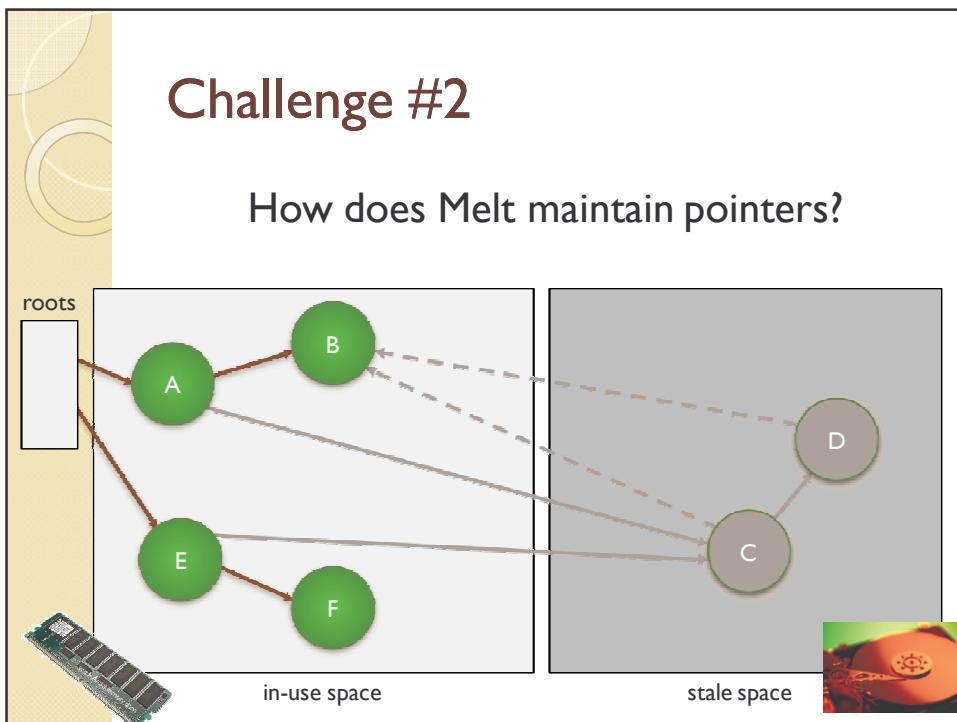
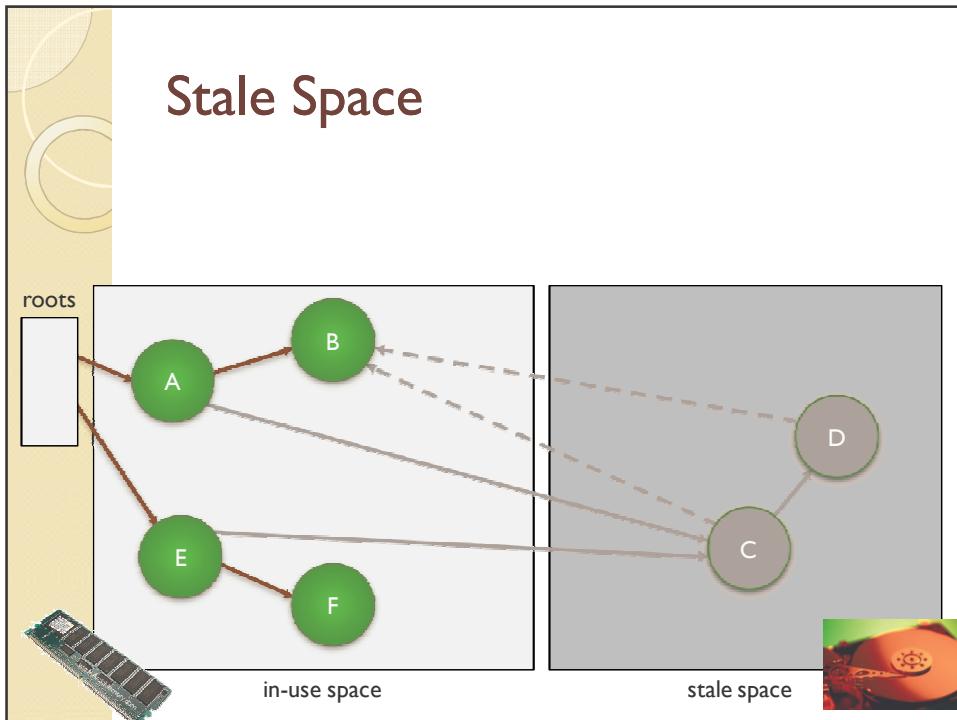
```

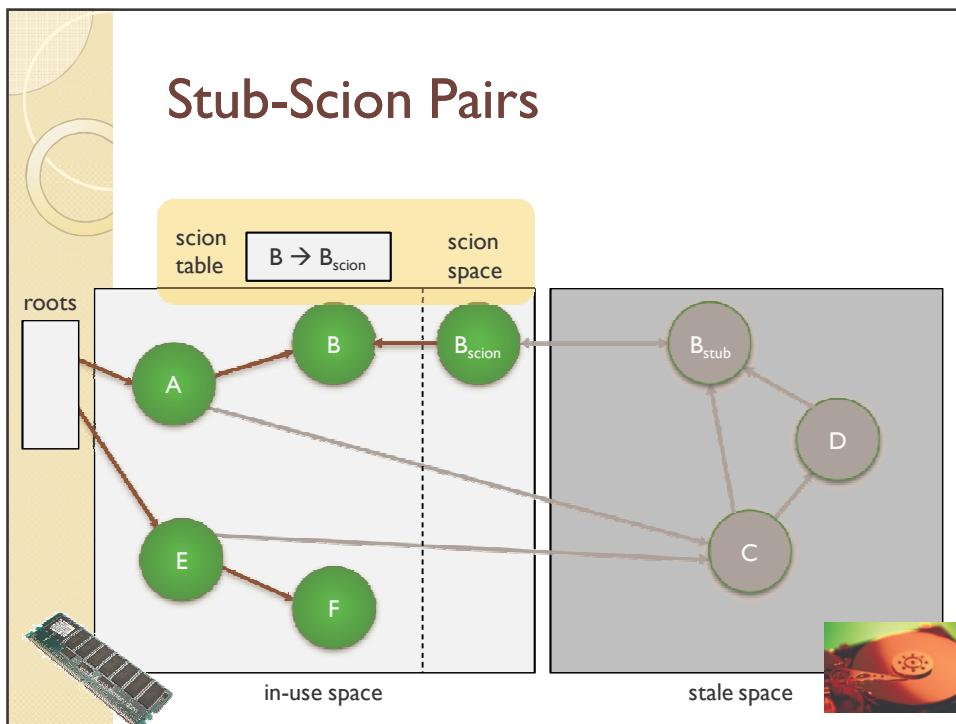
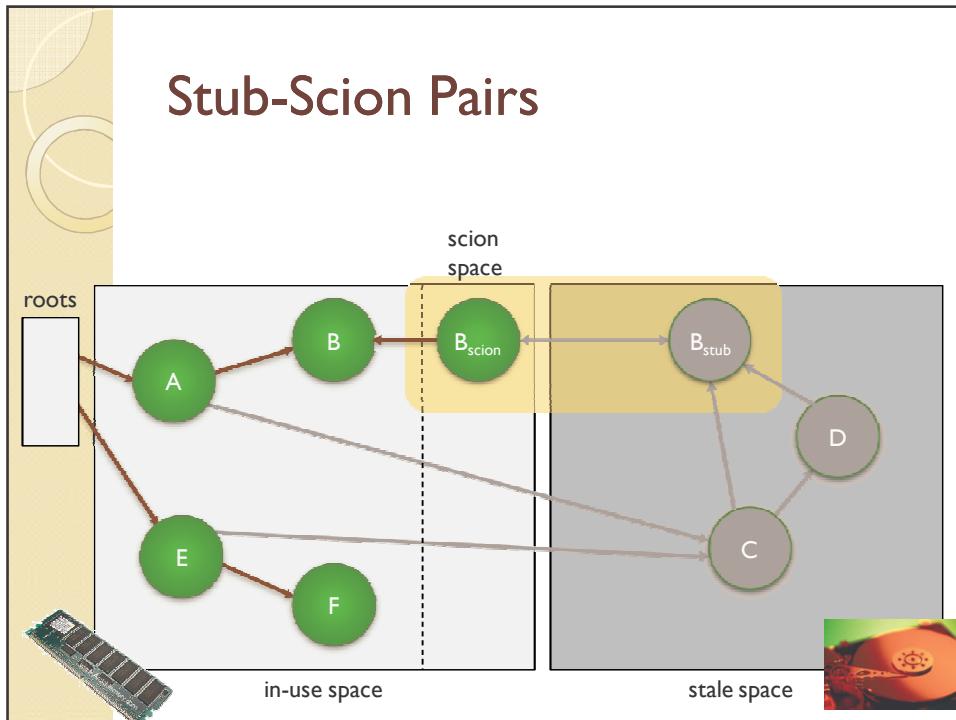
Stale Space

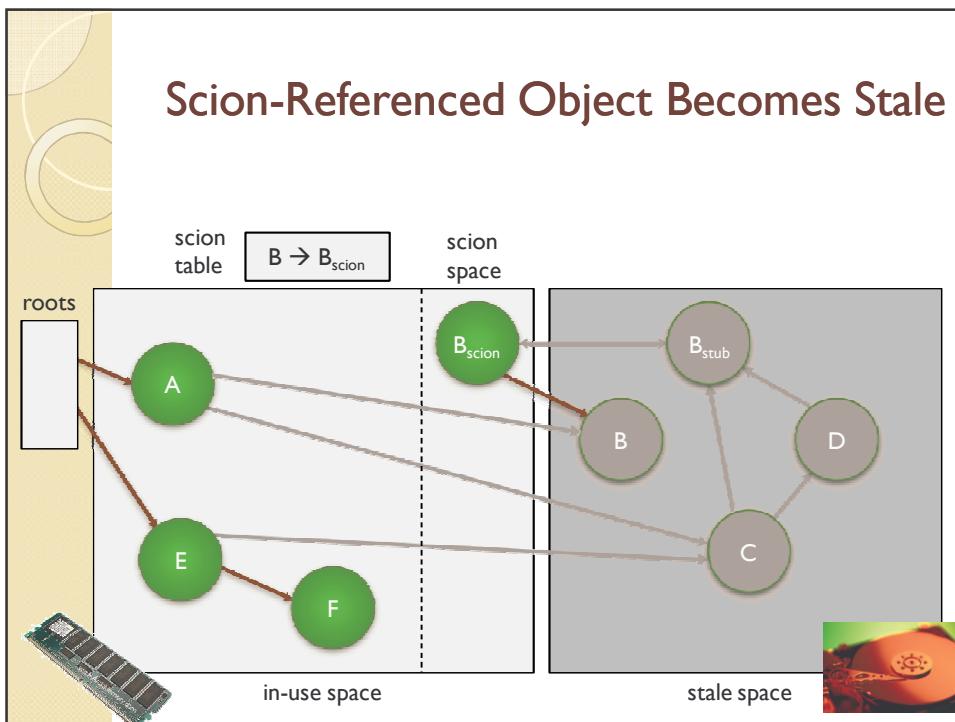
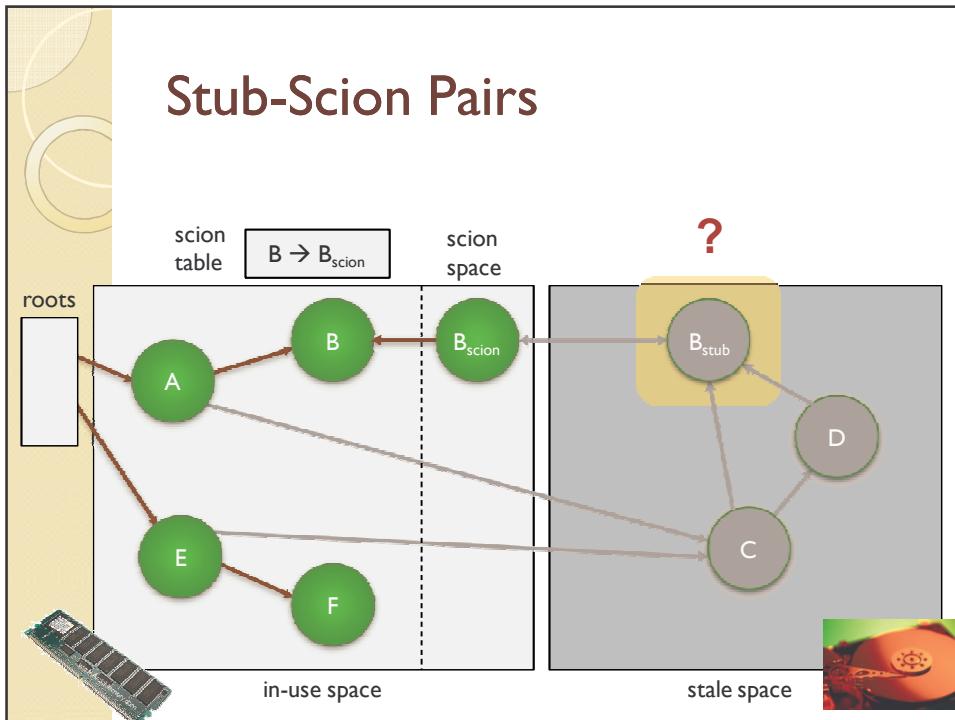
Heap nearly full → move stale objects to disk

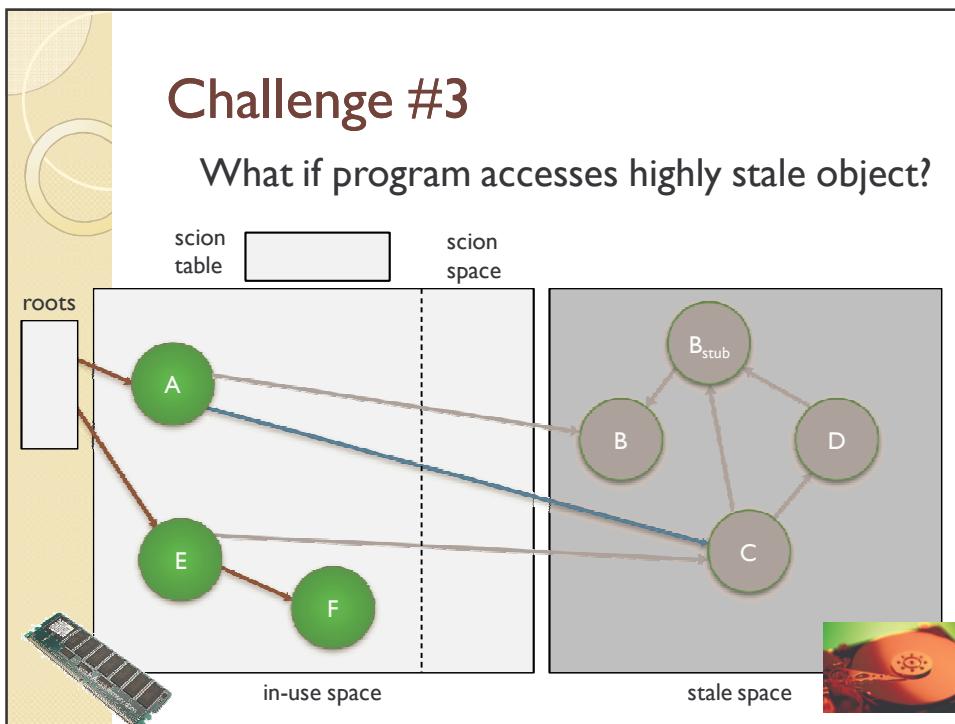
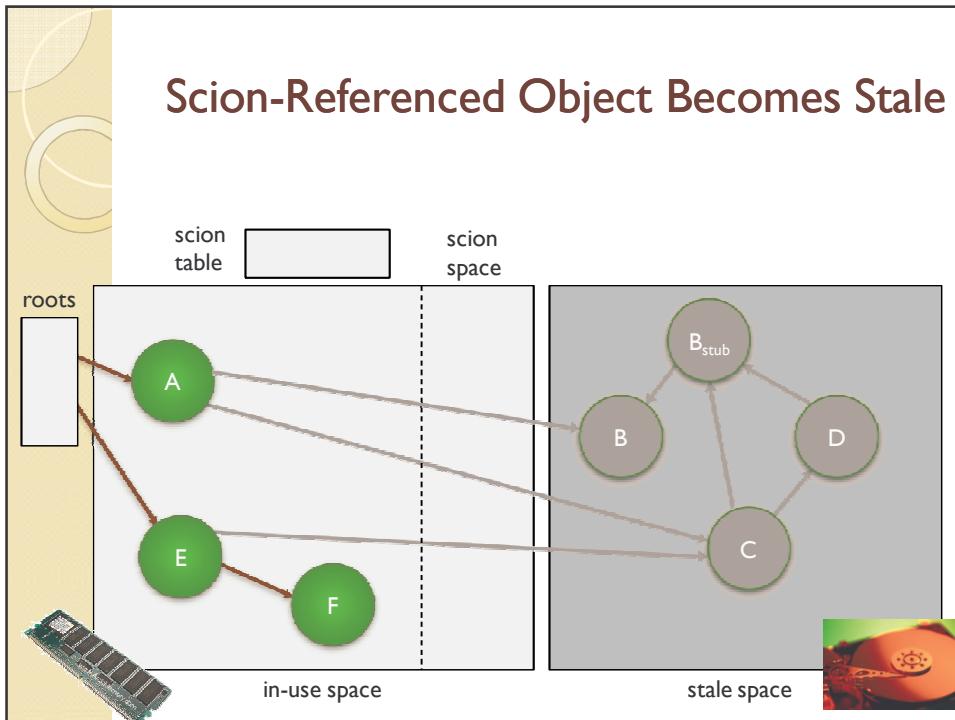
roots

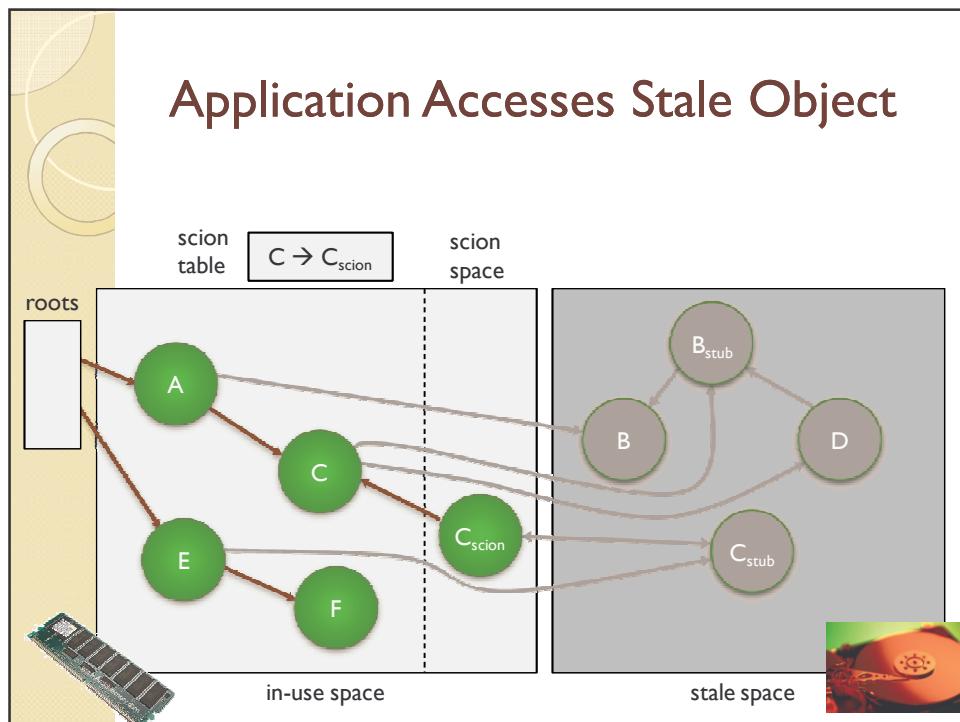
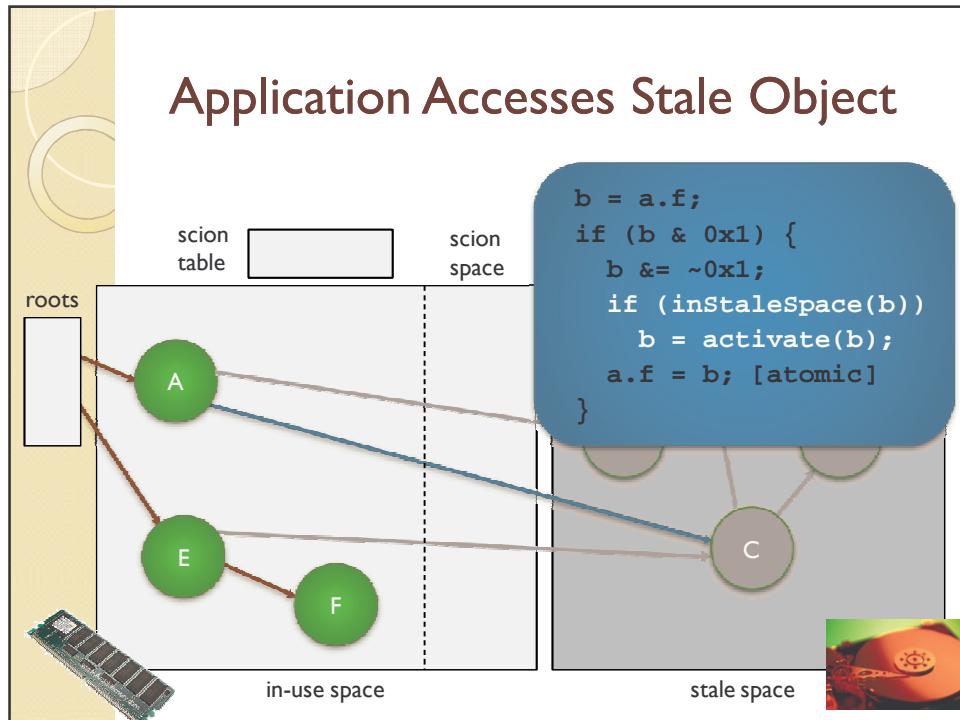
in-use space stale space

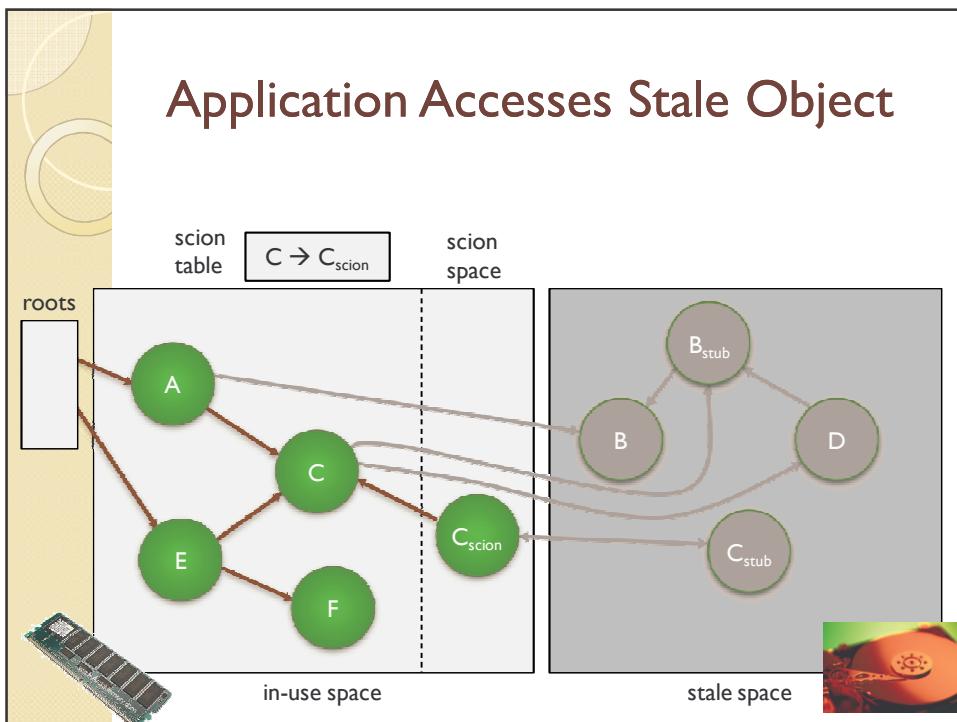
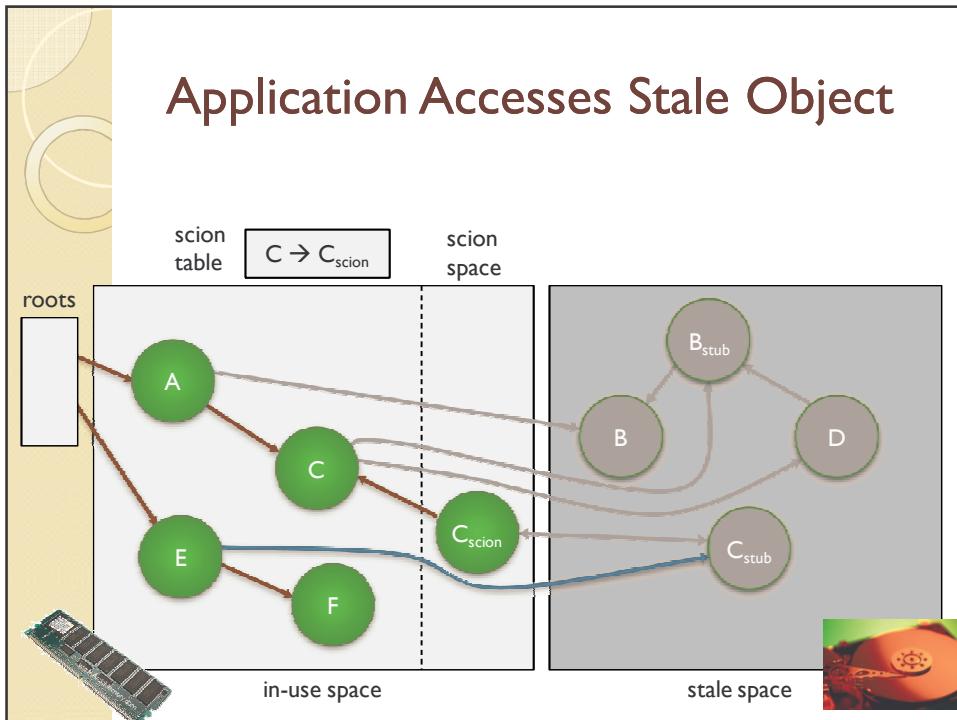










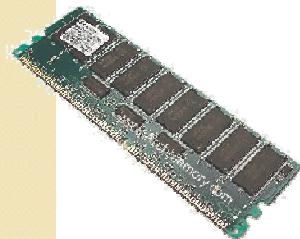


Implementation

- Integrated into Jikes RVM 2.9.2
 - Works with any tracing collector
 - Evaluation uses generational copying collector

Implementation

- Integrated into Jikes RVM 2.9.2
 - Works with any tracing collector
 - Evaluation uses generational copying collector



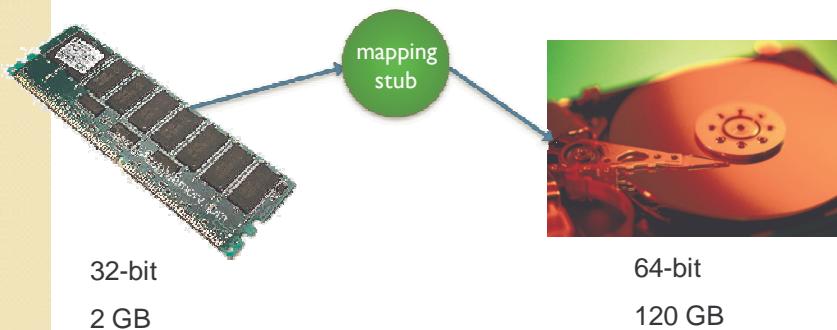
32-bit
2 GB



64-bit
120 GB

Implementation

- Integrated into Jikes RVM 2.9.2
 - Works with any tracing collector
 - Evaluation uses generational copying collector



Performance Evaluation

- Methodology
 - DaCapo, SPECjbb2000, SPECjvm98
 - Dual-core Pentium 4
 - Deterministic execution (replay)
- Results
 - 6% overhead (read barriers)
 - Stress test: still 6% overhead
 - Speedups in tight heaps (reduced GC workload)

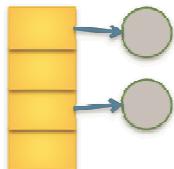
Leak	Melt's effect
Eclipse "Diff"	Tolerates until 24-hr limit (1,000X longer)
Eclipse "Copy-Paste"	Tolerates until 24-hr limit (194X longer)
JbbMod	Tolerates until 20-hr crash (19X longer)
ListLeak	Tolerates until disk full (200X longer)
SwapLeak	Tolerates until disk full (1,000X longer)
MySQL	Some highly stale but in-use (74X longer)
Delaunay Mesh	Short-running
DualLeak	Heap growth is in-use (2X longer)
SPECjbb2000	Heap growth is mostly in-use (2X longer)
Mckoi Database	Thread leak: extra support needed (2X longer)

Leak	Melt's effect
Eclipse "Diff"	Tolerates until 24-hr limit (1,000X longer)
Eclipse "Copy-Paste"	Tolerates until 24-hr limit (194X longer)
JbbMod	Tolerates until 20-hr crash (19X longer)
ListLeak	Tolerates until disk full (200X longer)
SwapLeak	Tolerates until disk full (1,000X longer)
MySQL	Some highly stale but in-use (74X longer)
Delaunay Mesh	Short-running
DualLeak	Heap growth is in-use (2X longer)
SPECjbb2000	Heap growth is mostly in-use (2X longer)
Mckoi Database	Thread leak: extra support needed (2X longer)

Tolerating Leaks

Leak	Melt's effect
Eclipse "Diff"	Tolerates until 24-hr limit (1,000X longer)
Eclipse "Copy-Paste"	Tolerates until 24-hr limit (194X longer)
JbbMod	Tolerates until 20-hr crash (19X longer)
ListLeak	Tolerates until disk full (200X longer)
SwapLeak	Tolerates until disk full (1,000X longer)
MySQL	Some highly stale but in-use (74X longer)
Delaunay Mesh	Short-running
DualLeak	Heap growth is in-use (2X longer)
SPECjbb2000	Heap growth is mostly in-use (2X longer)
Mckoi Database	Thread leak: extra support needed (2X longer)

Tolerating Leaks



MySQL	Some highly stale but in-use (74X longer)
Delaunay Mesh	Short-running
DualLeak	Heap growth is in-use (2X longer)
SPECjbb2000	Heap growth is mostly in-use (2X longer)
Mckoi Database	Thread leak: extra support needed (2X longer)

Tolerating Leaks



MySQL	Some highly stale but in-use (74X longer)
Delaunay Mesh	Short-running
DualLeak	Heap growth is in-use (2X longer)
SPECjbb2000	Heap growth is mostly in-use (2X longer)
Mckoi Database	Thread leak: extra support needed (2X longer)

Tolerating Leaks

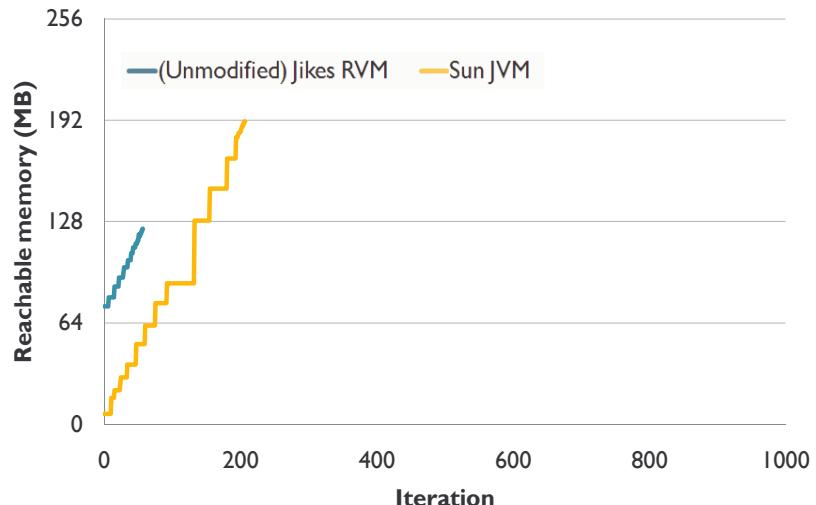


MySQL	Some highly stale but in-use (74X longer)
Delaunay Mesh	Short-running
DualLeak	Heap growth is in-use (2X longer)
SPECjbb2000	Heap growth is mostly in-use (2X longer)
Mckoi Database	Thread leak: extra support needed (2X longer)

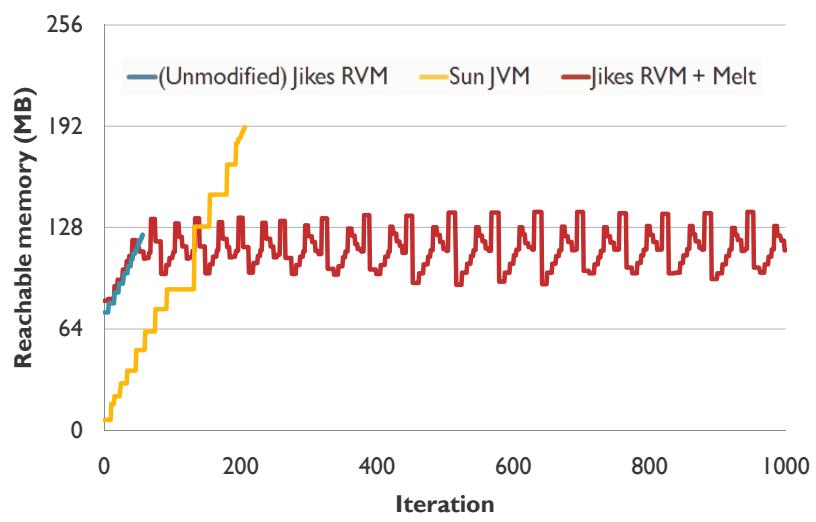
Tolerating Leaks	
MySQL	Some highly stale but in-use (74X longer)
Delaunay Mesh	Short-running
DualLeak	Heap growth is in-use (2X longer)
SPECjbb2000	Heap growth is mostly in-use (2X longer)
Mckoi Database	Thread leak: extra support needed (2X longer)

Tolerating Leaks	
Leak	Melt's effect
Eclipse "Diff"	Tolerates until 24-hr limit (1,000X longer)
Eclipse "Copy-Paste"	Tolerates until 24-hr limit (194X longer)
JbbMod	Tolerates until 20-hr crash (19X longer)
ListLeak	Tolerates until disk full (200X longer)
SwapLeak	Tolerates until disk full (1,000X longer)
MySQL	Some highly stale but in-use (74X longer)
Delaunay Mesh	Short-running
DualLeak	Heap growth is in-use (2X longer)
SPECjbb2000	Heap growth is mostly in-use (2X longer)
Mckoi Database	Thread leak: extra support needed (2X longer)

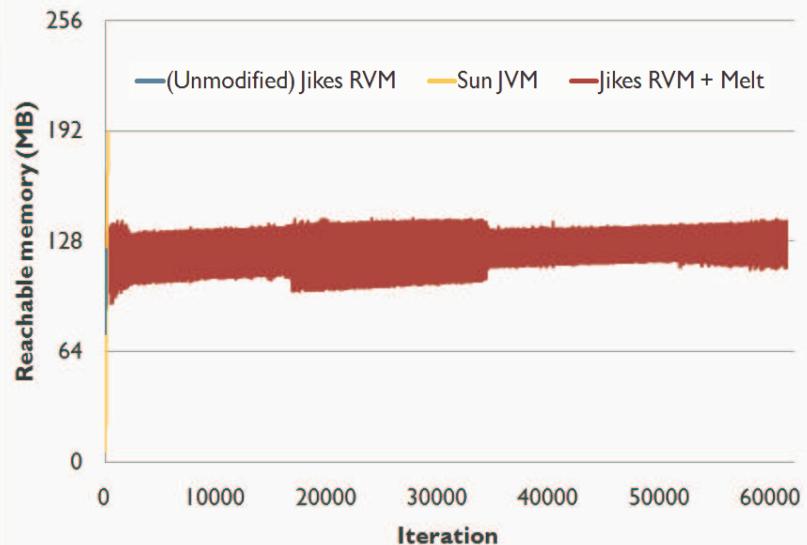
Eclipse Diff: Reachable Memory



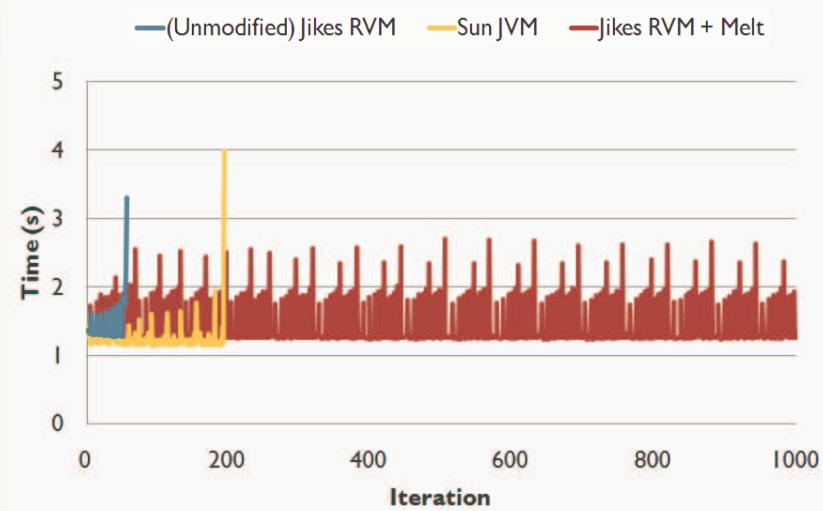
Eclipse Diff: Reachable Memory

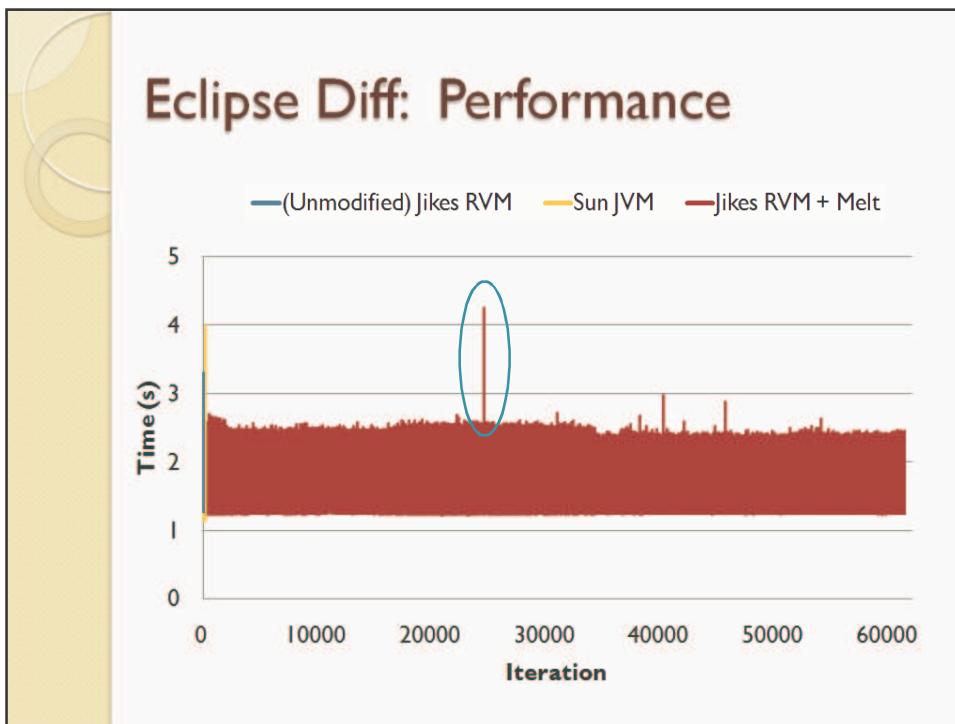
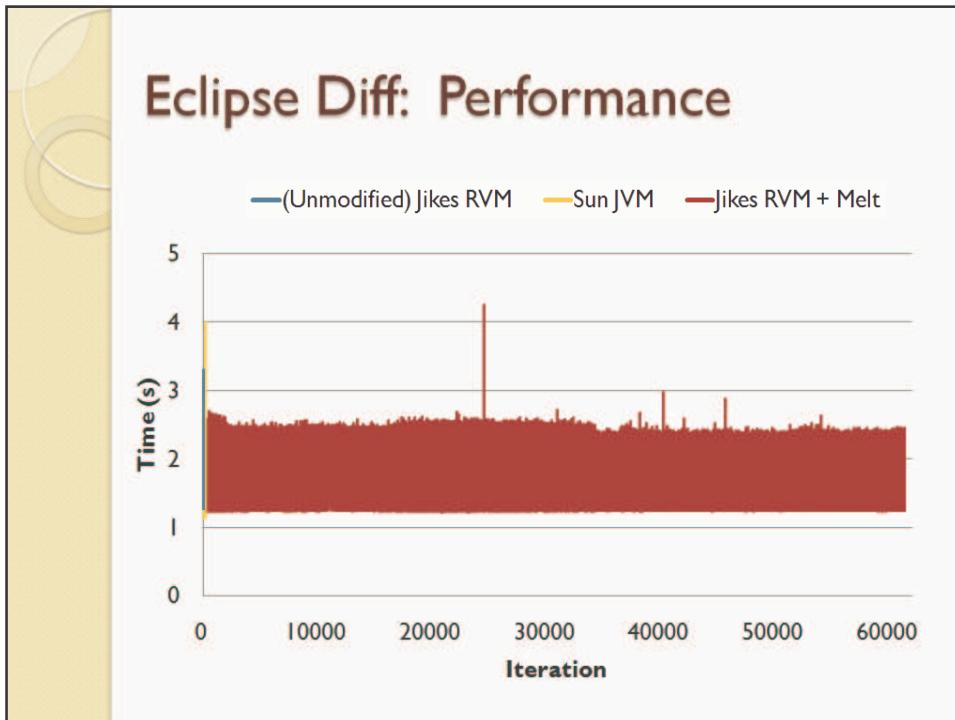


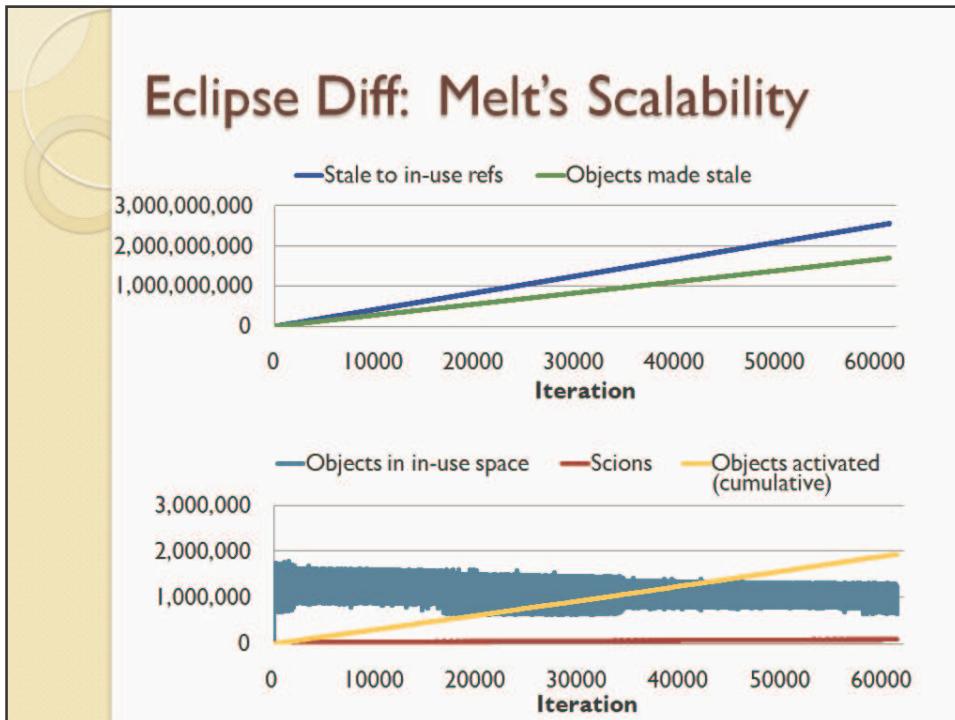
Eclipse Diff: Reachable Memory



Eclipse Diff: Performance







- ## Related Work
- Managed [LeakSurvivor, Tang et al. '08]
[Panacea, Goldstein et al. '07, Breitgand et al. '07]
 - Don't guarantee time & space proportional to live memory
 - Native [Cyclic memory allocation, Nguyen & Rinard '07]
[Plug, Novark et al. '08]
 - Different challenges & opportunities
 - Less coverage or change semantics
 - Orthogonal persistence & distributed GC
 - Barriers, swizzling, object faulting, stub-scion pairs

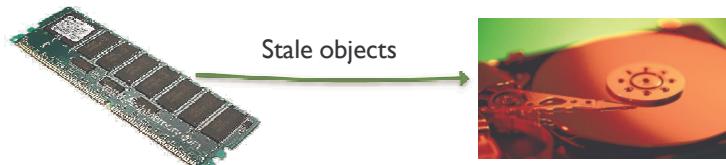
Conclusion

- Finding bugs before deployment is hard



Conclusion

- Finding bugs before deployment is hard
- Online diagnosis helps developers
- Help **users** in meantime
- Tolerate leaks with Melt: **illusion of fix**



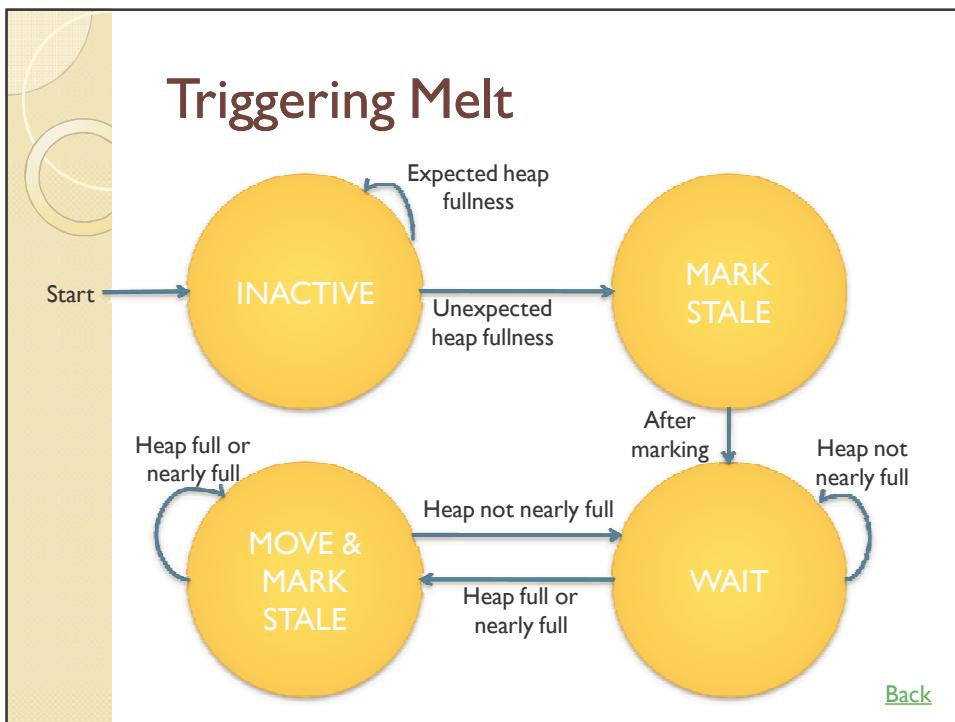
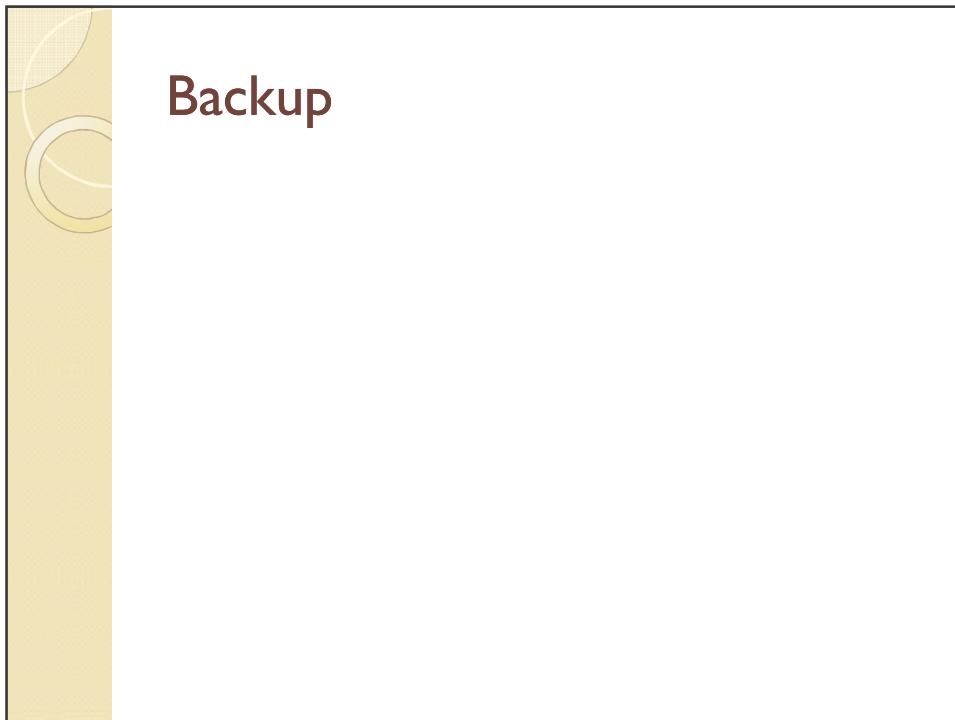
Conclusion

- Finding bugs before deployment is hard
- Online diagnosis helps developers
- Help **users** in meantime
- Tolerate leaks with Melt: **illusion of fix**
 - Time & space proportional to live memory
 - Preserve semantics

Conclusion

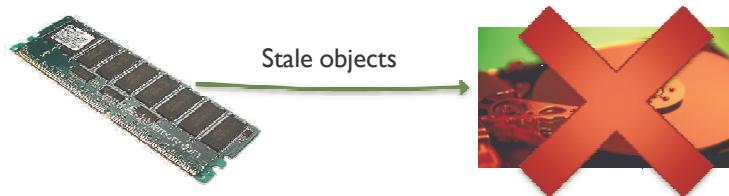
- Finding bugs before deployment is hard
- Online diagnosis helps developers
- Help **users** in meantime
- Tolerate leaks with Melt: **illusion of fix**
 - Time & space proportional to live memory
 - Preserve semantics
- Buys developers time to fix leaks

Thank you!



Conclusion

- Finding bugs before deployment is hard
- Online diagnosis helps developers
- To help **users** in meantime, tolerate bugs
- Tolerate leaks with Melt: **illusion of fix**



Related Work: Tolerating Bugs

- Nondeterministic errors
 - [Atom-Aid] [DieHard] [Grace] [Rx]
 - Memory corruption: perturb layout
 - Concurrency bugs: perturb scheduling
- General bugs
 - Ignore failing operations [FOC]
 - Need higher level, more proactive approaches

