

Prescient Memory: Exposing Weak Memory Model Behavior by Looking into the Future

MAN CAO

JAKE ROEMER

ARITRA SENGUPTA


MICHAEL D. BOND



Parallel Programming is Hard


Parallel Programming is Hard

- Shared-memory



Parallel Programming is Hard


- Shared-memory



- Difficult to be both correct and scalable

Parallel Programming is Hard


- Shared-memory



- Difficult to be both correct and scalable
 - Data race

Parallel Programming is Hard

- Shared-memory



- Difficult to be both correct and scalable

- Data race

- Fundamentally, lacks strong **semantic** guarantees

Example #1: Weak Semantics

Foo data = null;
boolean flag= false;

T1

data = new Foo();
flag = true;

T2

if (**flag**)
 data.bar();

Example #1: Weak Semantics


Foo data = null;
boolean flag= false;

T1

data = new Foo();
flag = true;

T2

if (**flag**)
 data.bar();



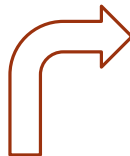
Null pointer
exception!

Example #1: Weak Semantics

Foo data = null;
boolean flag= false;

T1

No data
dependence

 **data** = new Foo();
flag = true;

T2

if (**flag**)
 data.bar();

Null pointer
exception!

Exposing Behaviors of Data Races

- Existing Approaches
 - Dynamic analyses
 - Model checkers

Exposing Behaviors of Data Races

- Existing Approaches
 - Dynamic analyses
 - Limitation: coverage
 - Model checkers
 - Limitation: scalability

Exposing Behaviors of Data Races

- Existing Approaches
 - Dynamic analyses
 - Limitation: coverage
 - Model checkers
 - Limitation: scalability

- Prescient Memory (PM)

Dynamic analysis with better coverage

Outline

- Memory Models and Behaviors of Data Races
- Design
 - Prescient Memory (PM)
 - PM-profiler
 - PM Workflow
- Evaluation

Memory Model

- Defines possible **values** that a **load** can return

Memory Model

- Defines possible **values** that a **load** can return

Strong

- Sequential Consistency (SC)
- Impractical to enforce

Memory Model

- Defines possible **values** that a **load** can return

Strong

- Sequential Consistency (SC)
- Impractical to enforce

Weak

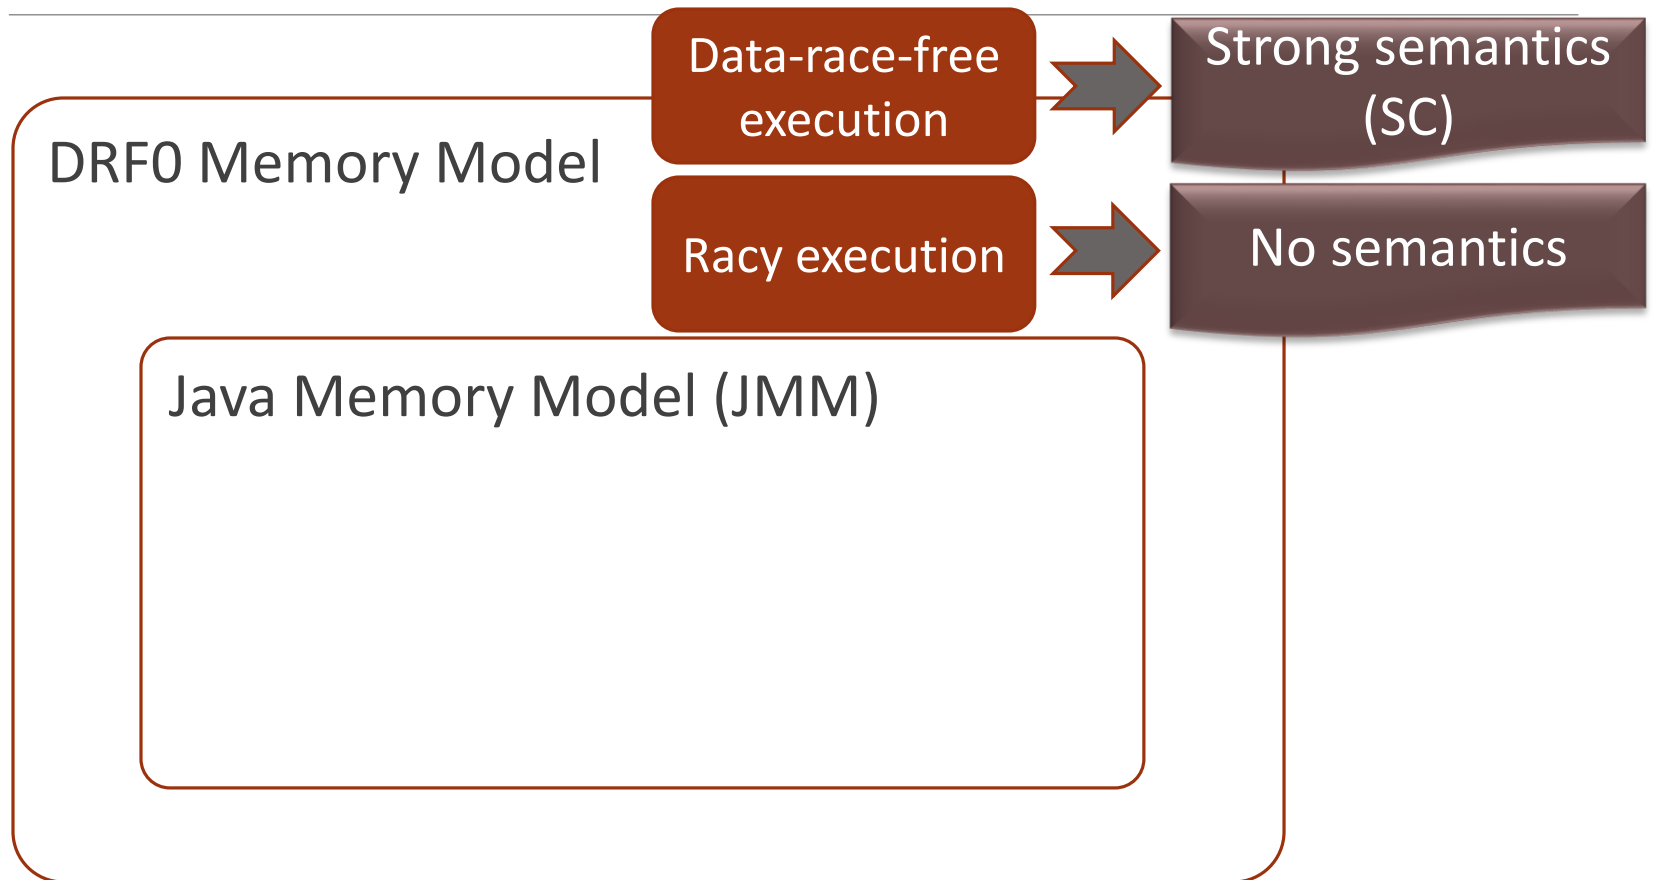
- Enables compiler & hardware optimizations
- DRF0, C++11, Java

Behaviors Allowed by Memory Models

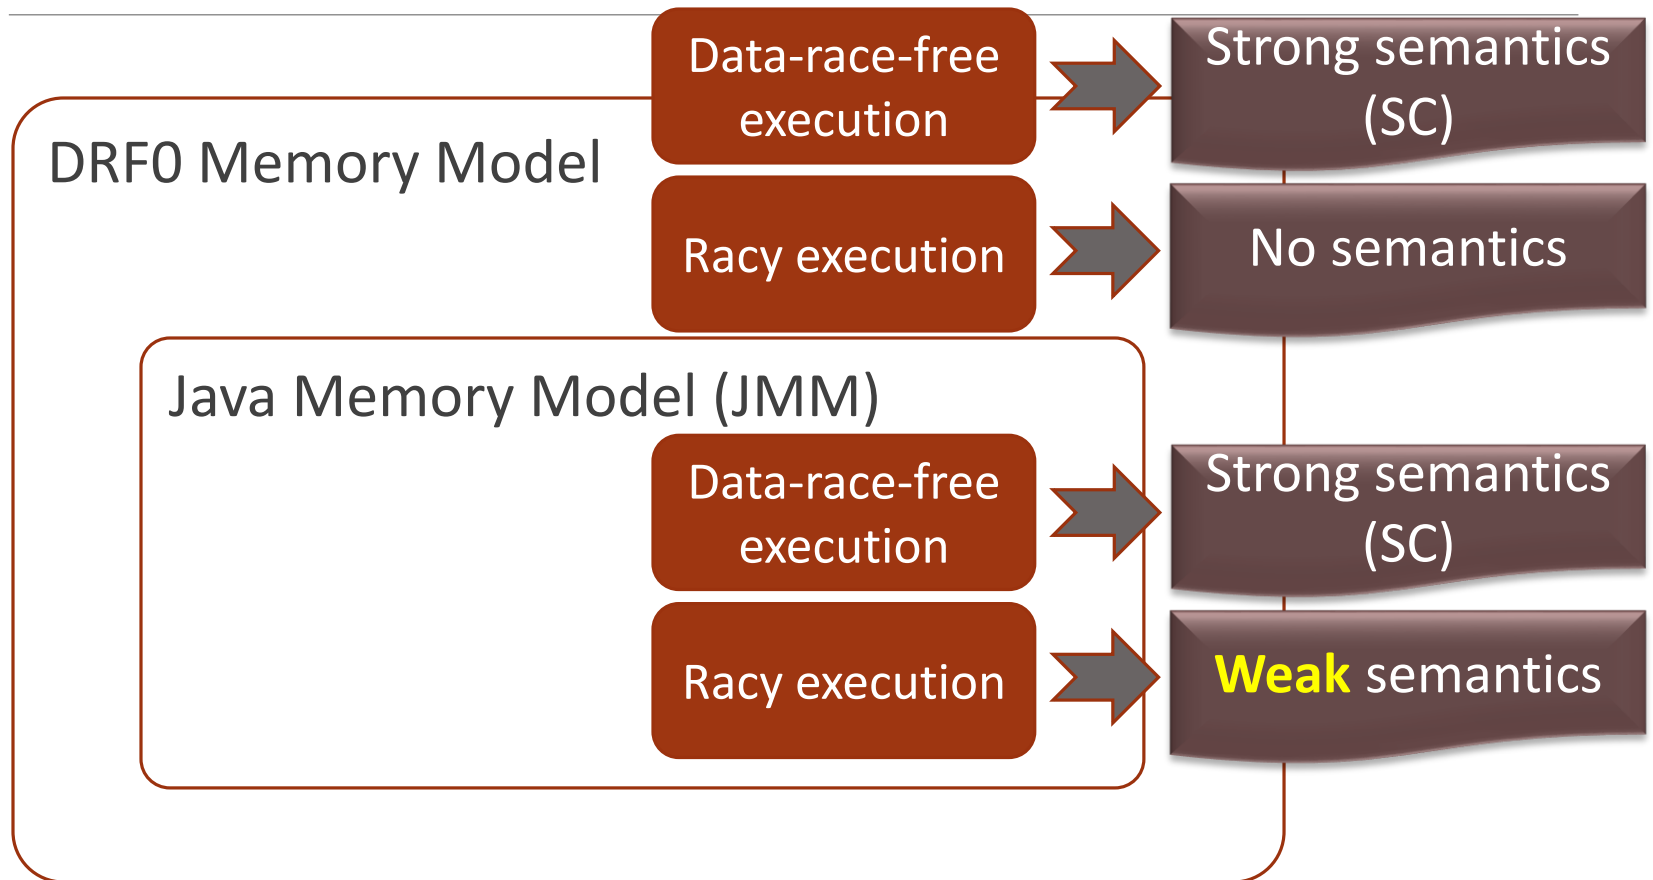
DRF0 Memory Model

Java Memory Model (JMM)

Behaviors Allowed by Memory Models



Behaviors Allowed by Memory Models



Behaviors Allowed by Memory Models

DRFO Memory Model

Data-race-free
execution

Strong semantics
(SC)

Racy execution

No semantics

Java Memory Model (JMM)

Data-race-free
execution

Strong semantics
(SC)

Racy execution

Weak semantics

Racy execution
can still lead to
surprising
behaviors!

Behaviors Allowed in JMM #1: Revisit

Foo data = null;
boolean flag= false;


T1

data = new Foo();
flag = true;


T2

if (**flag**)
 data.bar();

Behaviors Allowed in JMM #1: Revisit



Behaviors Allowed in JMM #1: Revisit



Null pointer
exception!

Behaviors Allowed in JMM #1: Revisit

Foo data = null;
boolean flag= false;

T1

data = new Foo();
flag = true;

T2

if (**flag**)
 data.bar();

Returning **stale value** can trigger the exception

Behaviors Allowed in JMM #2

```
int data = flag = 0;
```

T1


```
r = data;  
flag = 1;
```

```
assert r == 0;
```

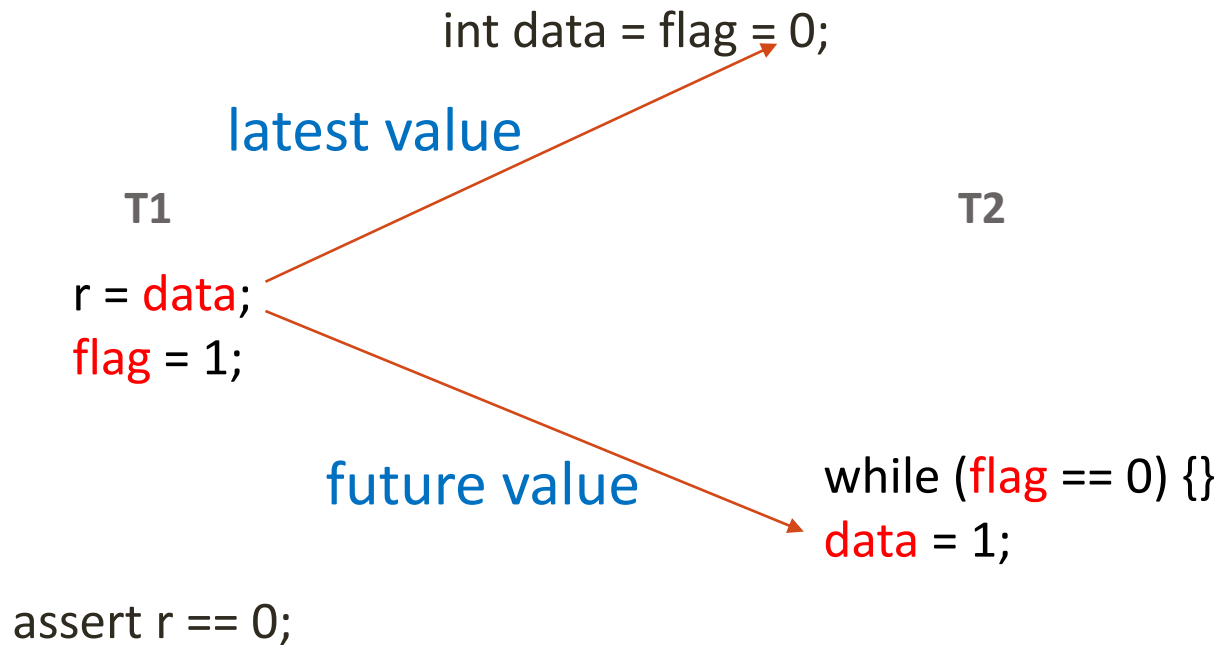
T2

```
while (flag == 0) {}  
data = 1;
```

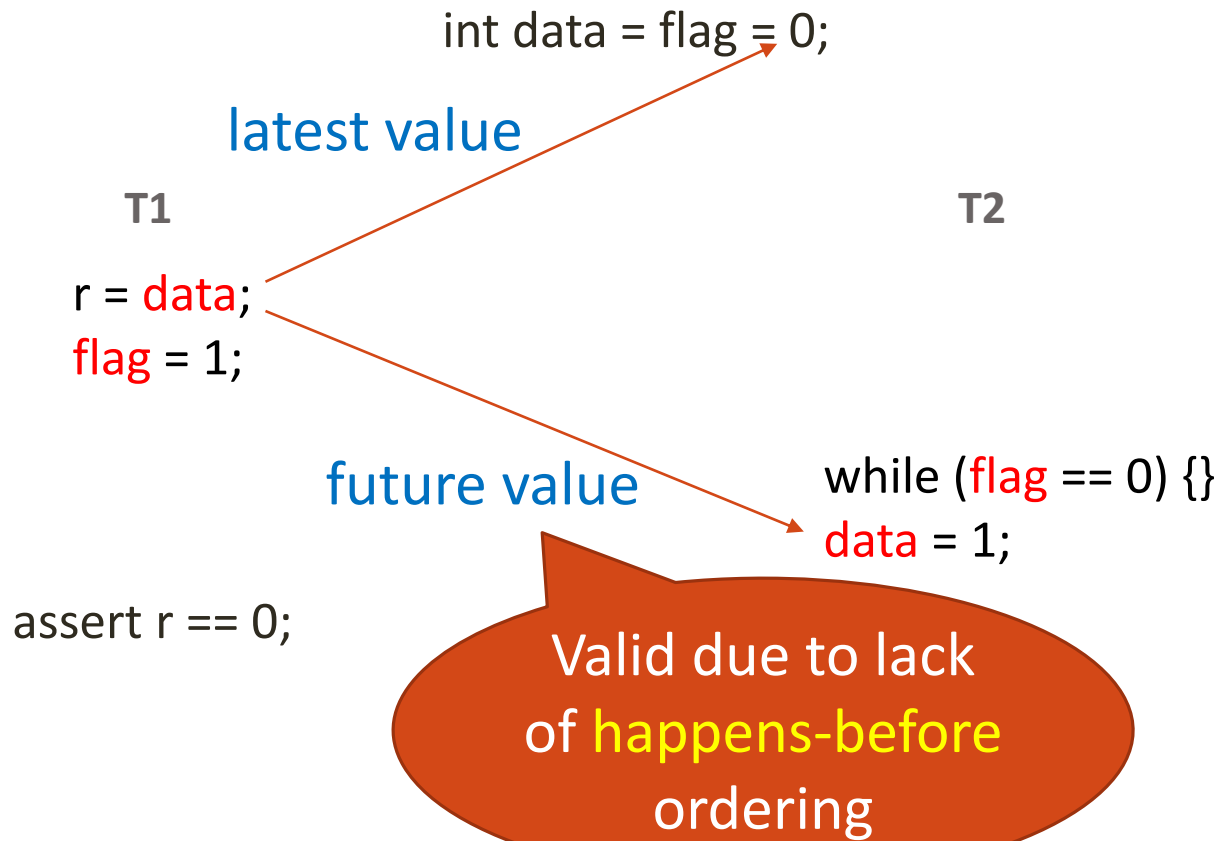
Behaviors Allowed in JMM #2



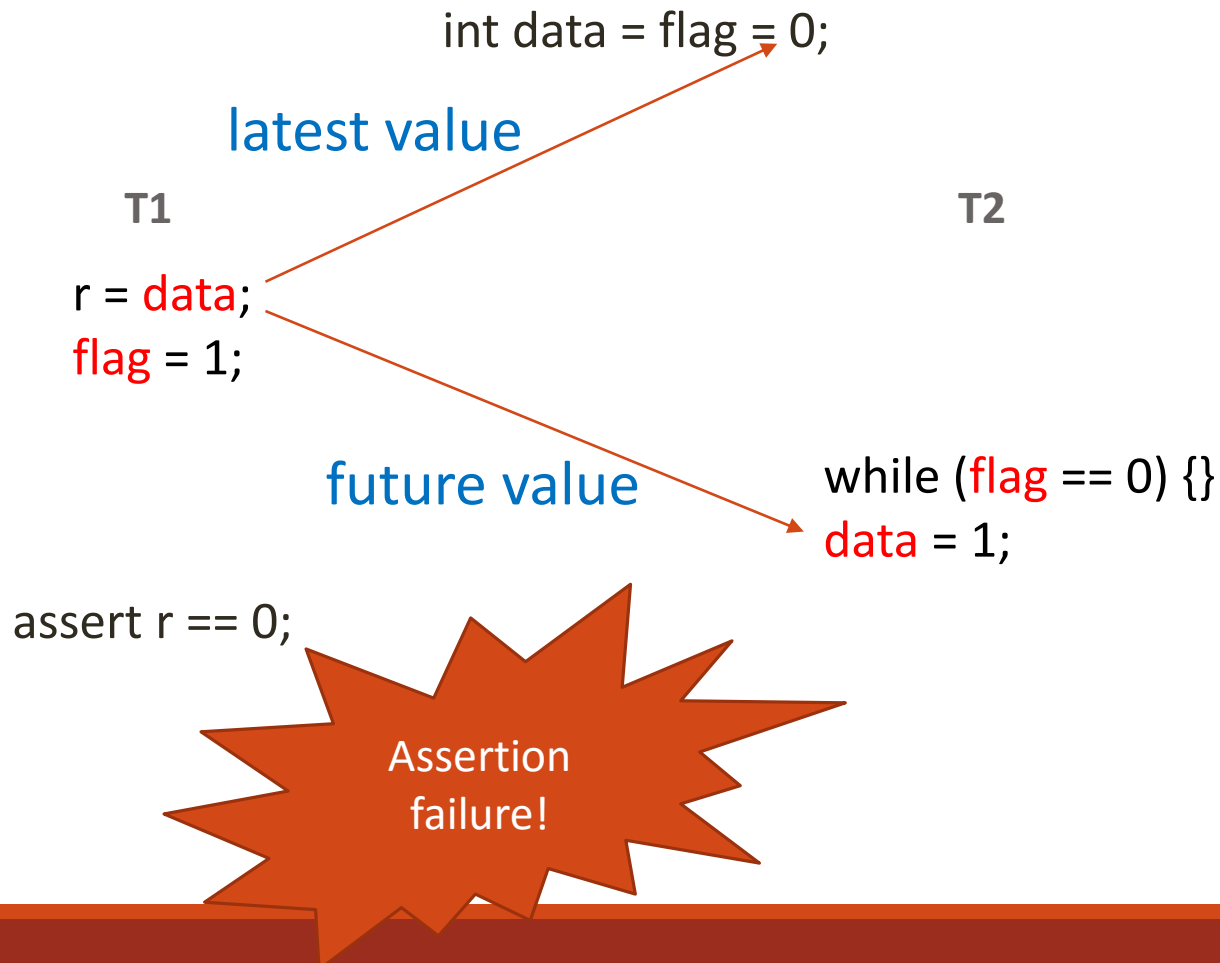
Behaviors Allowed in JMM #2



Behaviors Allowed in JMM #2




Behaviors Allowed in JMM #2




Behaviors Allowed in JMM #2

```
int data = flag = 0;
```

T1



```
r = data;  
flag = 1;
```



T2

```
while (flag == 0) {}  
data = 1;
```

```
assert r == 0;
```



Assertion
failure!

Behaviors Allowed in JMM #2

```
int data = flag = 0;
```

T1

```
r = data;  
flag = 1;  
assert r == 0;
```

T2

```
while (flag == 0) {}  
data = 1;
```

Requires returning future value or
reordering to trigger the assertion failure

Example #3

```
int x = y = 0;
```

T1

```
r1 = x;  
y = r1;
```

T2

```
r2 = y;  
if (r2 == 1) {  
    r3 = y;  
    x = r3;  
} else x = 1;
```

```
assert r2 == 0;
```


Example #3

int x = y = 0;

T1

r1 = x;
y = r1;

T2

r2 = y;
if (r2 == 1) {
 r3 = y;
 x = r3;
} else x = 1;

JMM disallows
r2 == 1 because
of causality
requirements

assert r2 == 0;

– Ševčík and Aspinall, ECOOP, 2008

Example #3

int x = y = 0;

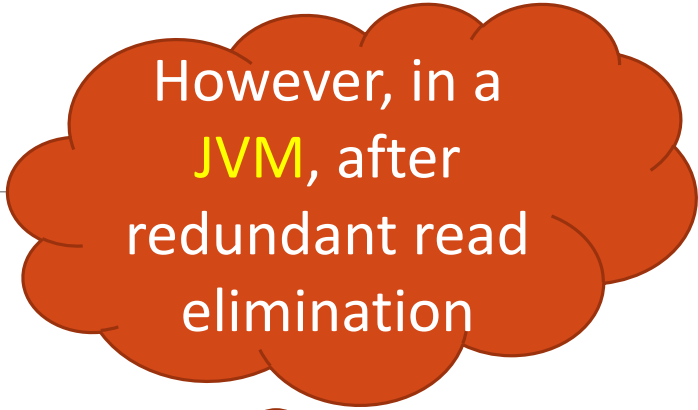
T1

r1 = x;
y = r1;

T2

r2 = y;
if (r2 == 1) {
 r3 = r2;
 x = r3;
} else x = 1;

assert r2 == 0;



However, in a
JVM, after
redundant read
elimination

Example #3

int x = y = 0;

T1

r1 = x;
y = r1;

T2

r2 = y;
if (r2 == 1) {
 r3 = r2;
 x = r3;
} else x = 1;

r2 = y;
if (r2 == 1)
 x = r2;
else x = 1;

assert r2 == 0;

However, in a
JVM, after
redundant read
elimination

Example #3

int x = y = 0;

T1

r1 = x;
y = r1;

However, in a
JVM, after
redundant read
elimination

T2

r2 = y;
if (r2 == 1) {
 r3 = r2;
 x = r3;
} else x = 1;

assert r2 == 0;

r2 = y;
if (r2 == 1)
 x = r2;
else x = 1;

r2 = y;
x = 1;

Example #3

```
int x = y = 0;
```

T1

```
r1 = x;  
y = r1;
```

T2

```
r2 = y;  
if (r2 == 1) {  
  r3 = r2;  
  x = r3;  
} else x = 1;
```

However, in a
JVM, after
redundant read
elimination

```
r2 = y;  
if (r2 == 1) {  
  x = r2;  
} else x = 1;
```

```
r2 = y;  
x = 1;
```

```
assert r2 == 0;
```

Assertion
failure
possible!

Behaviors Allowed by Memory Models and JVMs

DRF0 Memory Model

Java Memory Model

Typical JVMs

Behaviors Allowed by Memory Models and JVMs

DRF0 Memory Model

Java Memory Model

Typical JVMs

Unsatisfactory,
impractical to enforce

Exposing Behaviors of Example #3

int x = y = 0;

T1

r1 = x;
y = r1;

T2

r2 = y;
if (r2 == 1) {
 r3 = y;
 x = r3;
} else x = 1;

assert r2 == 0;



Consider
future value

Exposing Behaviors of Example #3

```
int x = y = 0;
```

T1

```
r1 = x; // r1 = 1  
y = r1; // y = 1
```

T2

```
r2 = y; // r2 = 1  
if (r2 == 1) {  
    r3 = y; // r3 = 1  
    x = r3; // x = 1  
} else x = 1;
```

```
assert r2 == 0;
```



Consider
future value

Exposing Behaviors of Example #3

```
int x = y = 0;
```

Consider
future value

T1

```
r1 = x; // r1 = 1  
y = r1; // y = 1
```

r1 = 1
justified!

Assertion
failure!

T2

```
r2 = y; // r2 = 1  
if (r2 == 1) {  
  r3 = y; // r3 = 1  
  x = r3; // x = 1  
} else x = 1;
```

```
assert r2 == 0;
```

Exposing Behaviors of Example #3

int x = y = 0;

T1

r1 = x;
y = r1;

T2

r2 = y;
if (r2 == 1) {
 r3 = y;
 x = r3;
} else x = 1;
assert r2 == 0;

Requires returning **future value** or
compiler optimization and reordering to
trigger the assertion failure

Exposing Behaviors with Dynamic Analyses

- Typical approaches
 - Simulate weak memory models behaviors [1,2,3]
 - Explore multiple thread interleavings [4, 5]

-
1. Adversarial Memory, Flanagan & Freund, PLDI'09
 2. Relaxer, Burnim et al, ISSTA'11
 3. Portend+, Kasikci et al, TOPLAS'15
 4. Replay Analysis, Narayanasamy et al, PLDI'07
 5. RaceFuzzer, Sen, PLDI'08

Exposing Behaviors with Dynamic Analyses

- Typical approaches
 - Simulate weak memory models behaviors [1,2,3]
 - Explore multiple thread interleavings [4, 5]
- Coverage Limitation
 - Return **stale values only**, not future values
 - **Cannot** expose assertion failures in Examples **#2, #3**

-
1. Adversarial Memory, Flanagan & Freund, PLDI'09
 2. Relaxer, Burnim et al, ISSTA'11
 3. Portend+, Kasikci et al, TOPLAS'15
 4. Replay Analysis, Narayanasamy et al, PLDI'07
 5. RaceFuzzer, Sen, PLDI'08

Relationship among memory models and exposed behaviors

DRF0 Memory Model

Java Memory Model

Existing Dynamic Analyses

Typical JVMs

Relationship among memory models and exposed behaviors

DRF0 Memory Model

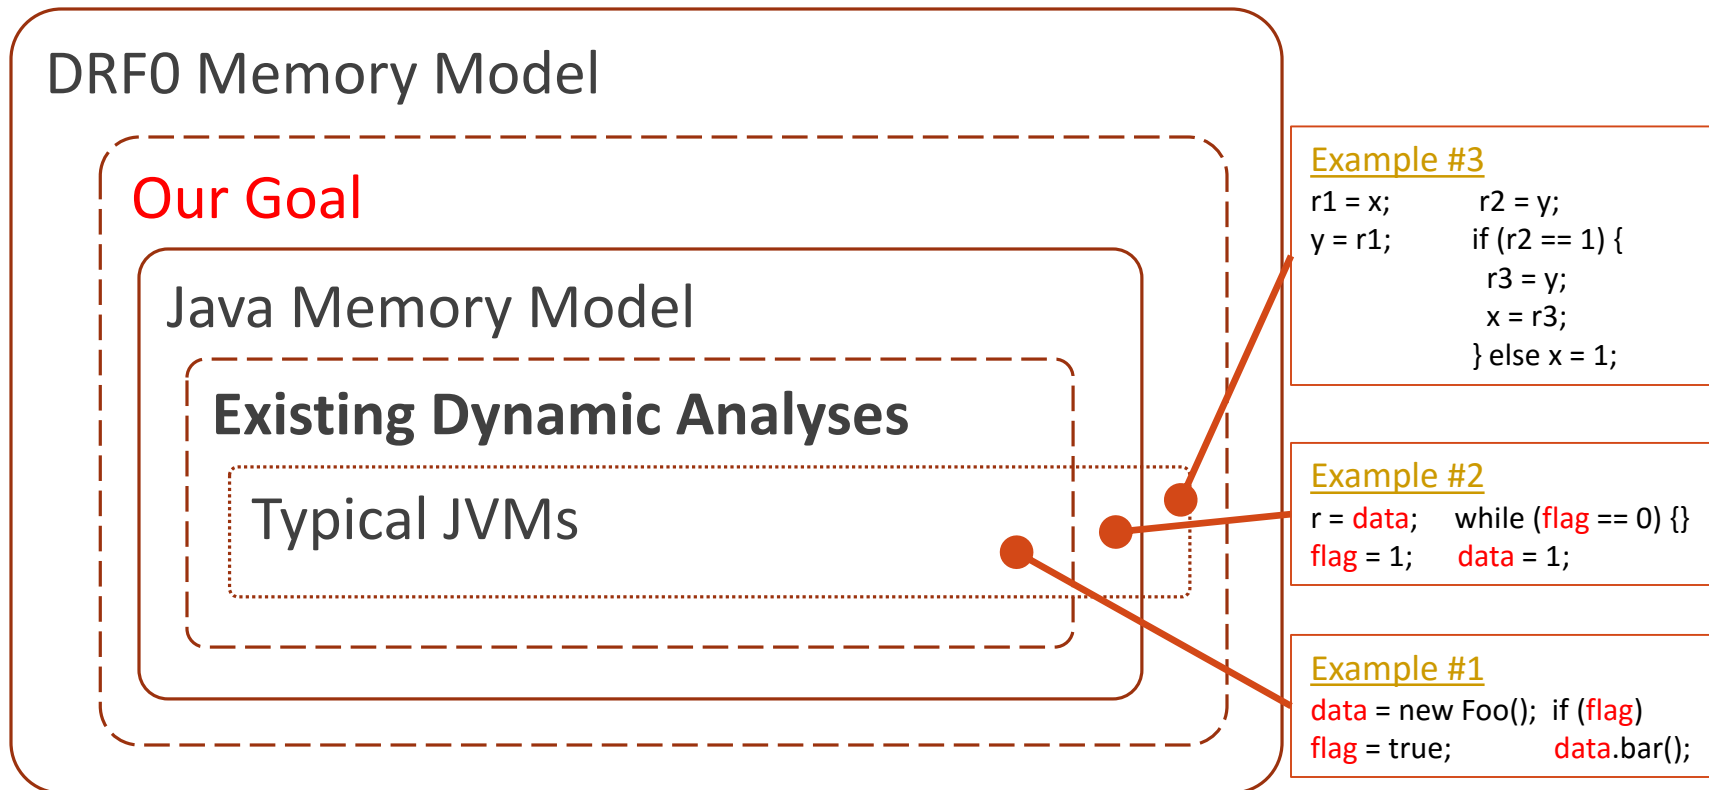
Our Goal

Java Memory Model

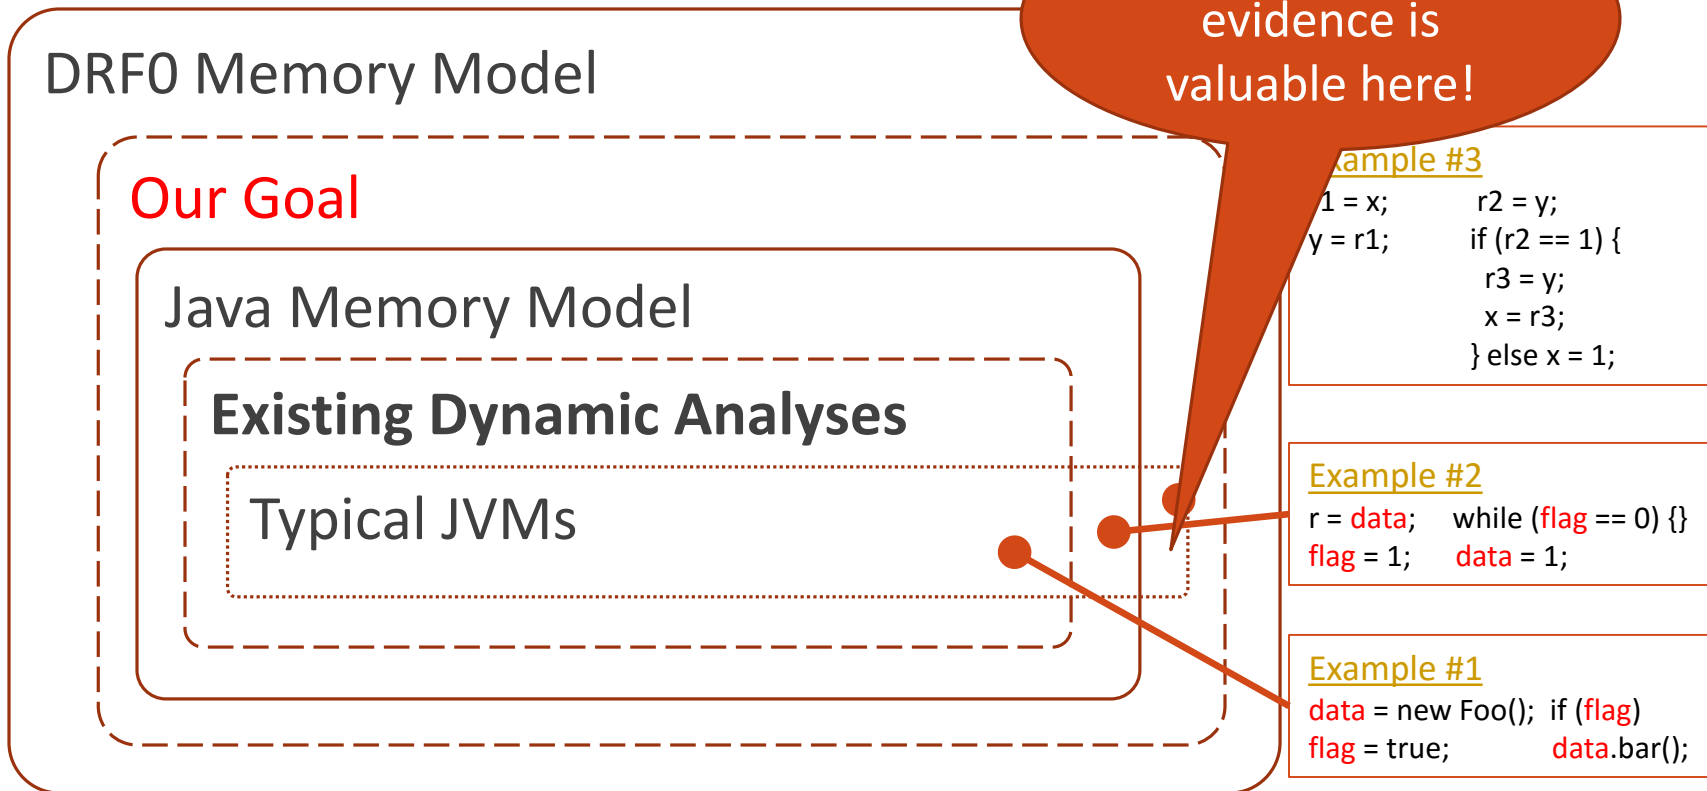
Existing Dynamic Analyses

Typical JVMs

Relationship among memory models and exposed behaviors



Relationship among memory models and exposed behaviors



Outline

- Memory Models and Behaviors of Data Races
- Design
 - Prescient Memory (PM)
 - PM-profiler
 - PM Workflow
- Evaluation

Prescient Memory: Key Idea

- *Speculatively* “guess” a future value at a **load**
- *Validate* the speculative value at a **later store**

Prescient Memory: Key Idea

- *Speculatively* “guess” a future value at a **load**
- **Validate** the speculative value at a **later store**

Returning Future Values is Tricky

```
int x = y = 0;
```

T1

```
r1 = x;  
y = r1;
```

T2

```
r2 = y;  
if (r2 == 0)  
    x = 1;
```

```
assert r1 == 0 || r2 == 0;
```

Returning Future Values is Tricky

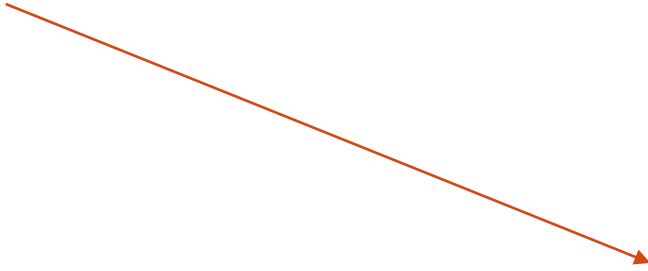
```
int x = y = 0;
```

T1

```
r1 = x;  
y = r1;
```

T2

```
r2 = y;  
if (r2 == 0)  
    x = 1;
```



```
assert r1 == 0 || r2 == 0;
```

Returning Future Values is Tricky

```
int x = y = 0;
```

T1

```
r1 = x; // r1 = 1  
y = r1; // y = 1
```

T2

```
r2 = y; // r2 = 1  
if (r2 == 0)  
    x = 1;
```

```
assert r1 == 0 || r2 == 0;
```

Returning Future Values is Tricky

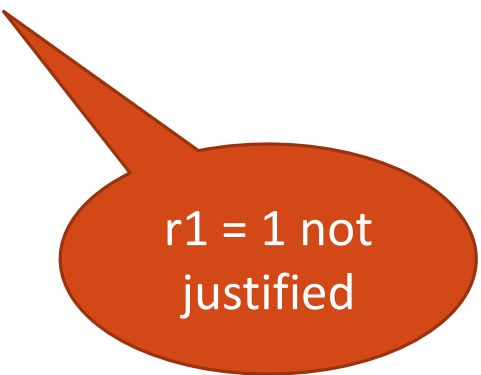
```
int x = y = 0;
```

T1

```
r1 = x; // r1 = 1  
y = r1; // y = 1
```

T2

```
r2 = y; // r2 = 1  
if (r2 == 0)  
    x = 1;
```



r1 = 1 not
justified

```
assert r1 == 0 || r2 == 0;
```


Returning Future Values is Tricky

```
int x = y = 0;
```

T1

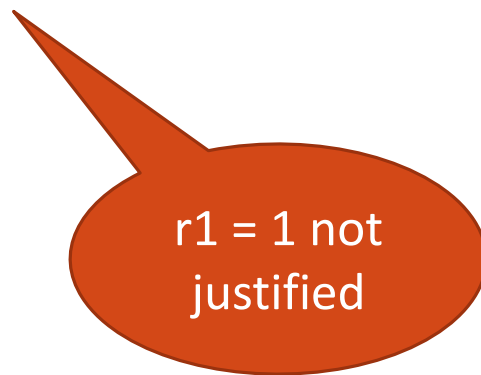
```
r1 = x; // r1 = 1  
y = r1; // y = 1
```



Invalid
execution!

T2

```
r2 = y; // r2 = 1  
if (r2 == 0)  
    x = 1;
```



r1 = 1 not
justified

```
assert r1 == 0 || r2 == 0;
```

Returning Future Values is Tricky

```
int x = y = 0;
```

T1

```
r1 = x;  
y = r1;
```

T2

```
r2 = y;  
if (r2 == 0)  
    x = 1;
```



Should
never fail!

```
assert r1 == 0 || r2 == 0;
```

Returning Future Values is Tricky

```
int x = y = 0;
```

T1

```
r1 = x;  
y = r1;
```

T2

```
r2 = y;  
if (r2 == 0)  
    x = 1;
```

```
assert r1 == 0 || r2 == 0;
```

Validating speculative values is necessary to
prevent nonsensical results


Prescient Memory: Key Idea

- *Speculatively* “guess” a future value at a **load**
- *Validate* the speculative value at a **later store**

Prescient Memory: Key Idea


- *Speculatively* “guess” a future value at a **load**

- *Validate* the speculative value at a **later store**



Prescient Memory: Key Idea

- *Speculatively* “guess” a future value at a **load**
 - Maintain a per-variable **speculative read history**
 - Records <logical timestamp, speculative value>
- *Validate* the speculative value at a **later store**



PM Example

```
int x = y = 0;  
S[x] = ∅
```

T1 Timestamp: K_1

```
1: r = x;  
2: y = 1;
```

T2 Timestamp: K_2

```
3: while (y == 0) {}
```

```
4: x = 1;
```

```
assert r == 0;
```

PM Example

```
int x = y = 0;  
S[x] =  $\emptyset$ 
```

T1 Timestamp: K_1

```
1: r = x;    1  $\leftarrow$  predict(...)    // guess value 1  
2: y = 1;    S[x] = {<K1, 1>}
```

T2 Timestamp: K_2

```
3: while (y == 0) {}
```

```
4: x = 1;
```

```
assert r == 0;
```


PM Example

```
int x = y = 0;  
S[x] =  $\emptyset$ 
```

T1 Timestamp: K_1

```
1: r = x; 1  $\leftarrow$  predict(...) // guess value 1  
2: y = 1; S[x] = {< $K_1$ , 1>}
```

T2 Timestamp: K_2

```
3: while (y == 0) {}
```

```
validate S[x]:  
   $K_1 \not\sqsubseteq K_2 \ \&\& \ 1 == 1$ 
```

```
4: x = 1;
```

```
assert r == 0;
```

1 is a valid future value!

Challenges

- How to guess a future value?



predict(...) ?

Challenges

- How to guess a future value?
 - Which *load* should return a future value?
 - What *value* should be returned?

Challenges

- How to guess a future value?
 - Which *load* should return a future value?
 - What *value* should be returned?
- Solution
 - *Profile* possible future values in a prior run

Profiling Future Values

Helper Dynamic Analysis: PM-profiler

- Maintains a per-variable **concrete read history**
- At a load, records:
 - <logical timestamp, instruction ID, set of visible values>

Profiling Future Values

Helper Dynamic Analysis: PM-profiler

- At a store, detects:


Potential **future value**
for a previous load




Store **paces** with the
previous load

Store writes a value
distinct from visible values
of the previous load

Prescient Memory Workflow

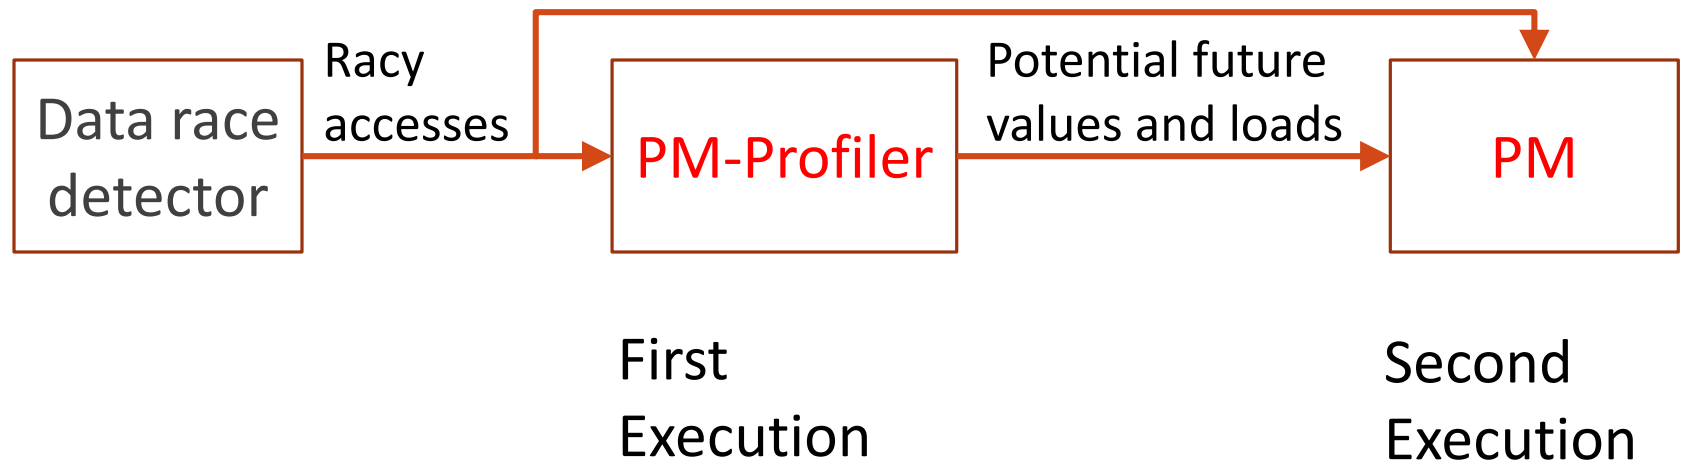


Prescient Memory Workflow



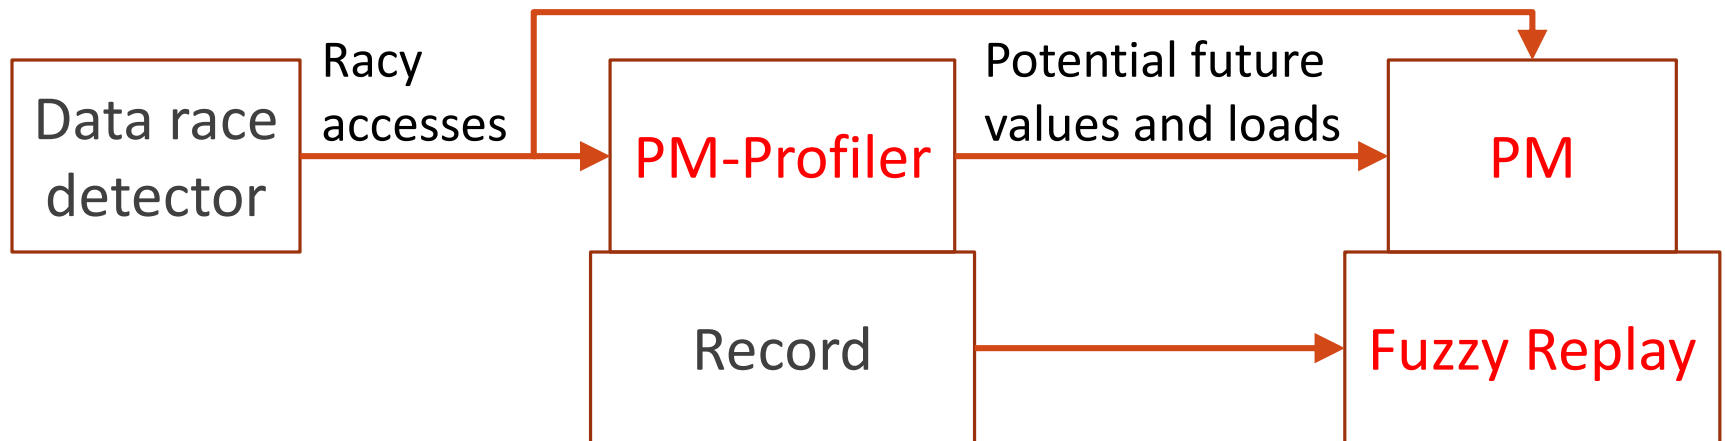
Run-to-run **nondeterminism** affects validatable future values

Prescient Memory Workflow




- Run-to-run **nondeterminism** affects validatable future values
- Solution: record and replay

Prescient Memory Workflow



Prescient Memory Workflow



Returning a future value could **diverge** from the record execution

- Best-effort, fuzzy replay

Outline

- Memory Models and Behaviors of Data Races
- Design
 - Prescient Memory (PM)
 - PM-profiler
 - PM Workflow
- Evaluation

Methodology and Implementation

- Compare with

Adversarial Memory (AM) [Flanagan & Freund, PLDI'09]: a dynamic analysis that only uses **stale** values

Methodology and Implementation

- Compare with

Adversarial Memory (AM) [Flanagan & Freund, PLDI'09]: a dynamic analysis that only uses **stale** values

- Platform

Jikes RVM 3.1.3

DaCapo Benchmark 2006, 2009 and SPEC JBB 2000 & 2005

4-Core Intel Core i5-2500

Record and Replay [Replay, Bond et al. PPPJ'15]

Methodology and Implementation

- Compare with

Adversarial Memory (AM) [Flanagan & Freund, PLDI'09]: a dynamic analysis that only uses **stale** values

- Platform

Jikes RVM 3.1.3

DaCapo Benchmark 2006, 2009 and SPEC JBB 2000 & 2005

4-Core Intel Core i5-2500

Record and Replay [Replay, Bond et al. PPPJ'15]

- Implementation limitation

Does not support reference-type fields

Exposed Erroneous Behaviors

Program	AM	PM
hsqldb	Non-termination	Data corruption
hsqldb	None	Performance bug
avroa	Data corruption	Data corruption
lusearch (GNU Classpath)	Performance bug	None
sunflow	Null ptr exception	Null ptr exception
jbb2000	Non-termination	Data corruption
jbb2000	Data corruption	Data corruption
jbb2005 (GNU Classpath)	Data corruption	Data corruption
jbb2005 (GNU Classpath)	Data corruption	None

Exposed Erroneous Behaviors

Program	AM	PM
hsqldb	Non-termination	Data corruption
hsqldb	None	Performance bug
avroa	Data corruption	Data corruption
lusearch (GNU Classpath)	Performance bug	None
sunflow	Null ptr exception	Null ptr exception
jbb2000	Non-termination	Data corruption
jbb2000	Data corruption	Data corruption
jbb2005 (GNU Classpath)	Data corruption	Data corruption
jbb2005 (GNU Classpath)	Data corruption	None

PM found 3 new erroneous behaviors!

Exposed Erroneous Behaviors

Program	AM	PM	
hsqldb	Non-termination	Data corruption	✓
hsqldb	None	Performance bug	
avroa	Data corruption	Data corruption	✓
lusearch (GNU Classpath)	Performance bug	None	✗
sunflow	Null ptr exception	Null ptr exception	✓
jbb2000	Non-termination	Data corruption	✓
jbb2000	Data corruption	Data corruption	✓
jbb2005 (GNU Classpath)	Data corruption	Data corruption	✓
jbb2005 (GNU Classpath)	Data corruption	None	✗

PM exposes most bugs that AM found.

Exposed Erroneous Behaviors

Program	AM	PM
hsqldb	Non-termination	Data corruption
hsqldb	None	Performance bug
avroa	Data corruption	Data corruption
lusearch (GNU Classpath)	Performance bug	None
sunflow	Null ptr exception	Null ptr exception
jbb2000	Non-termination	Data corruption
jbb2000	Data corruption	Data corruption
jbb2005 (GNU Classpath)	Data corruption	Data corruption
jbb2005 (GNU Classpath)	Data corruption	None

Paper contains detailed analysis of each bug.

Conclusion

- **First** dynamic analysis to expose legal behaviors due to **future values** in **large, real** programs
- **Successfully** found **new** harmful behaviors due to future values in real programs
- Reaffirms that “**benign**” races are **harmful**
- Helps future **revisions** to language **specifications** by finding evidence of **controversial** behaviors in real programs