

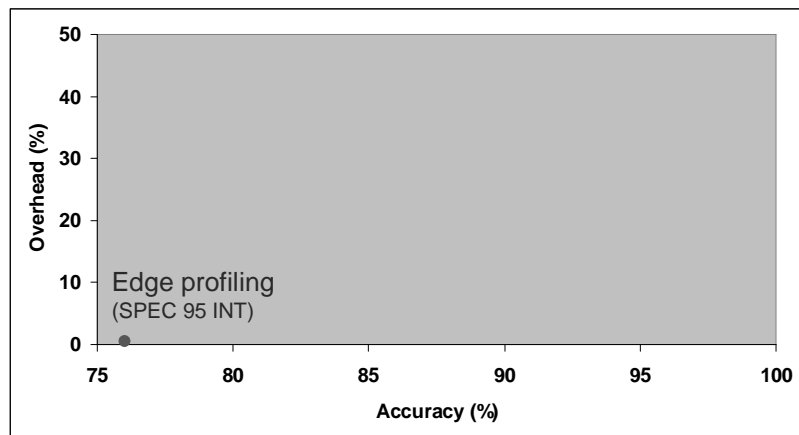
Targeted Path Profiling: Lower Overhead Path Profiling for Staged Dynamic Optimization Systems

Rahul Joshi, UIUC
Michael Bond*, UT Austin
Craig Zilles, UIUC

Path information is useful

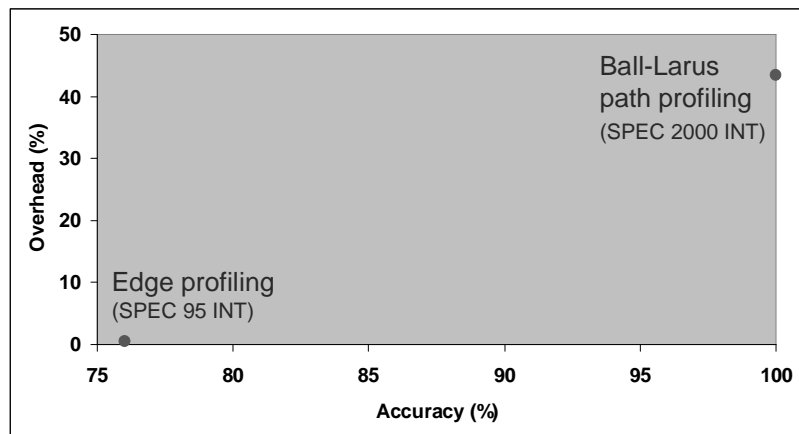
- Enlarges scope of optimizations
 - Superblock formation
 - Hyperblock formation
- Improves other optimizations
 - Code scheduling and register allocation
 - Dataflow analysis
 - Software pipelining
 - Code layout
 - Static branch prediction

Overhead vs. accuracy



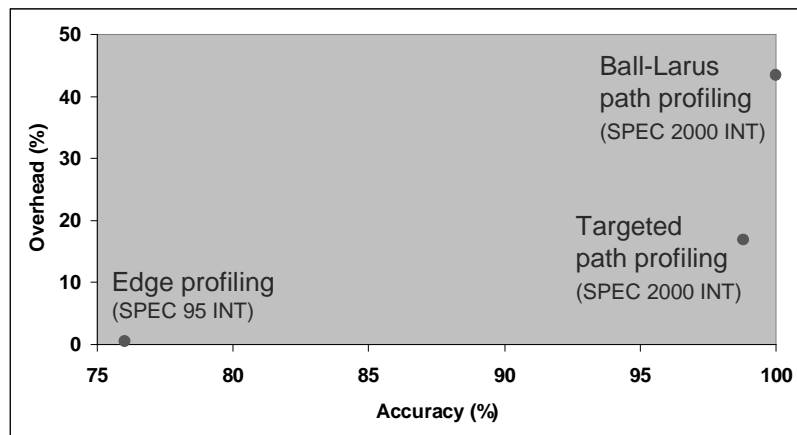
3

Overhead vs. accuracy




4

Overhead vs. accuracy



5

Overhead vs. accuracy



6

Outline

- Background
 - Staged dynamic optimization and profile-guided profiling
 - Ball-Larus path profiling
 - Opportunities for reducing overhead
- Targeted path profiling
- Results
 - Overhead and accuracy

7


Staged dynamic optimization

Stage 0

Static
optimizations

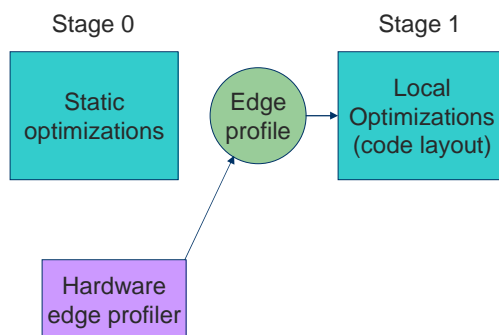
8

Staged dynamic optimization



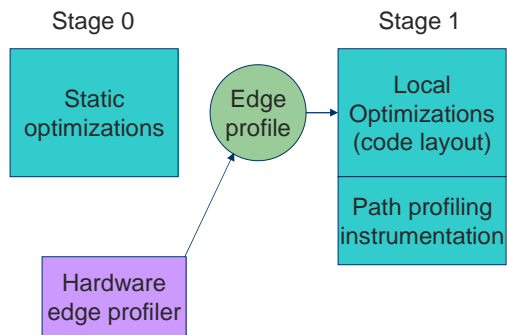
9

Staged dynamic optimization



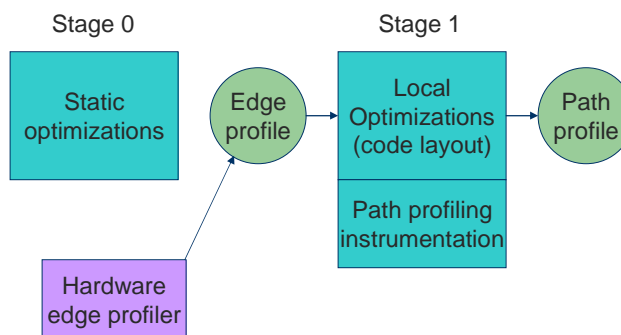
10

Staged dynamic optimization



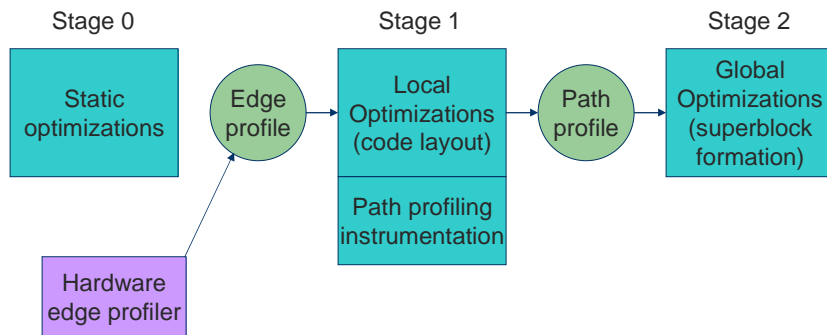
11

Staged dynamic optimization



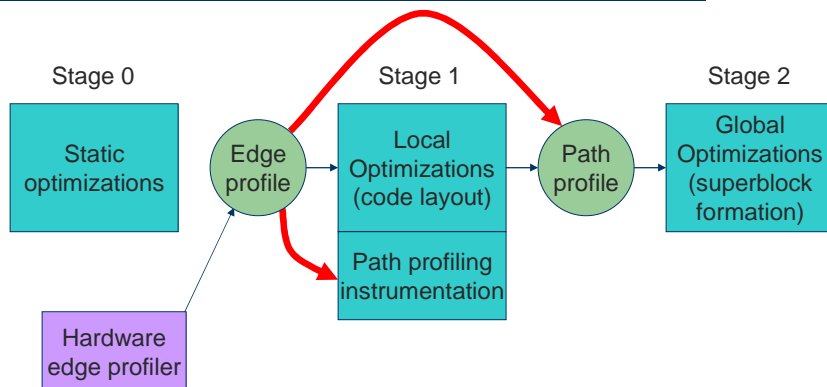
12

Staged dynamic optimization



13

Profile-guided profiling



14


Ball-Larus path profiling

- Acyclic, intraprocedural paths
- Handles cyclic CFGs
 - Paths end at loop back edges
- Each path computes unique integer

15

Ball-Larus path profiling


- 4 paths



16

Ball-Larus path profiling


- 4 paths
- Each path computes unique integer



17

Ball-Larus path profiling


- 4 paths
- Each path computes unique integer
- Path 0



18

Ball-Larus path profiling


- 4 paths
- Each path computes unique integer
- Path 0
- Path 1



19

Ball-Larus path profiling


- 4 paths
- Each path computes unique integer
- Path 0
- Path 1
- Path 2



20

Ball-Larus path profiling


- 4 paths
- Each path computes unique integer
- Path 0
- Path 1
- Path 2
- Path 3



21

Ball-Larus path profiling

- **r**: path register
- **count**: array of path frequencies



22

Overhead in Ball-Larus path profiling

	SPEC 95	SPEC 2000
gcc	96%	87%
INT Avg	41%	43%
FP Avg	12%	22%
Overall Avg	28%	37%

23

Overhead in Ball-Larus path profiling

	SPEC 95	SPEC 2000
gcc	96%	87%
INT Avg	41%	43%
FP Avg	12%	22%
Overall Avg	28%	37%

- Opportunities for reducing overhead?
 - When there are many paths
 - When edge profile gives perfect path profile


24

Routines with many paths

- Many *possible* paths
 - Exponential in number of edges
 - Can't use array of counters
- Number of *taken paths* small
 - Ball-Larus uses hash table
 - Hash function call expensive
 - Hashed path ~5 times overhead


25

Edge profile gives perfect path profile



26


Edge profile gives perfect path profile



27

Edge profile gives perfect path profile

- An *obvious path* contains an edge that is only on that path
 - Path uniquely identified by edge
 - Path freq = edge freq
- If all paths obvious, edge profile gives perfect path profile



28

Outline

- Background
 - Staged dynamic optimization and profile-guided profiling
 - Ball-Larus path profiling
 - Opportunities for reducing overhead
- Targeted path profiling
- Results
 - Overhead and accuracy

29

Targeted path profiling


- Profile-guided profiling
 - Use existing edge profile
- Exploits opportunities for reducing overhead
 - When there are many paths
 - Remove cold edges
 - When edge profile gives perfect path profile
 - Don't instrument obvious routines and loops

30

Removing cold edges

- Examine relative execution frequency of each branch

```
if (relFreq < threshold)  
    edge is cold
```





31

Removing cold edges

- Examine relative execution frequency of each branch

```
if (relFreq < threshold)  
    edge is cold
```





32

Removing cold edges

- Examine relative execution frequency of each branch


```
if (relFreq < threshold)
    edge is cold
```



33

Removing cold edges


- A path that contains a cold edge is a *cold path*
- Removing an edge may halve number of paths



34

Removing cold edges

- A path that contains a cold edge is a *cold path*
- Removing an edge may halve number of paths
- Number of paths: $16 \rightarrow 4$




35

Removing cold edges

- A path that contains a cold edge is a *cold path*
- Removing an edge may halve number of paths
- Number of paths: $16 \rightarrow 4$


Goal: hashed \rightarrow non-hashed



36

Removing cold edges


- Remaining paths potentially hot
- 4 paths $\rightarrow [0, 3]$



37


Removing cold edges

- Remaining paths potentially hot
- 4 paths $\rightarrow [0, 3]$




38

- What if cold edge taken?



Removing cold edges


- What if cold edge taken?
- Cold edges poison path



40

Removing cold edges

- What if cold edge taken?
- Cold edges poison path
- Instrumentation checks for poisoned path



41


Checking for poison

```
if (r poisoned)
  cold_counter++
else
  count[r]++
```

42

Obvious routines


- All paths obvious
- We don't instrument obvious routines
- Edge profile gives perfect path profile



43

Obvious loops


- Loop with obvious body
- Don't instrument obvious loops with high average trip counts
- Edge profile yields high-accuracy path profile



44

Obvious loops

- Loop with obvious body
- Don't instrument obvious loops with high average trip counts
- Edge profile yields high-accuracy path profile



45

Summary of our techniques

- Remove cold edges
 - Eliminates many cold paths
 - Count paths with array (instead of hash table)
- Don't instrument obvious routines and loops
 - Edge profile derives path profile

46

Outline

- Background
 - Staged dynamic optimization and profile-guided profiling
 - Ball-Larus path profiling
 - Opportunities for reducing overhead
- Targeted path profiling
- Results
 - Overhead and accuracy


47

Implementation

- Static profiling
- **PP**: tool for path profiling
- **TPP**: tool for targeted path profiling
- Tools instrument native SPARC executables
 - SPEC 95 **ref**
 - SPEC 2000 **ref**

48

Results: SPEC 2000 INT



49

Where does benefit come from?

- Cold path elimination alone: 60%
- Add obvious path elimination: + 40%
- Little benefit from obvious path elimination alone

50

Related work

- Dynamo [Bala et al. '00]
 - Successful online path-guided optimization
 - “Bails out” when no dominant path
- Instrumentation sampling [Arnold & Ryder '01]
 - Orthogonal to targeted path profiling
- Selective path profiling [Apiwattanapong & Harrold '02]
 - Useful when only a few paths of interest

51

Summary

- Profile-guided profiling in a staged dynamic optimization system
- Two synergistic techniques
 - Remove cold paths
 - Don't instrument obvious routines and loops
- Reduces overhead by half (SPEC 95) to two-thirds (SPEC 2000)
- High accuracy: ~99%

52